

Introduction to Programming (C/C++)

09: Performance Fun!

Huanchen Zhang



清华大学
Tsinghua University



交叉信息研究院
Institute for Interdisciplinary
Information Sciences

Measuring Time

Turn On Compiler Optimization

-O2/-O3

Command-Line Arguments

```
$ ./myProg 10 happy
```

Command-Line Arguments

```
$ ./myProg 10 happy
```

```
int argc
```



3

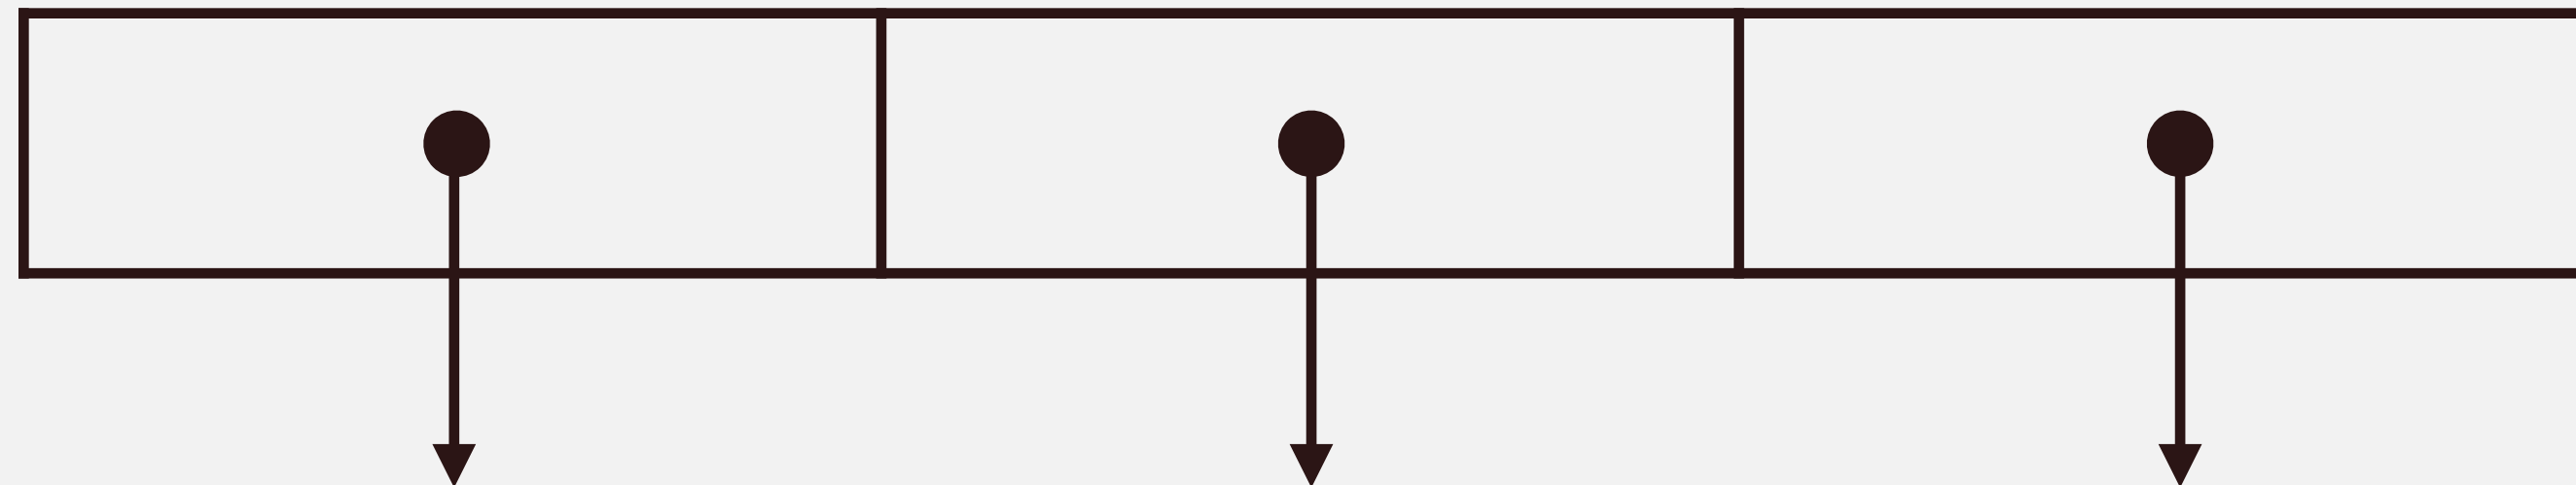
Command-Line Arguments

```
$ ./myProg 10 happy
```

`int argc`

3

`char *argv[]`



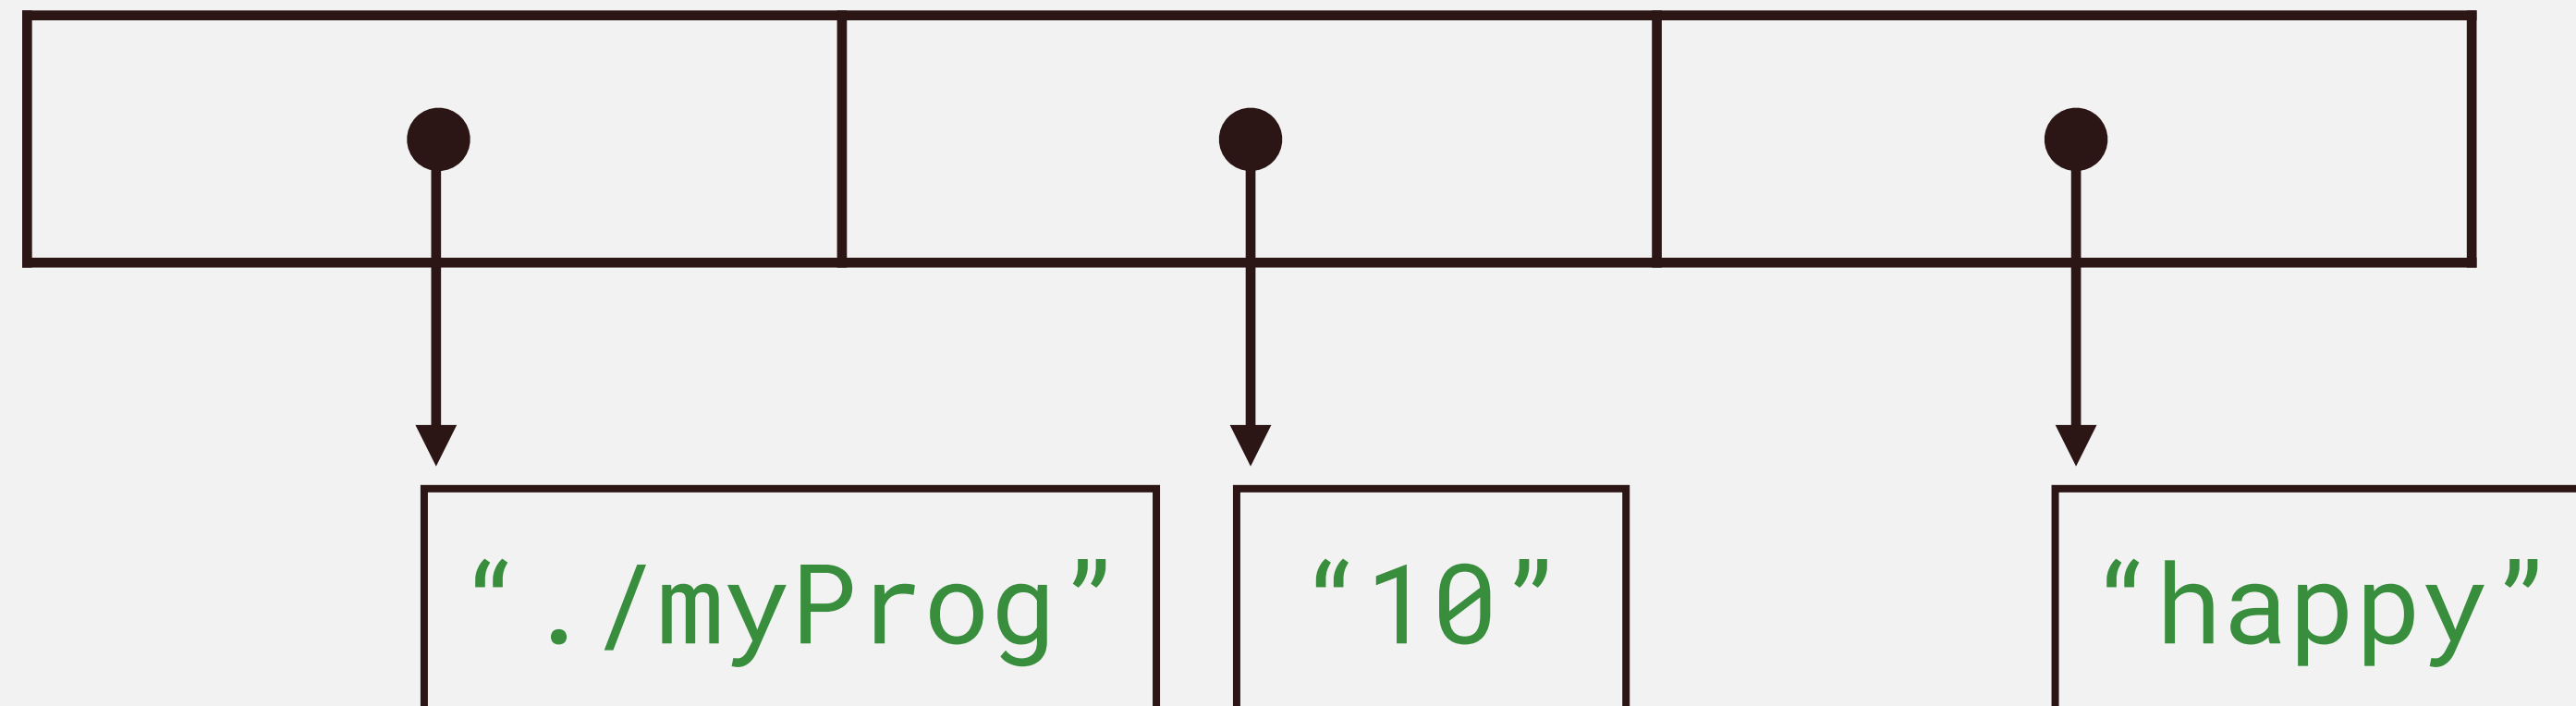
Command-Line Arguments

```
$ ./myProg 10 happy
```

`int argc`

3

`char *argv[]`



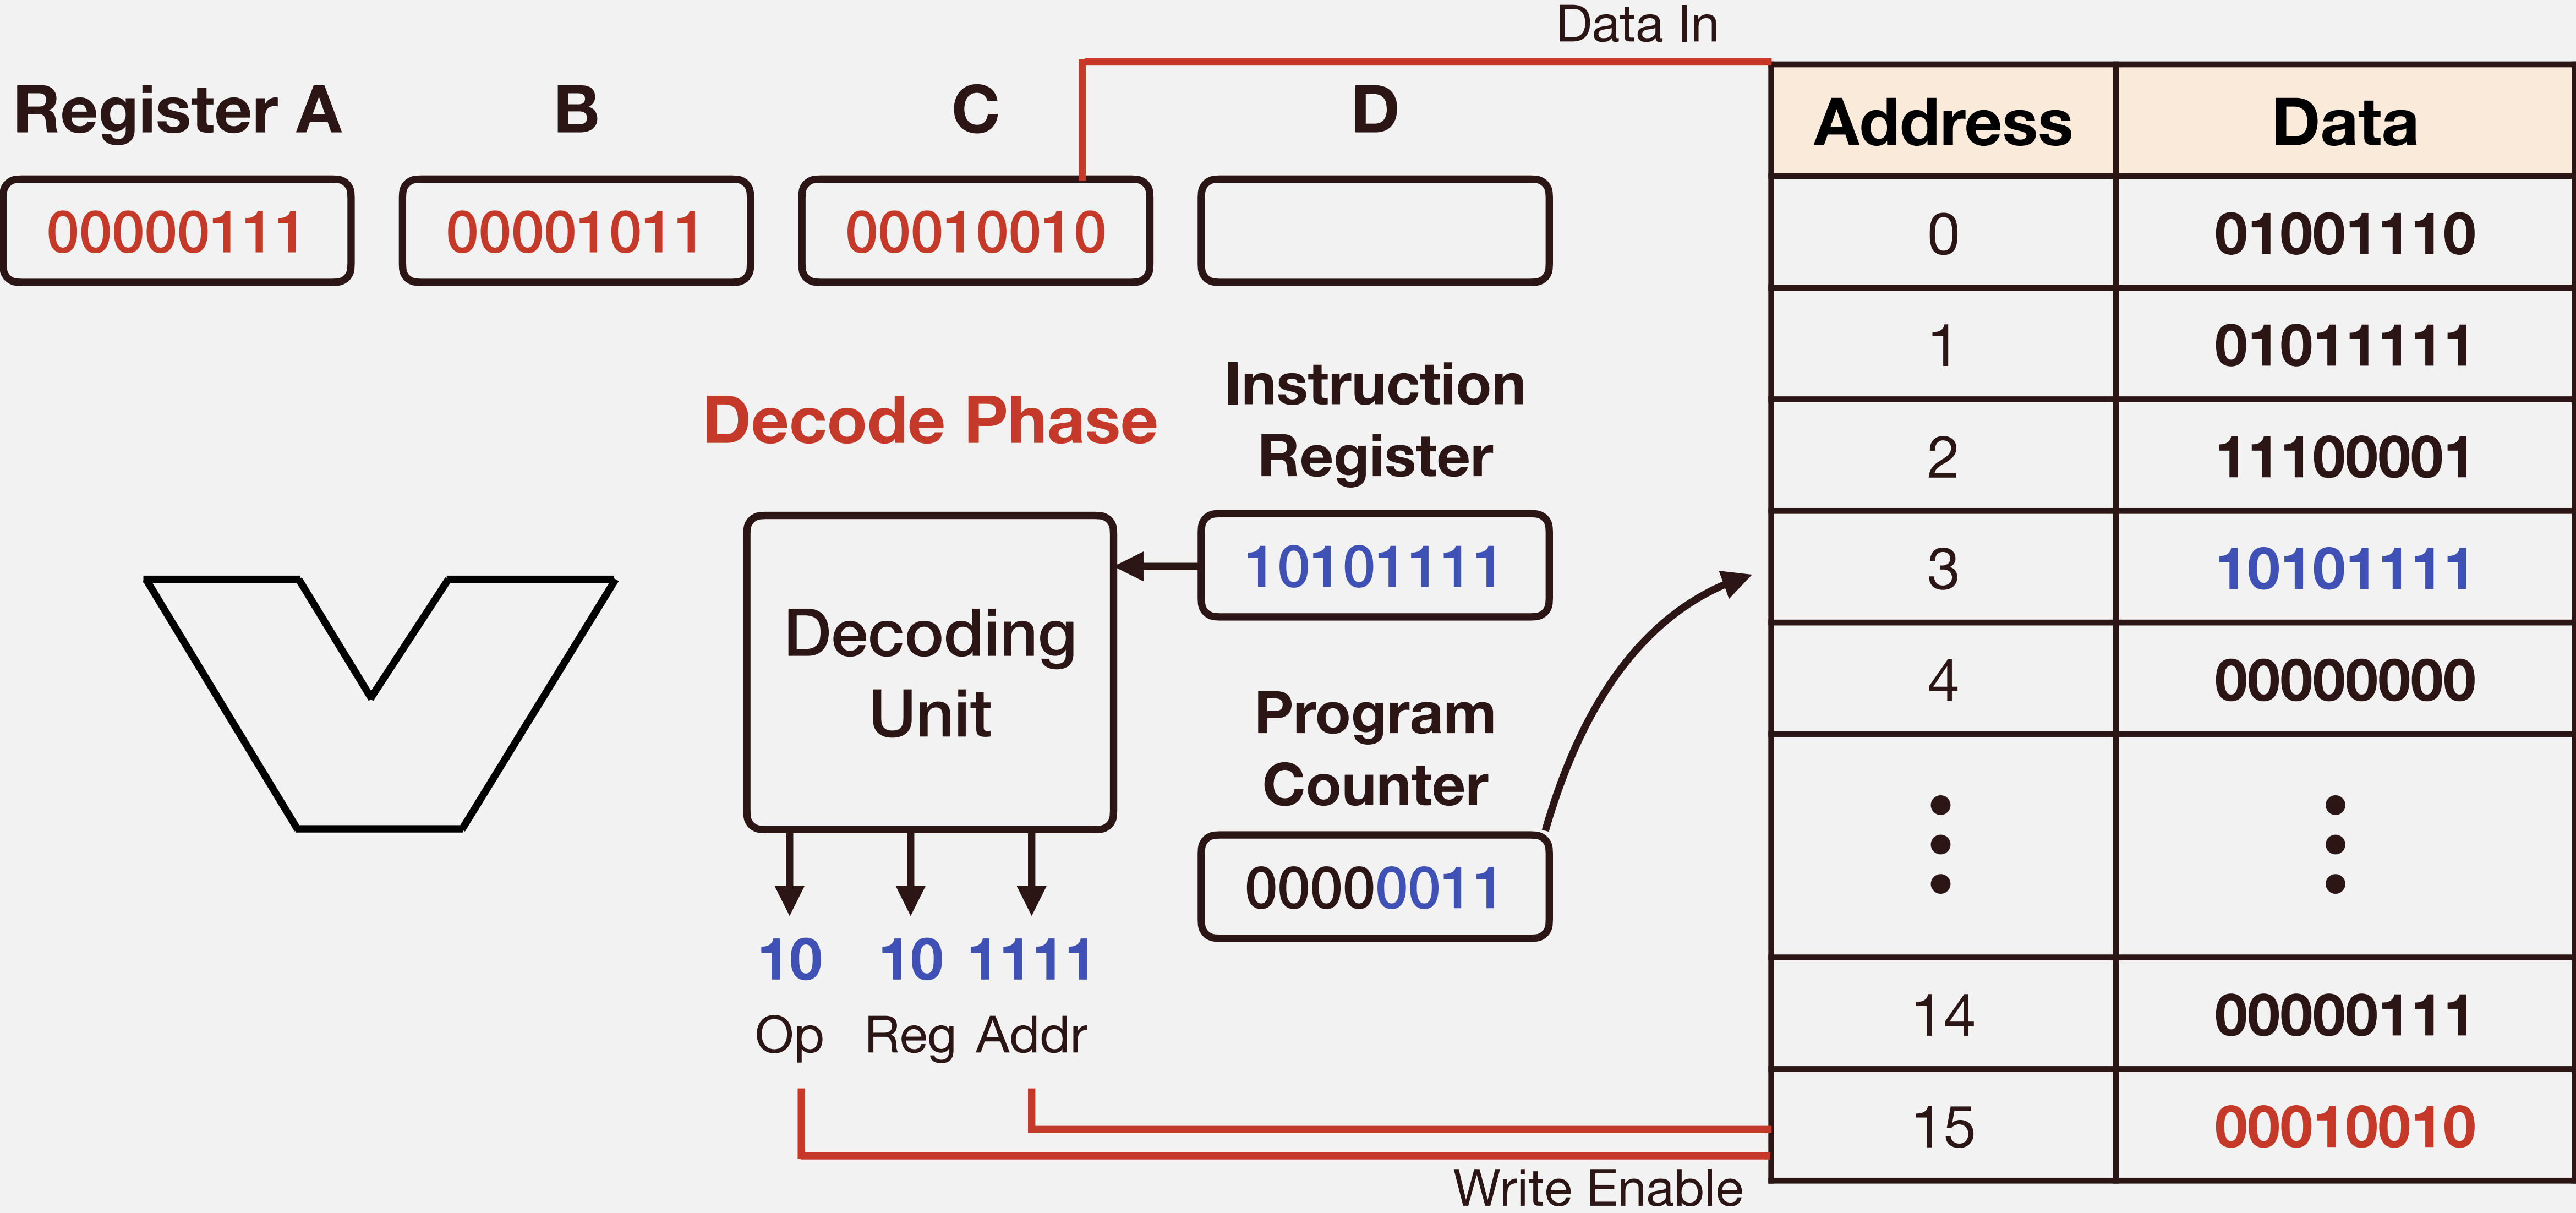
Algorithms Matter a Lot

Algorithms Matter a Lot

At Large Scale

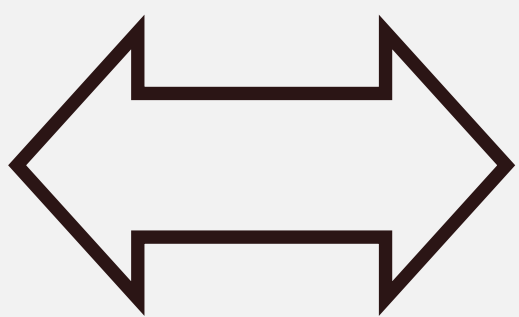
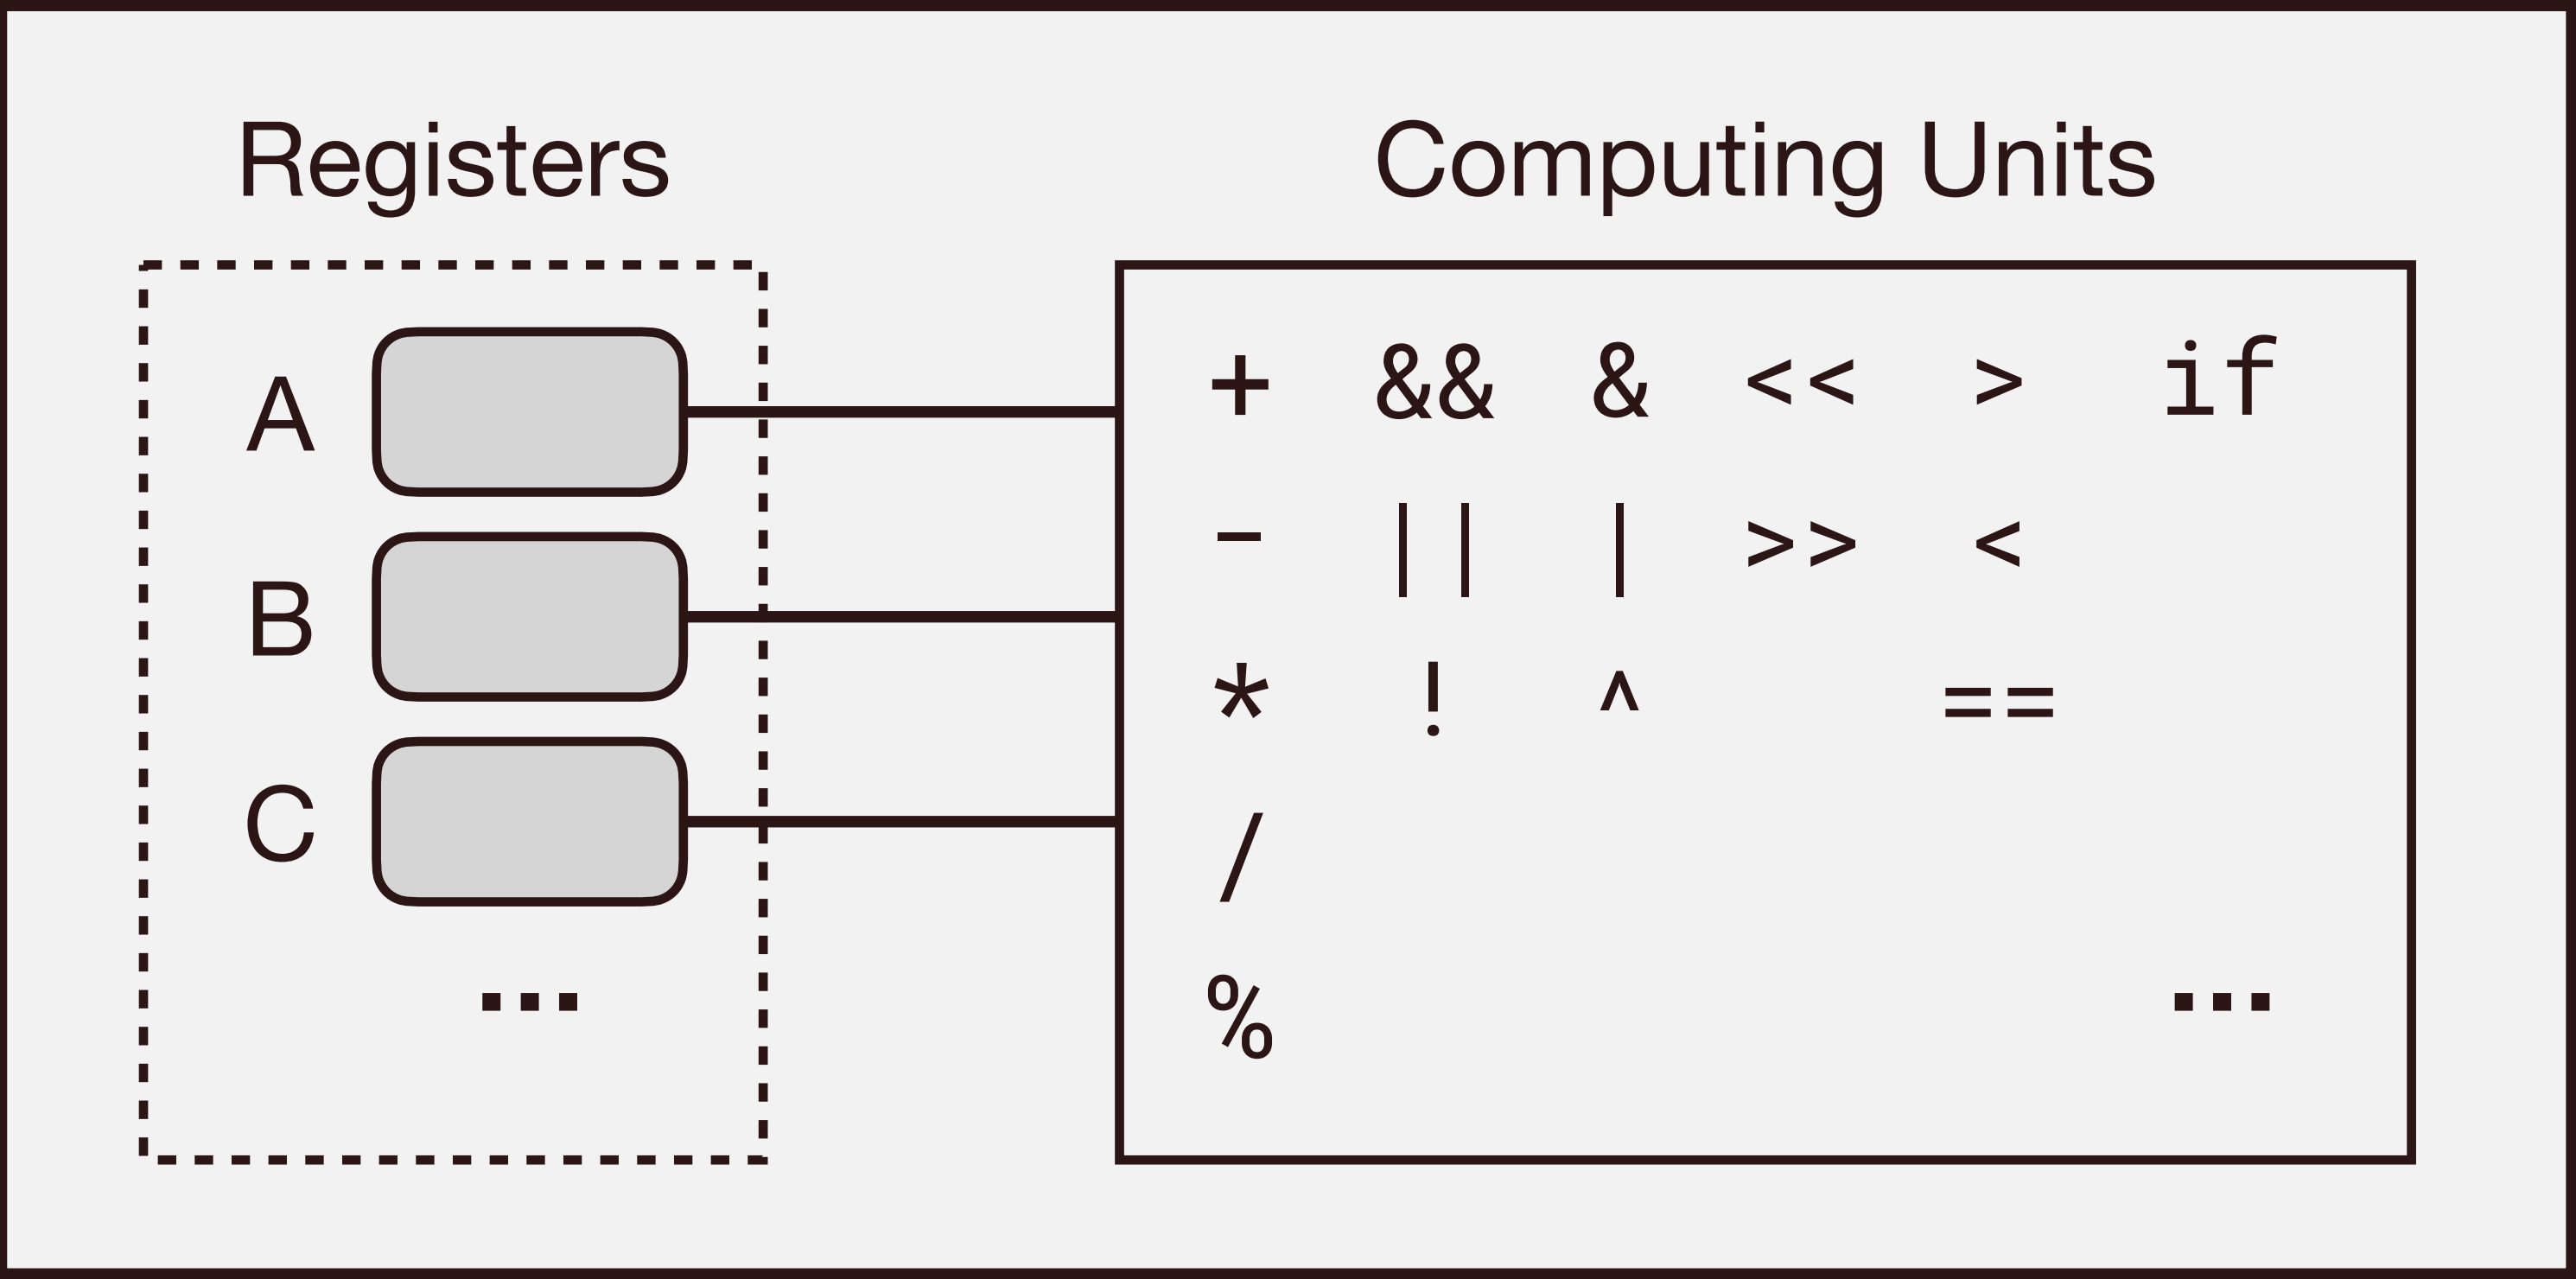
**Don't Do More Work Than
You Have To**

Building a Central Processing Unit (CPU)



Know Your Architecture

CPU

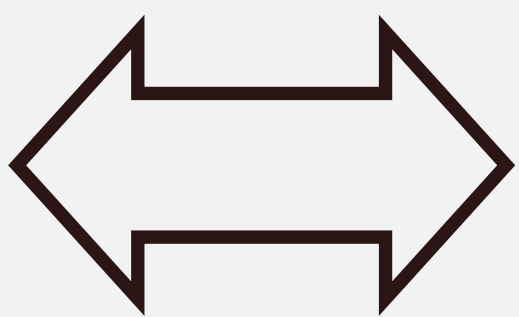
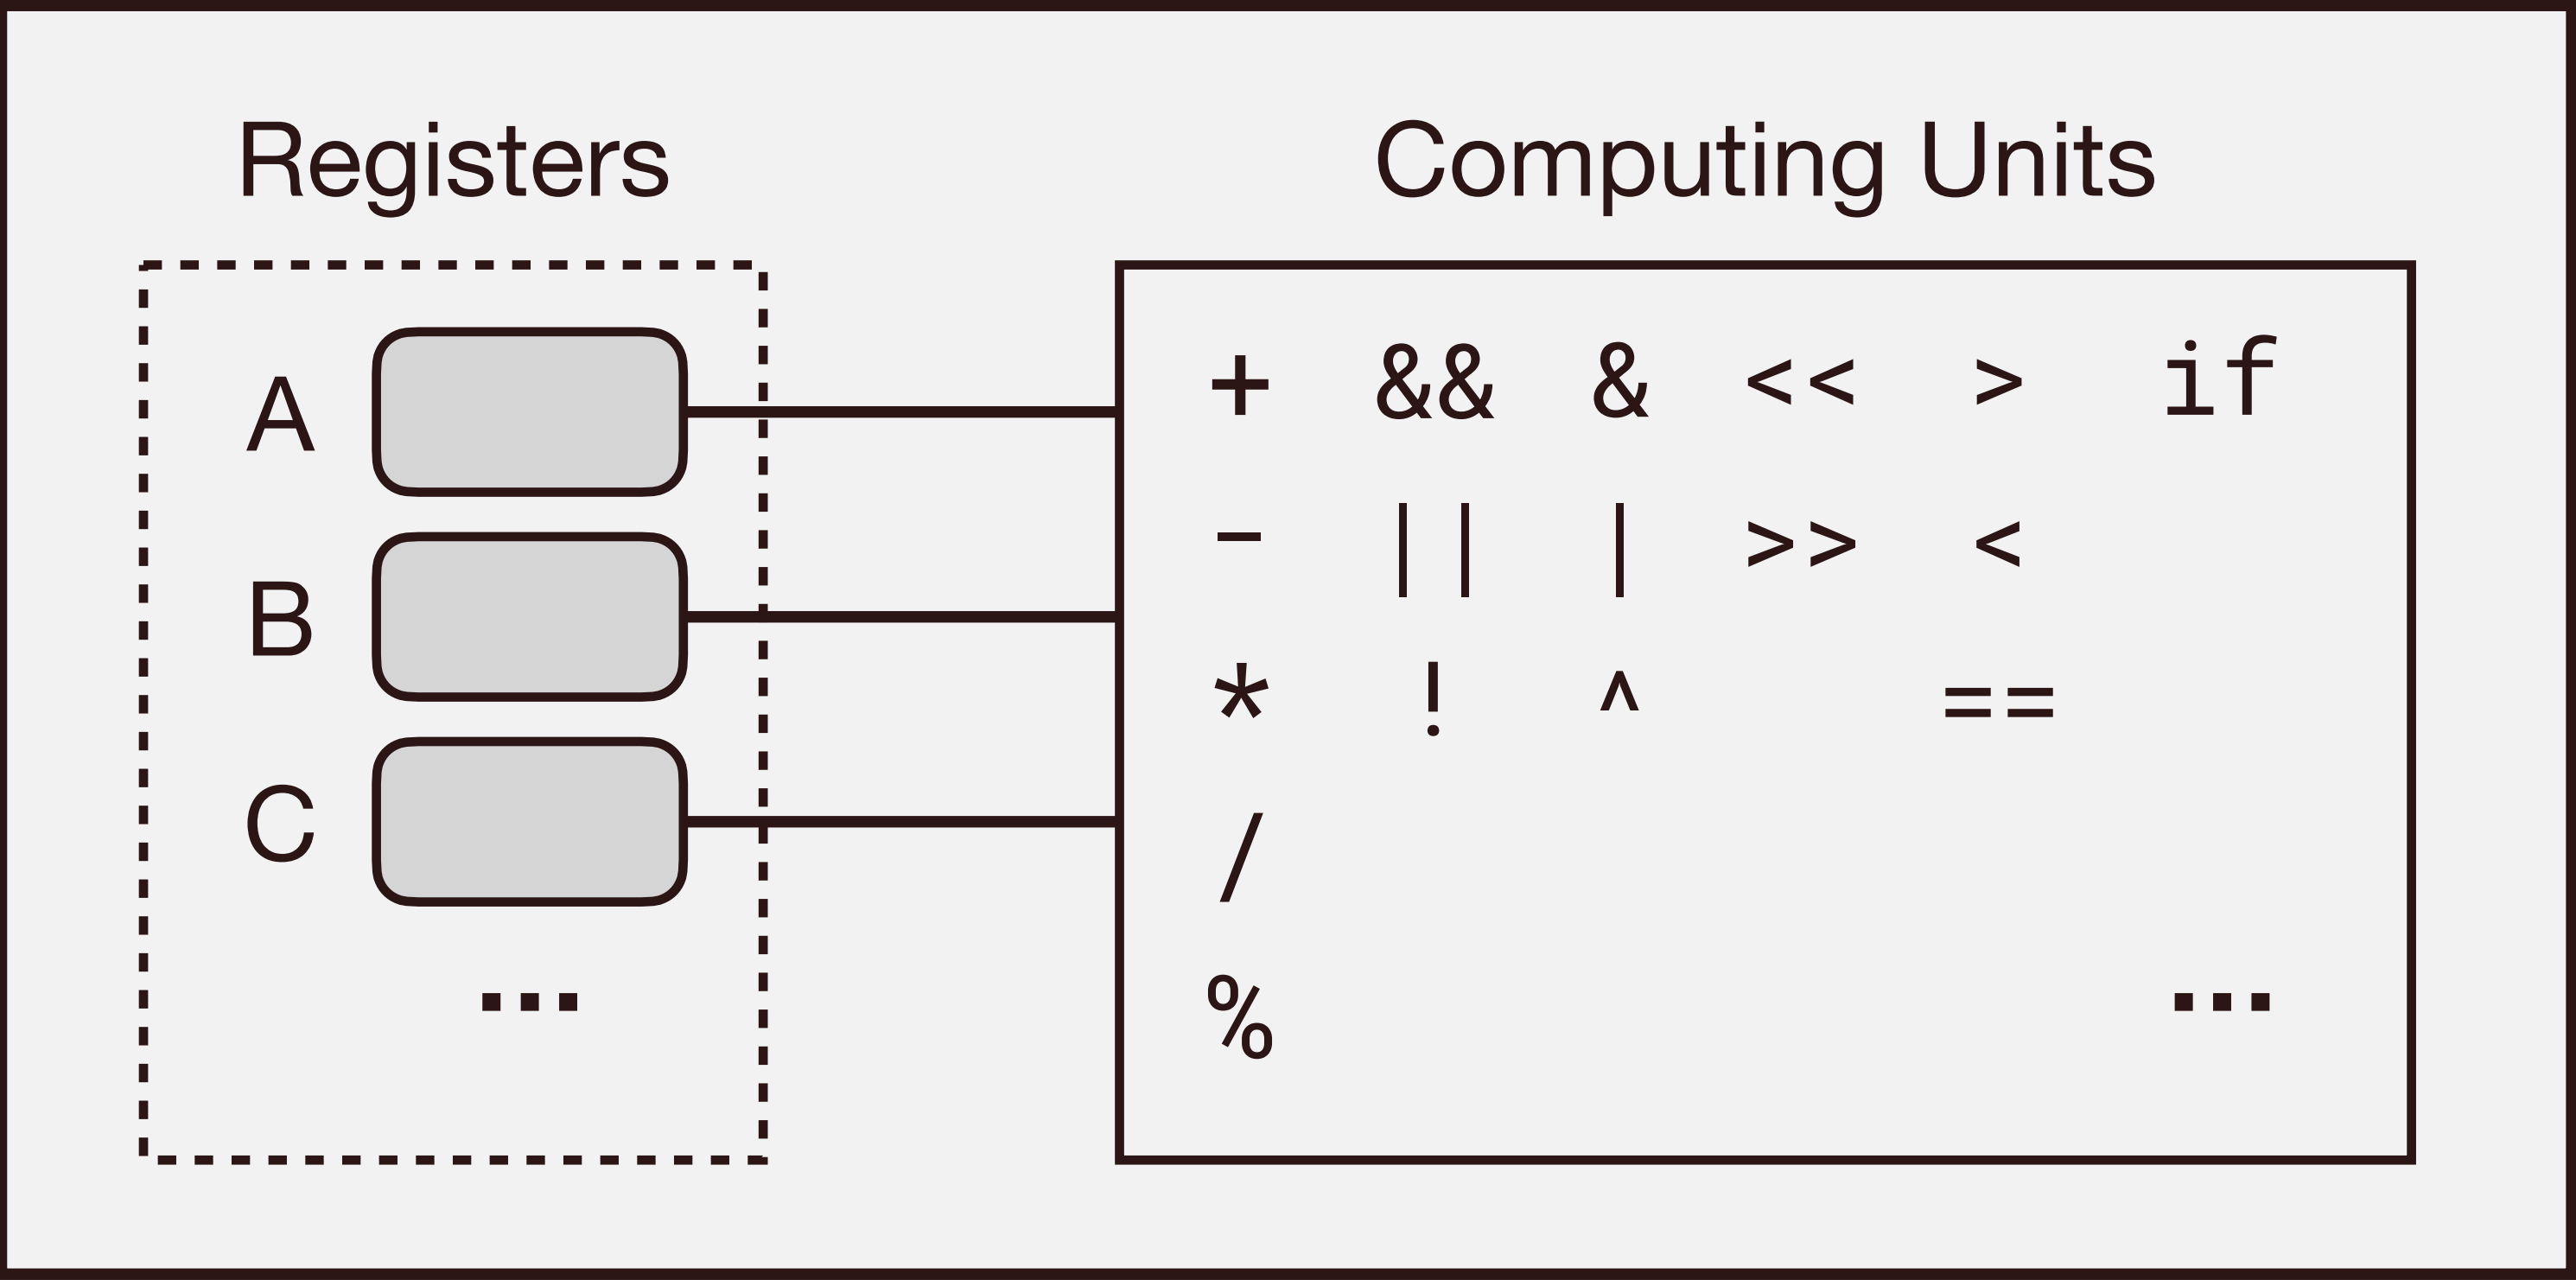


Memory



Know Your Architecture

CPU



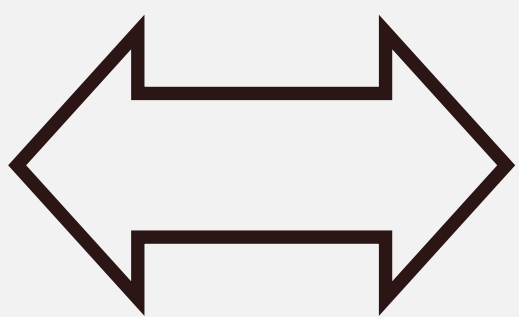
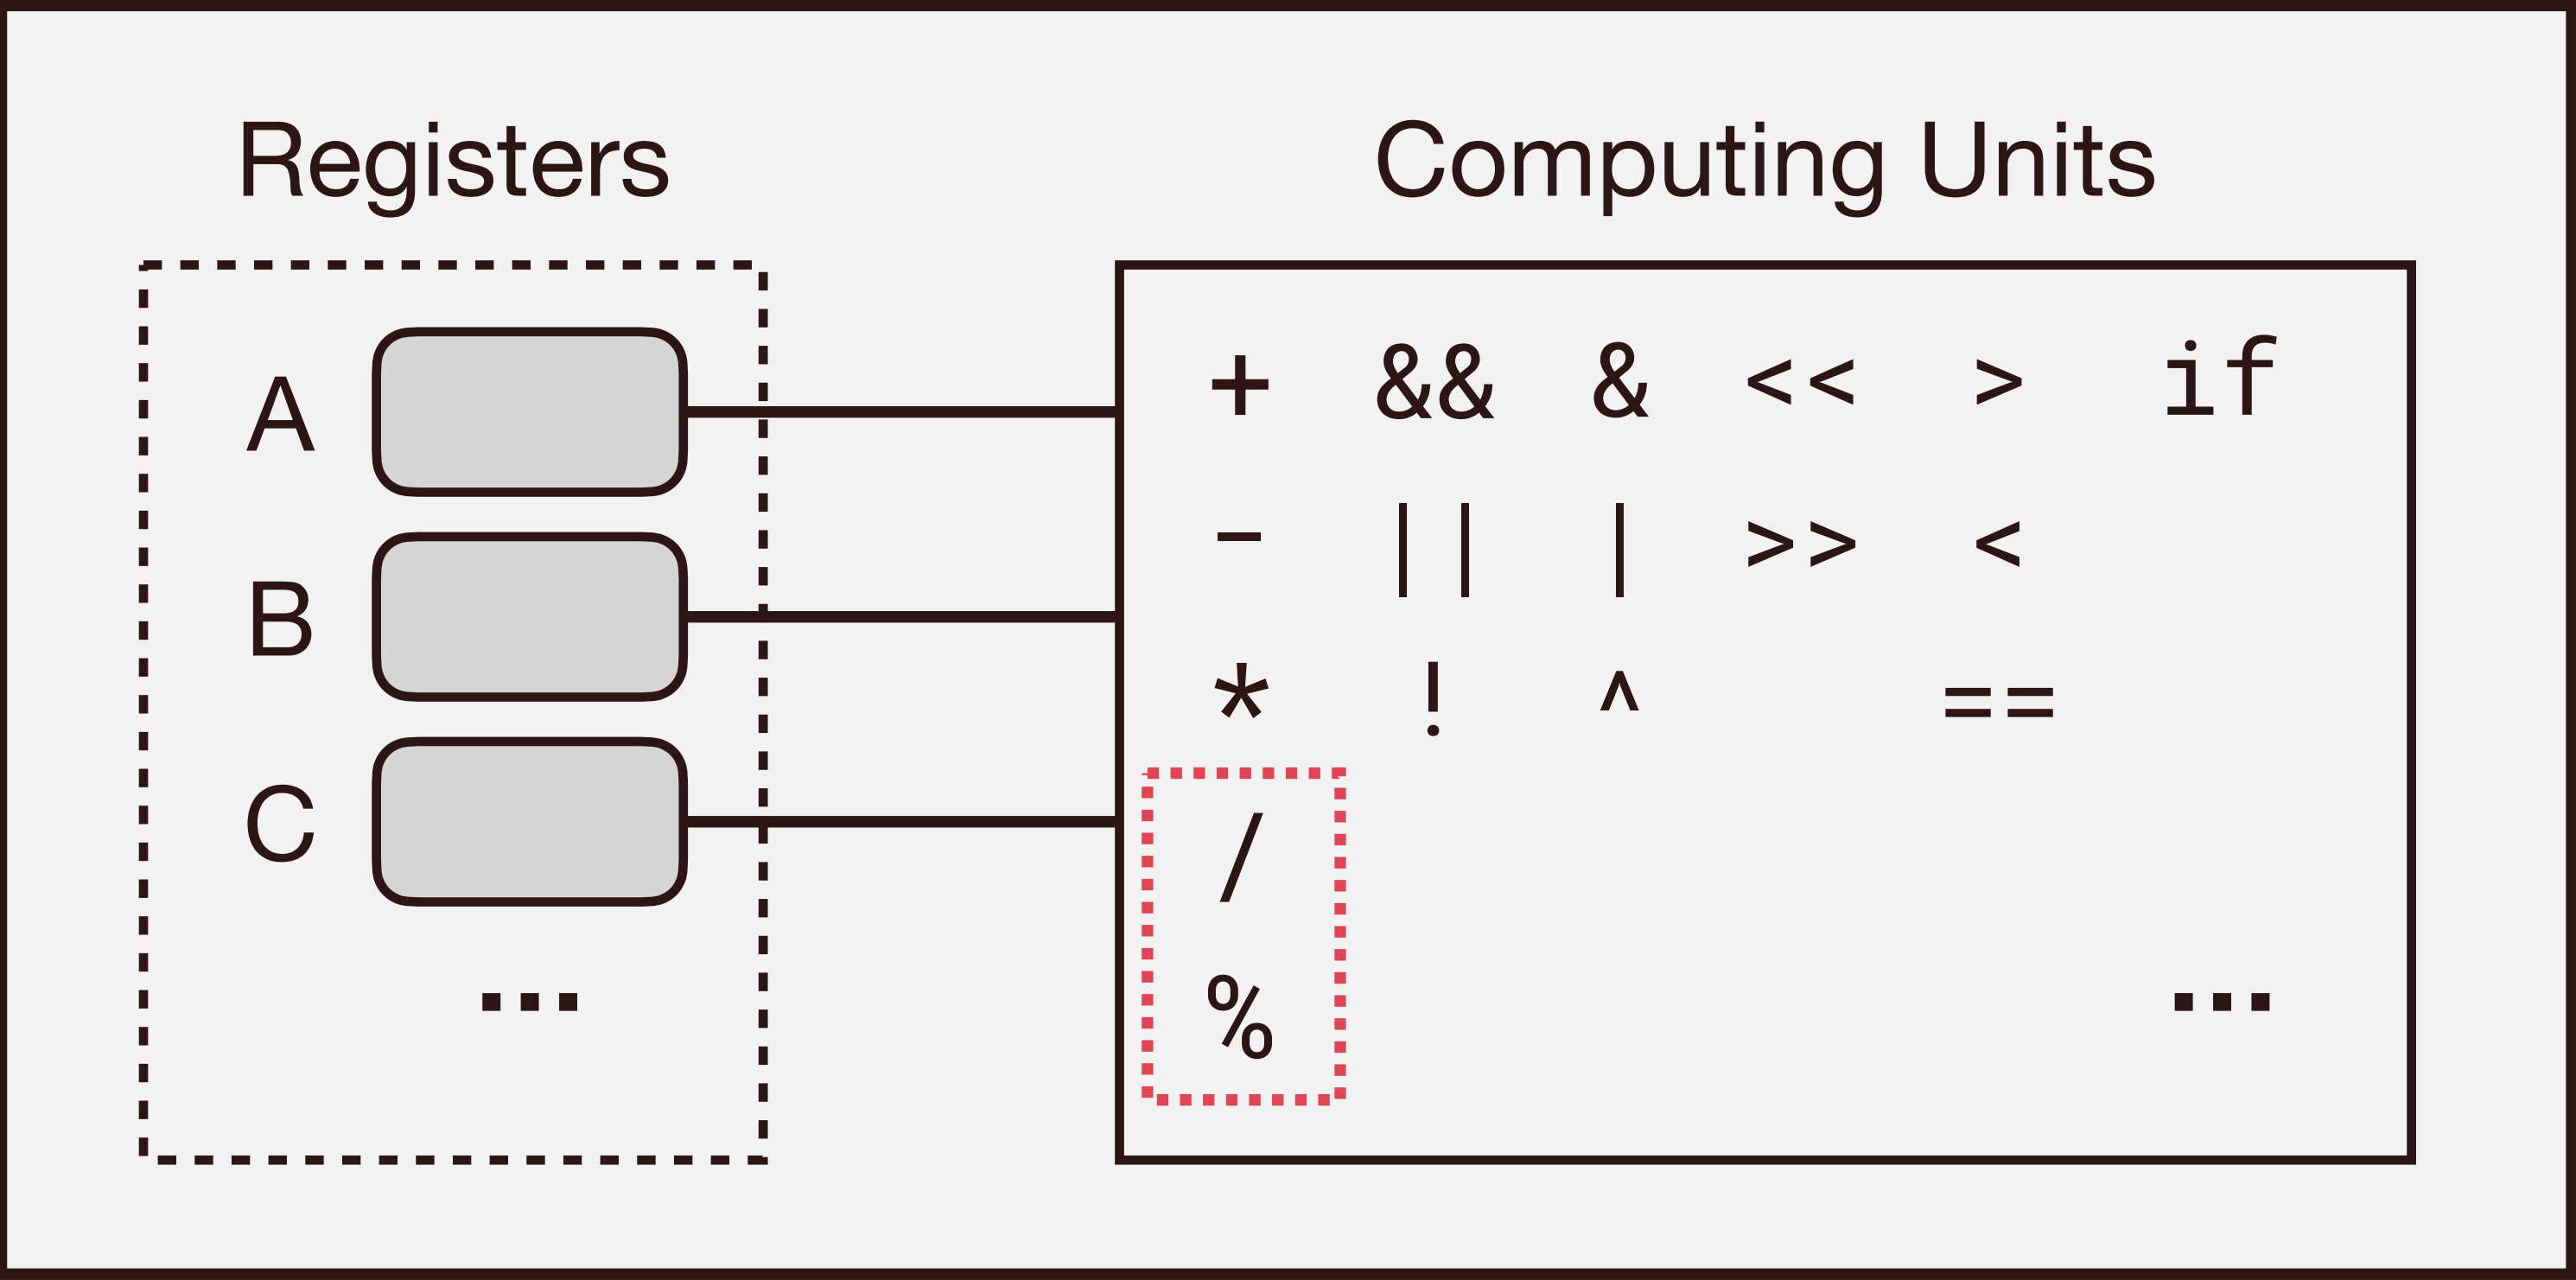
Memory



They are NOT equally fast!

Know Your Architecture

CPU



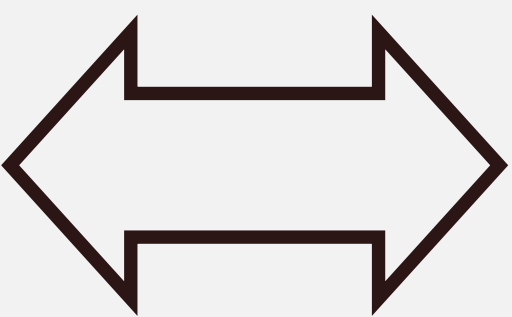
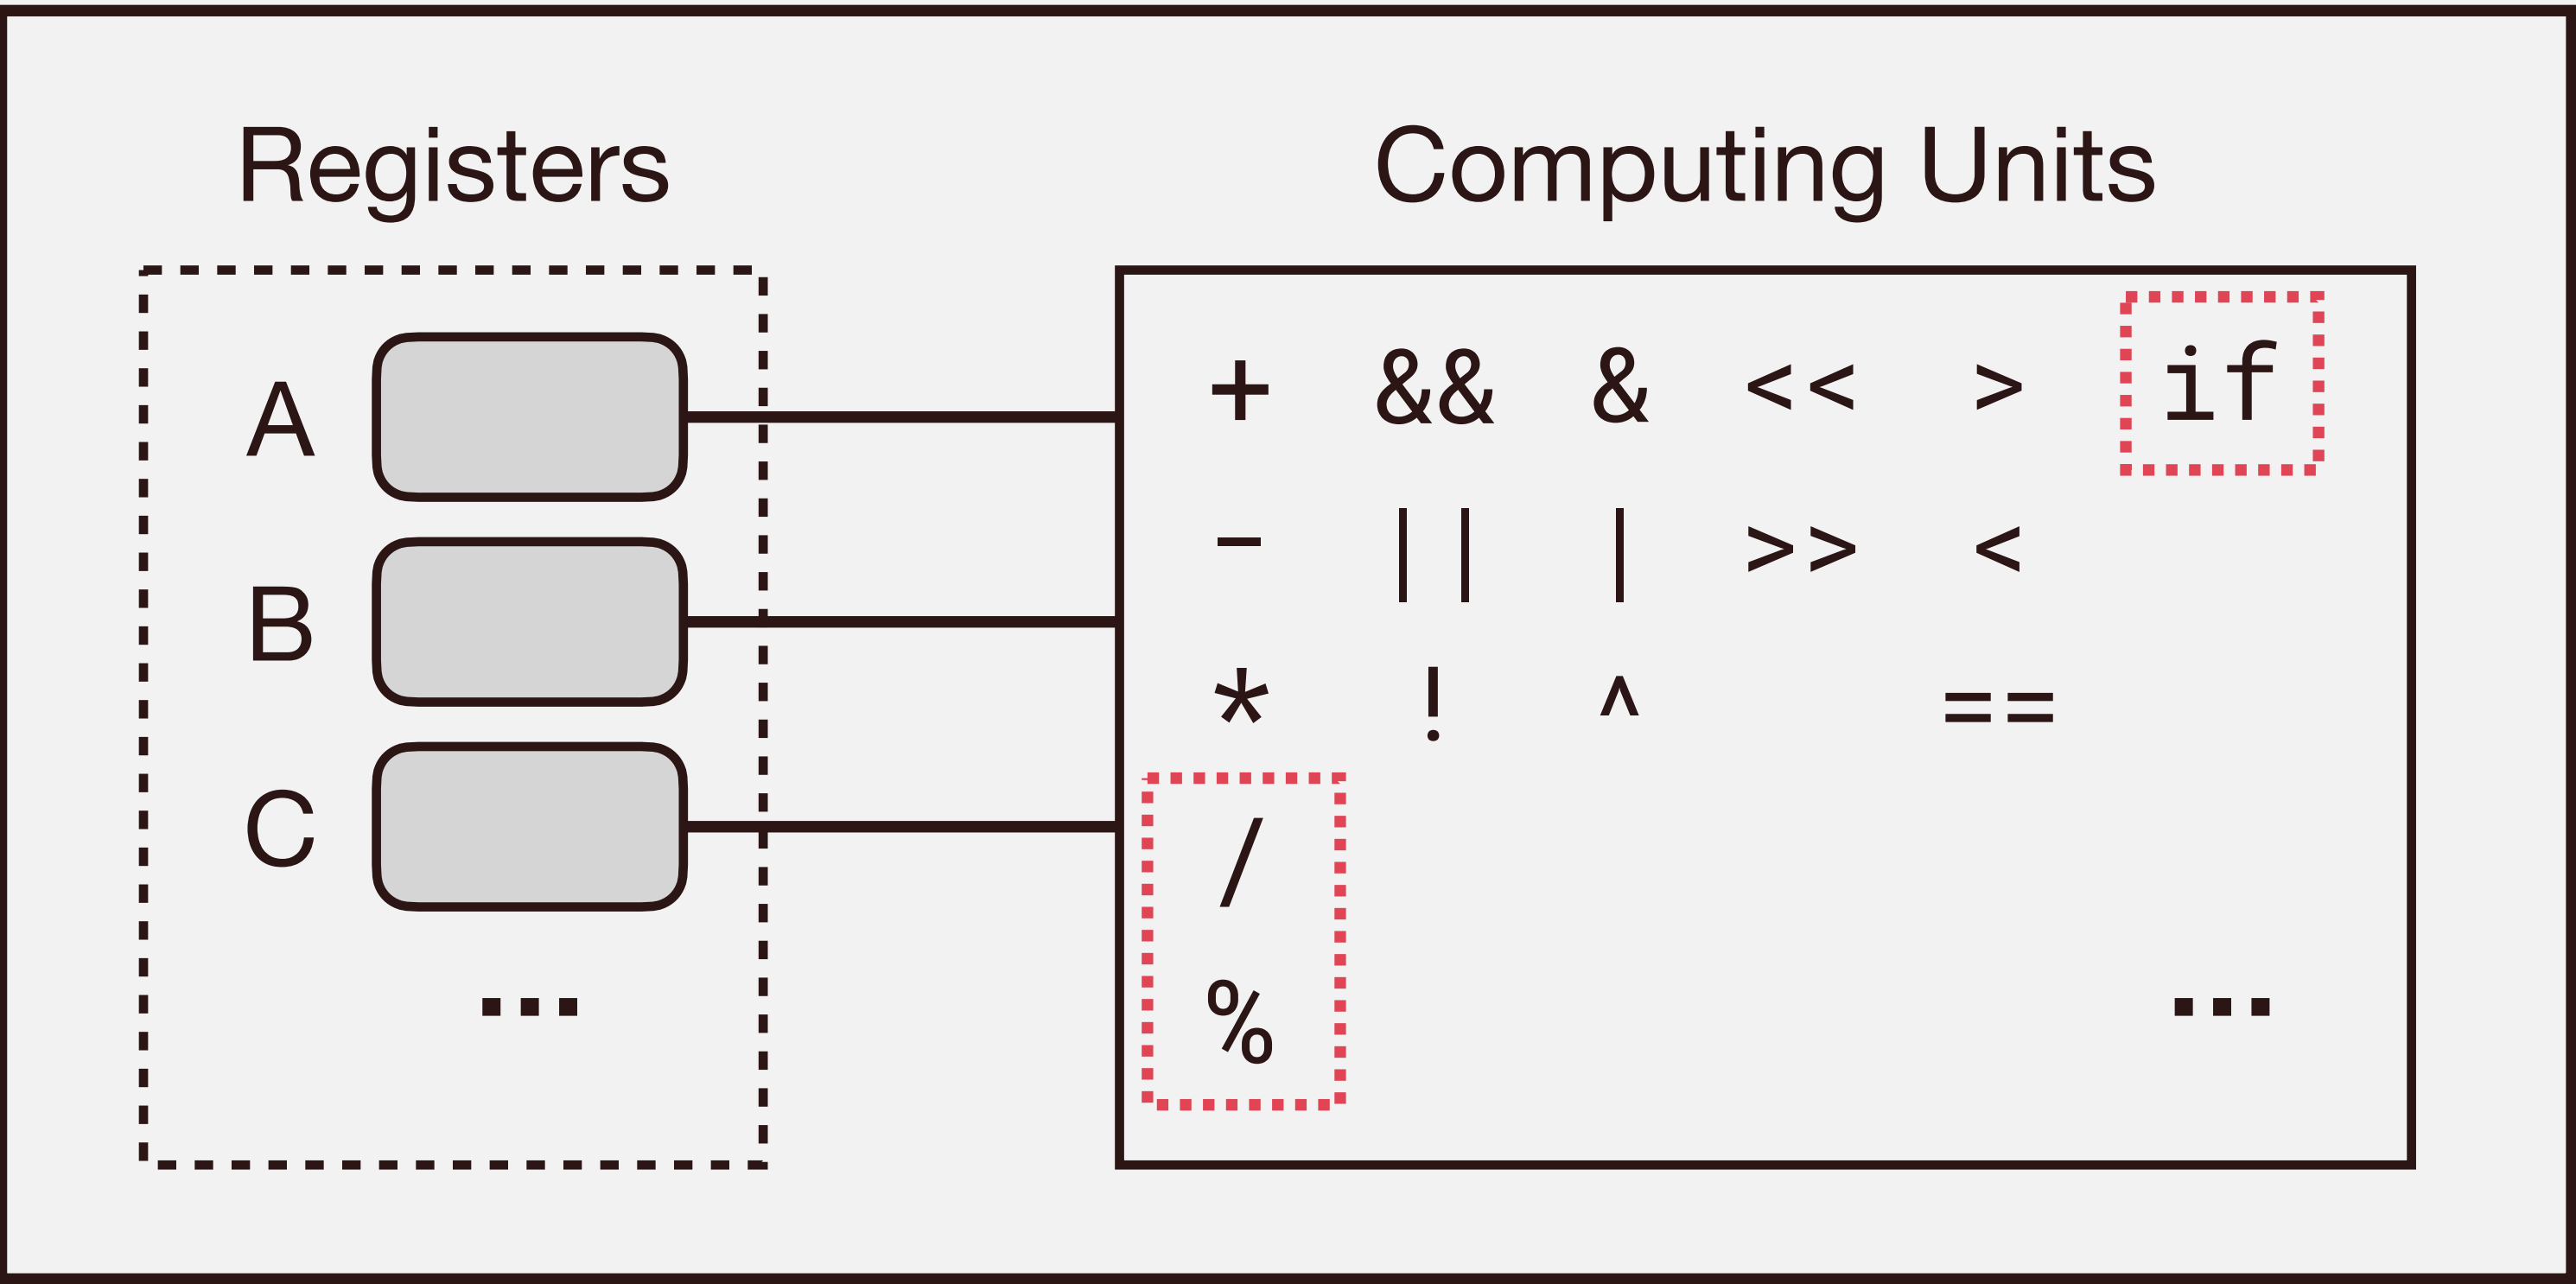
Memory



They are NOT equally fast!

Know Your Architecture

CPU

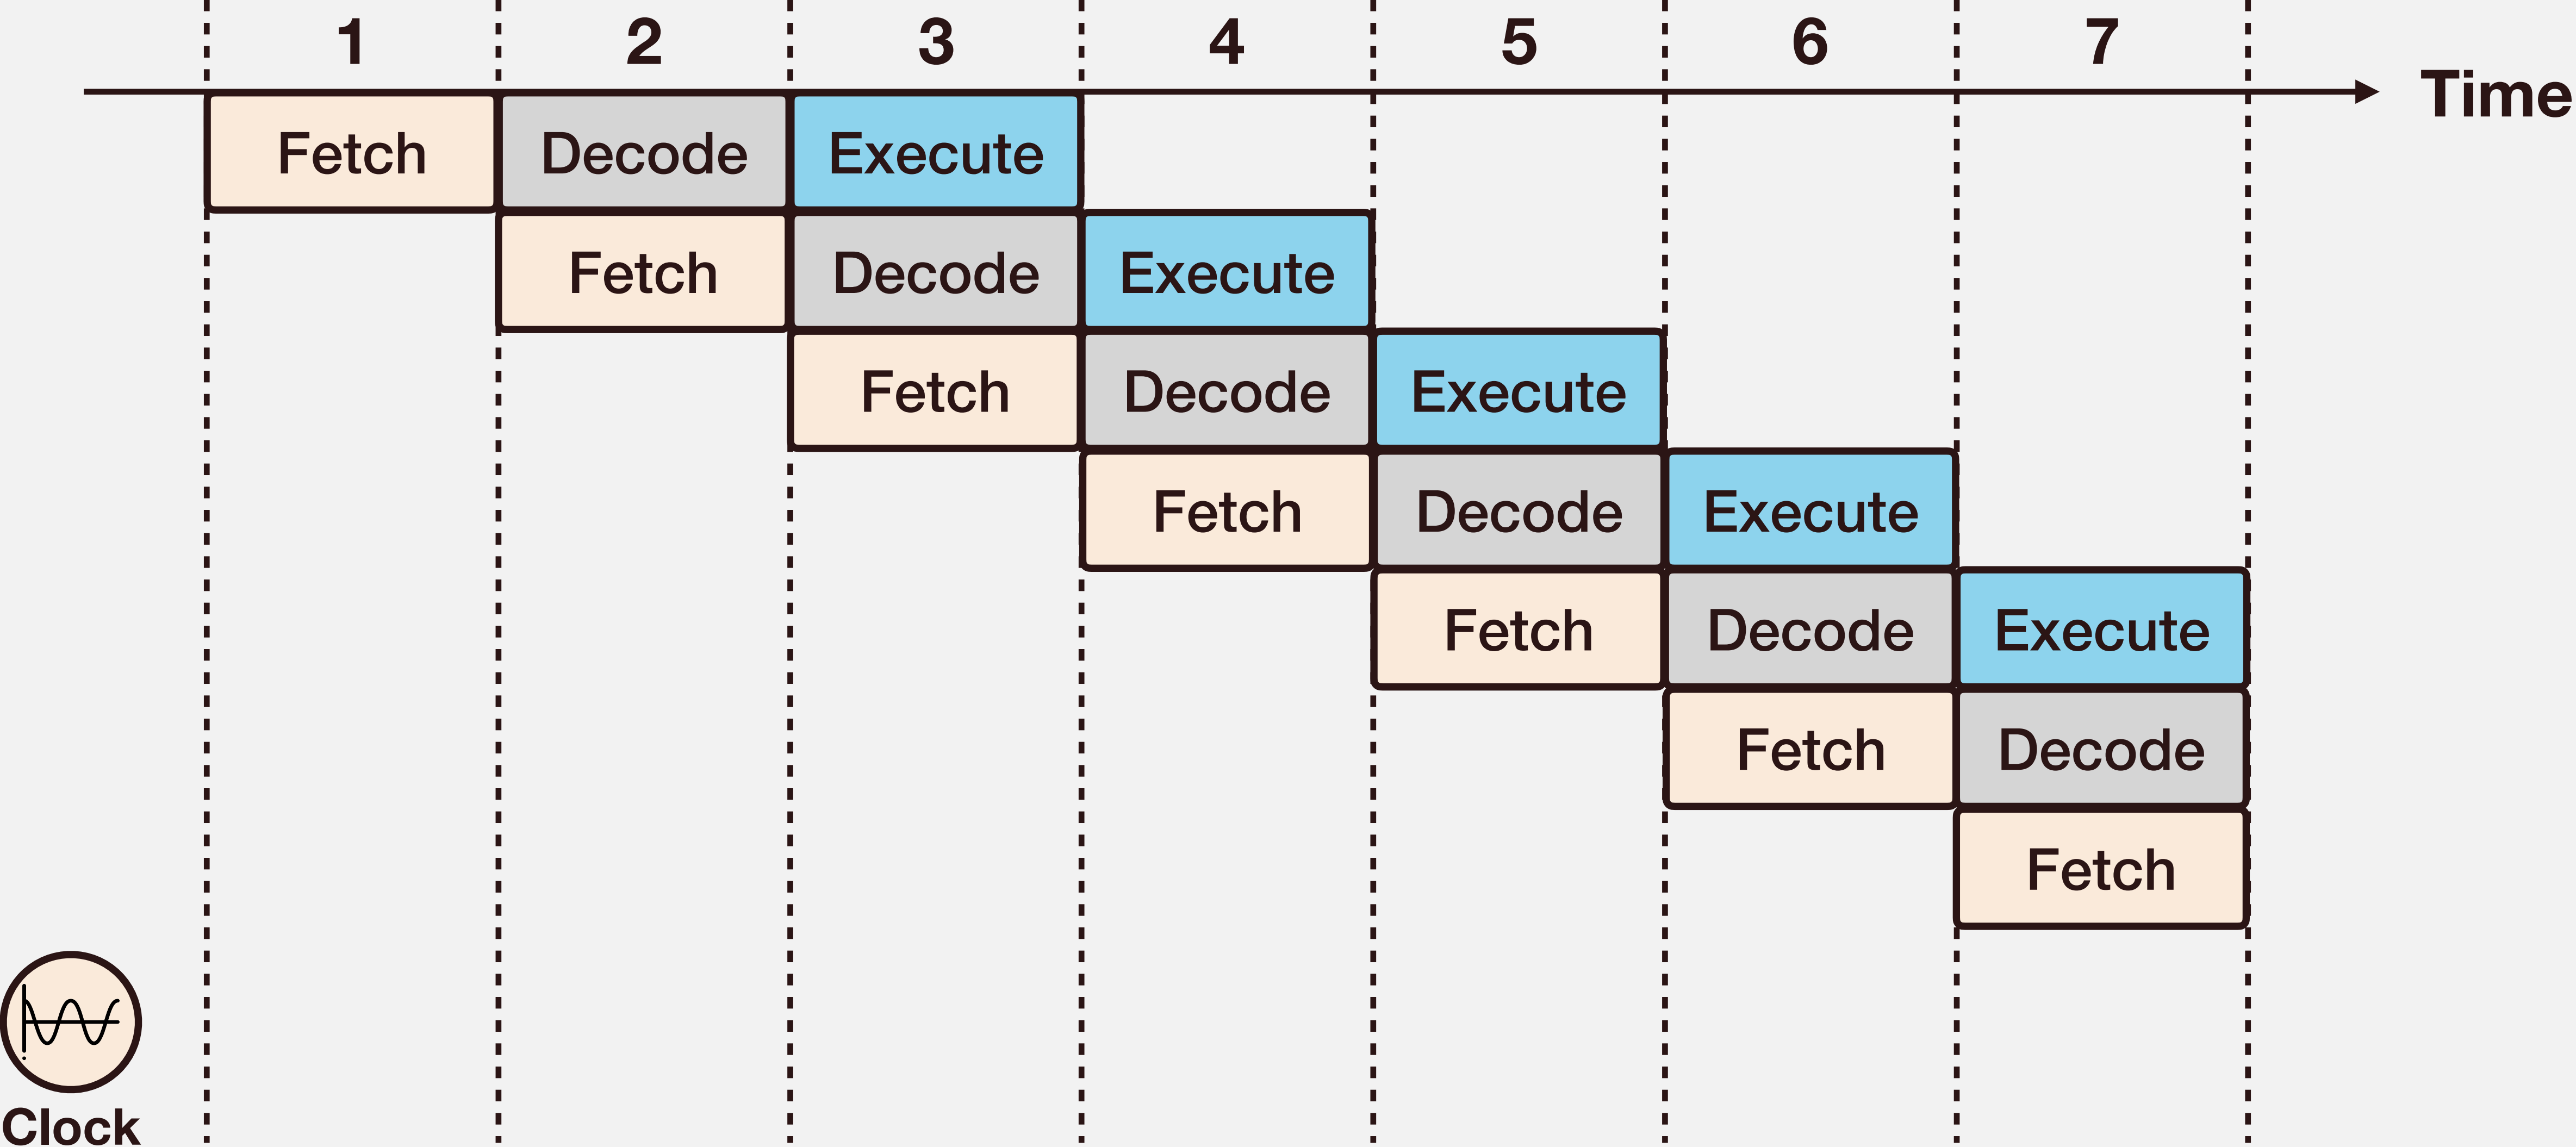


Memory



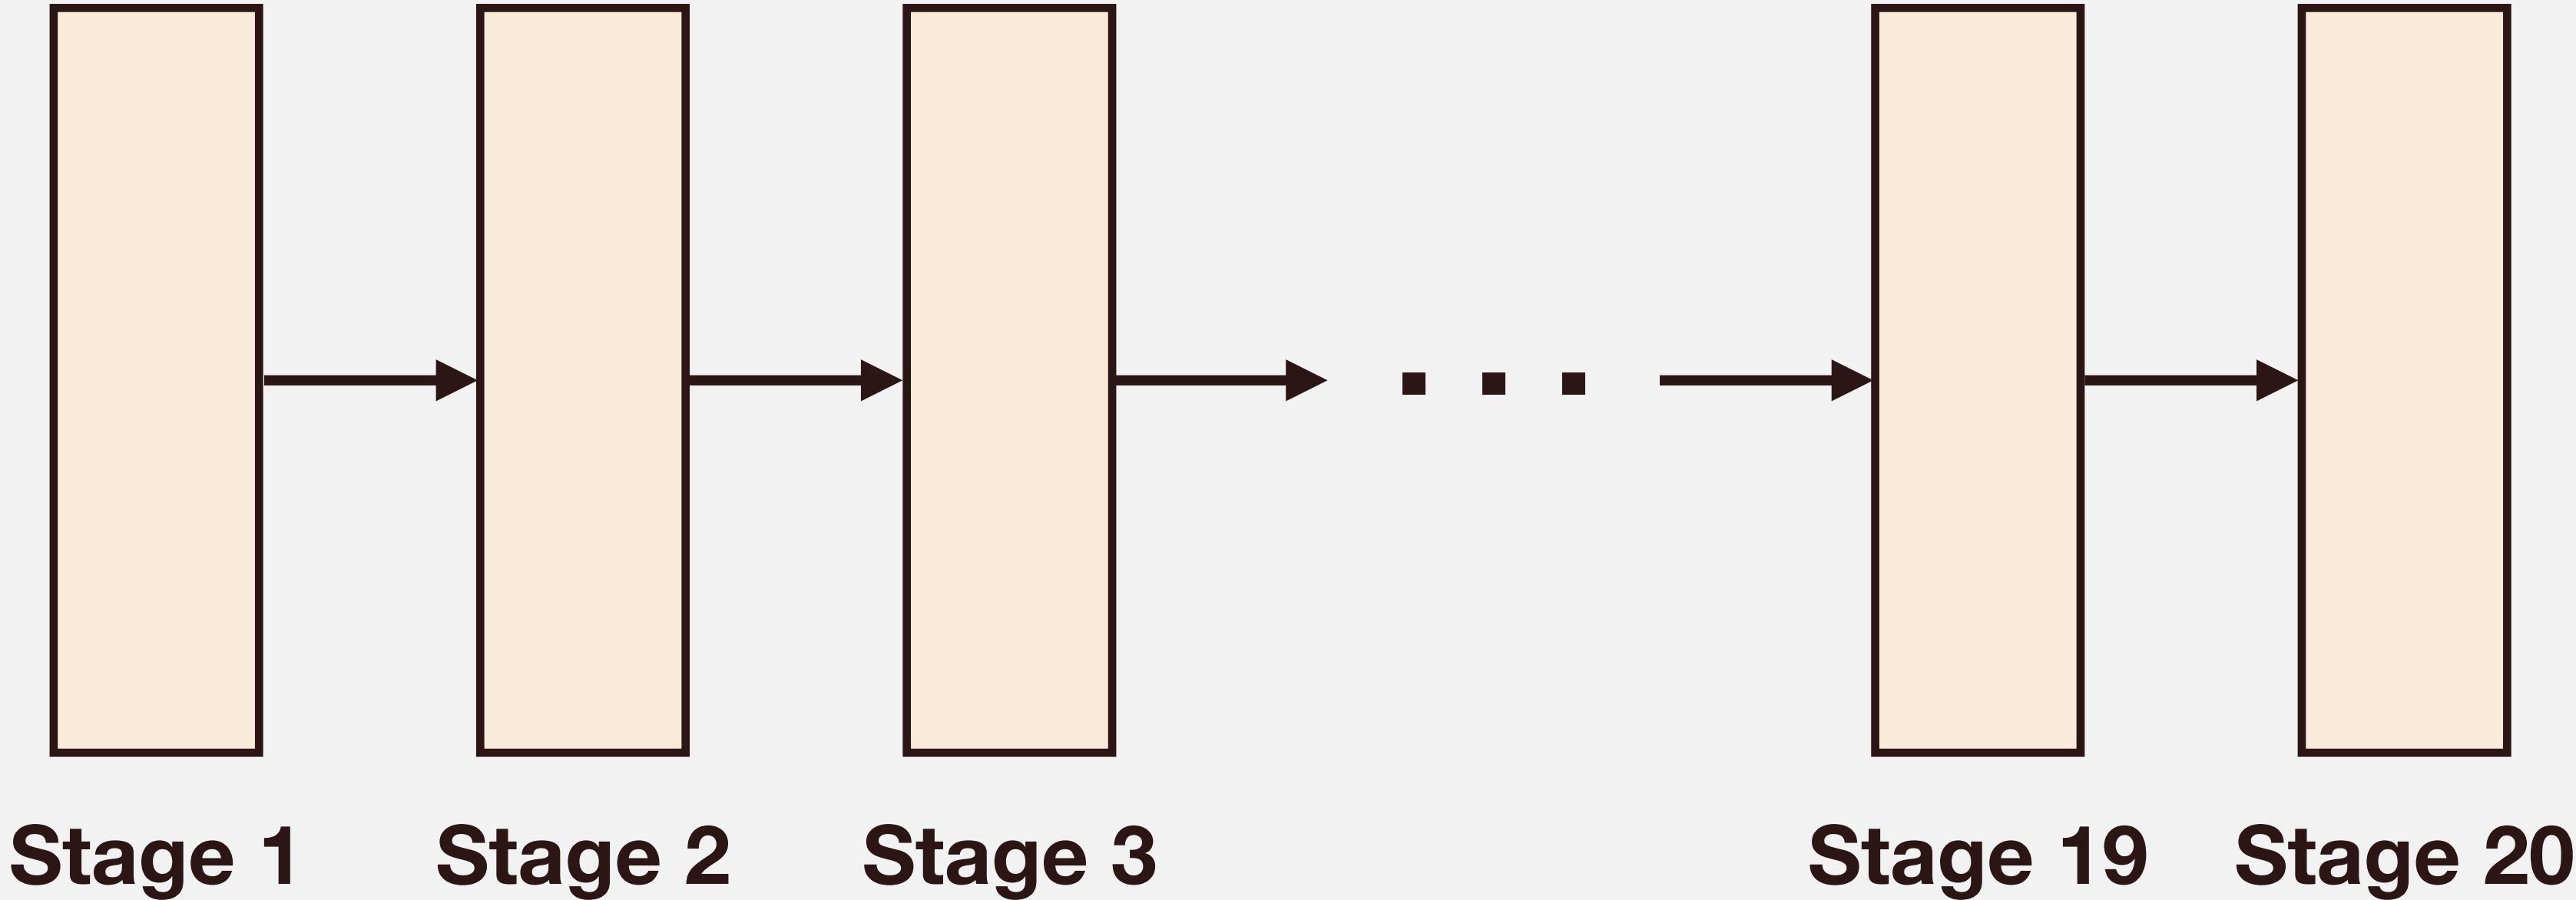
They are NOT equally fast!

CPU Pipeline



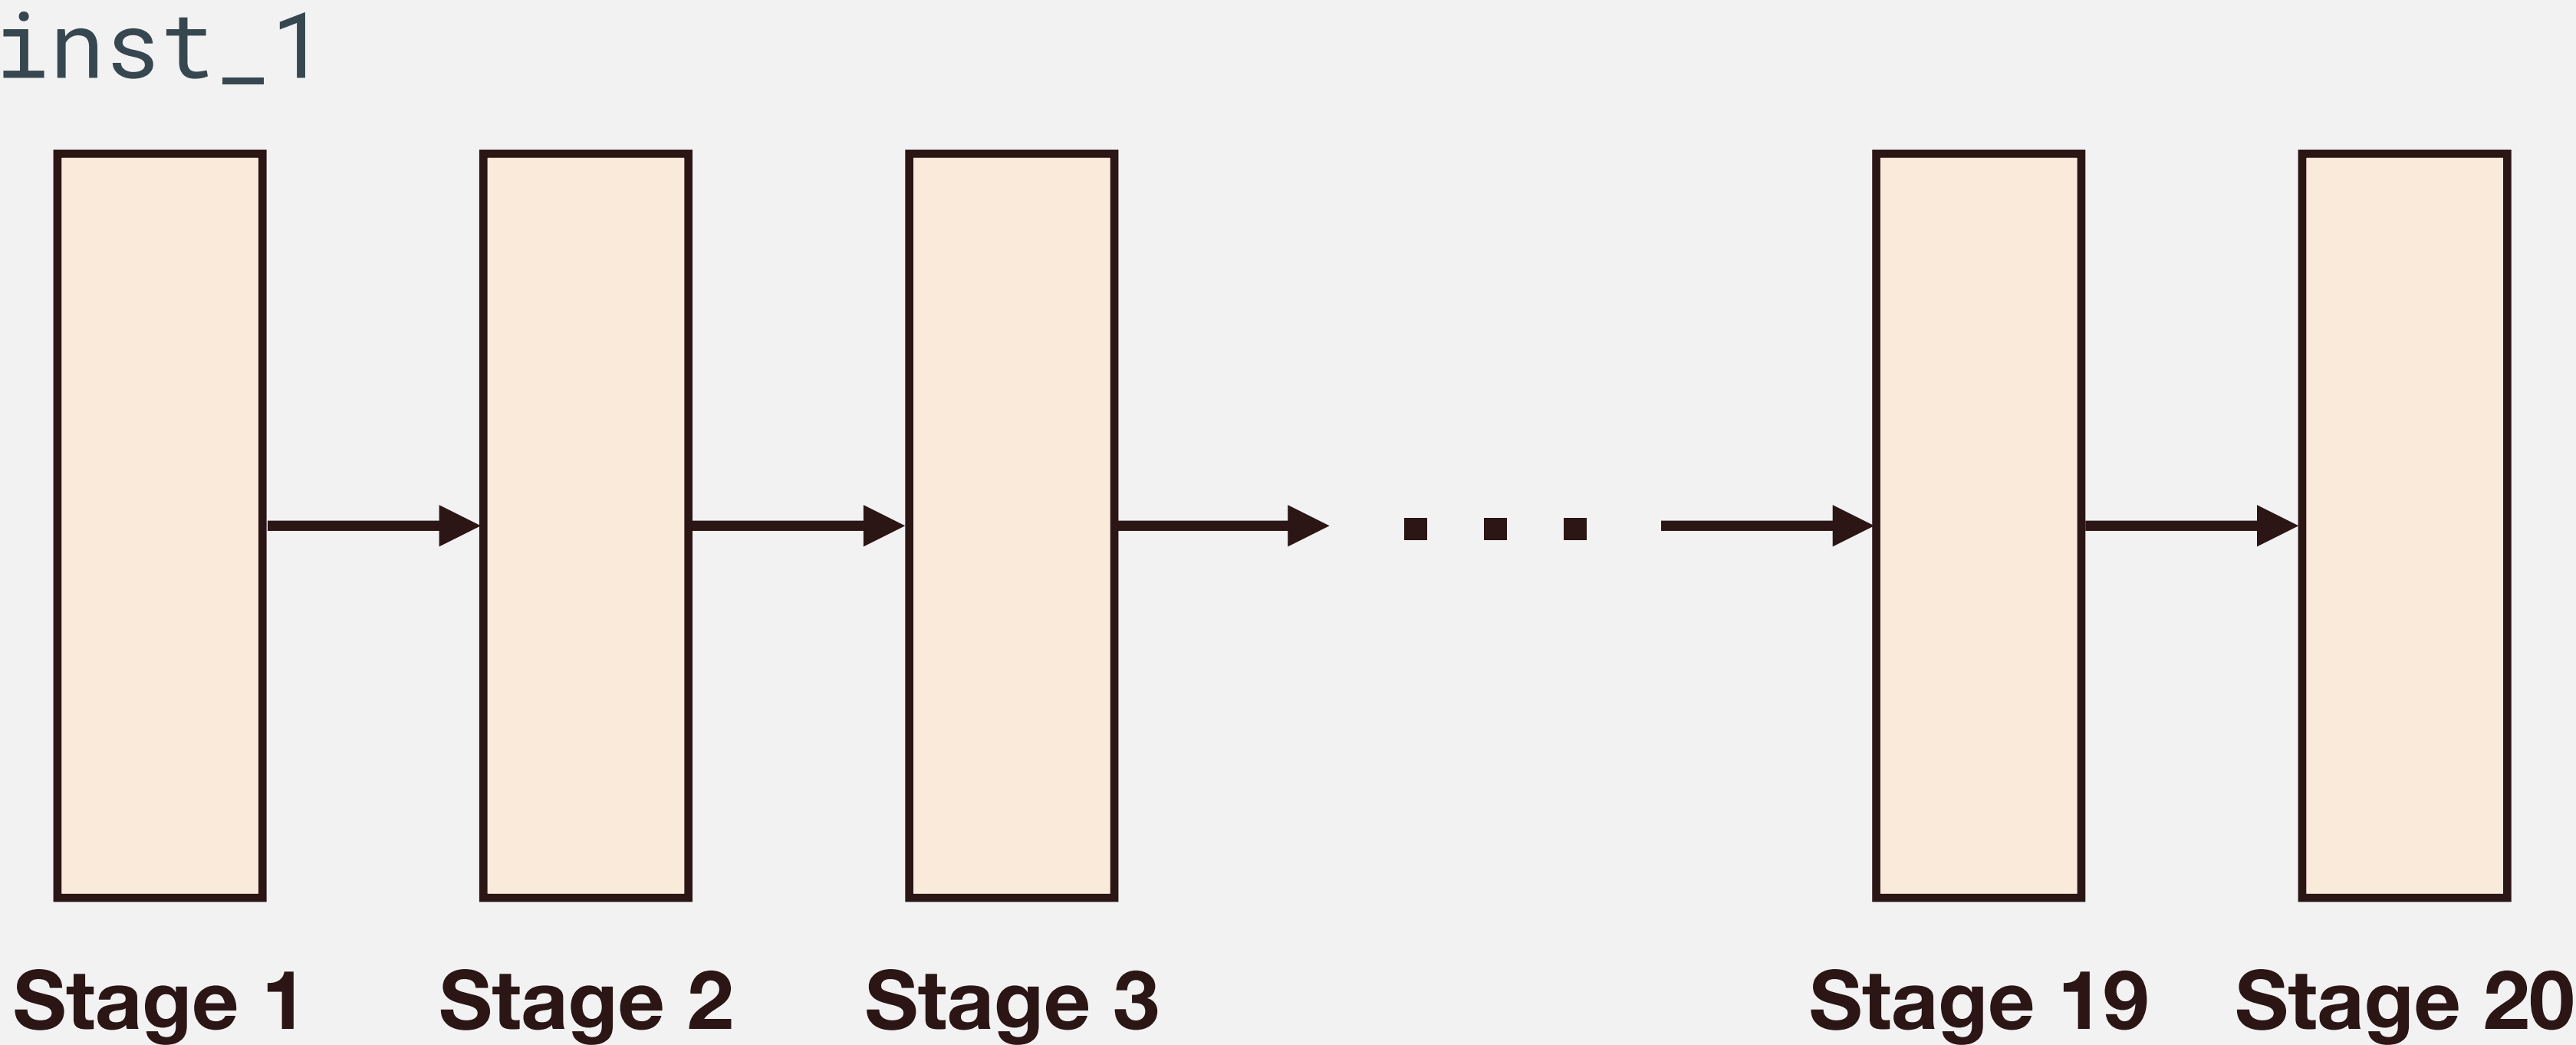
CPU Pipeline

inst_1
inst_2
inst_3
...
inst_N



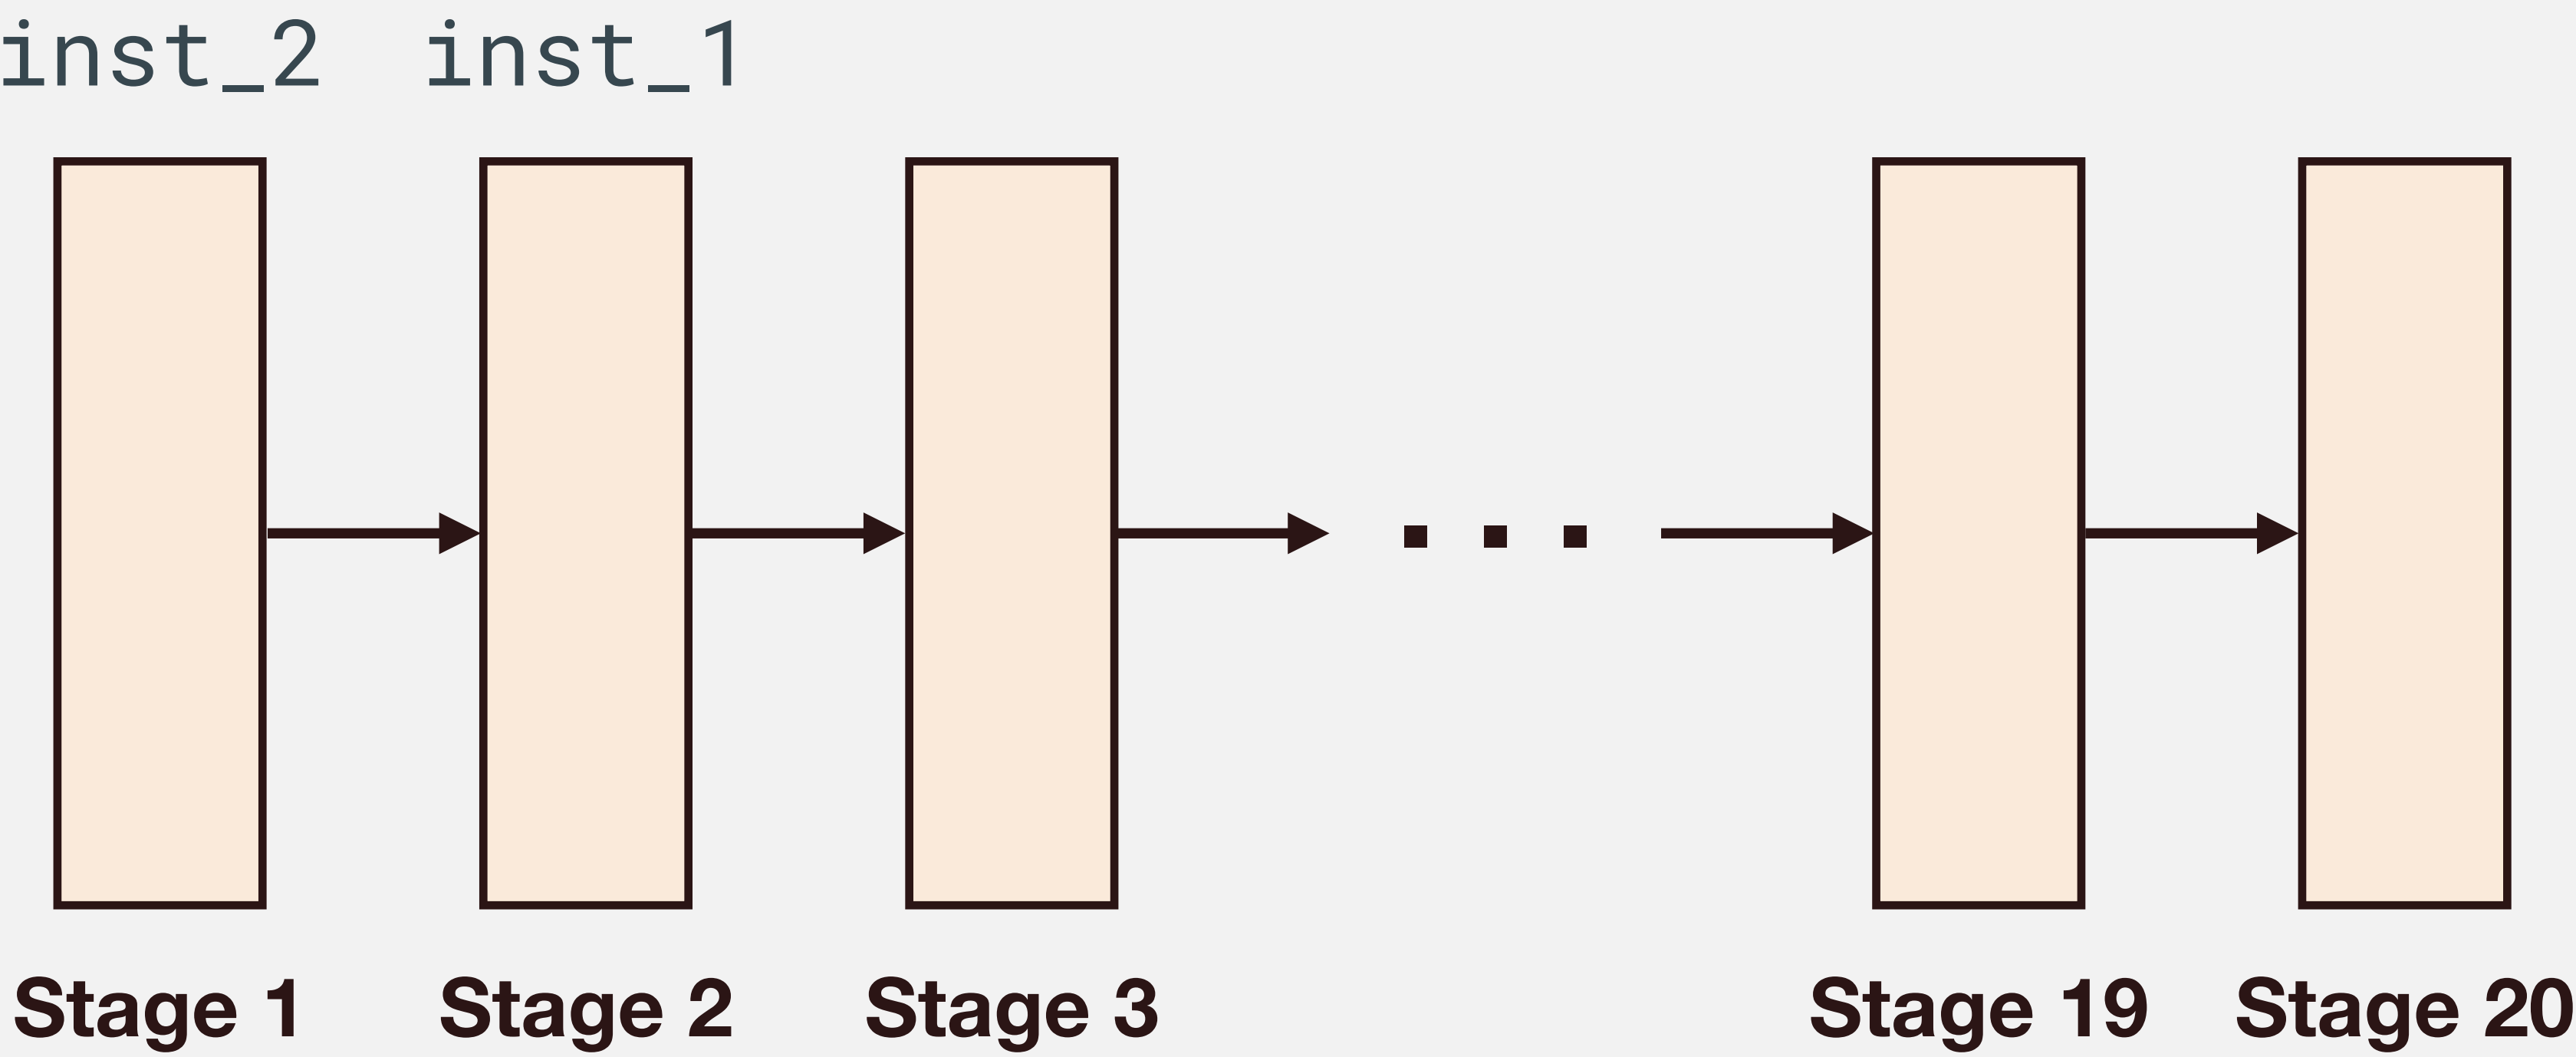
CPU Pipeline

inst_1
inst_2
inst_3
...
inst_N



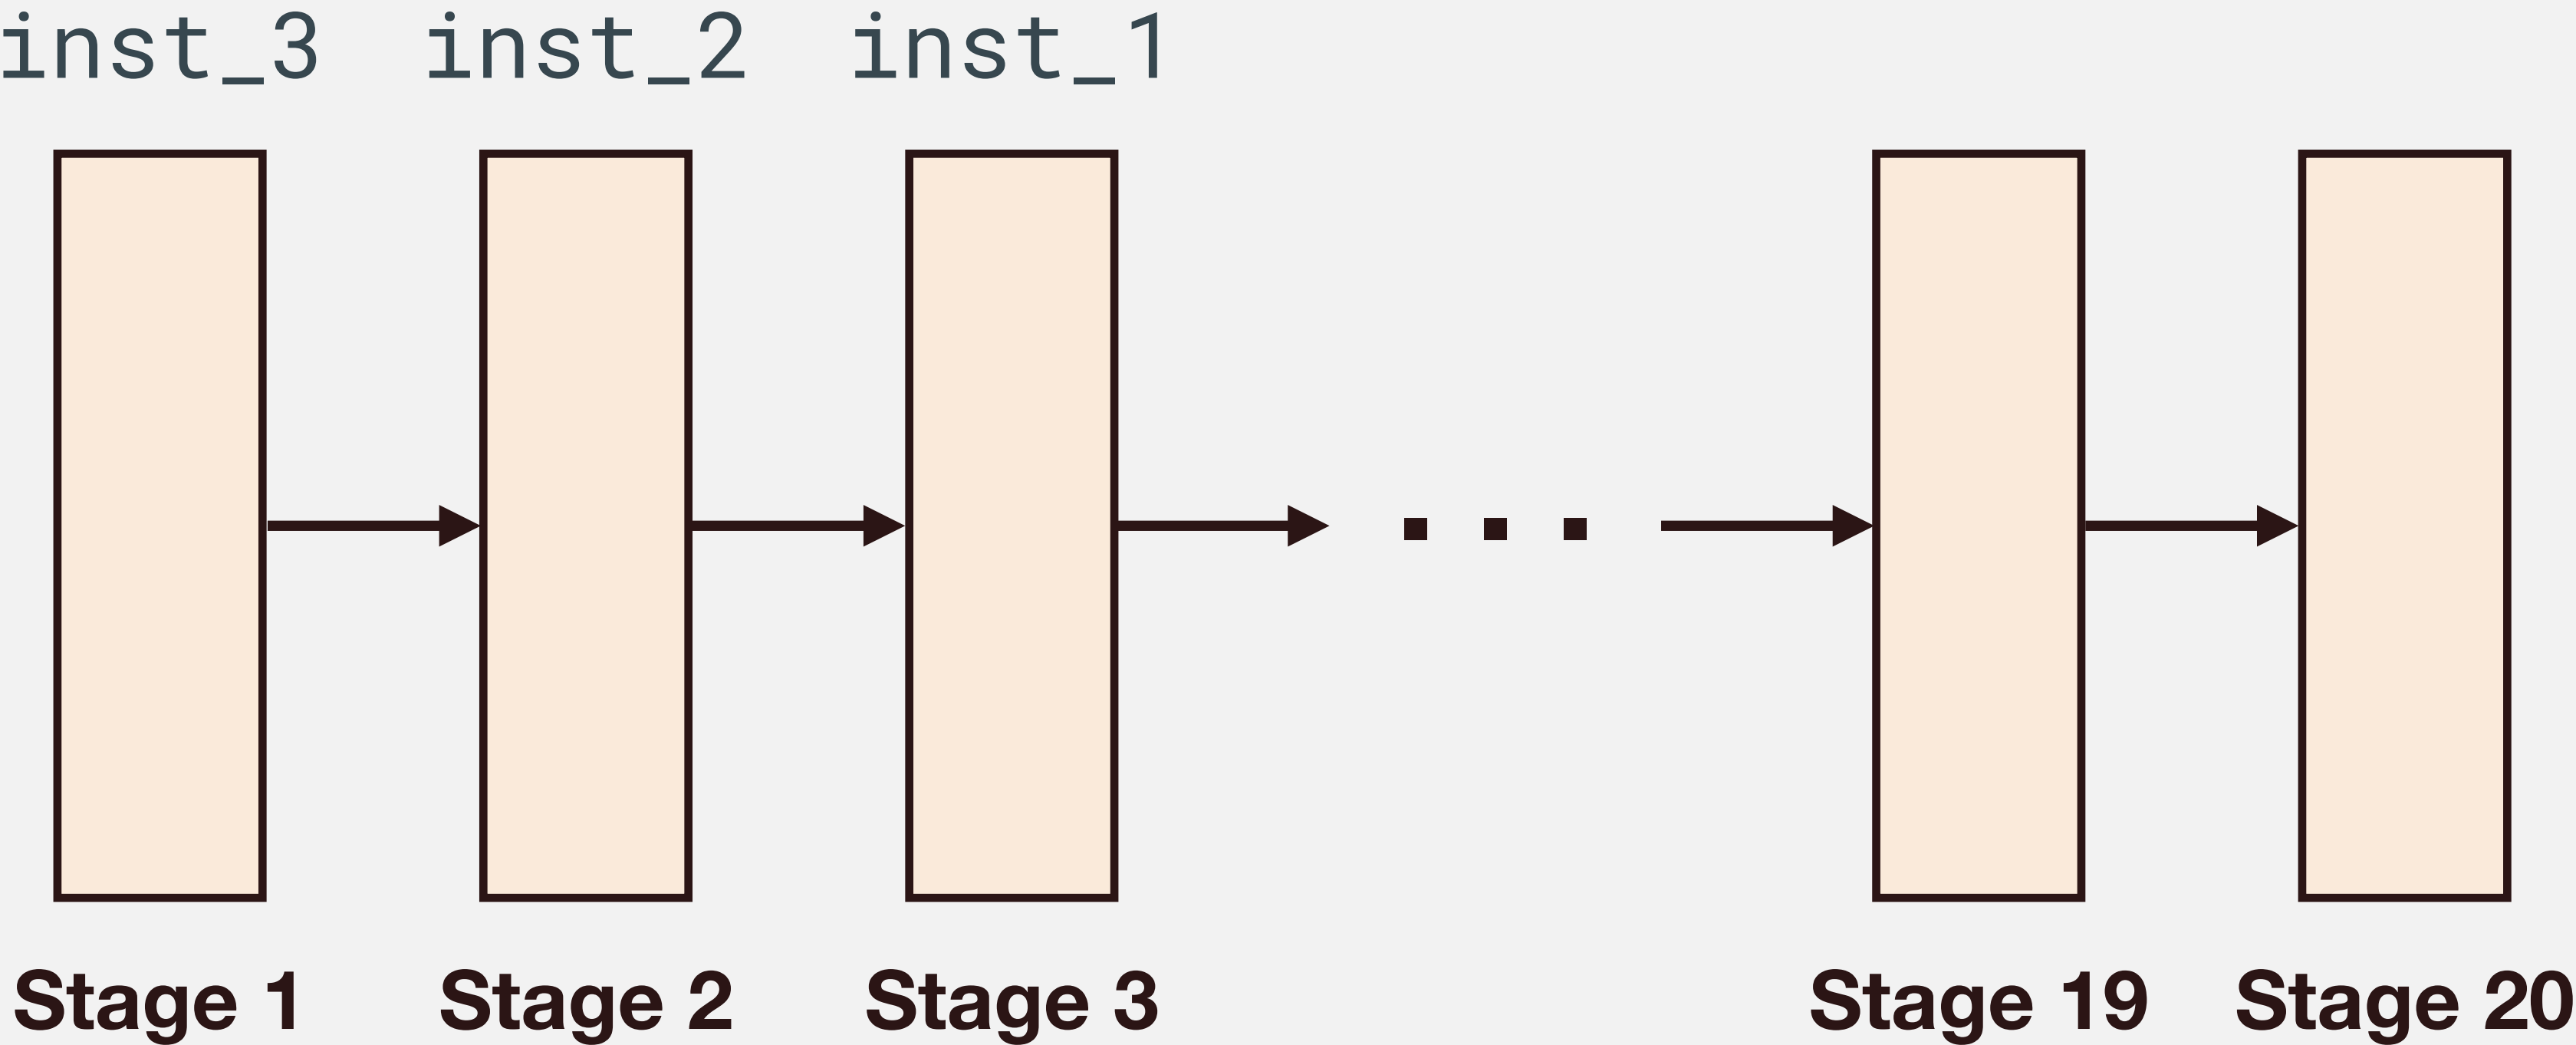
CPU Pipeline

inst_1
inst_2
inst_3
...
inst_N



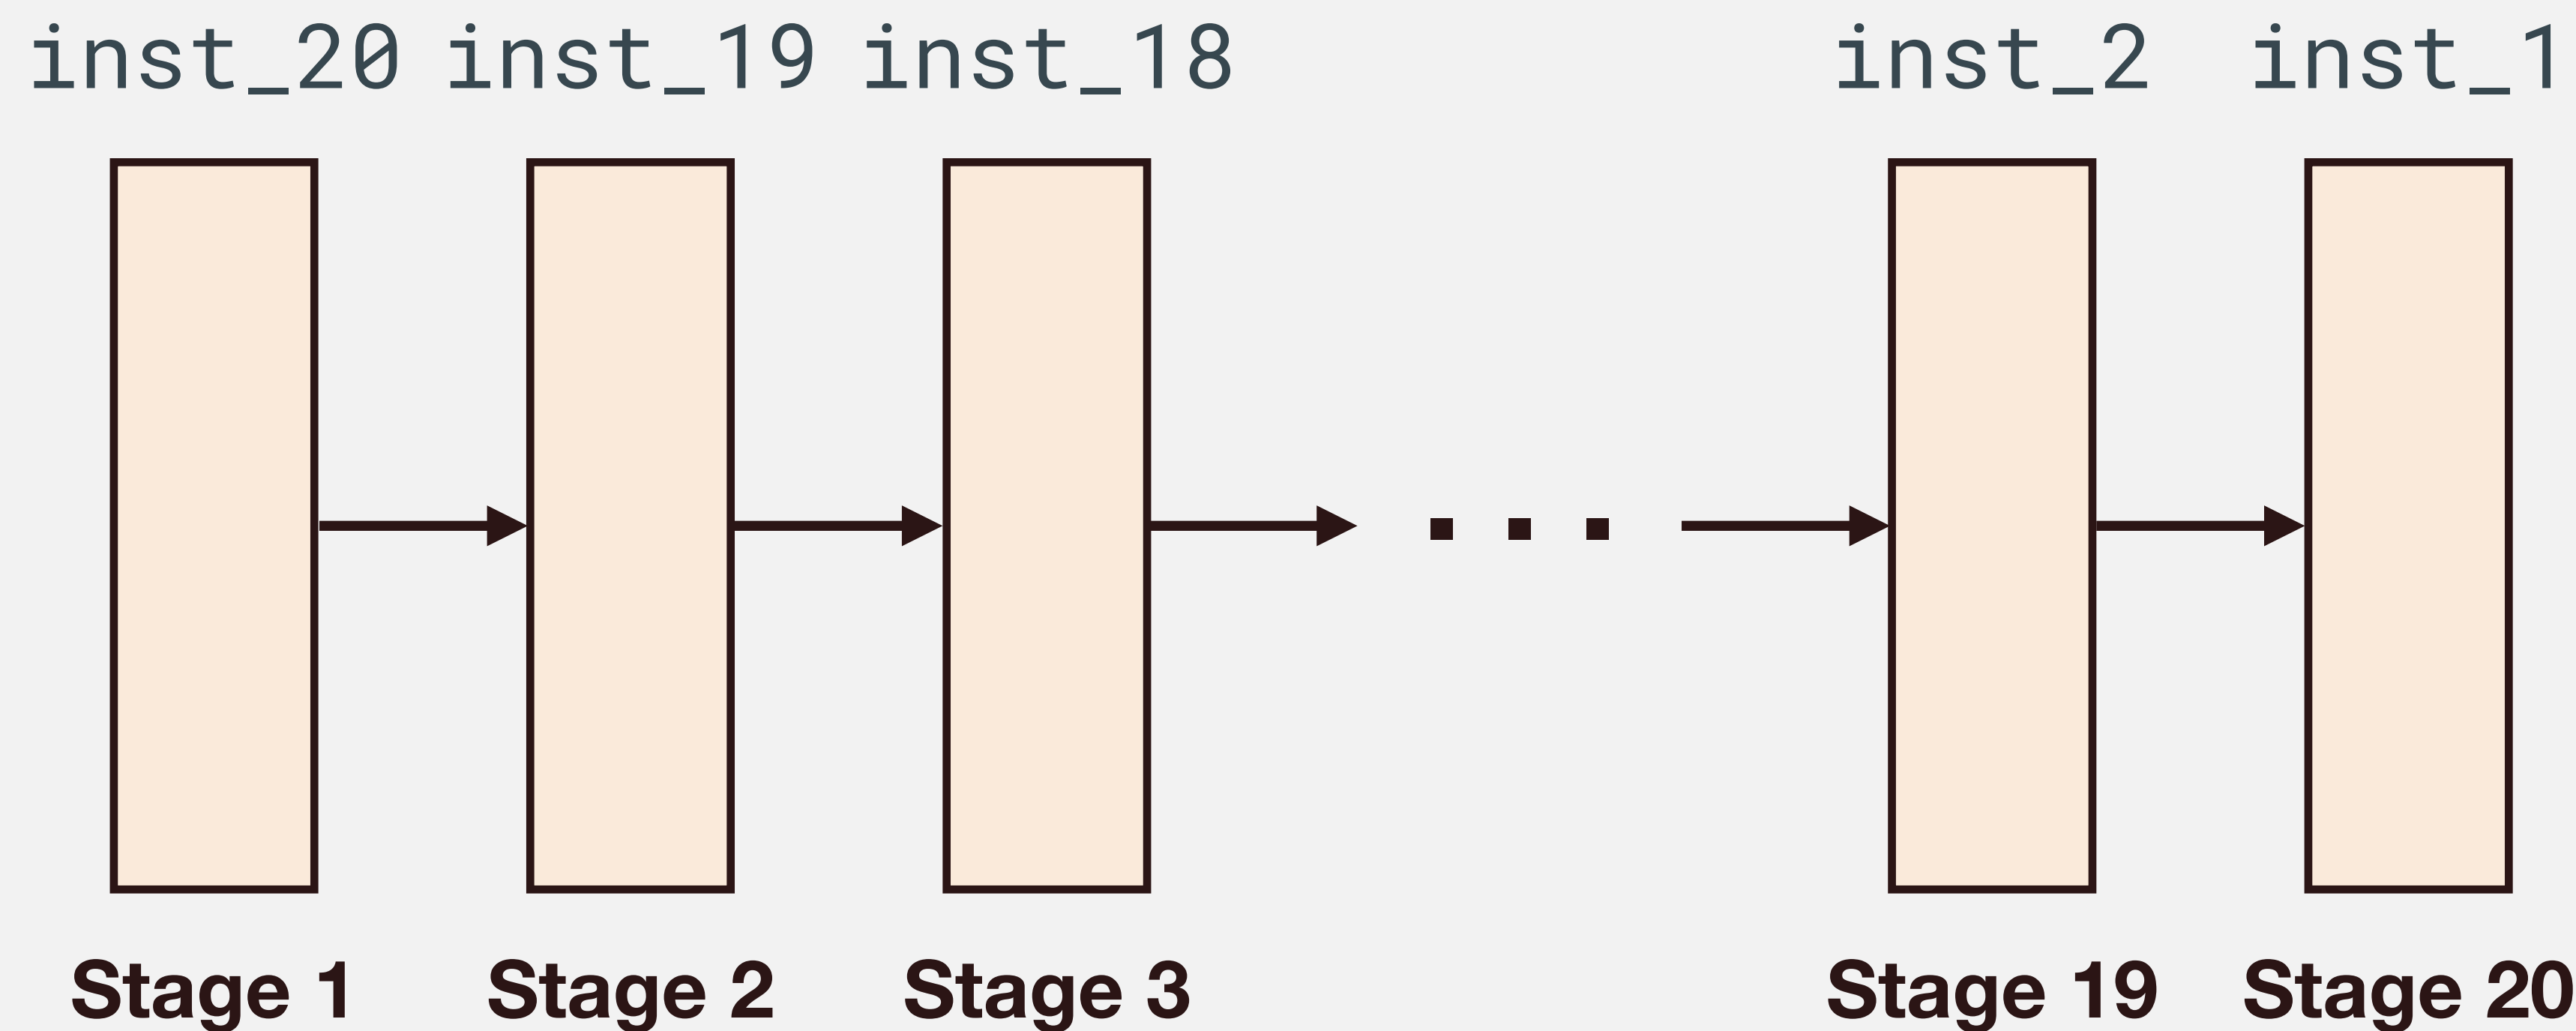
CPU Pipeline

inst_1
inst_2
inst_3
...
inst_N

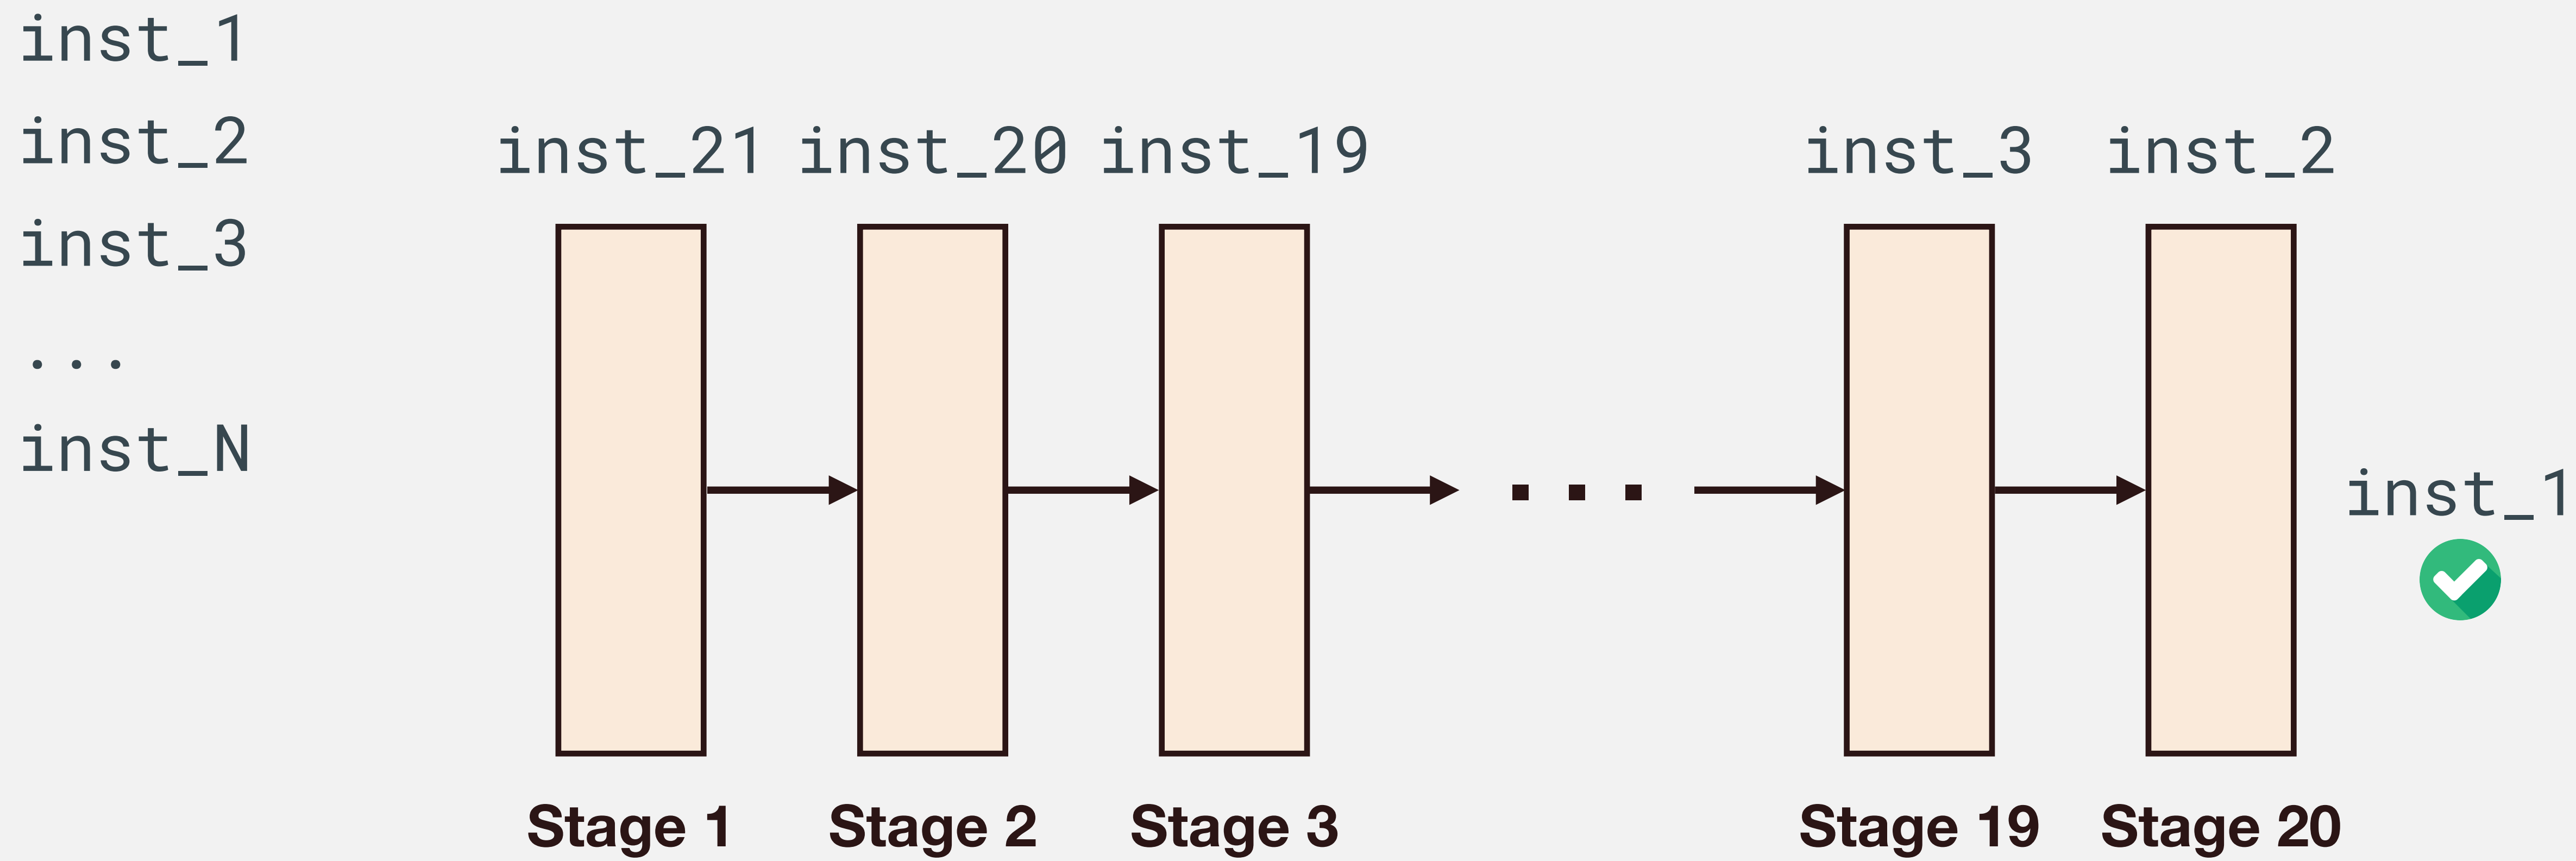


CPU Pipeline

inst_1
inst_2
inst_3
...
inst_N

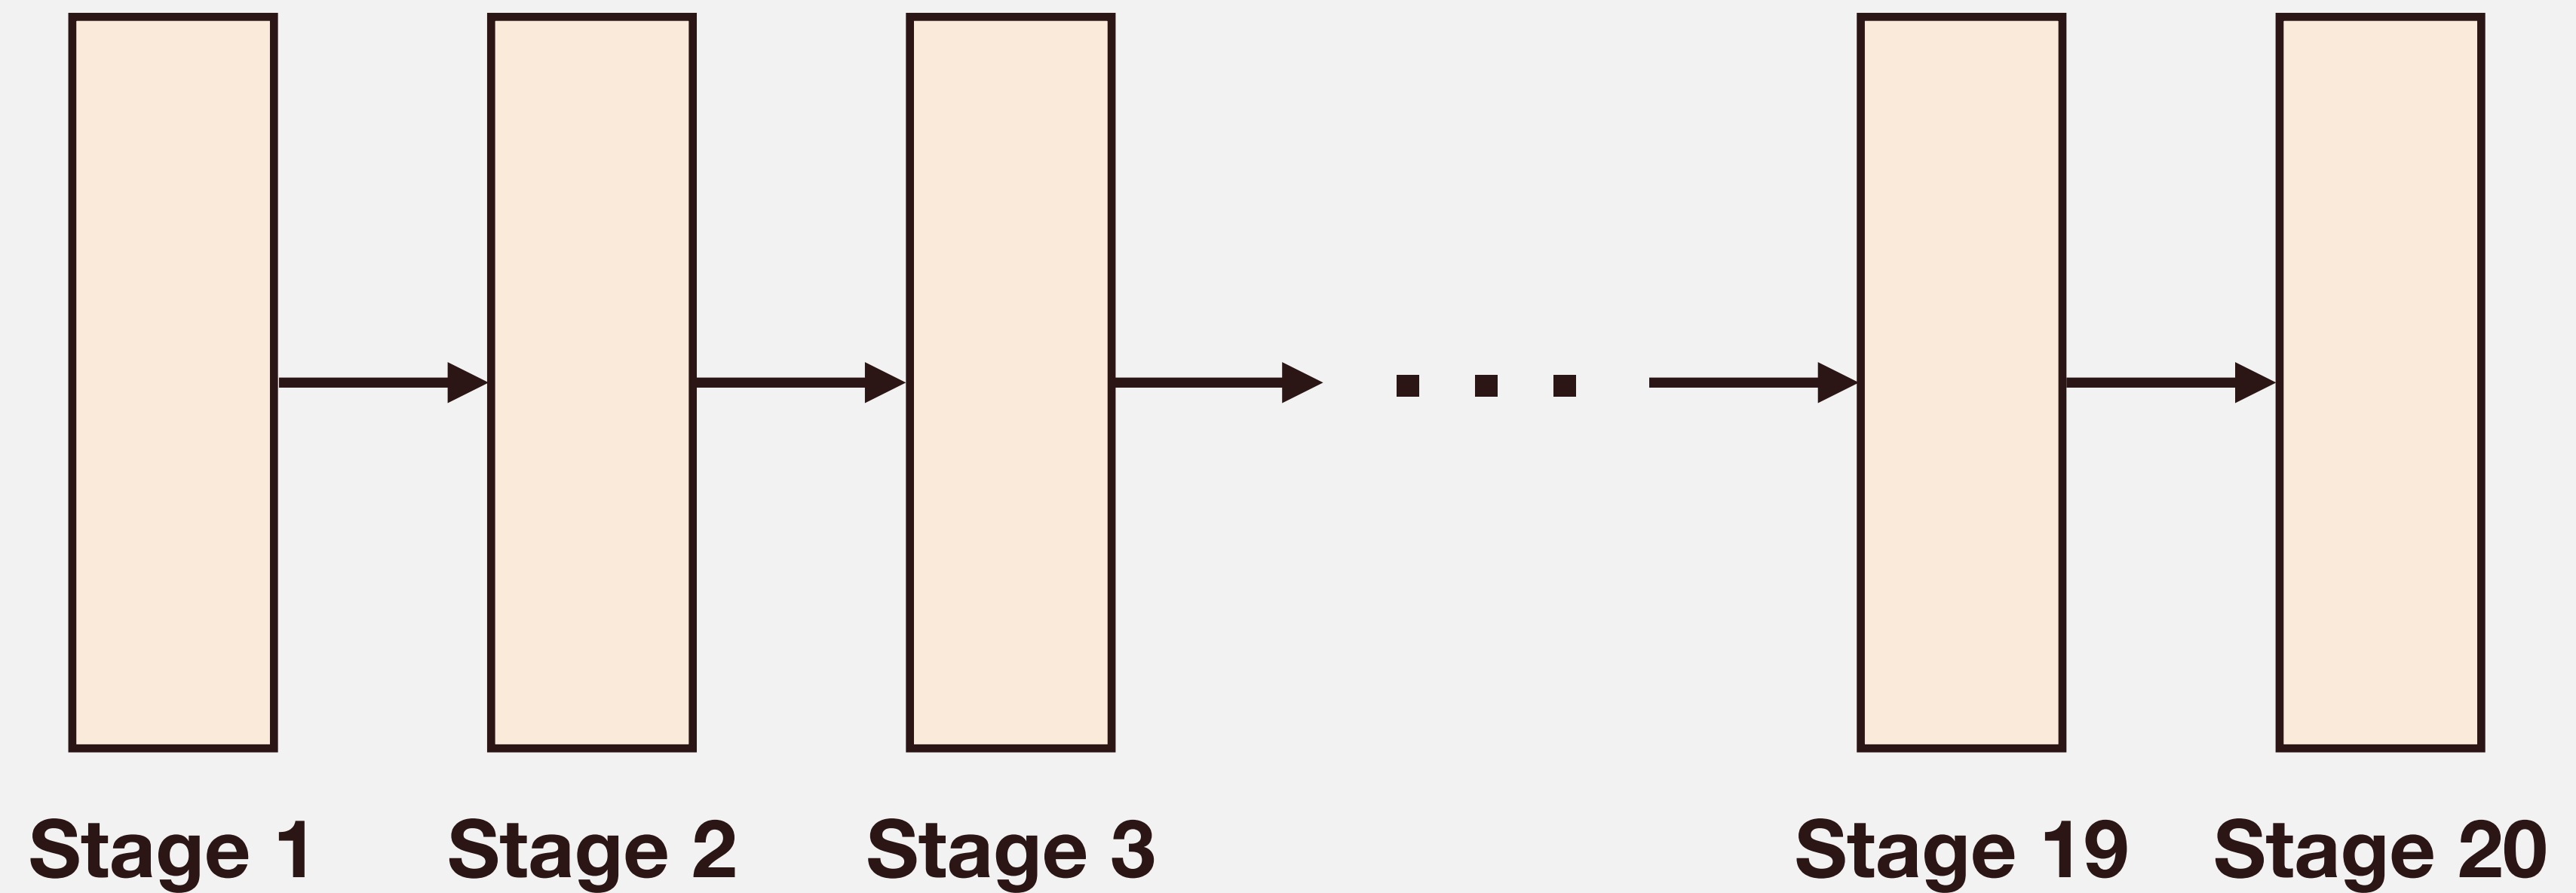


CPU Pipeline



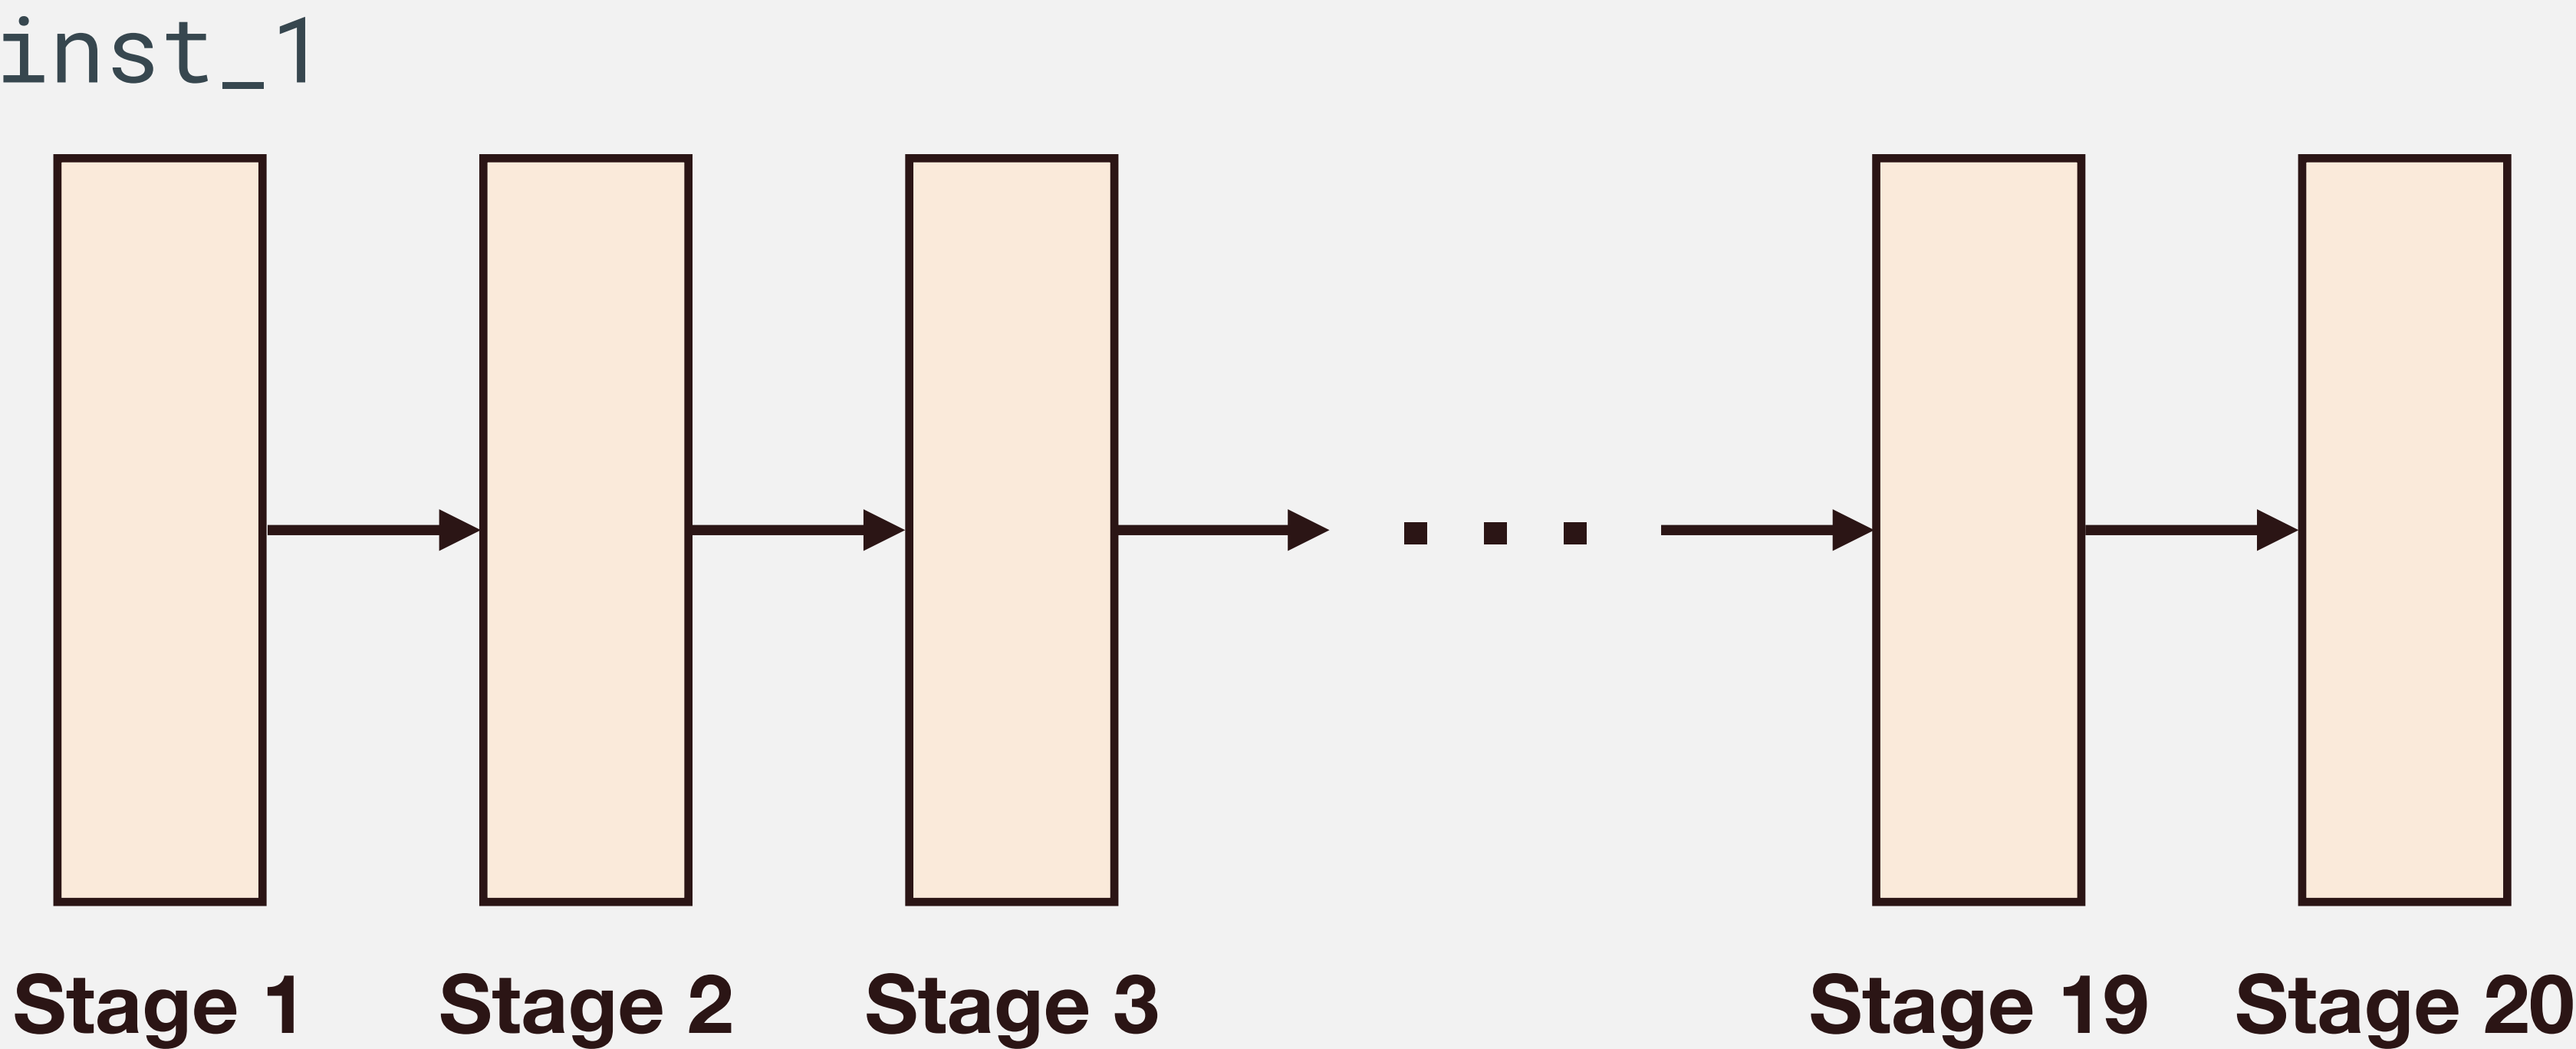
CPU Pipeline

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



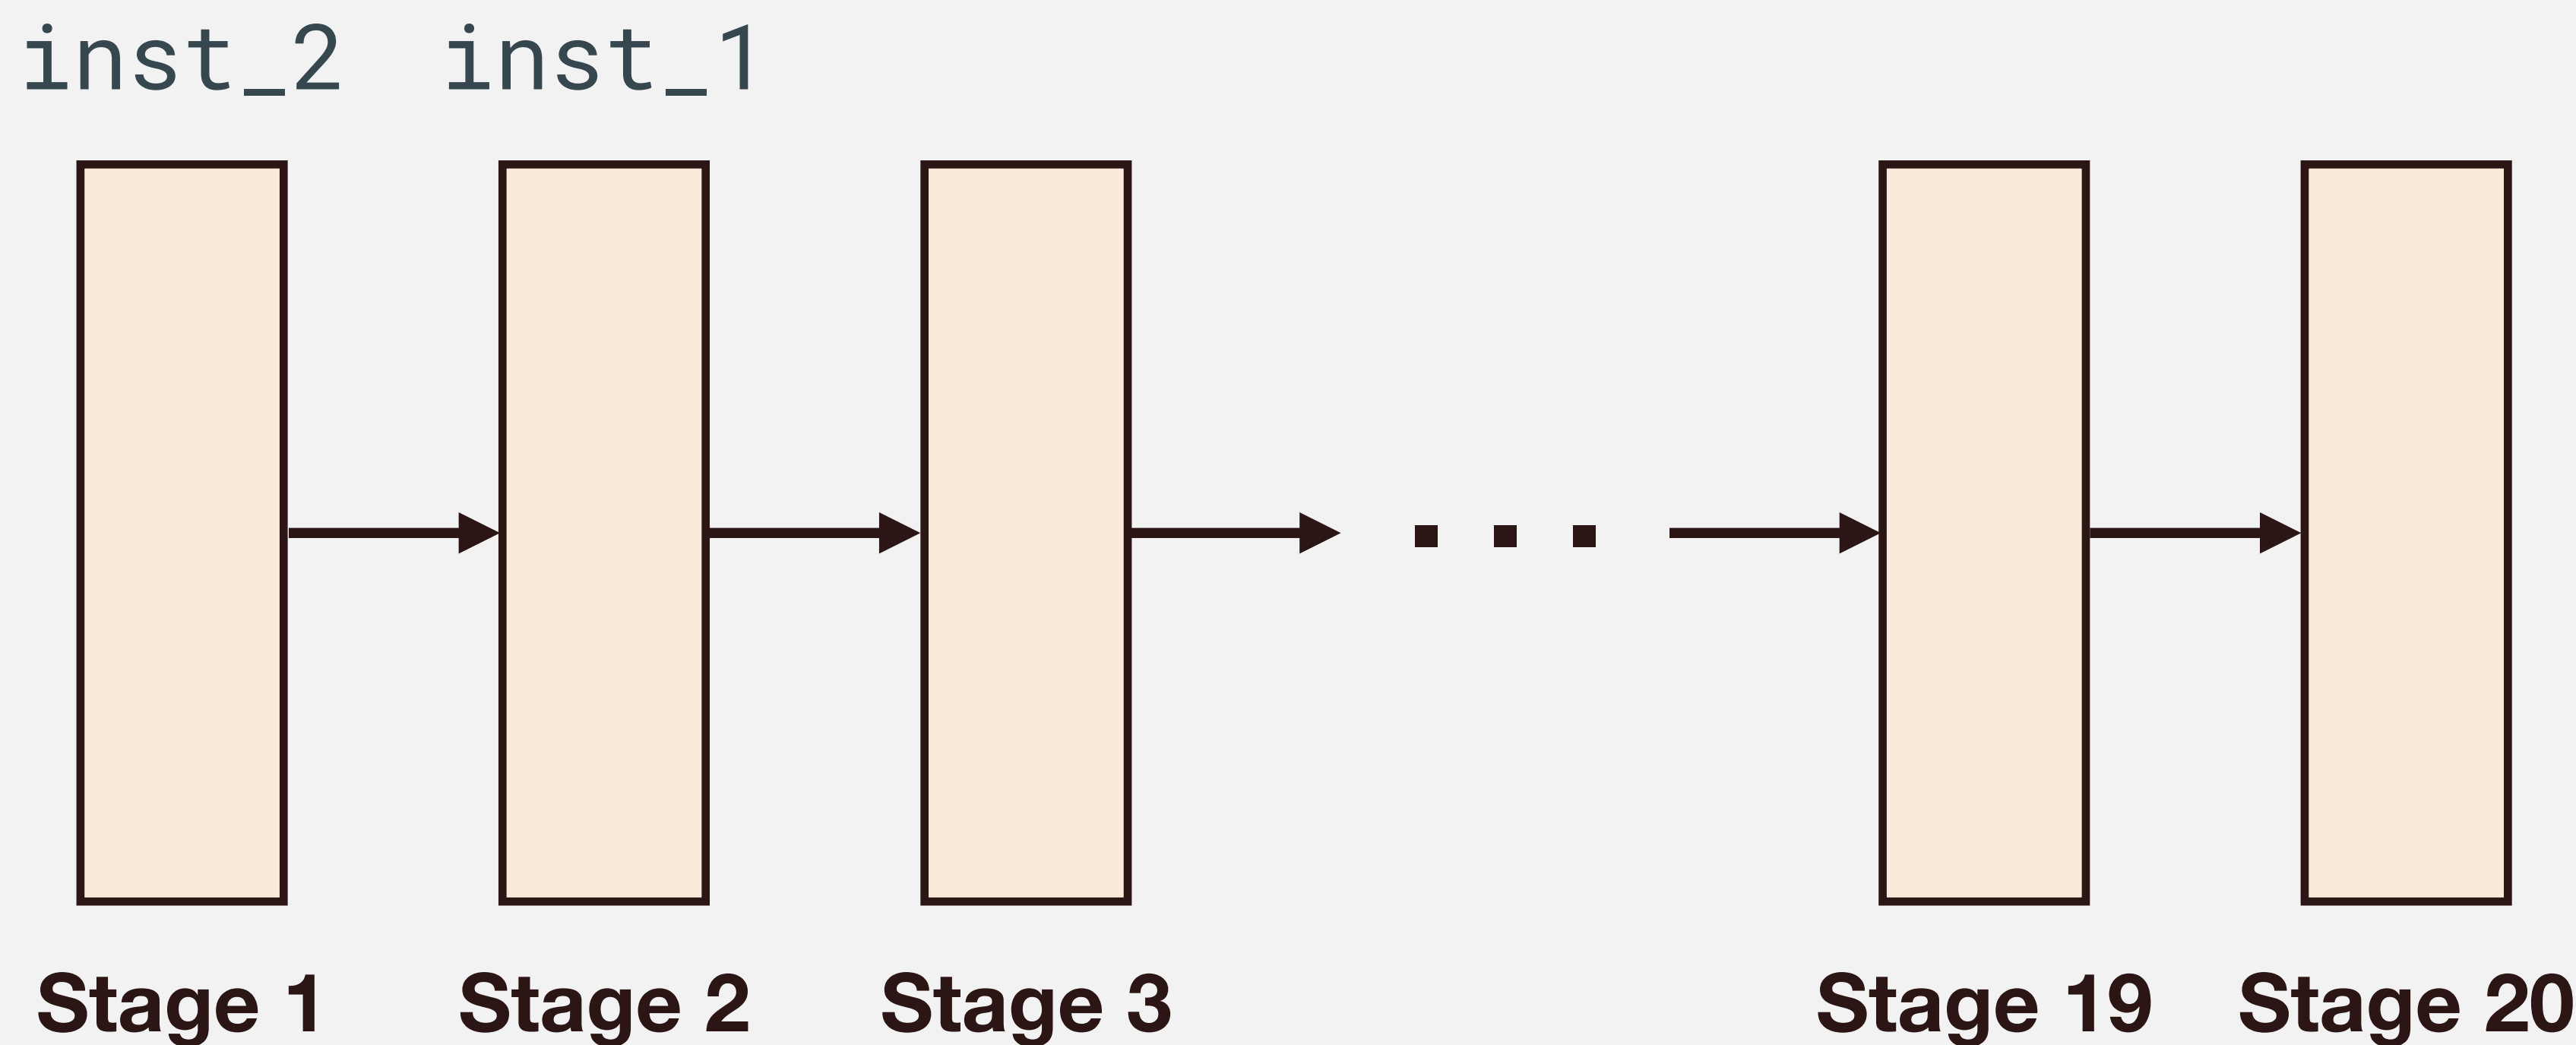
CPU Pipeline

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



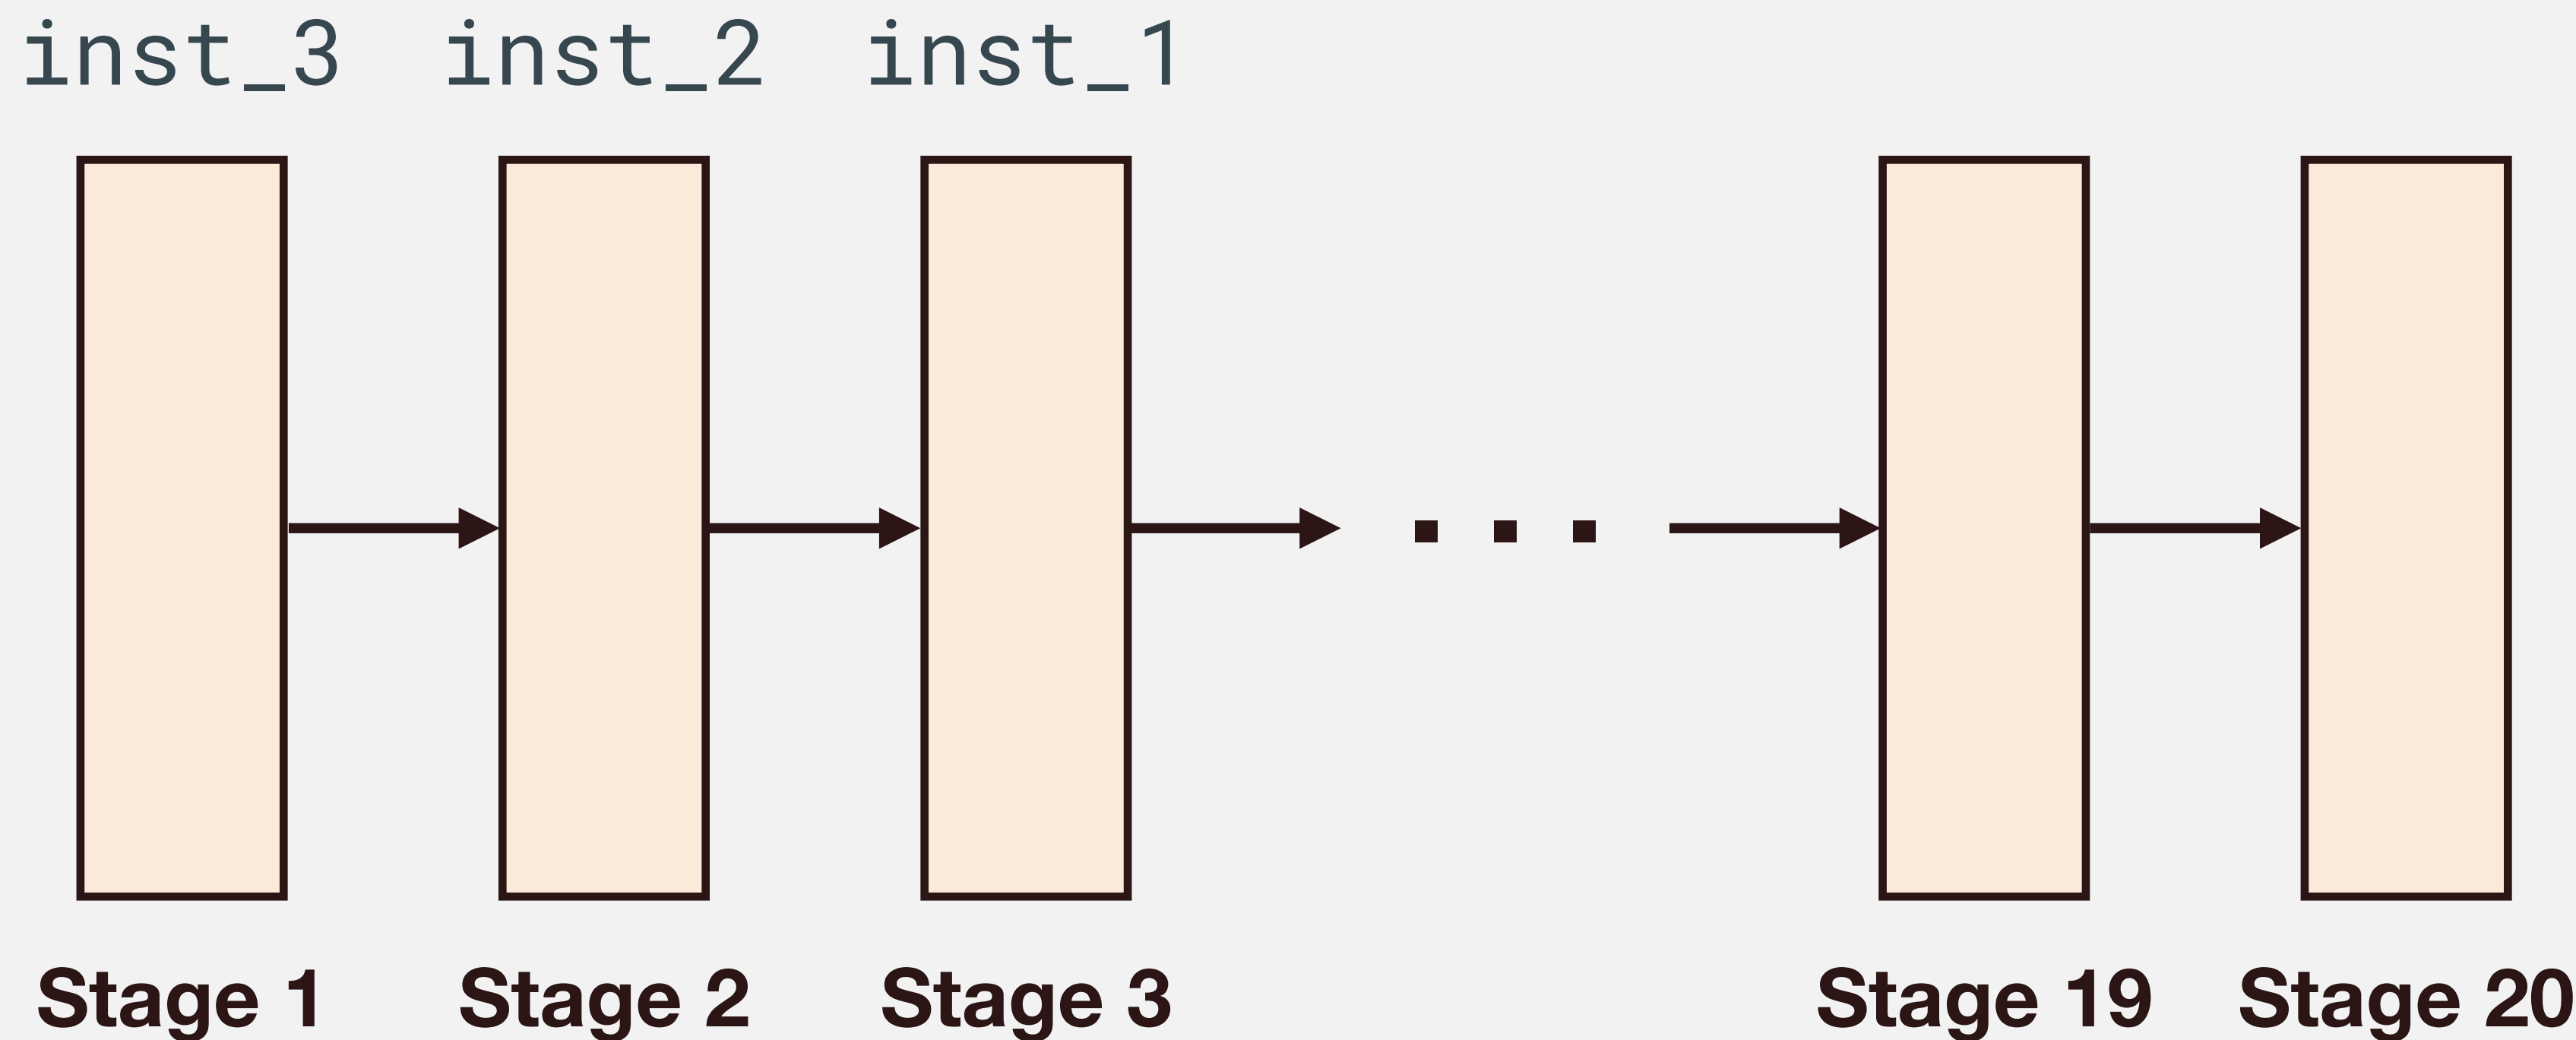
CPU Pipeline

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



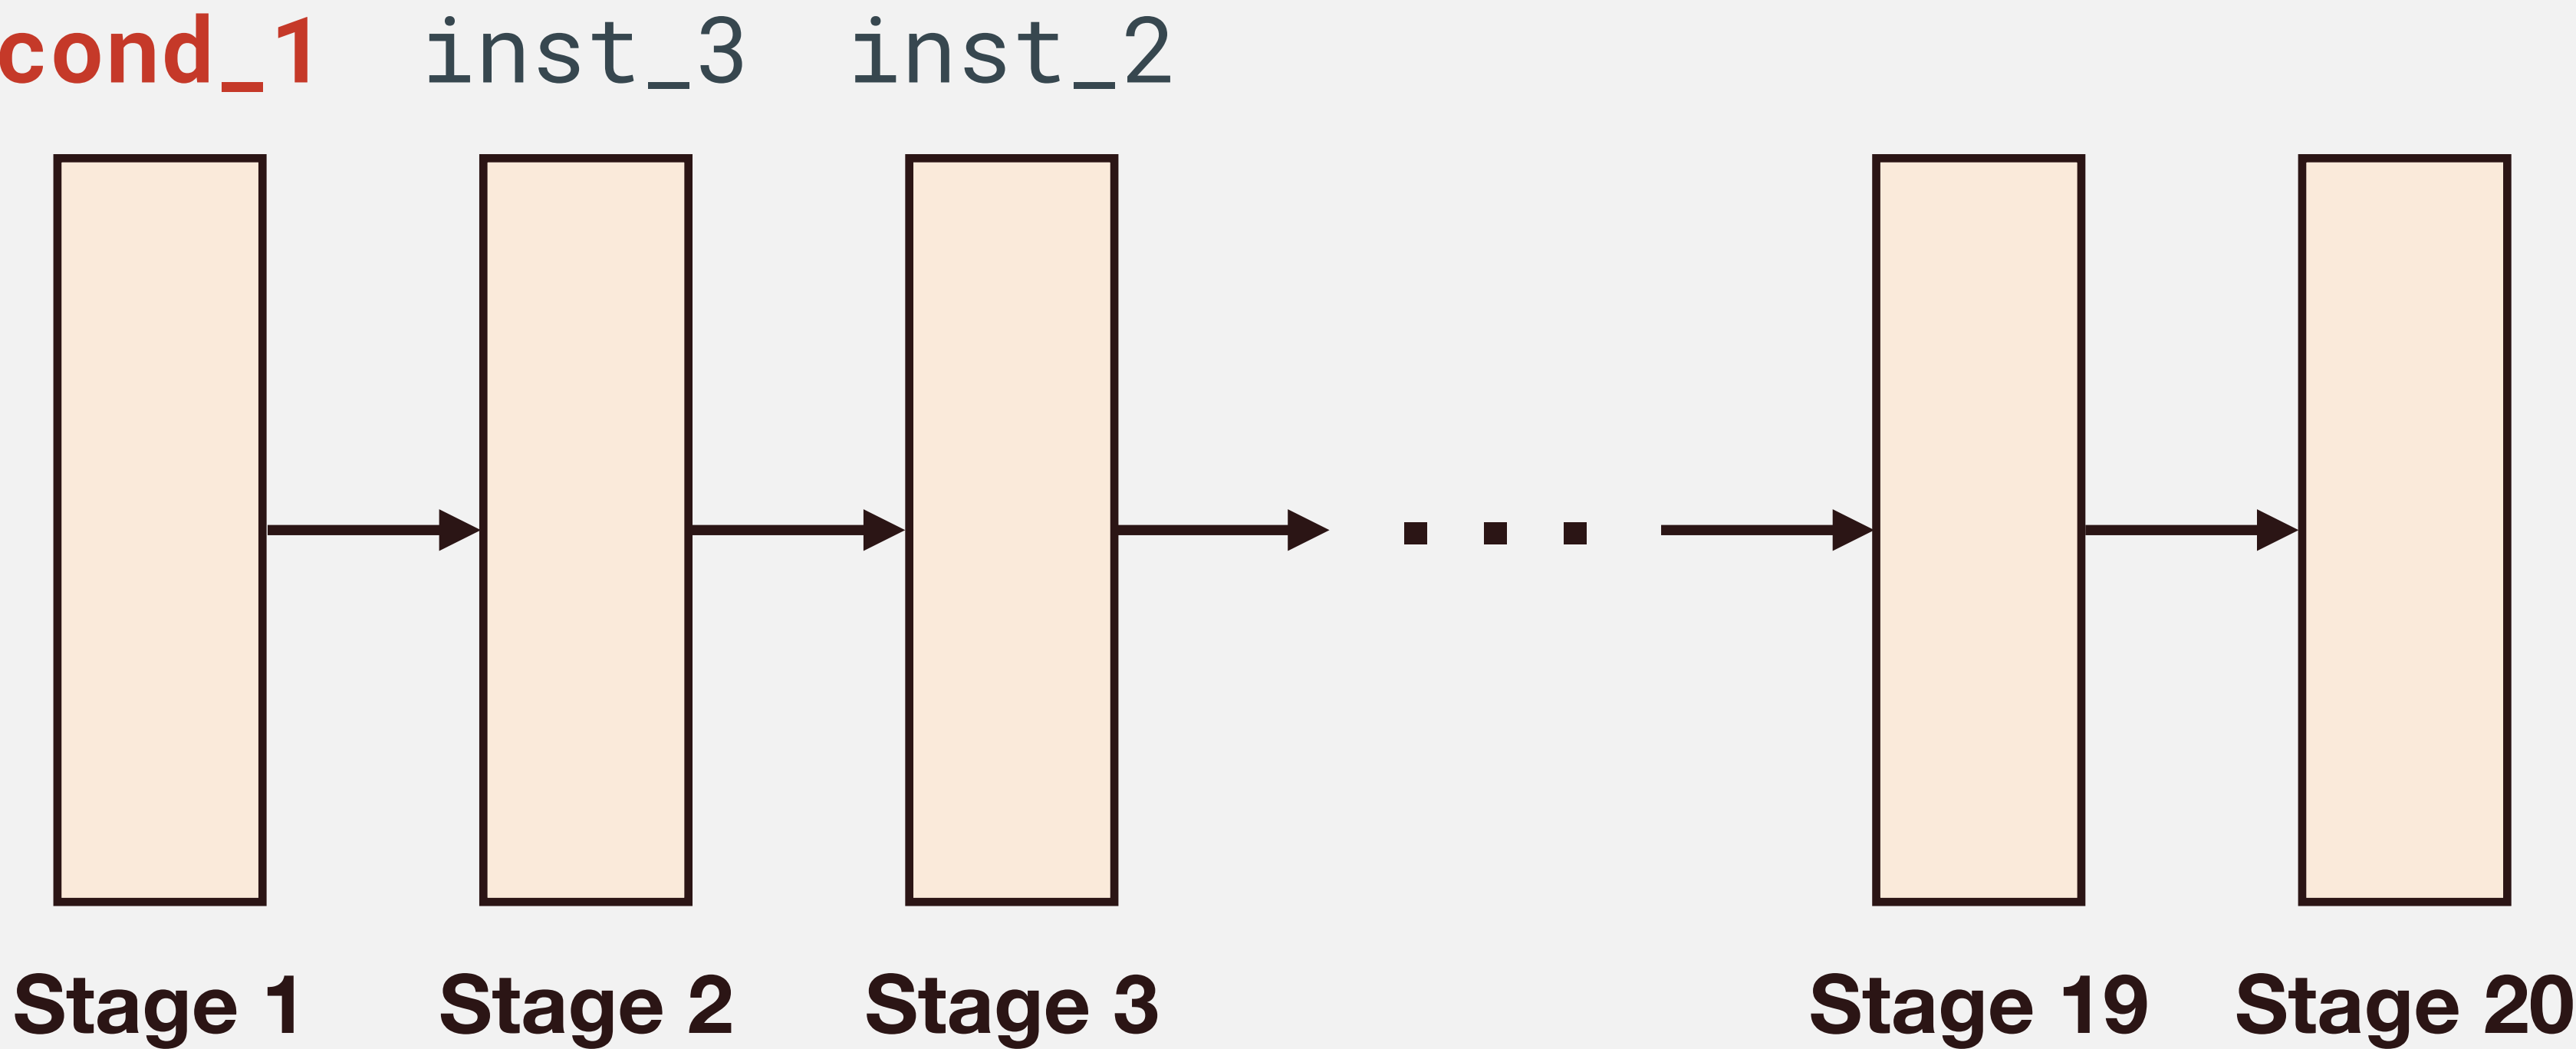
CPU Pipeline

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



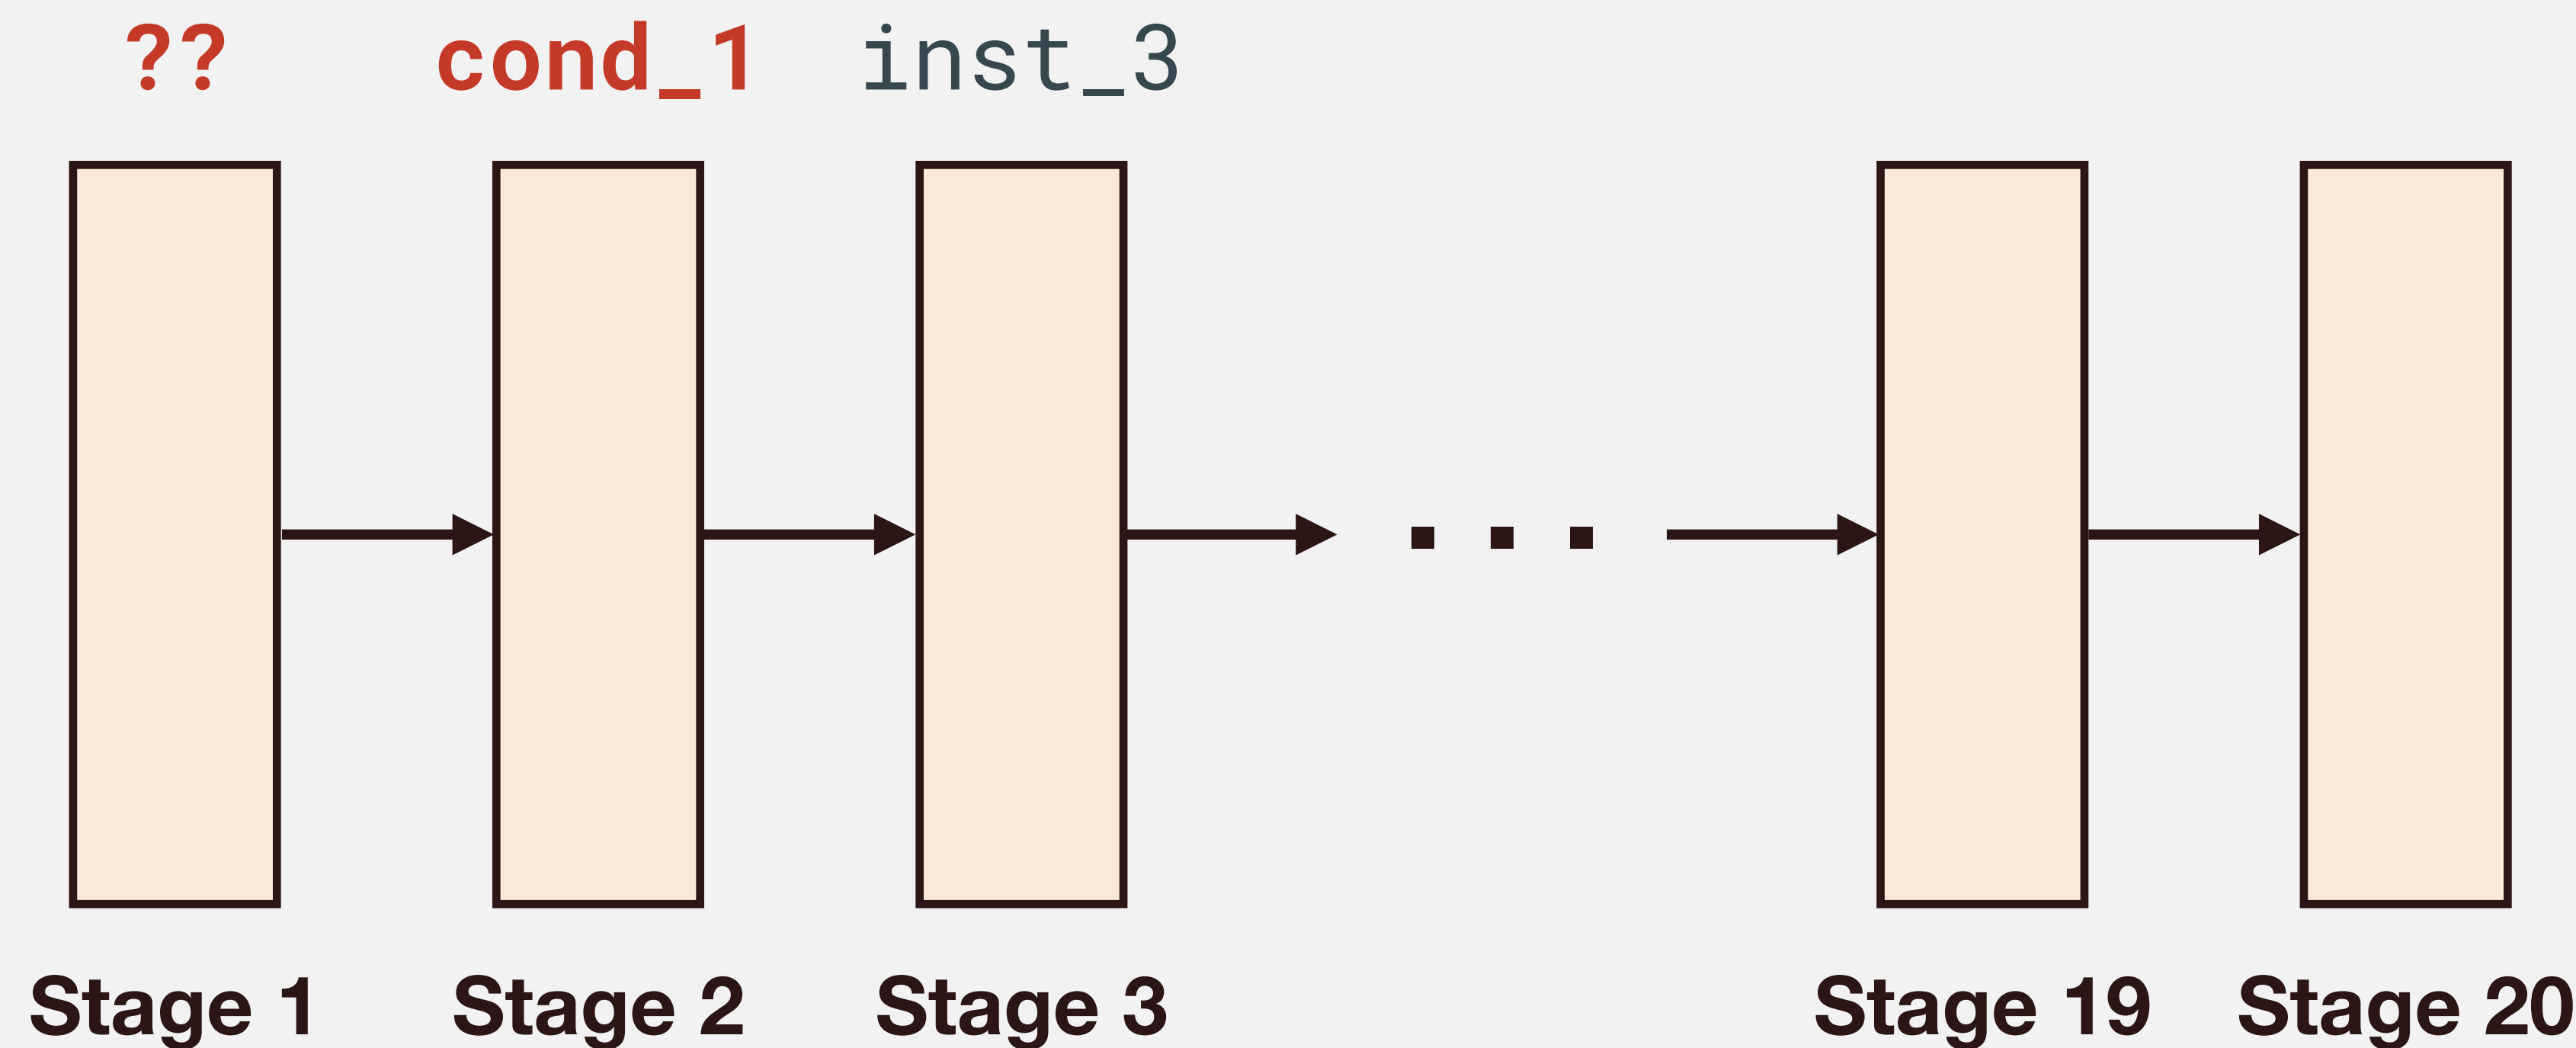
CPU Pipeline

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



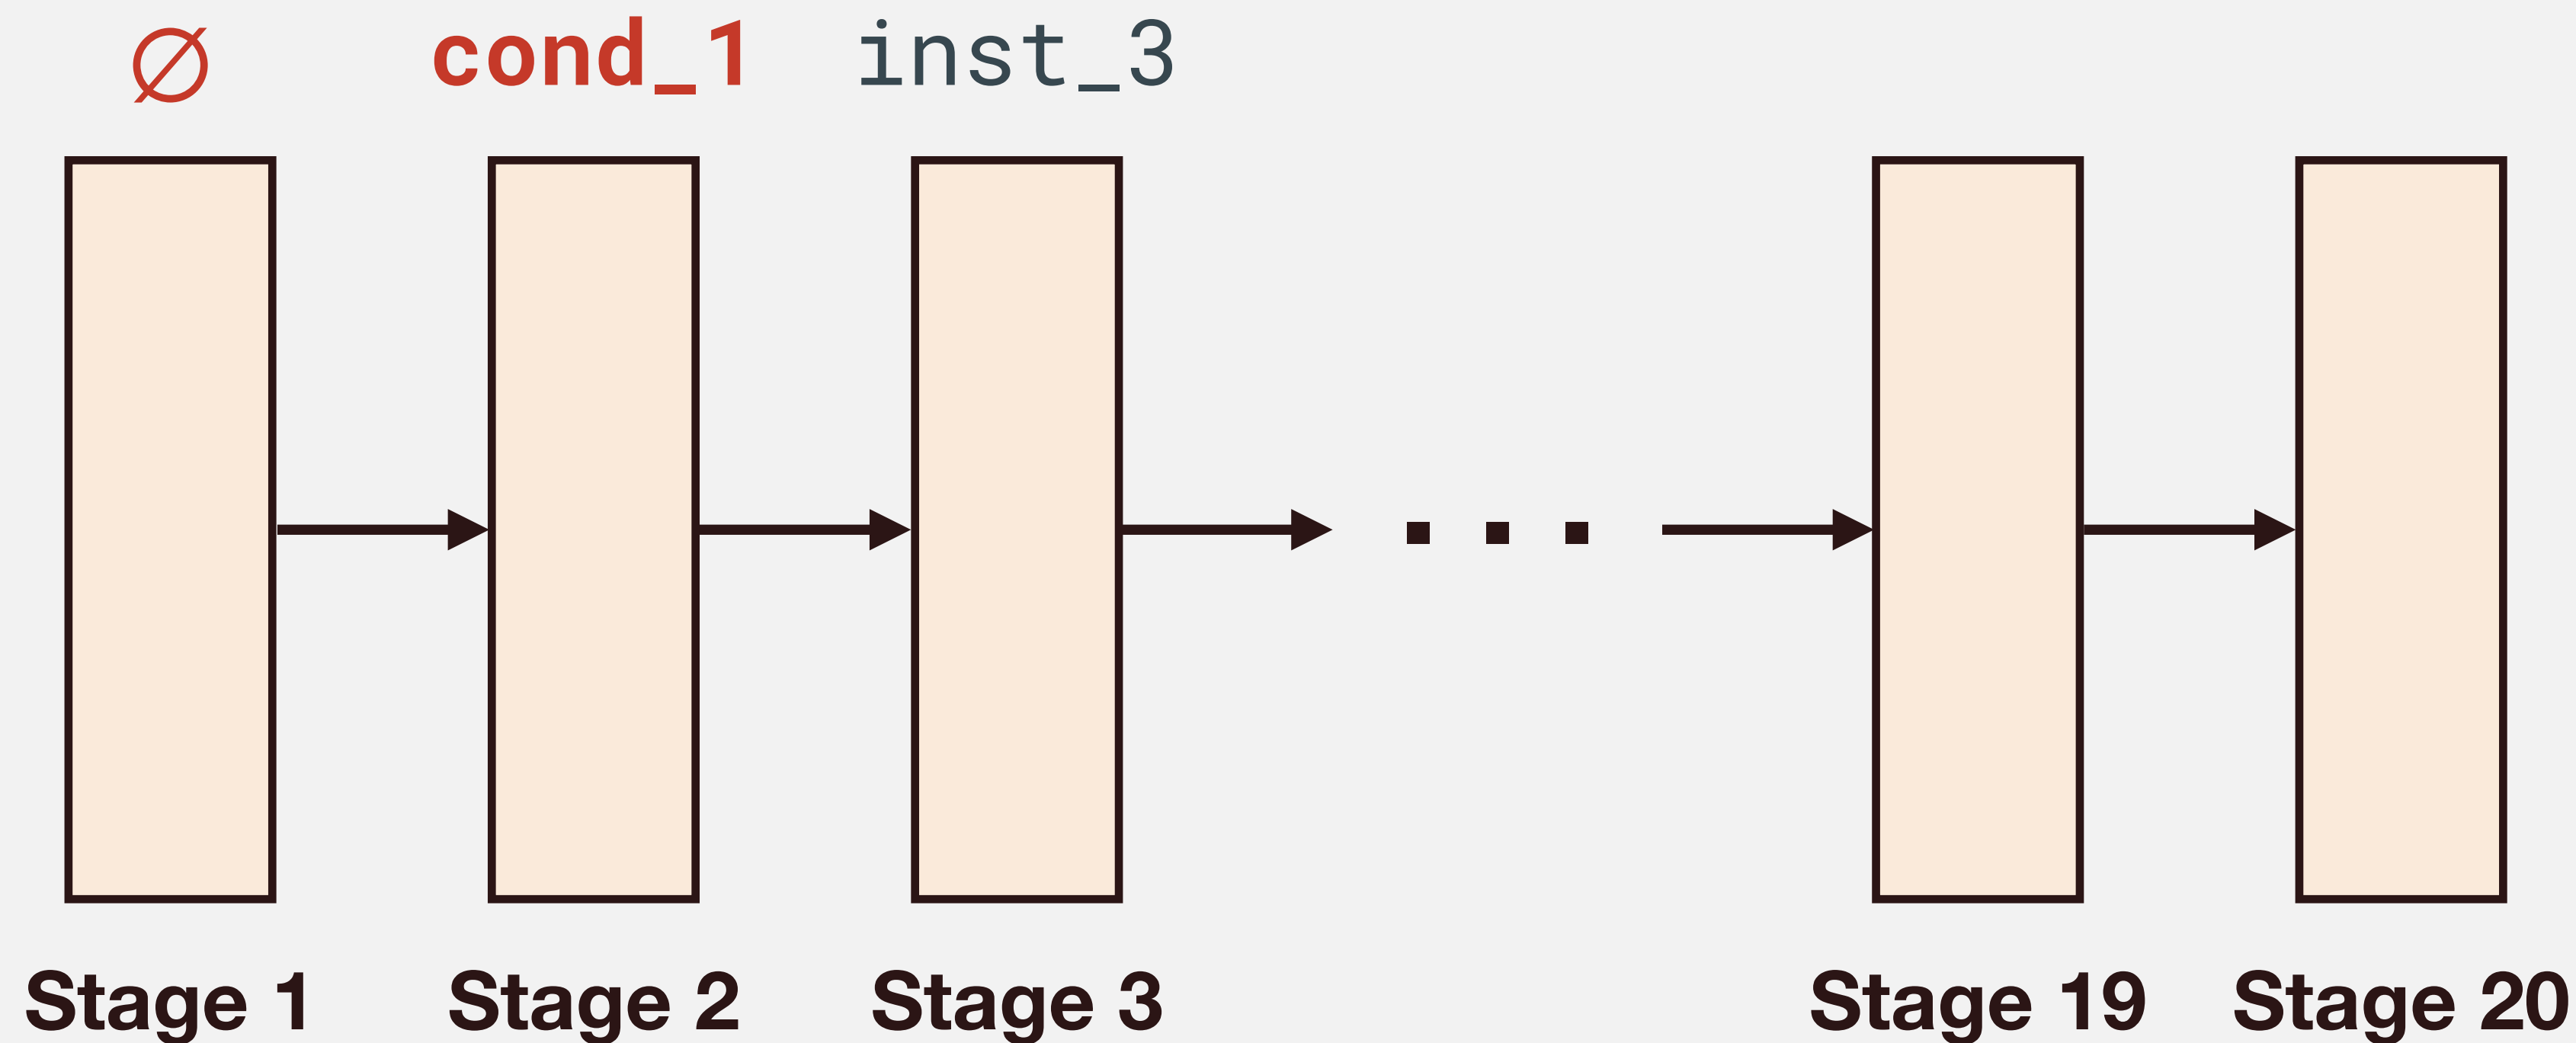
CPU Pipeline

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



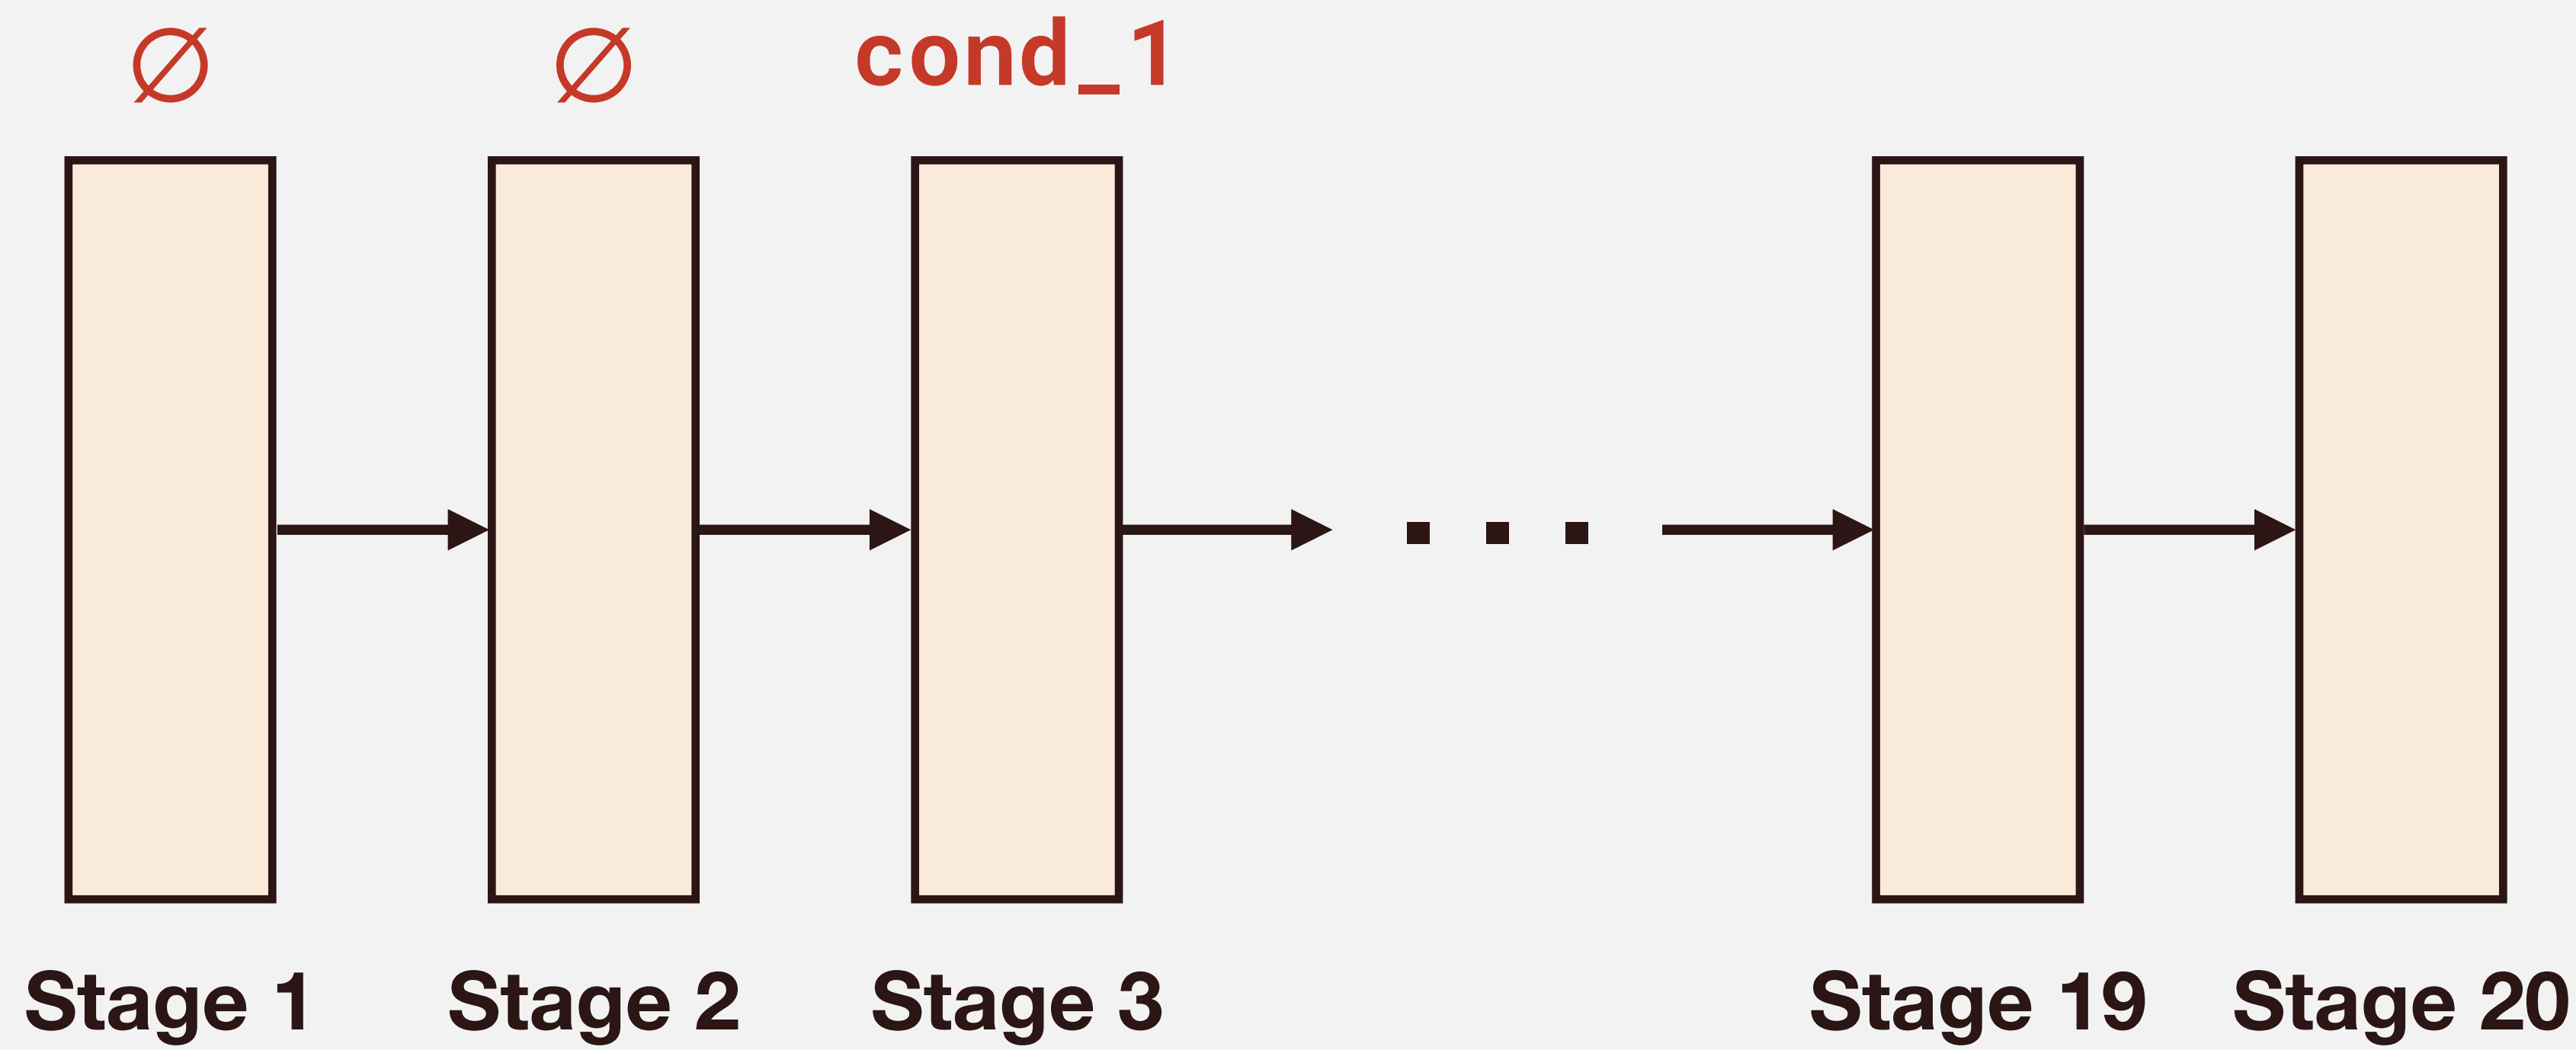
CPU Pipeline

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



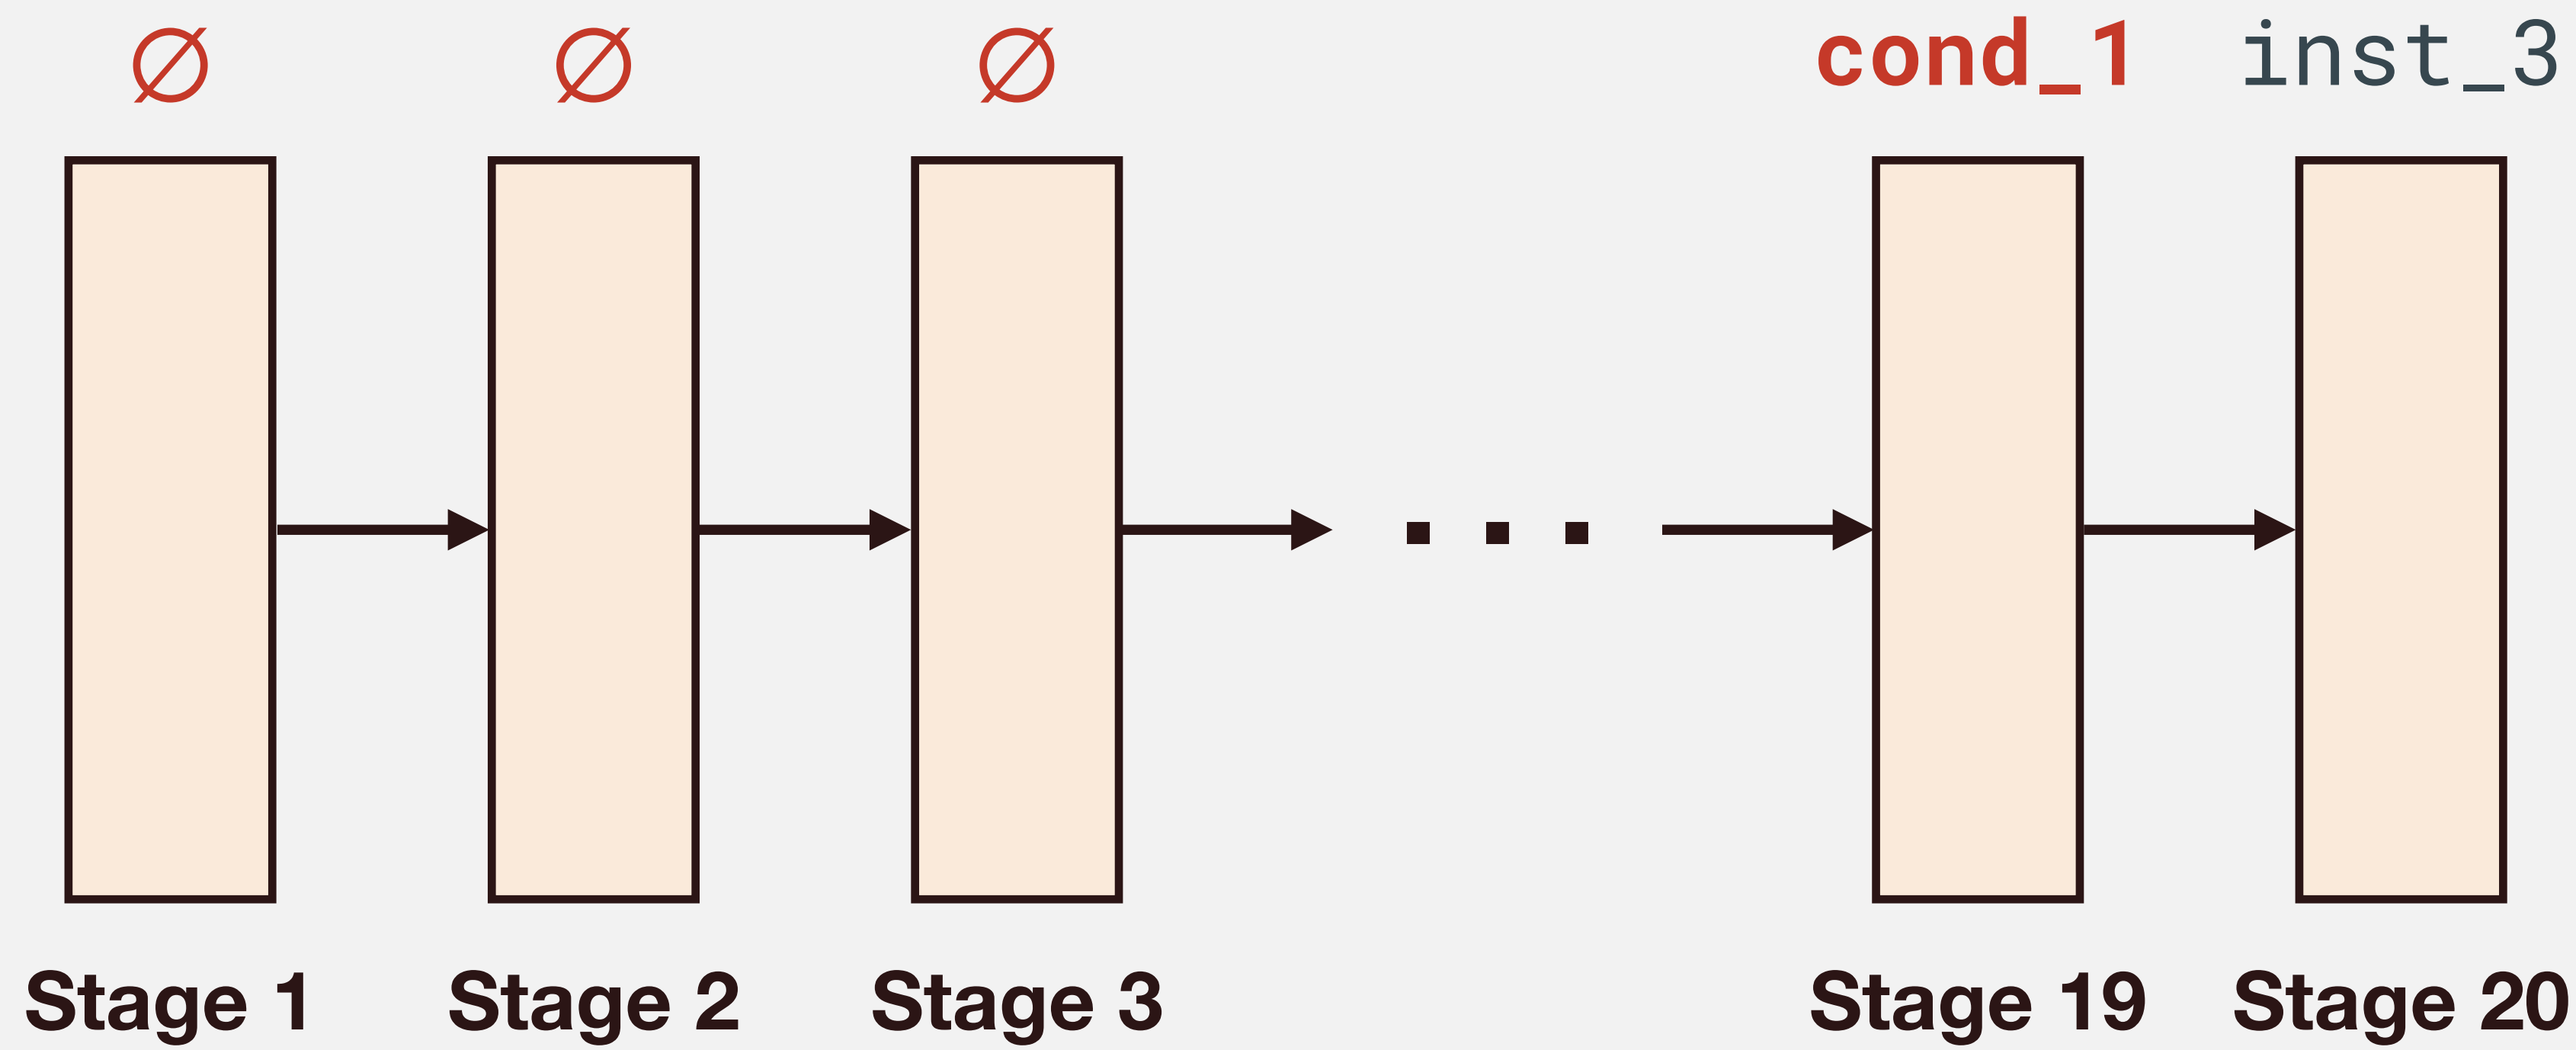
CPU Pipeline

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



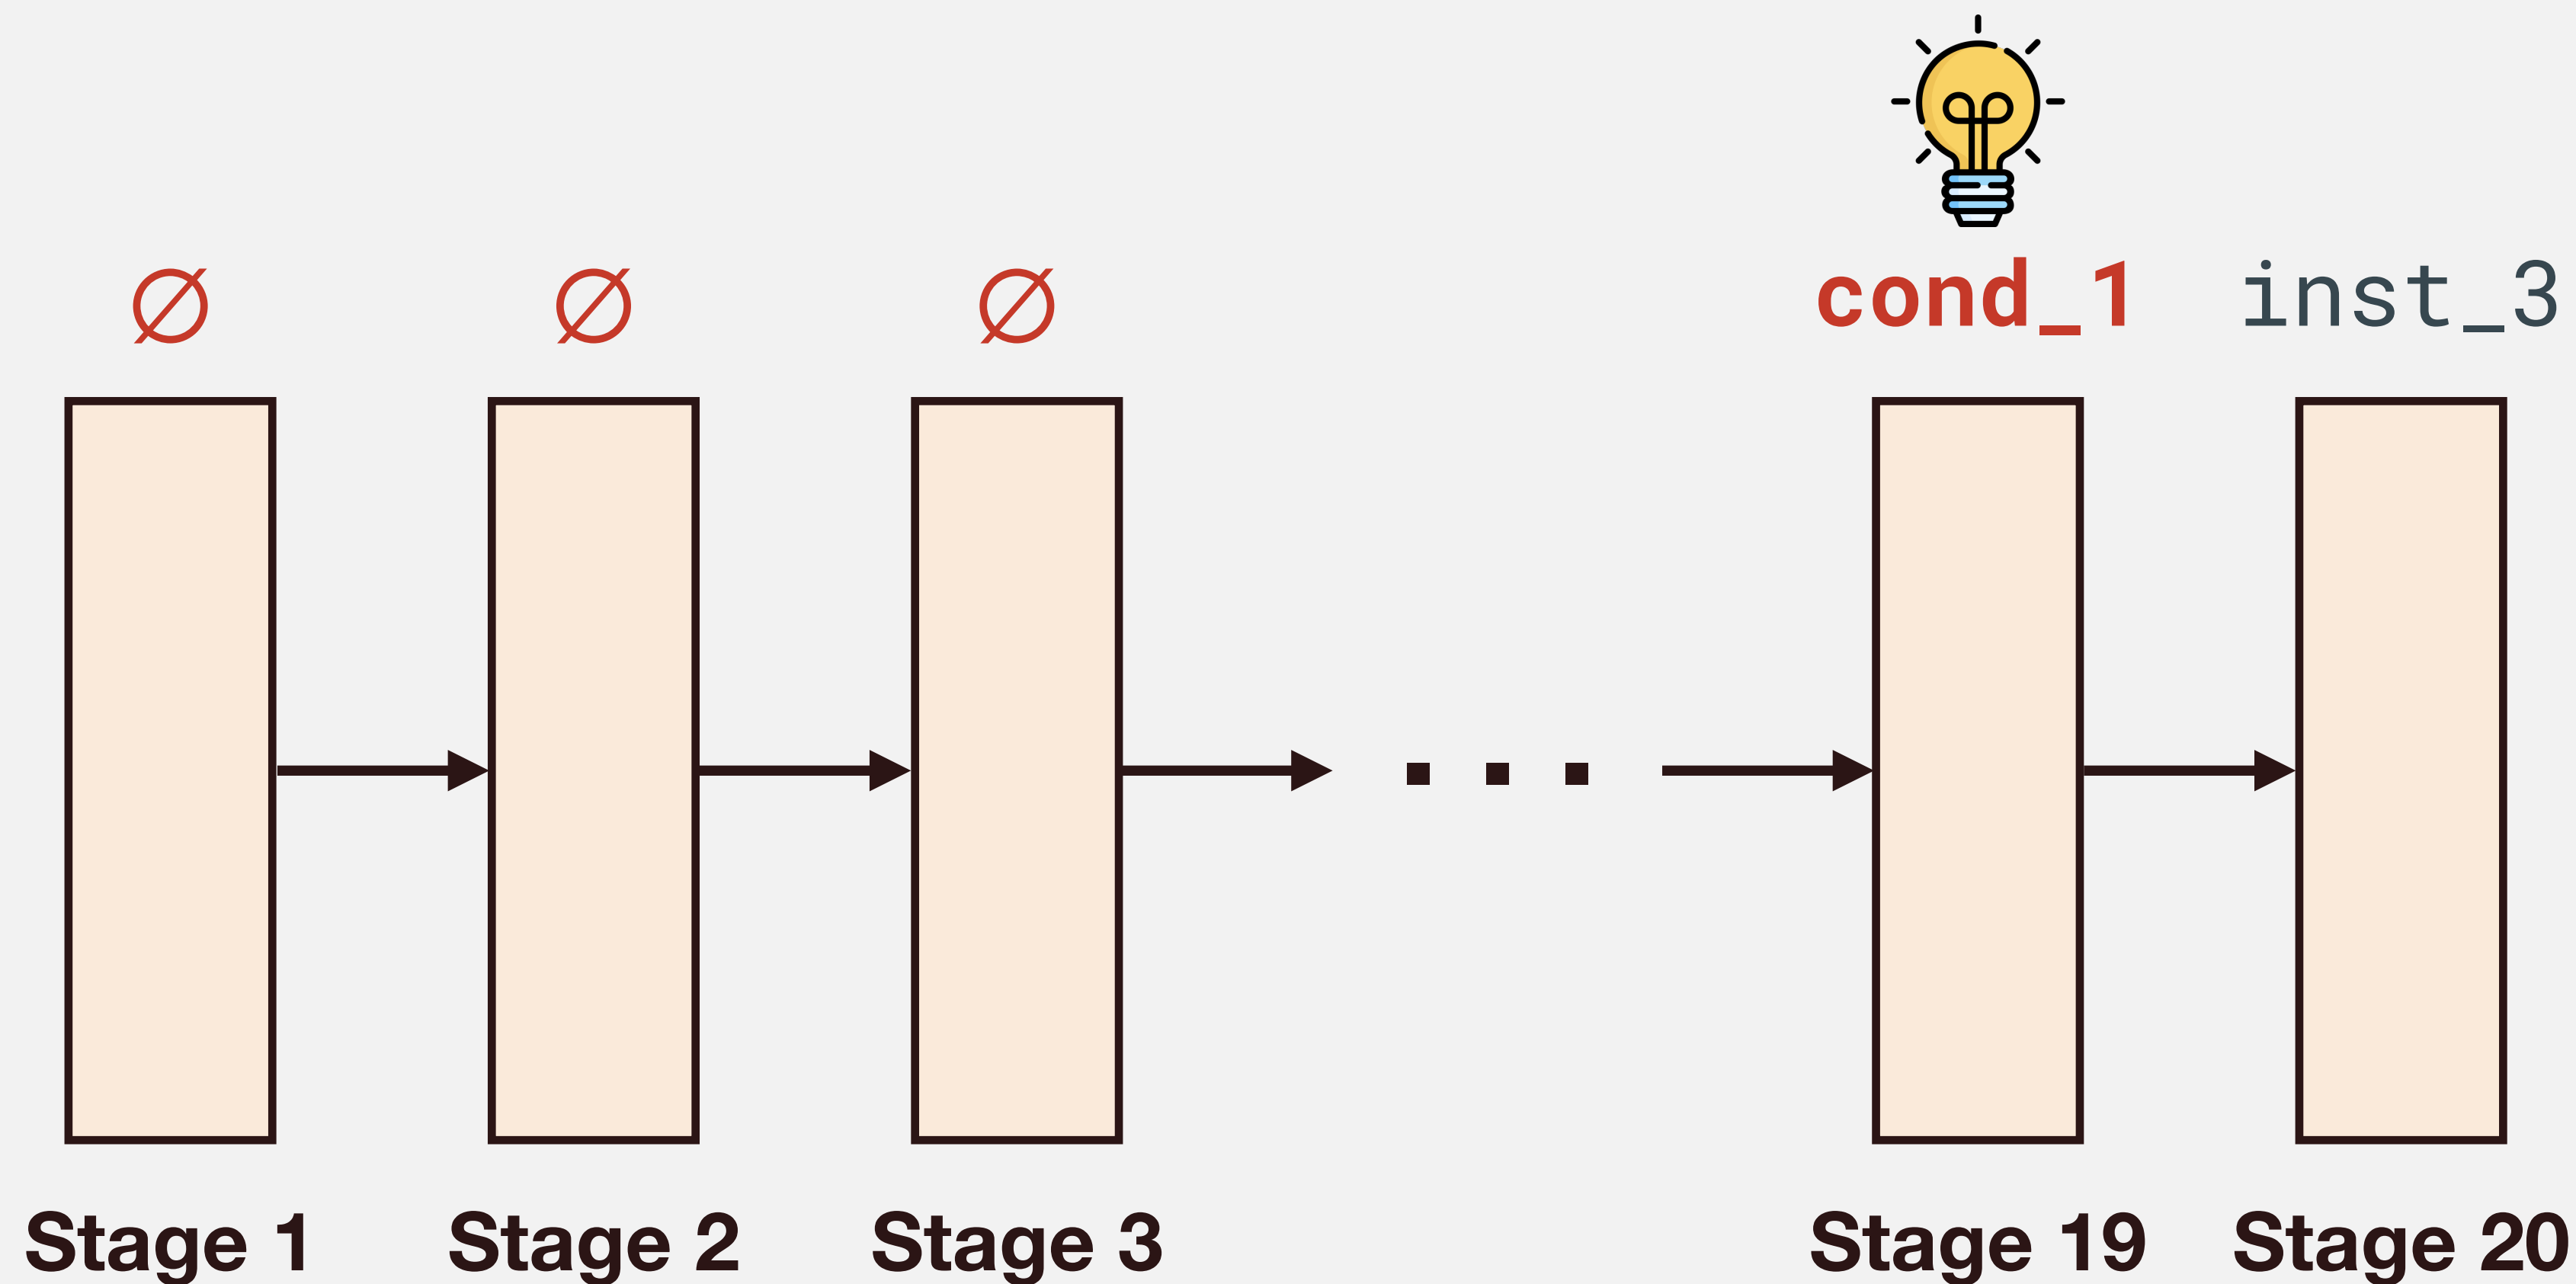
CPU Pipeline

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



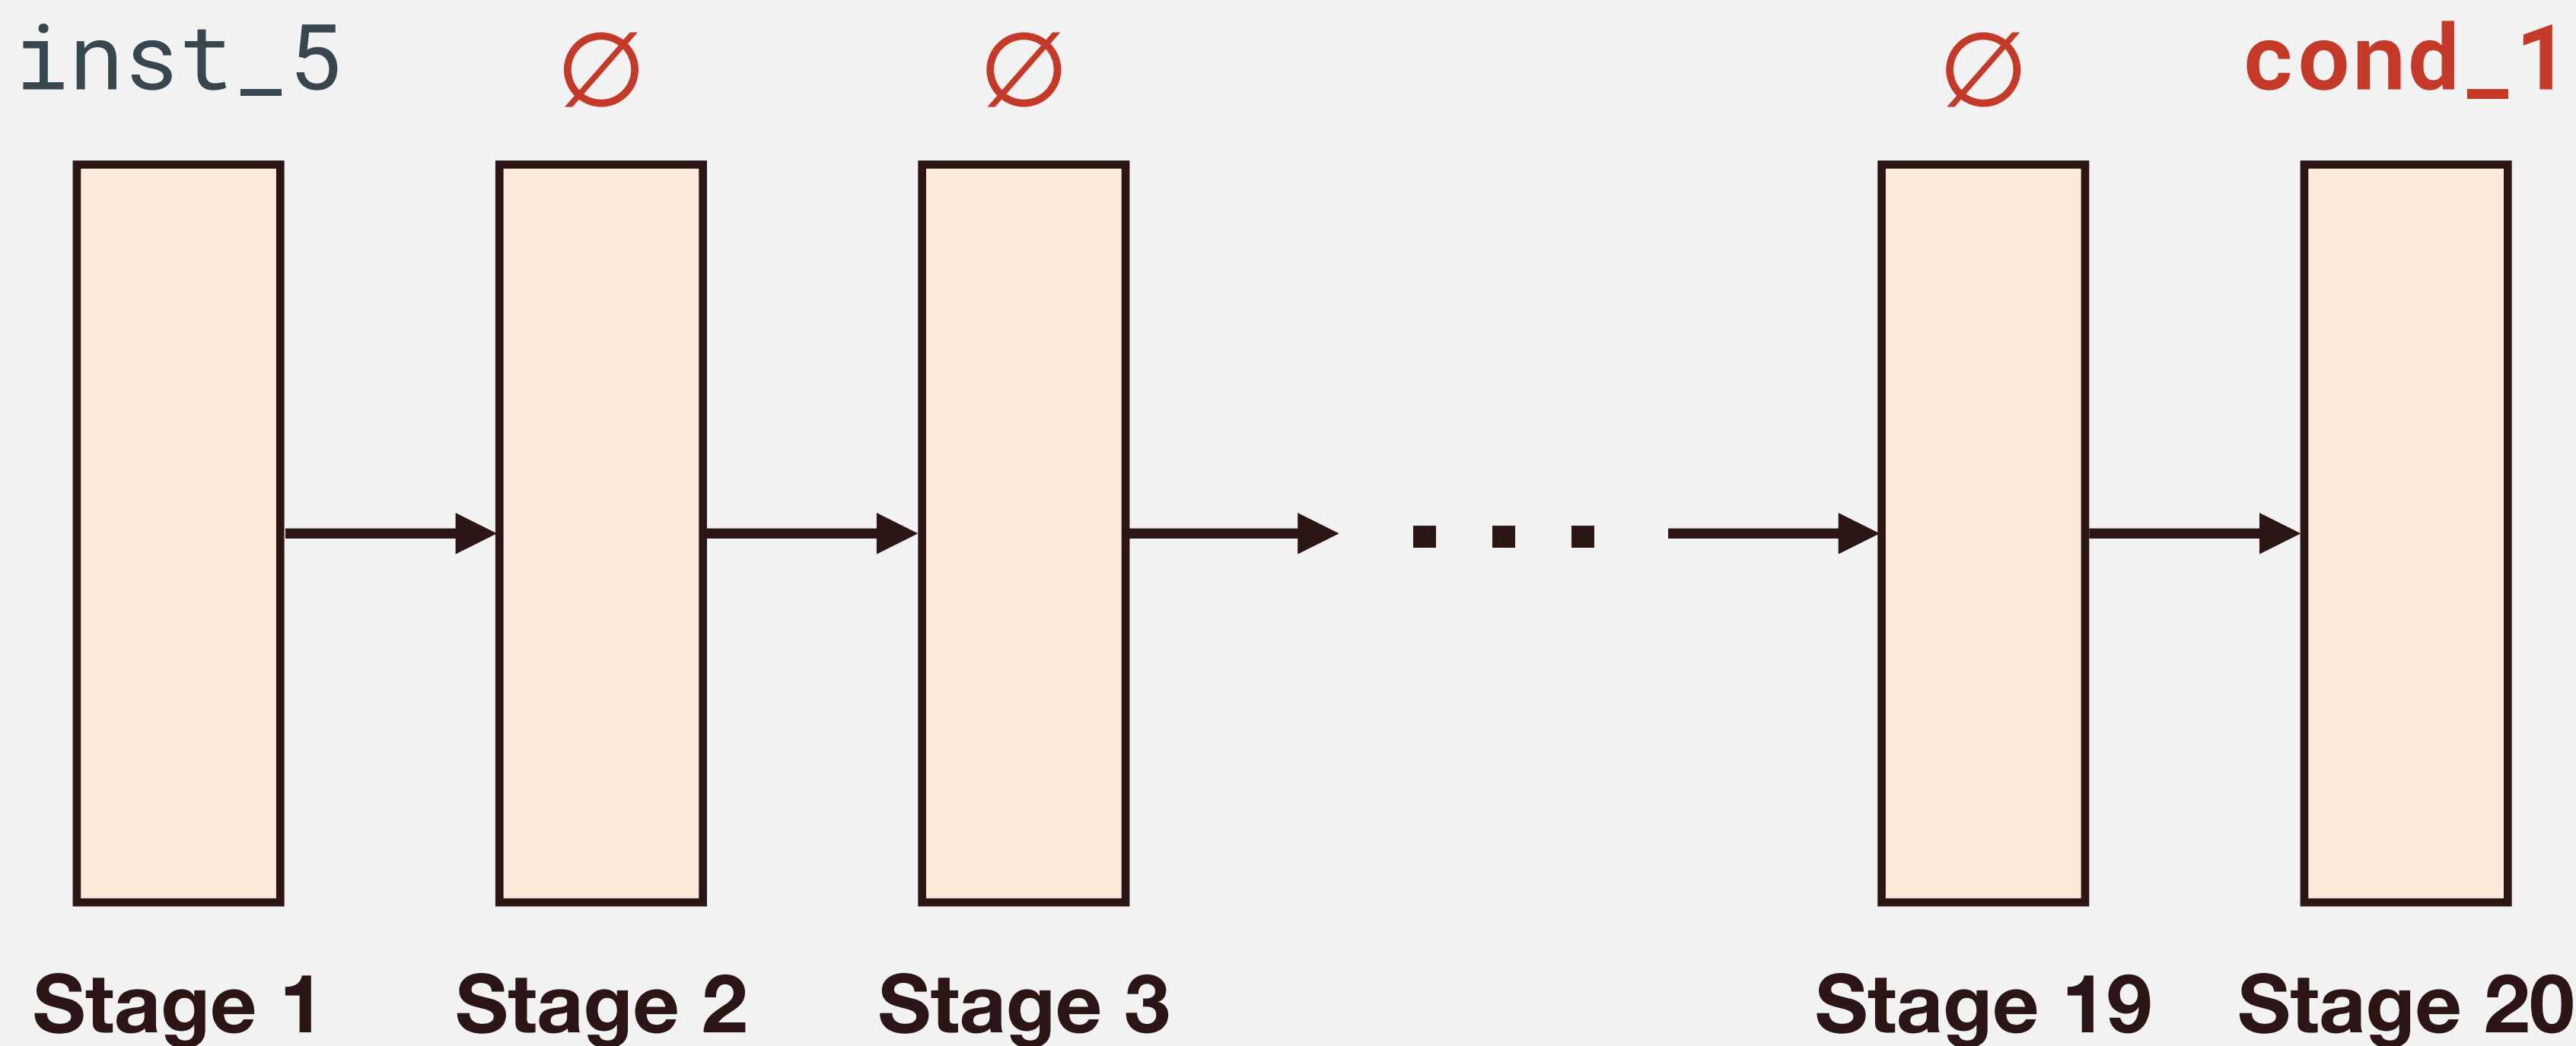
CPU Pipeline

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



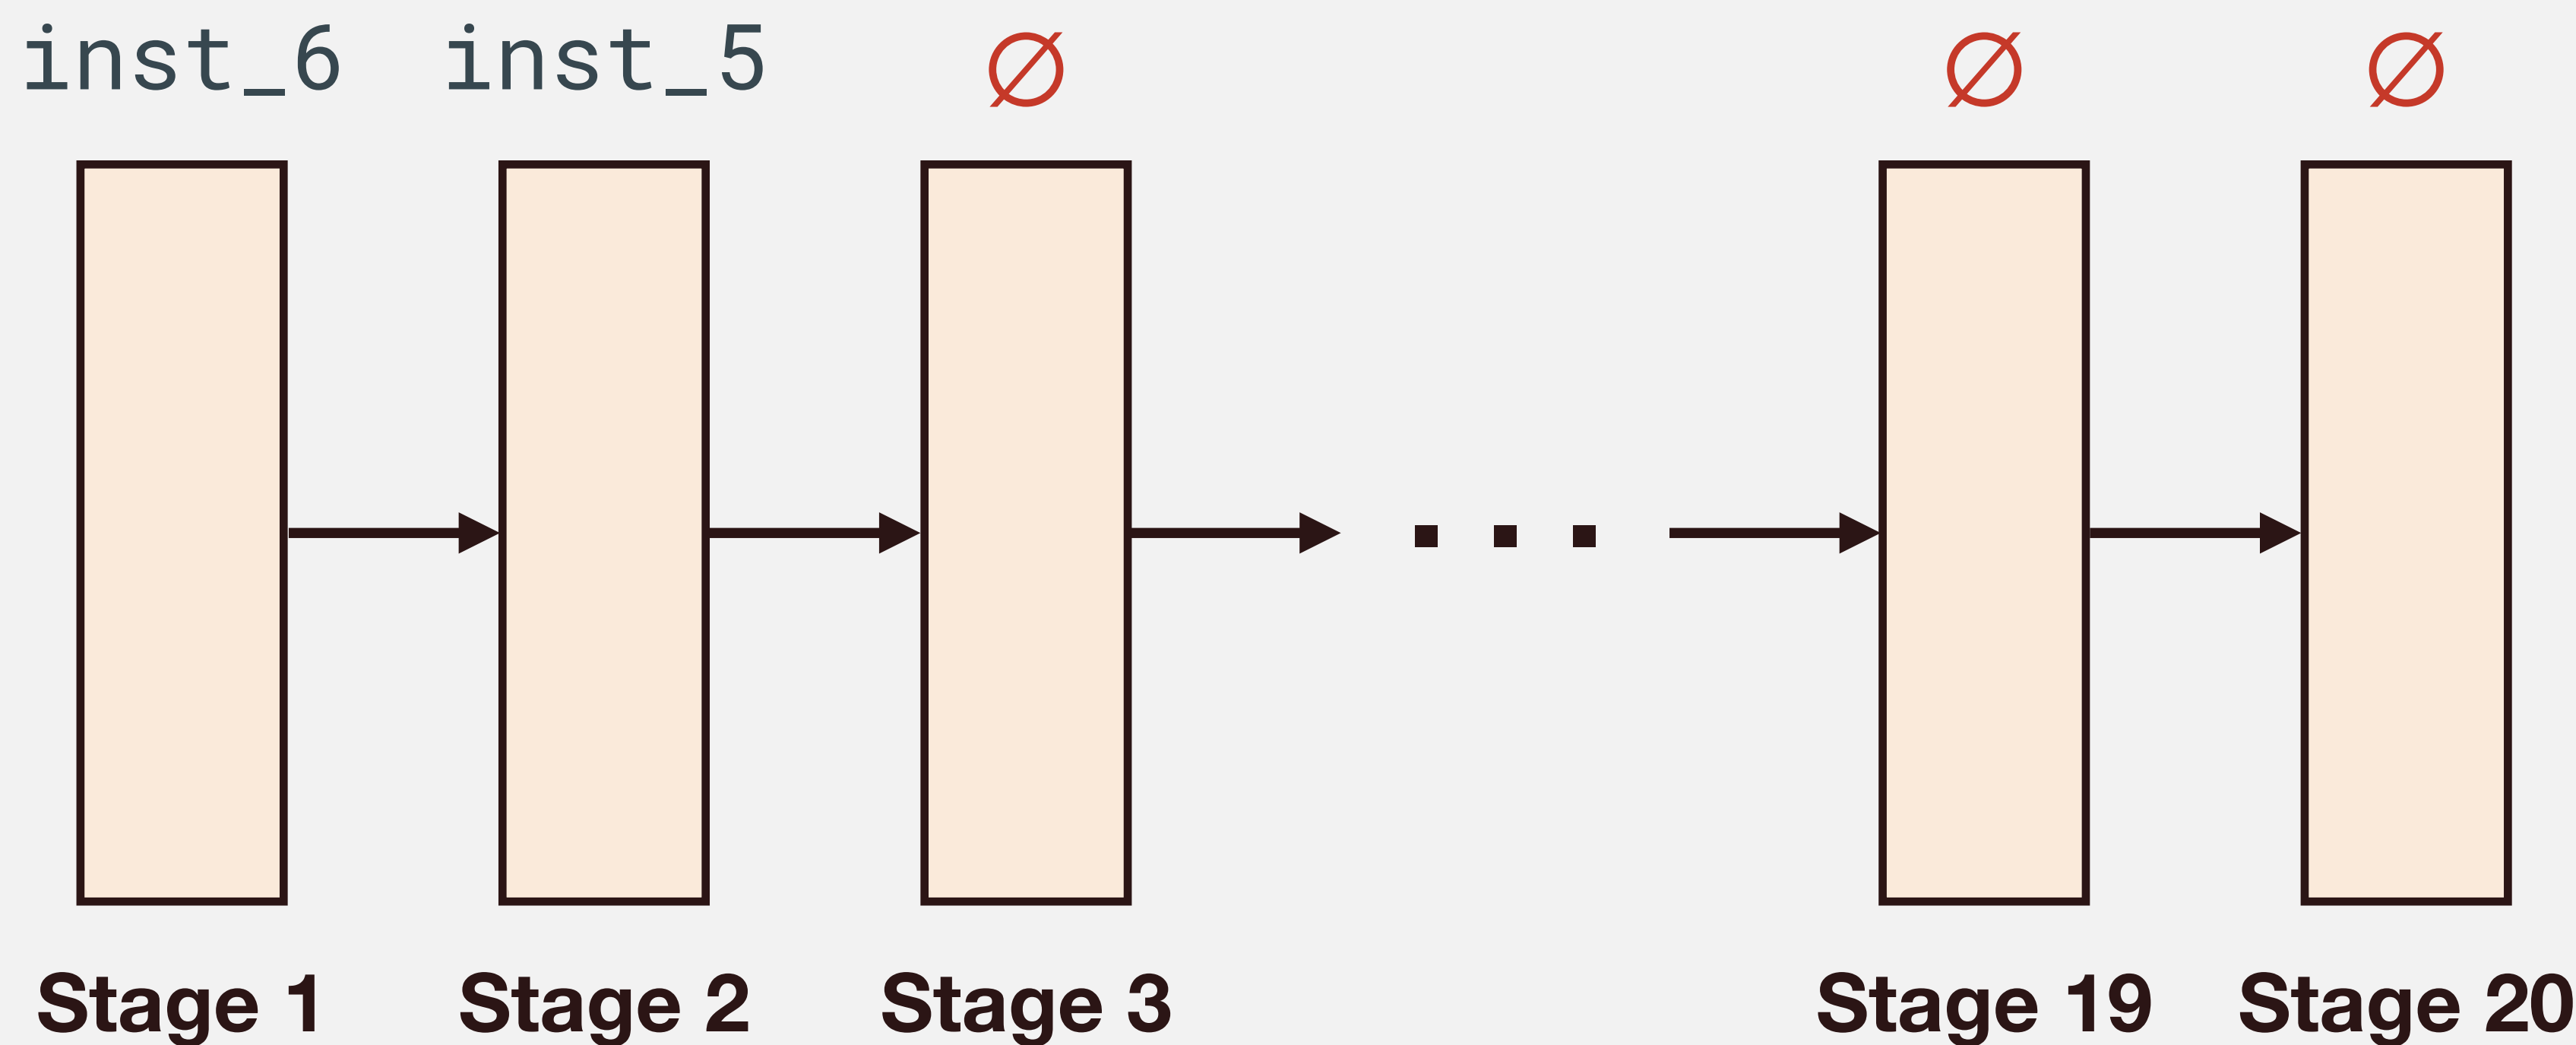
CPU Pipeline

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



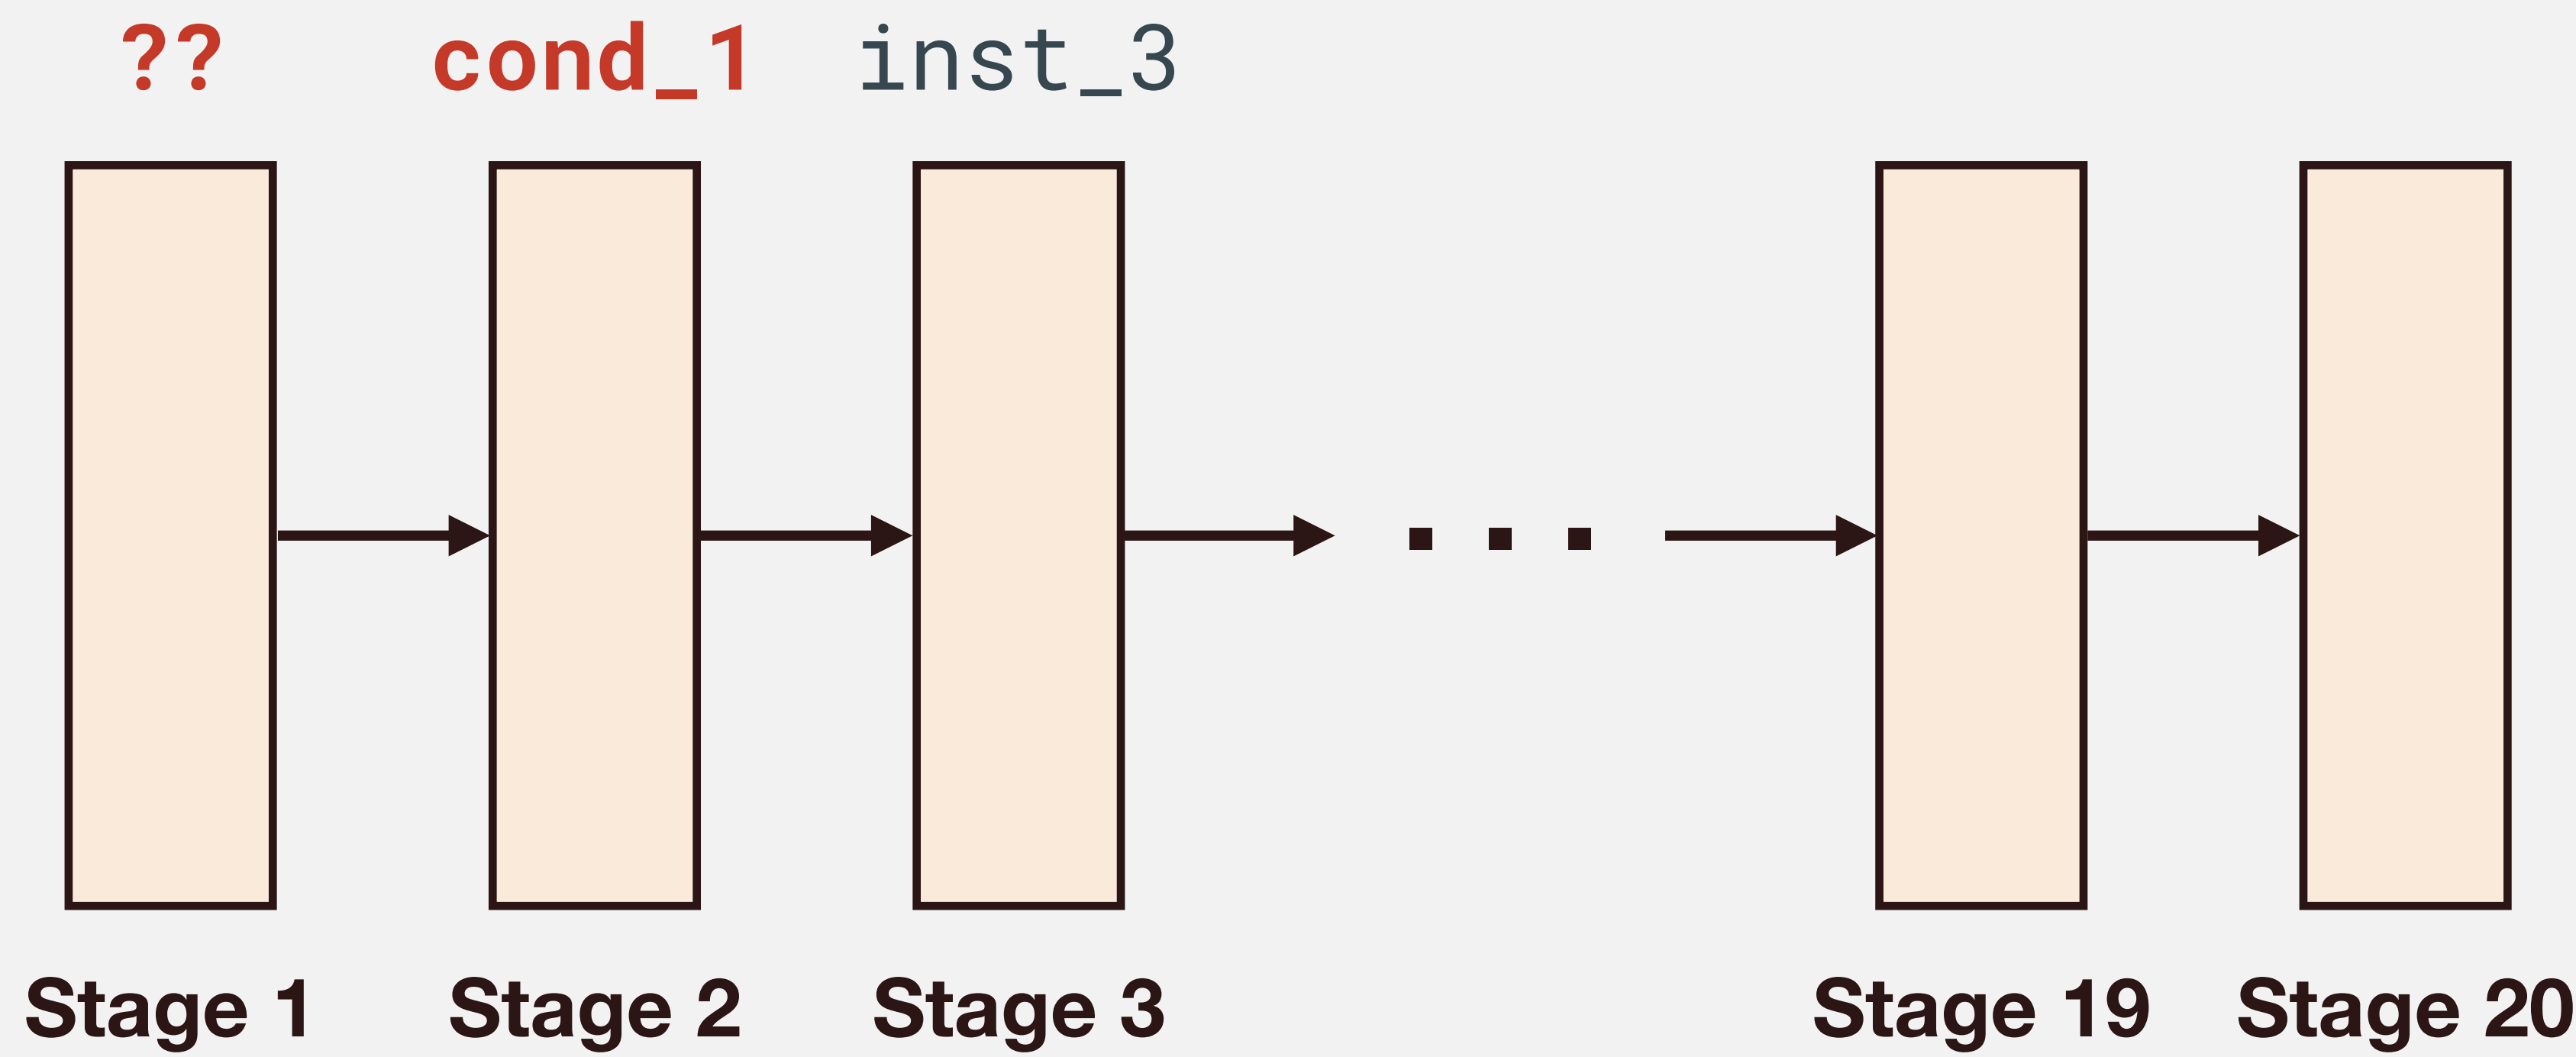
CPU Pipeline

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



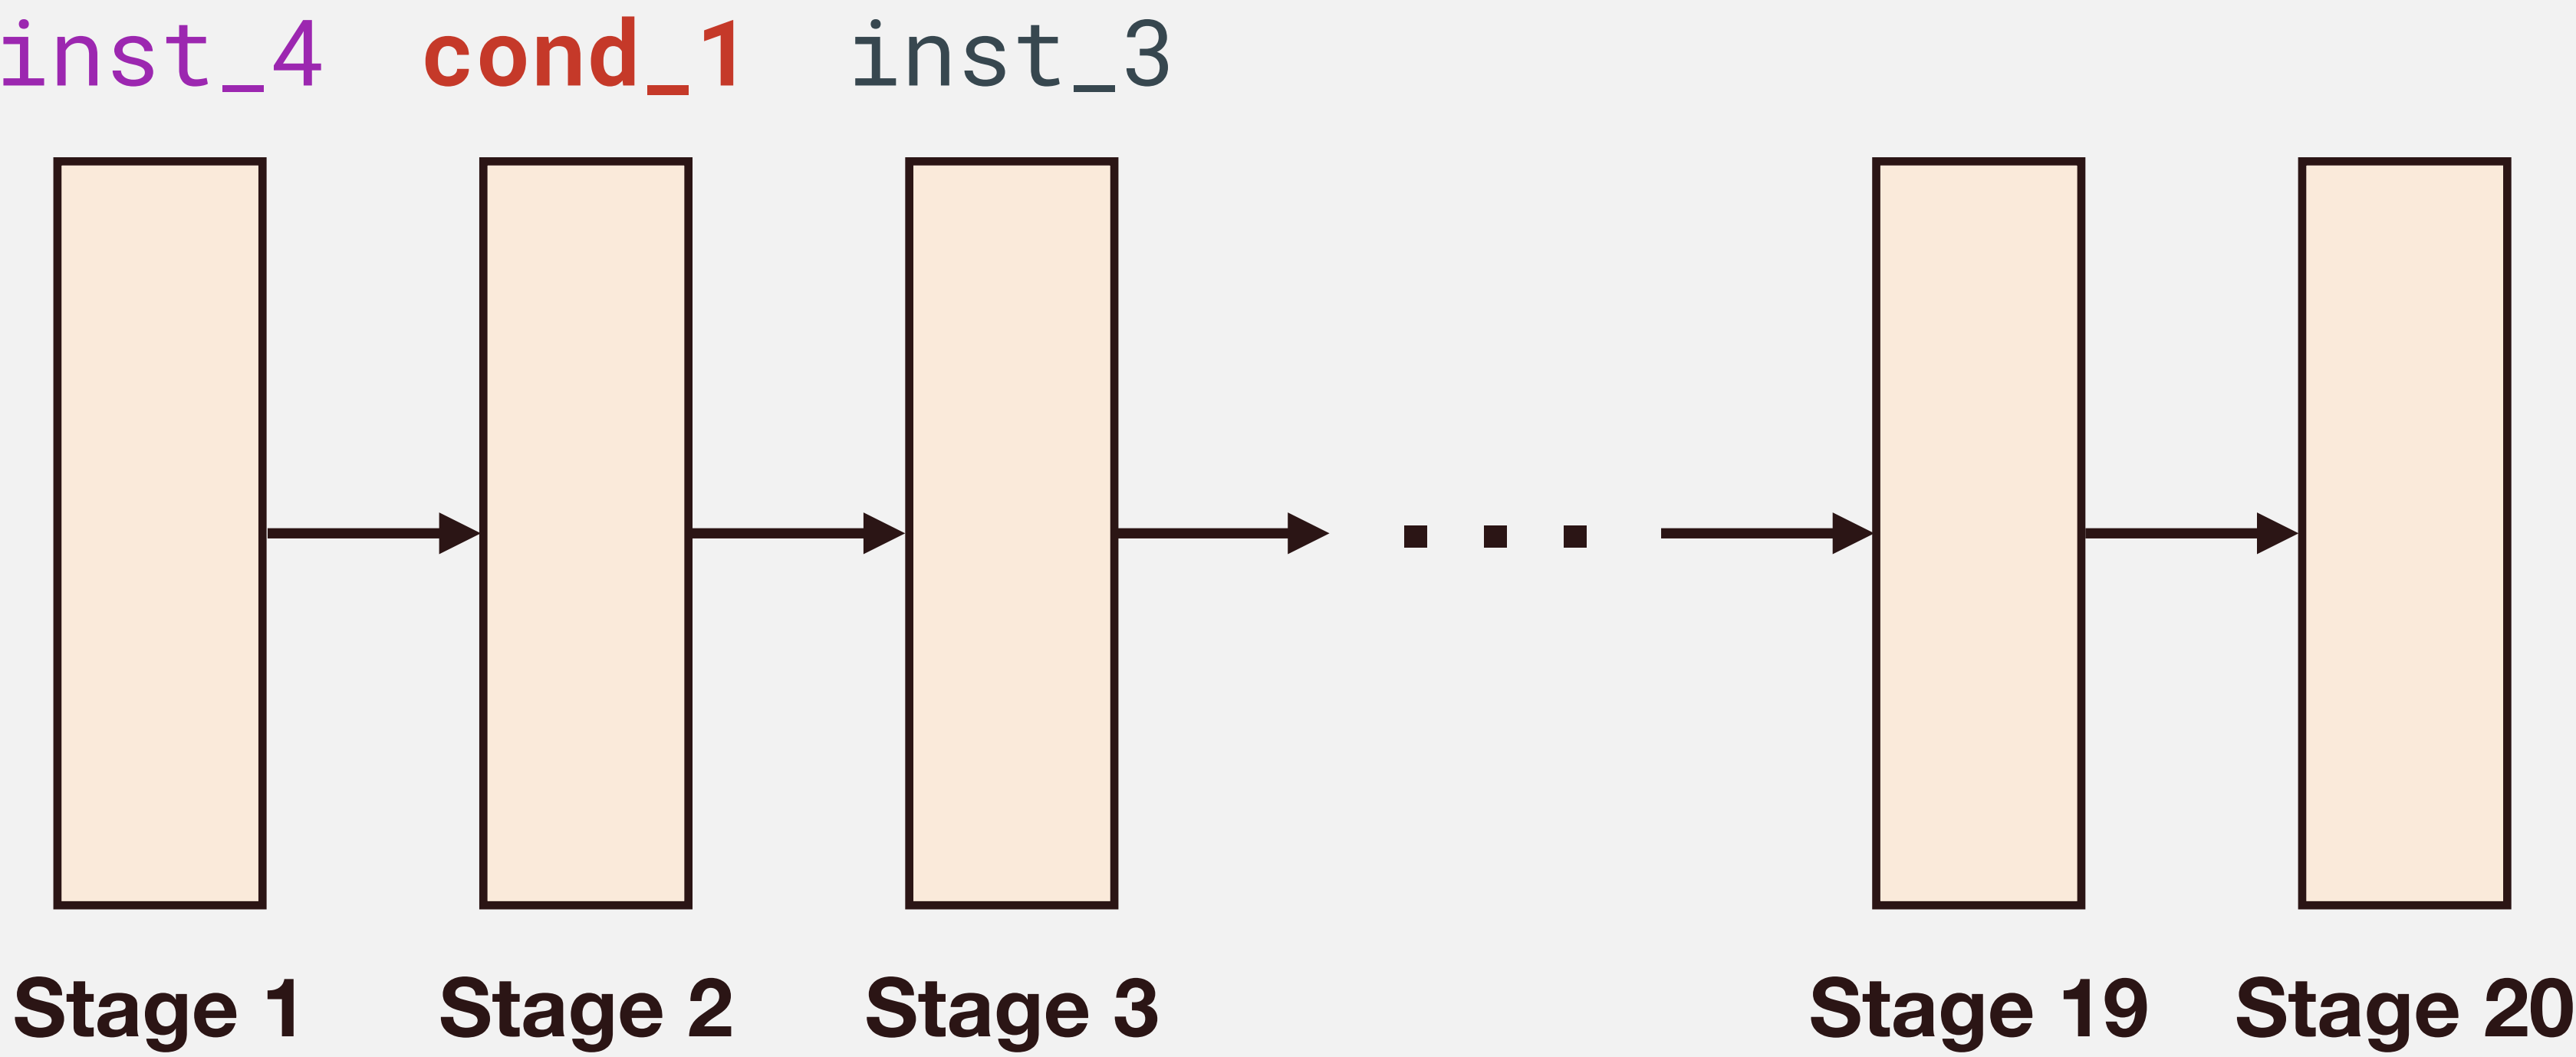
Branch Prediction

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



Branch Prediction

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```

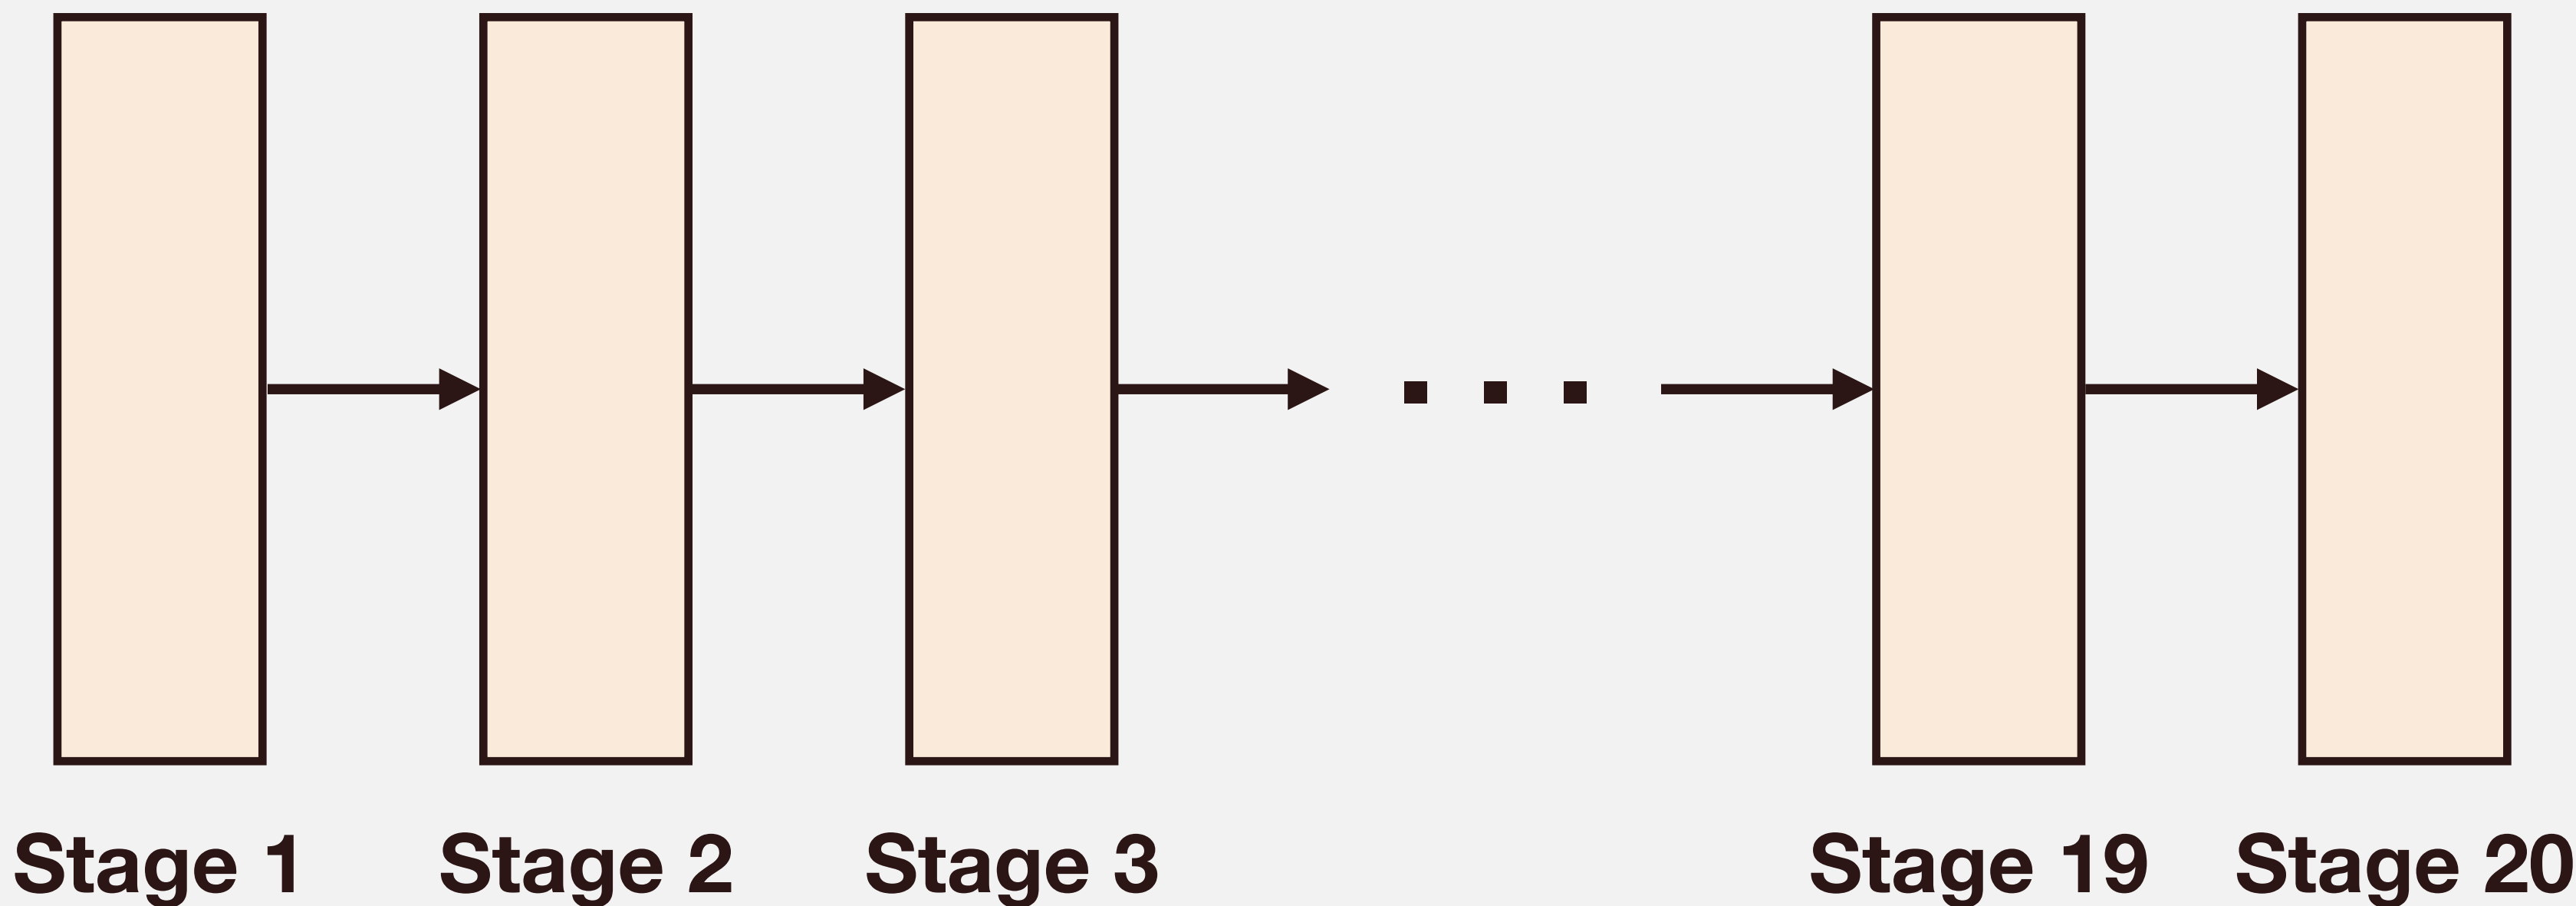


Branch Prediction

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```

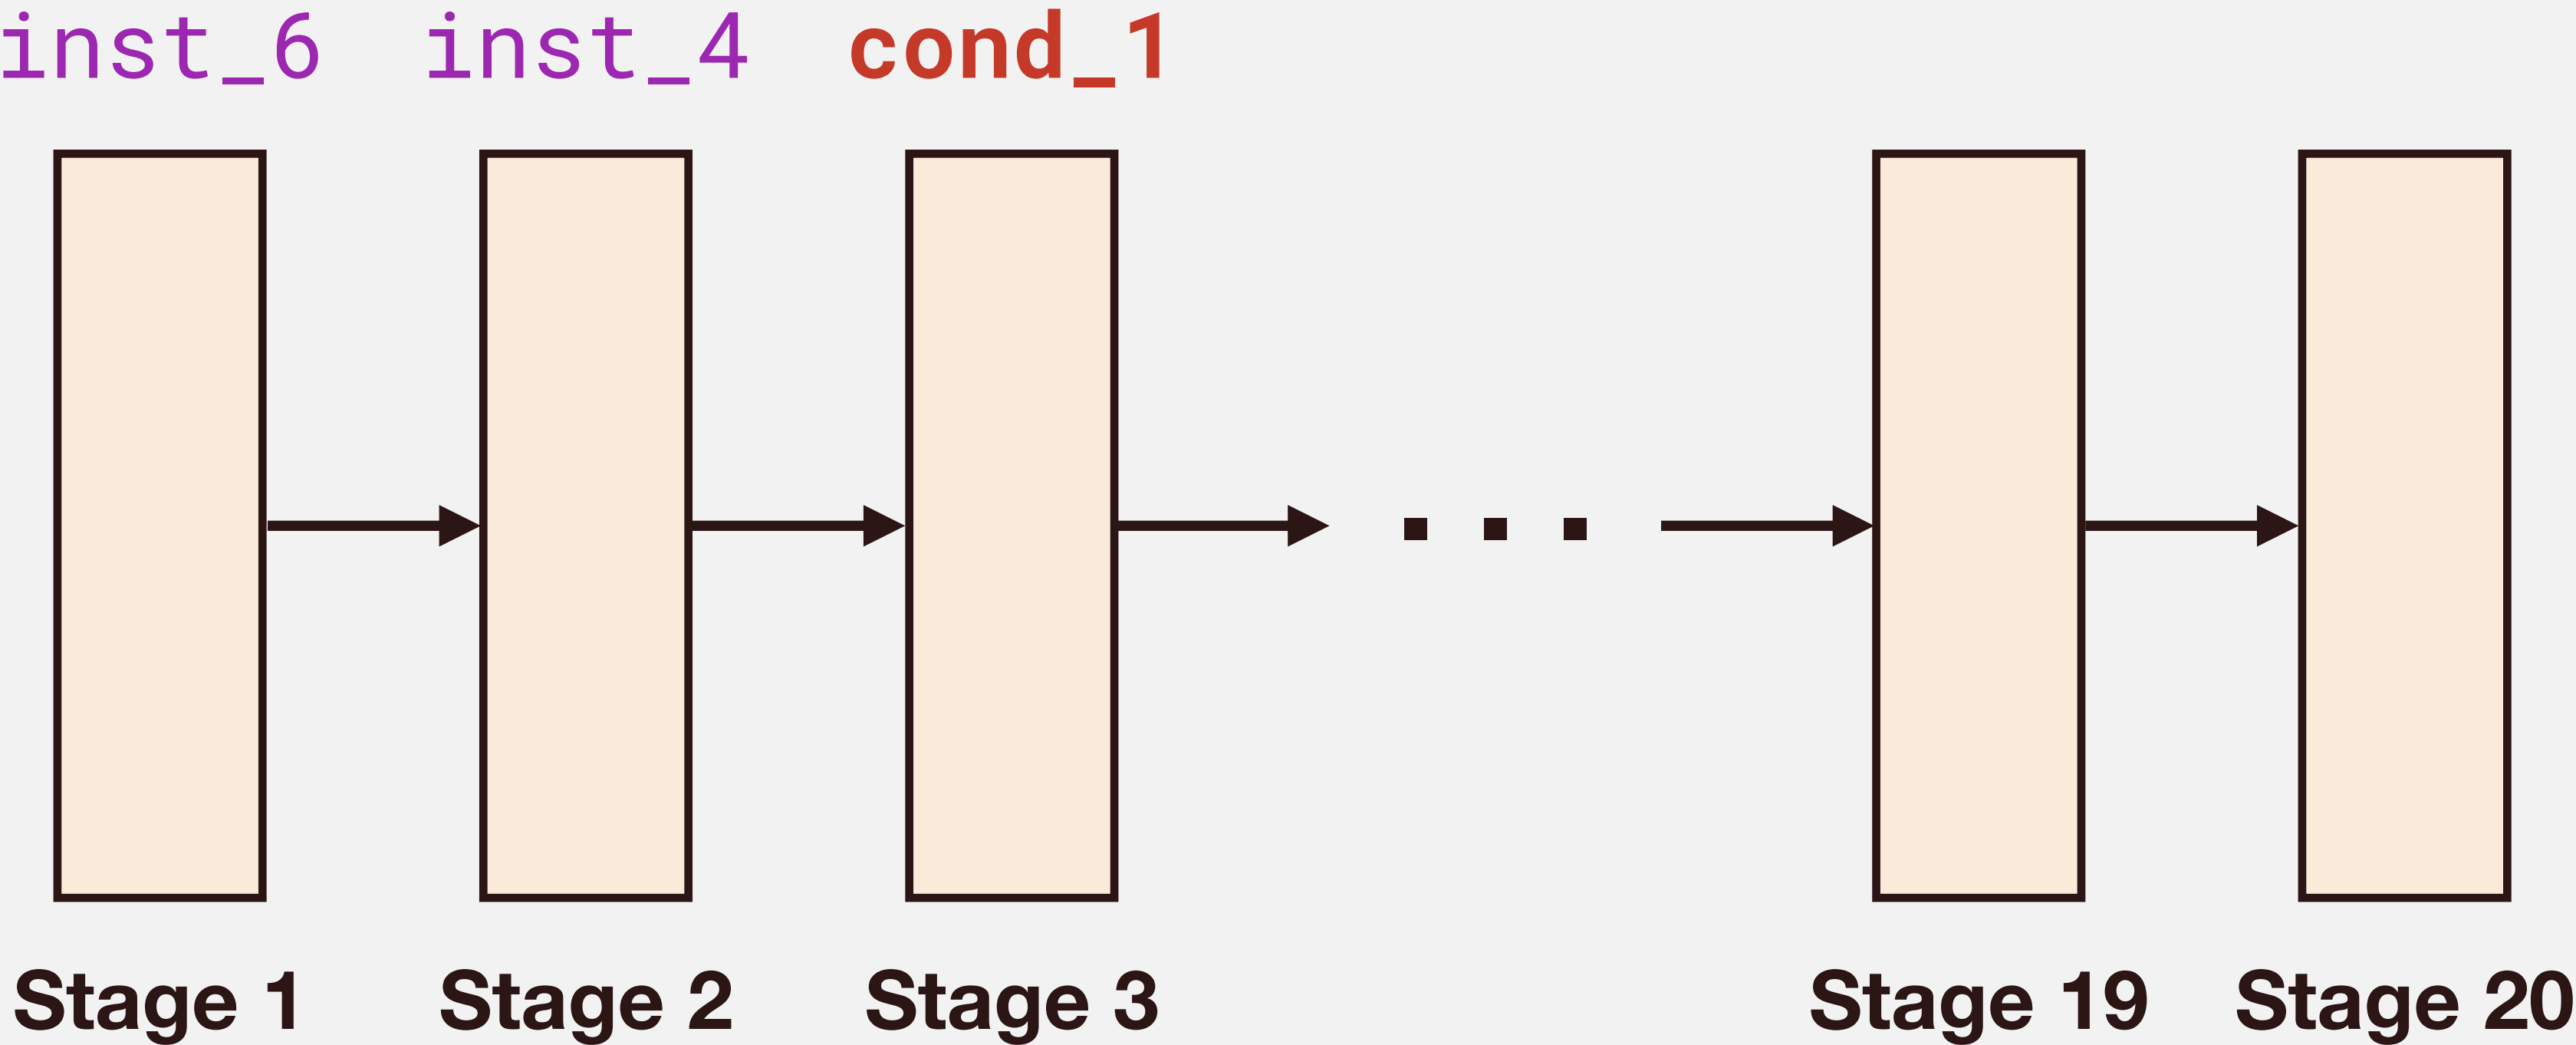
Speculative Execution

inst_4 cond_1 inst_3



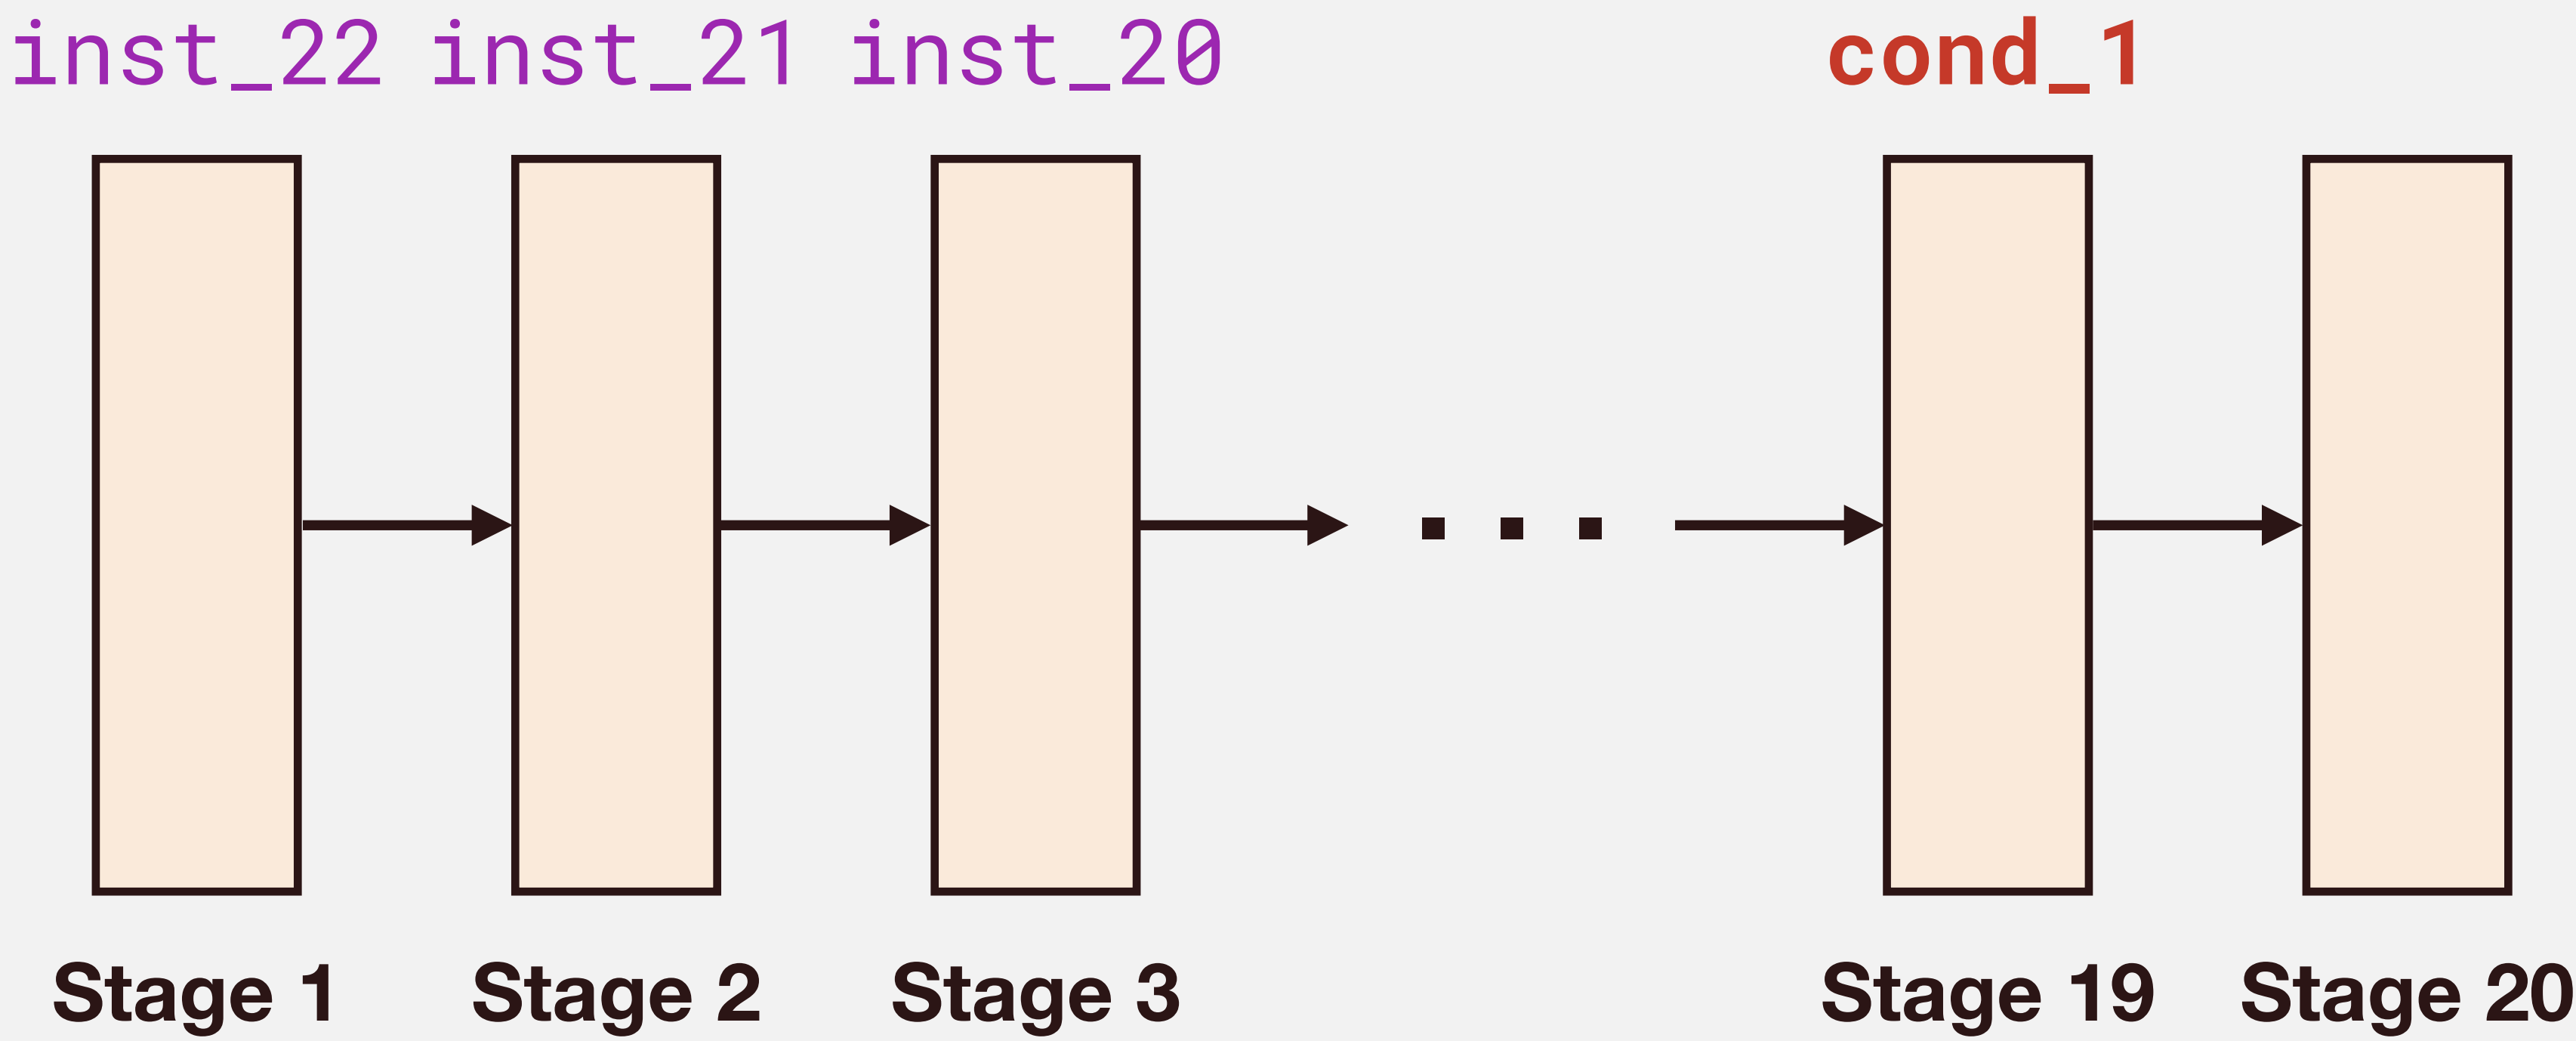
Branch Prediction

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



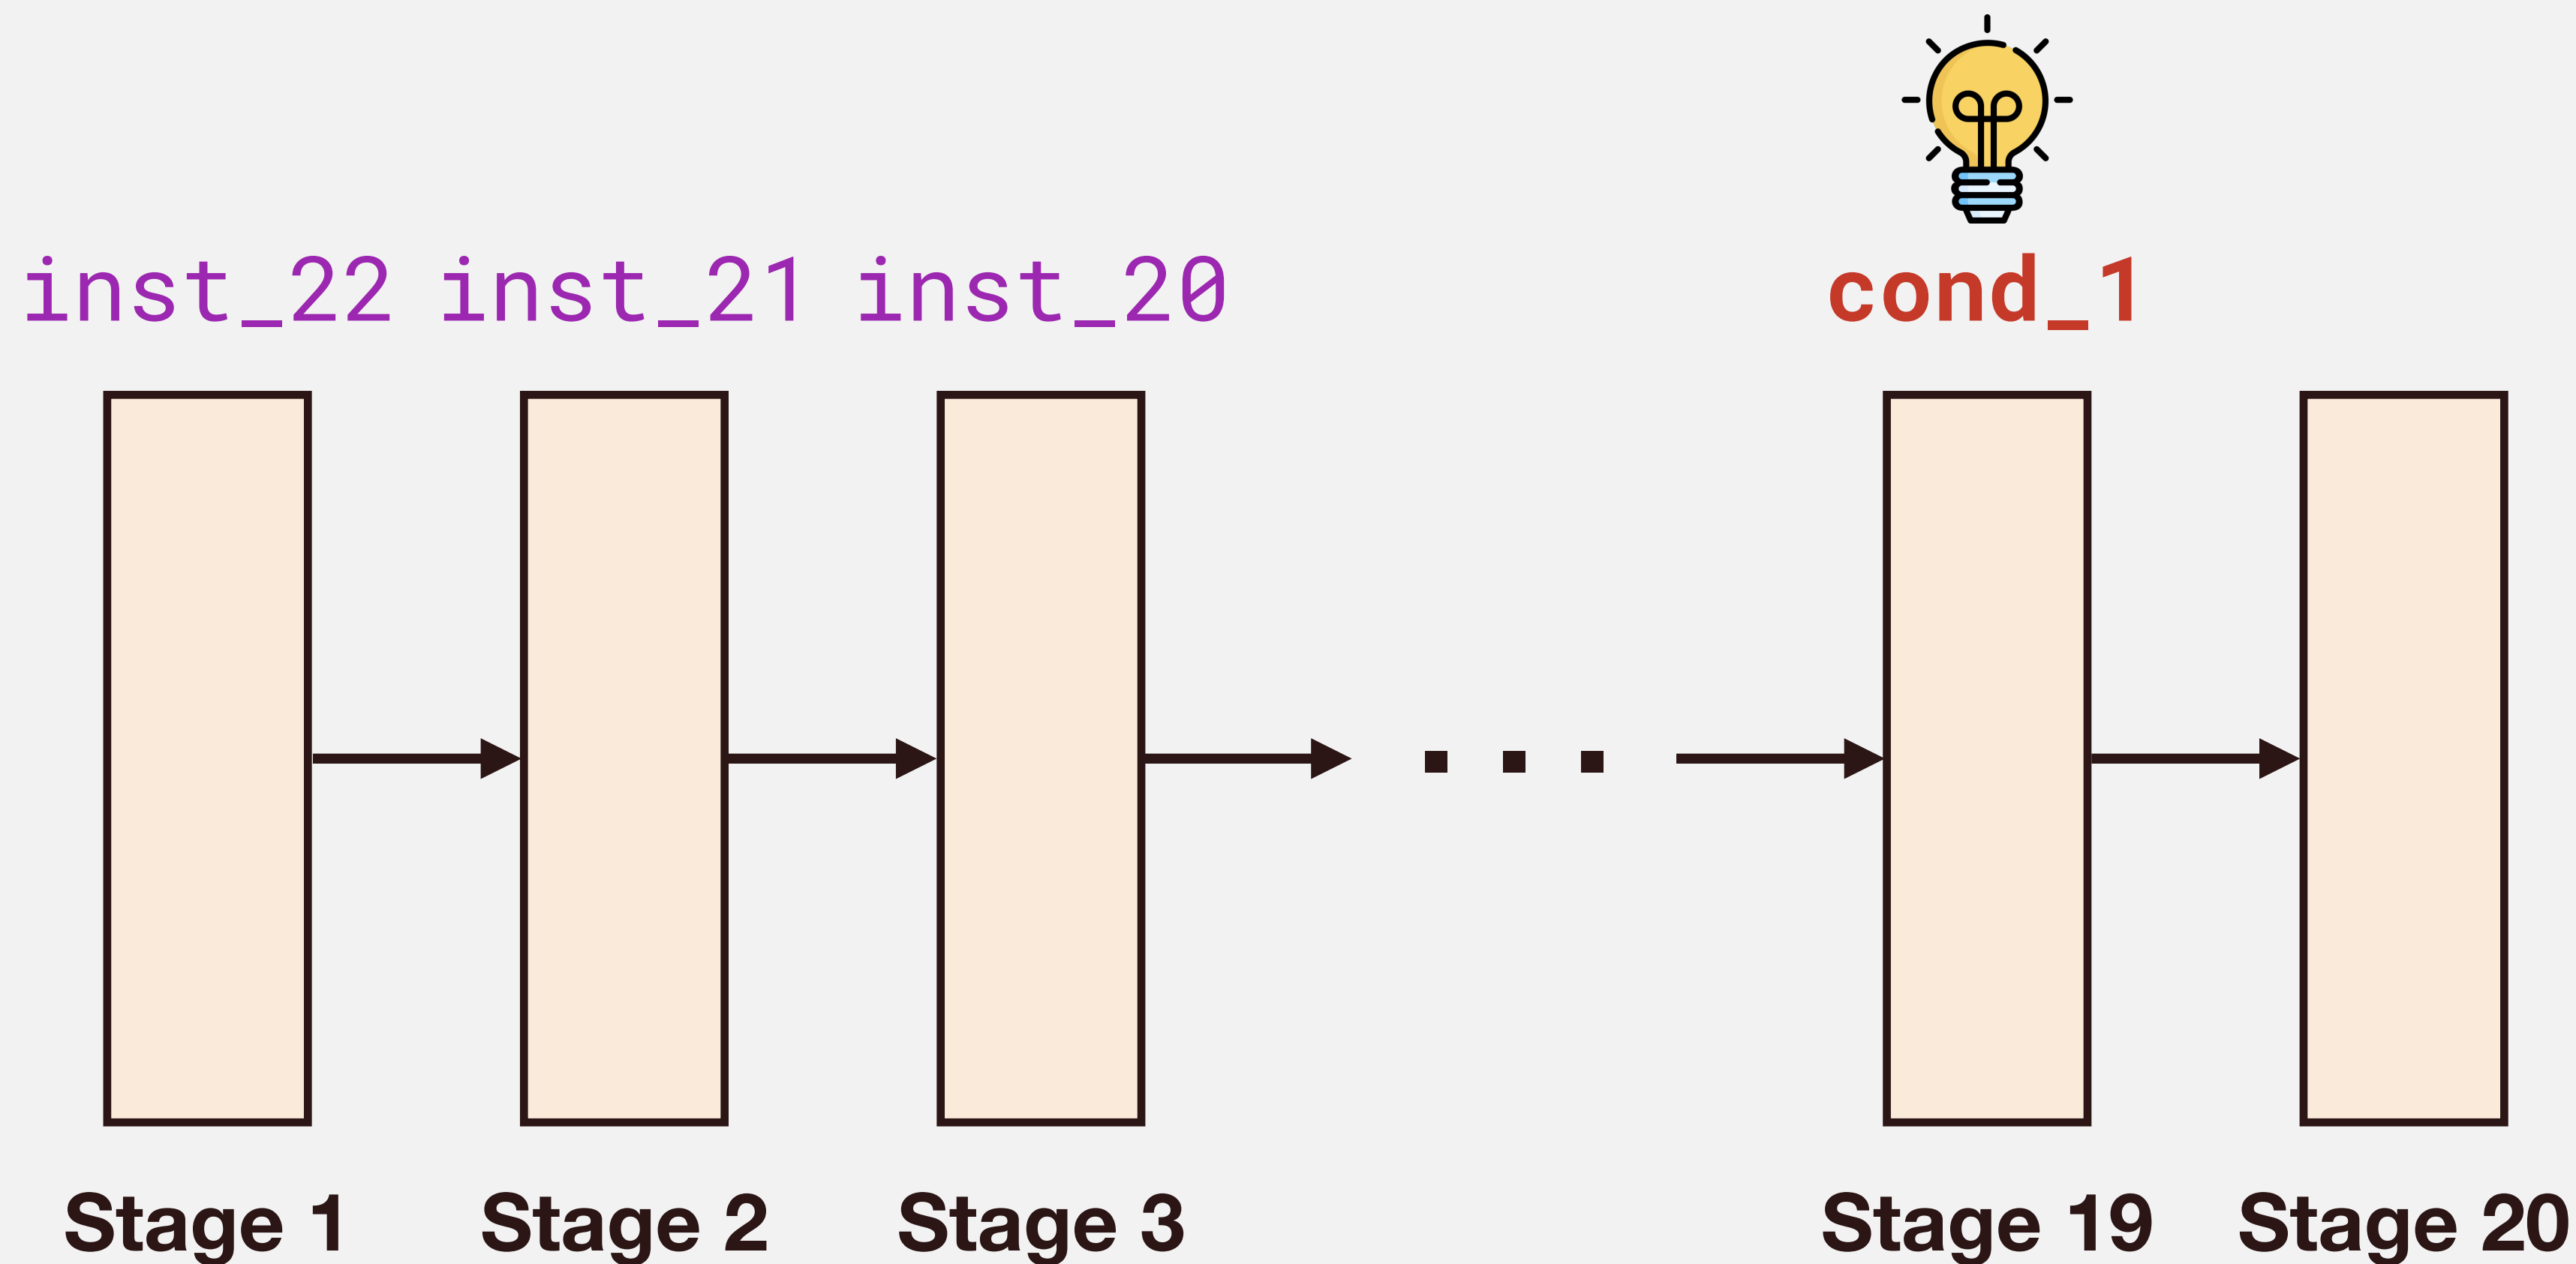
Branch Prediction

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



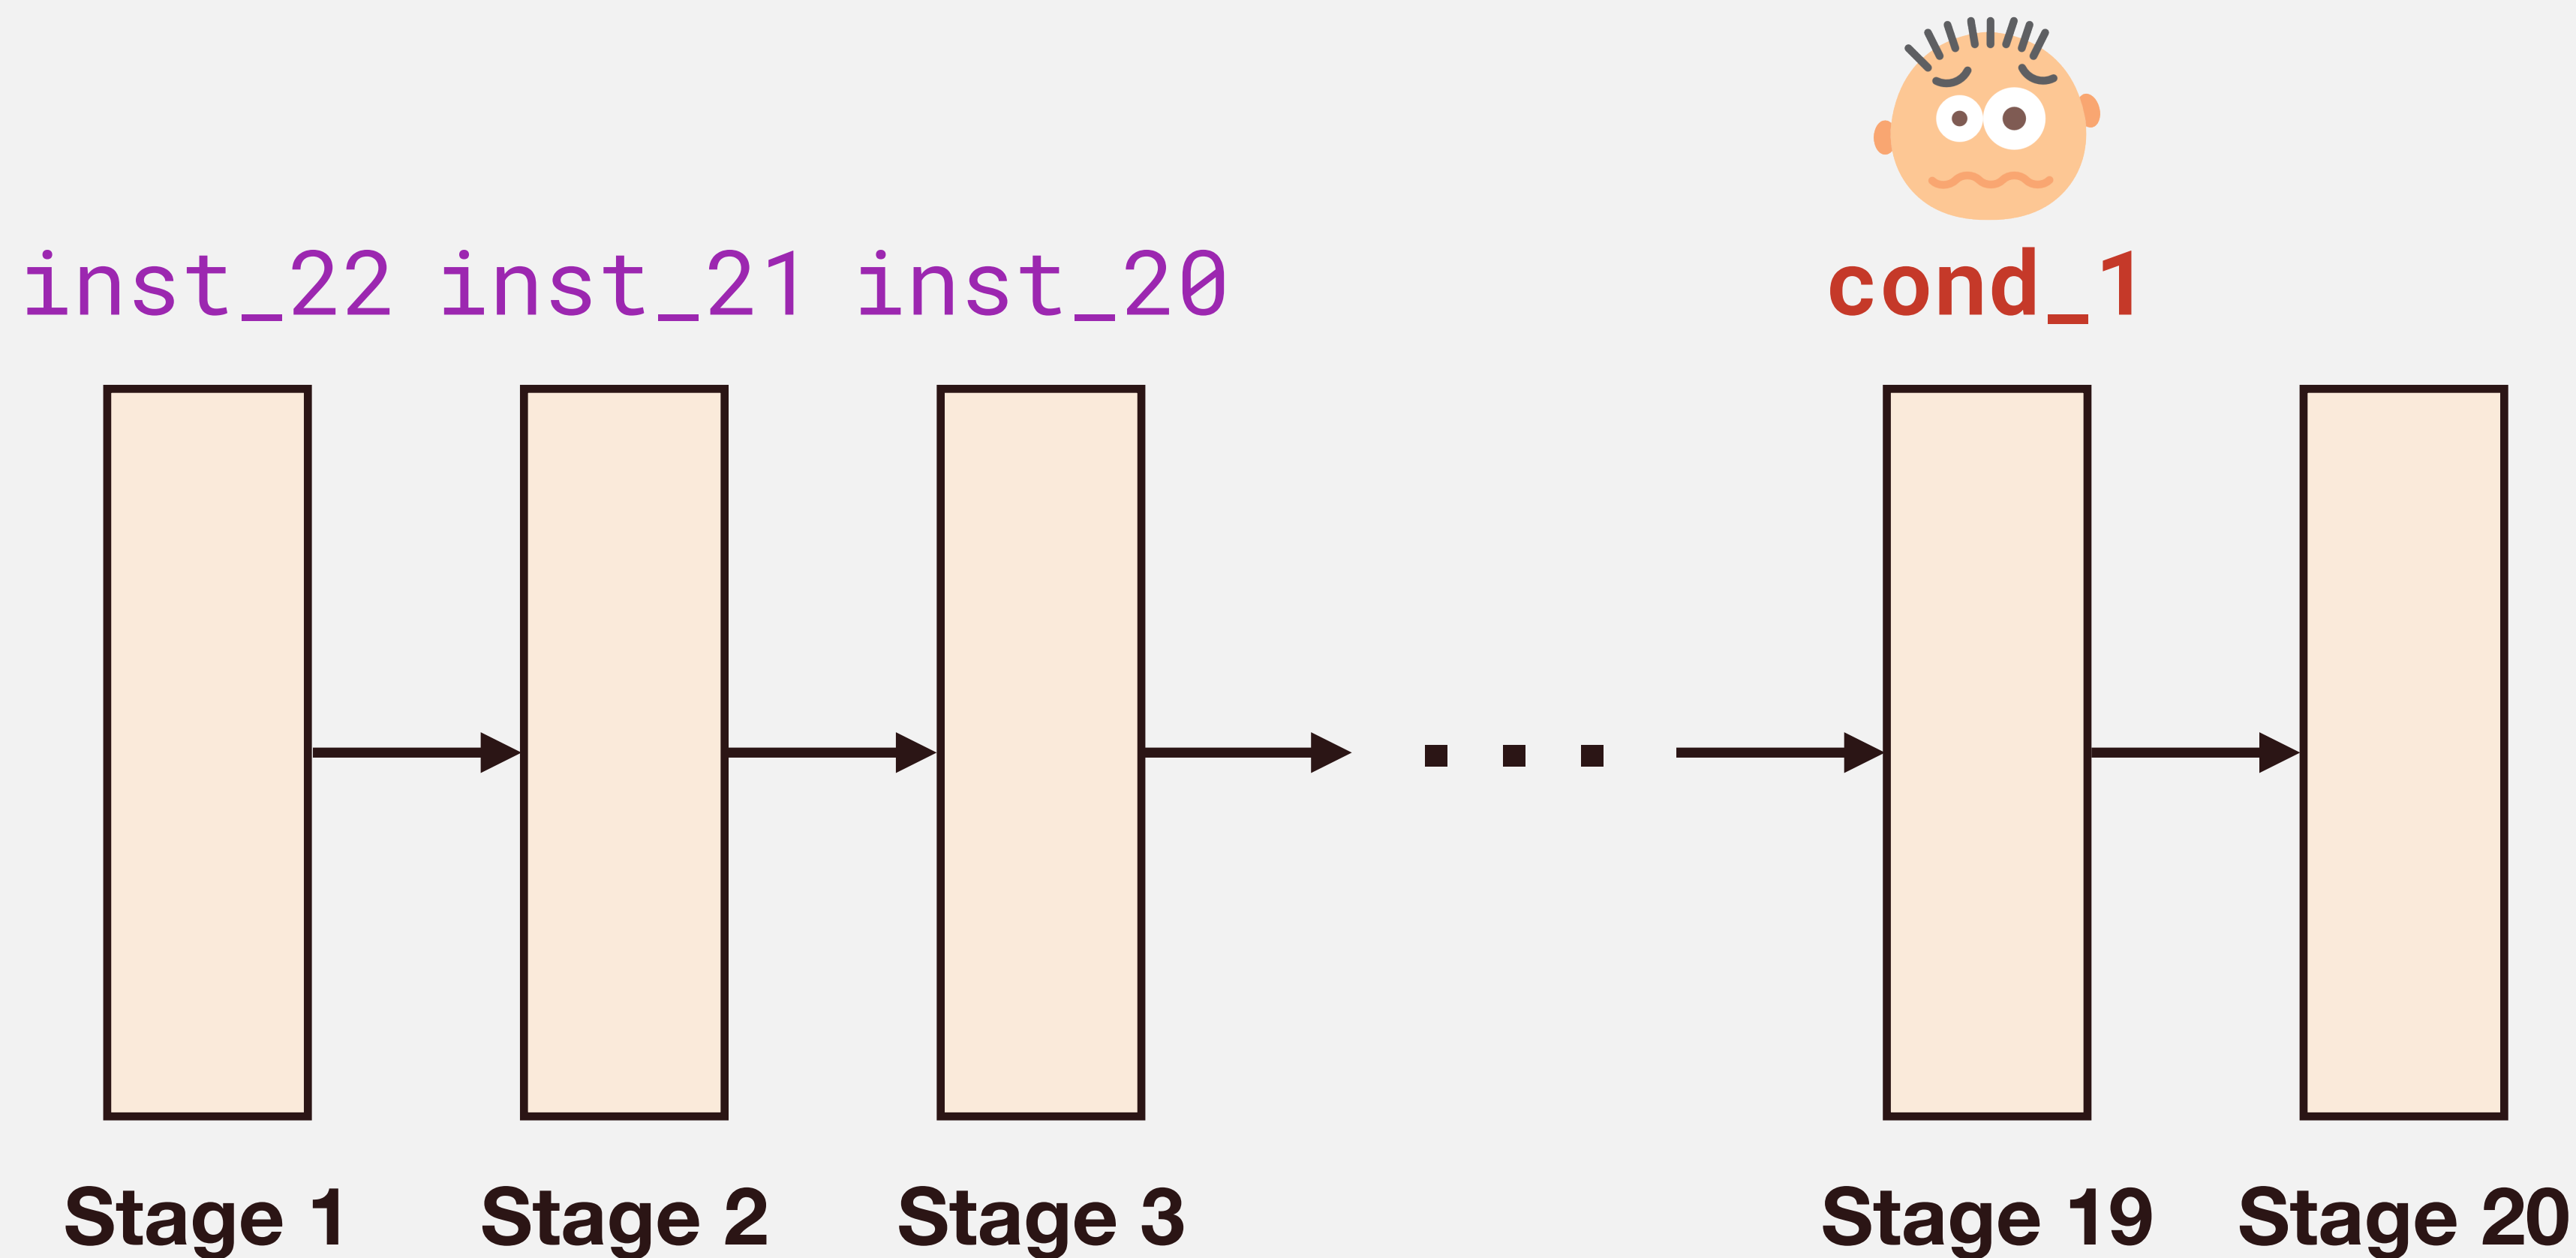
Branch Prediction

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



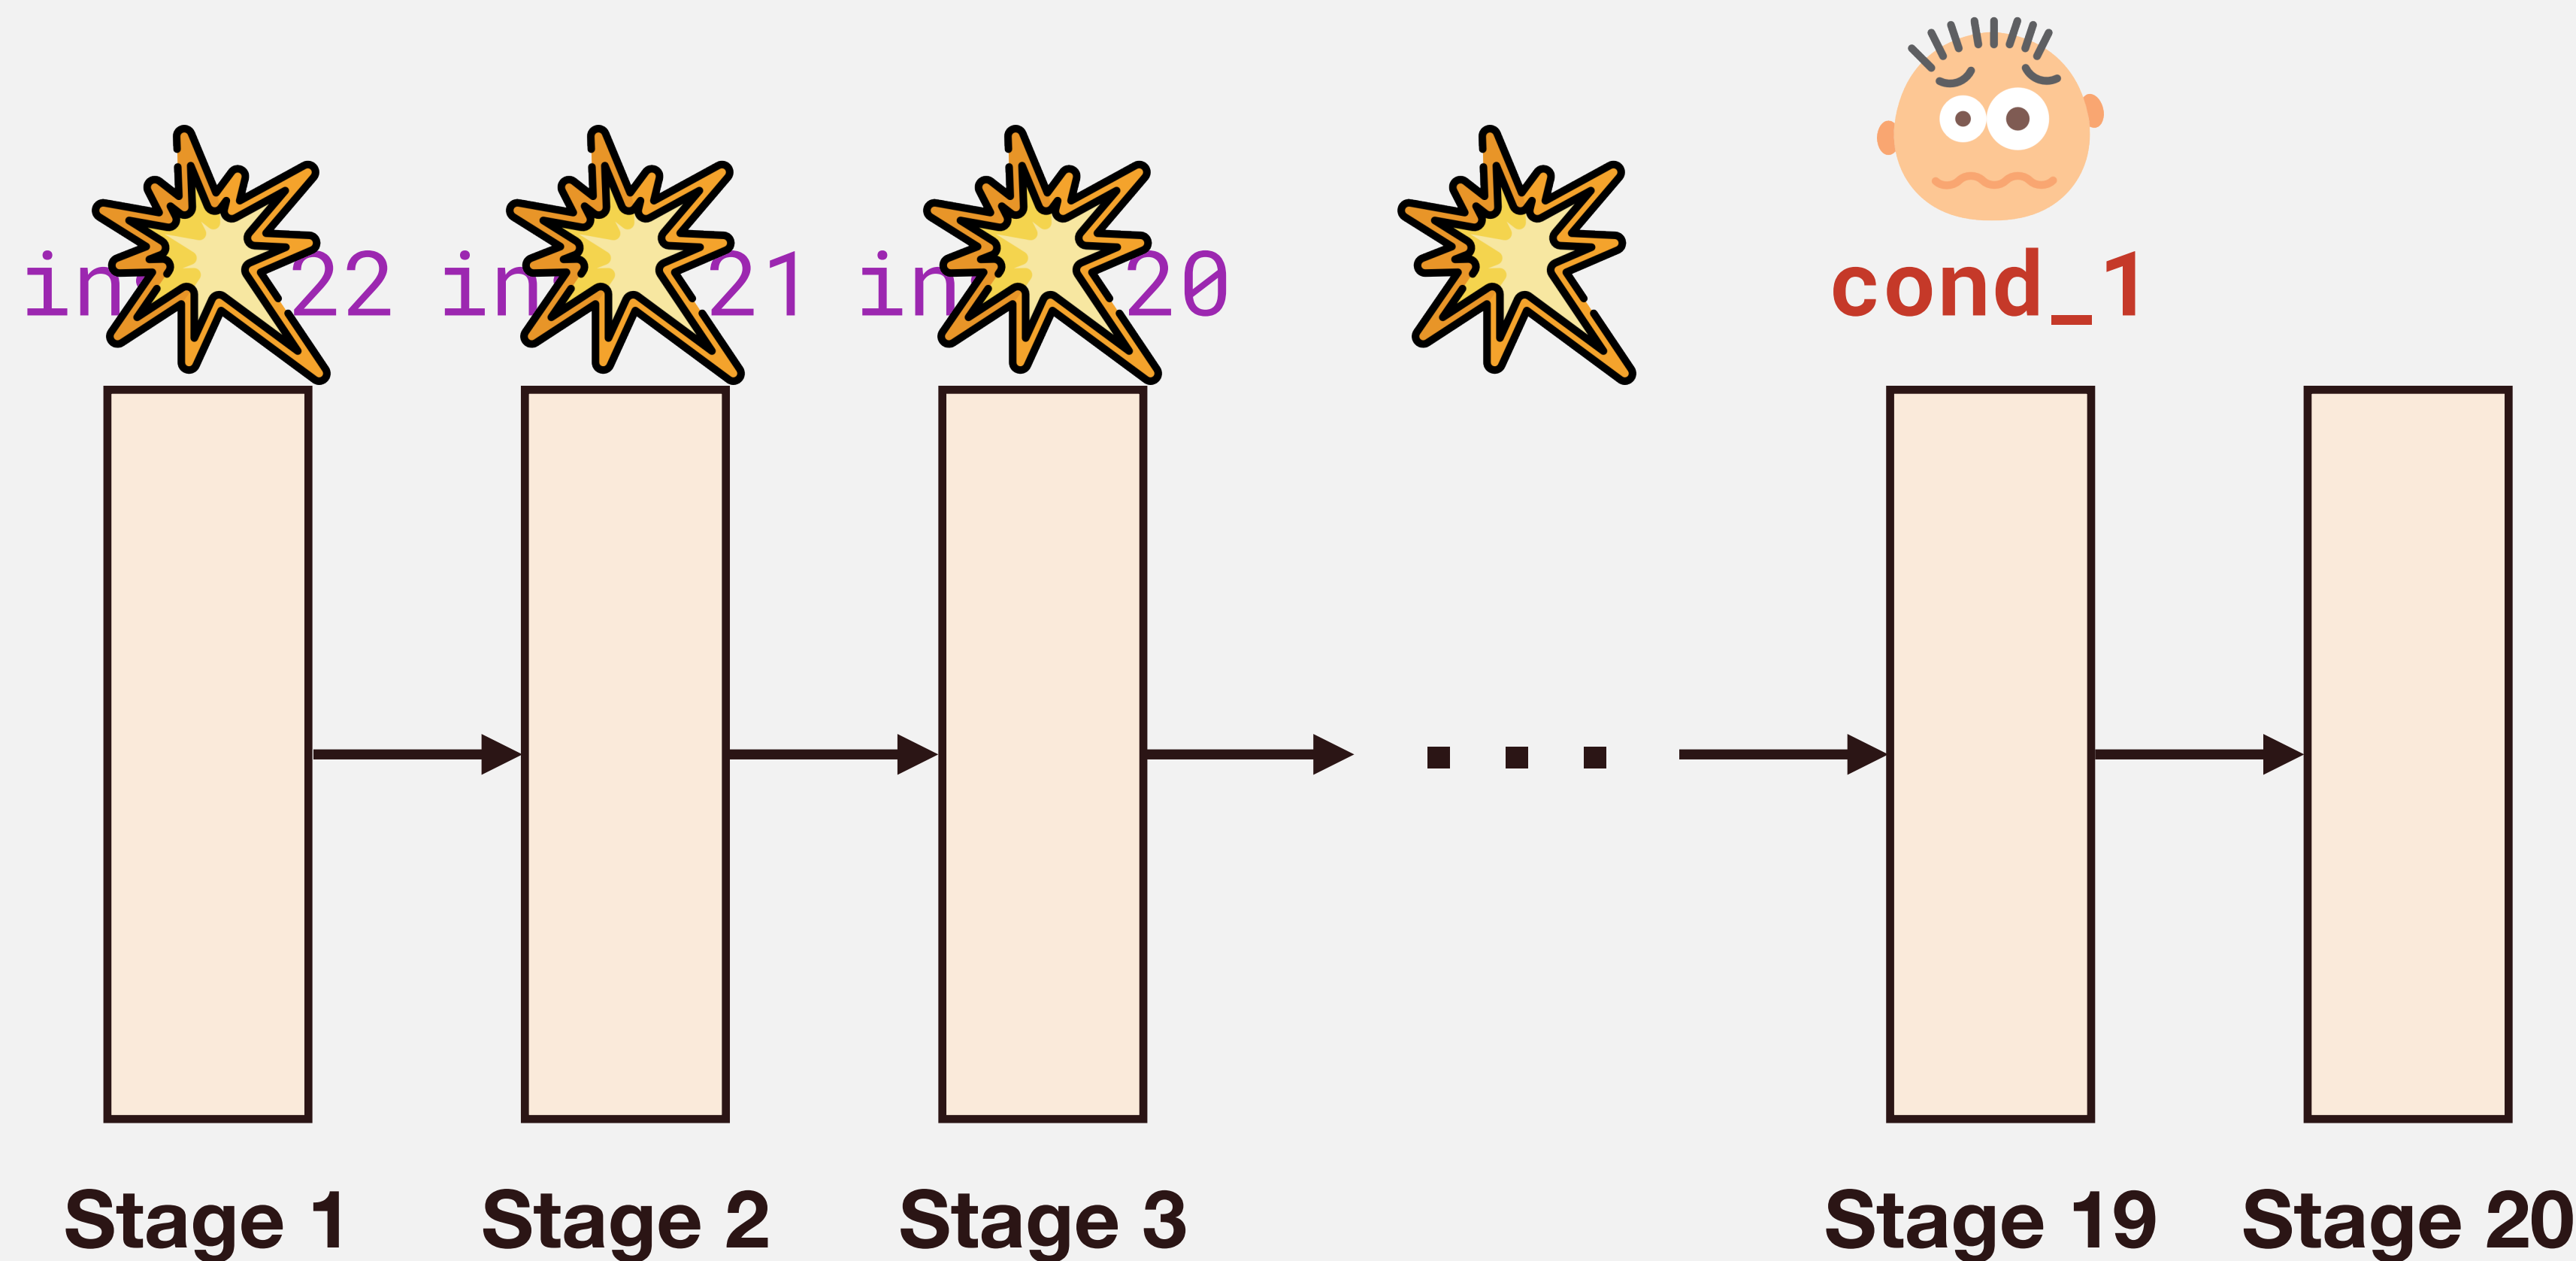
Branch Prediction

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



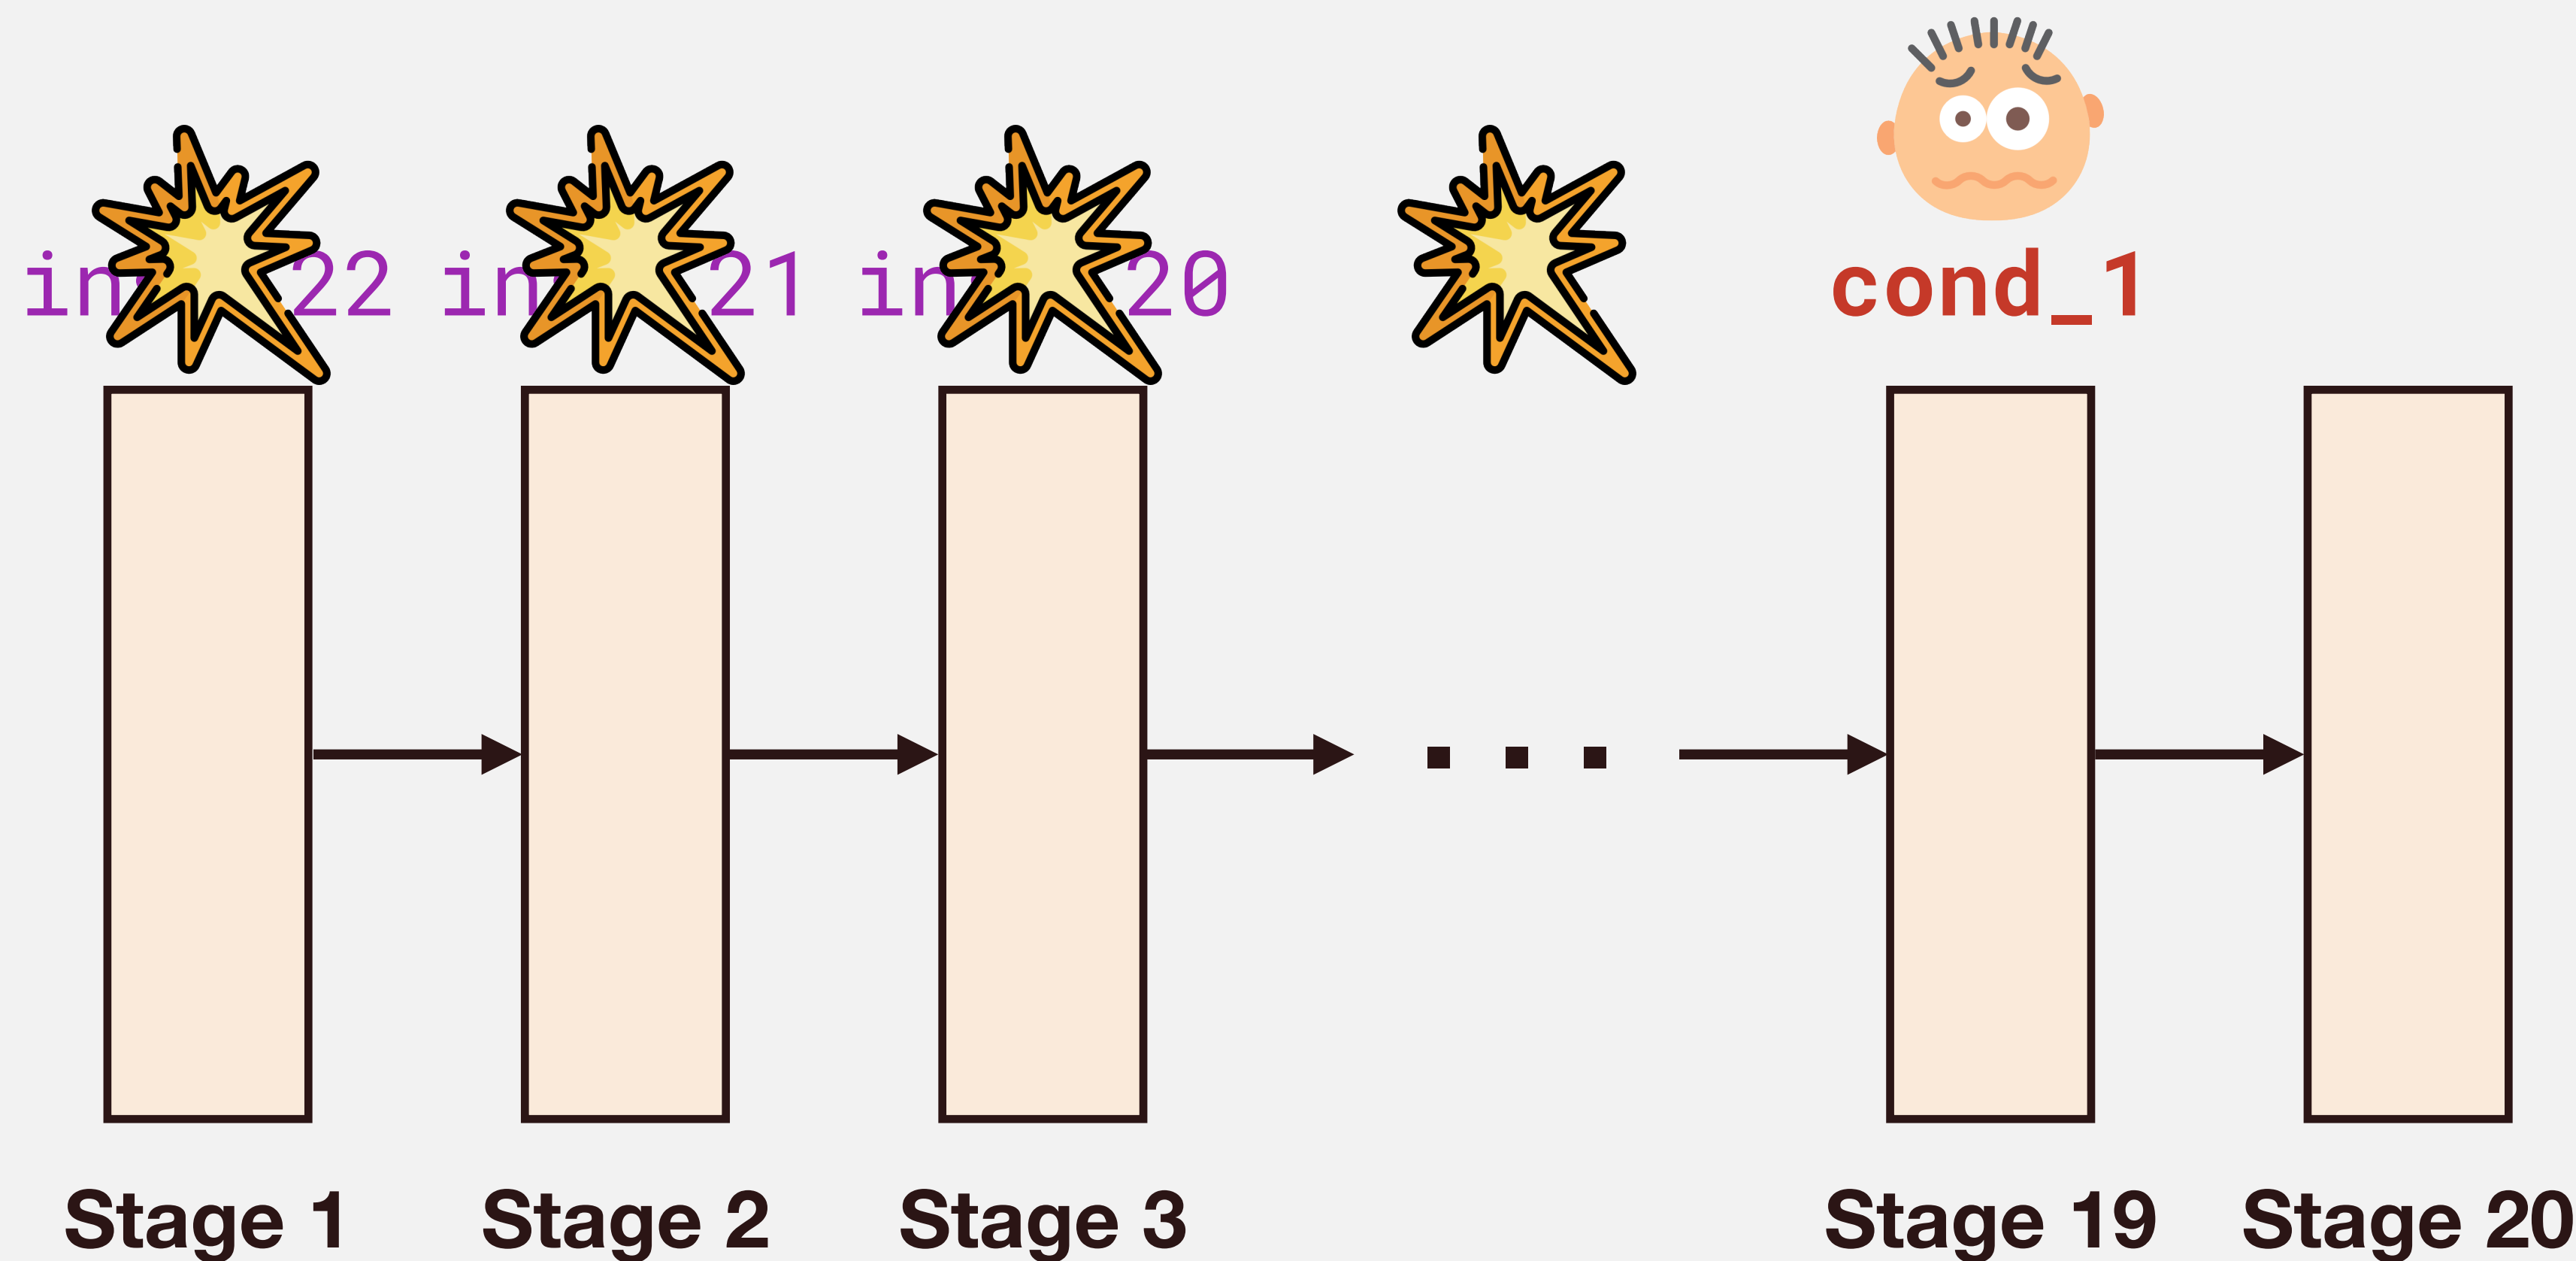
Branch Prediction

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```



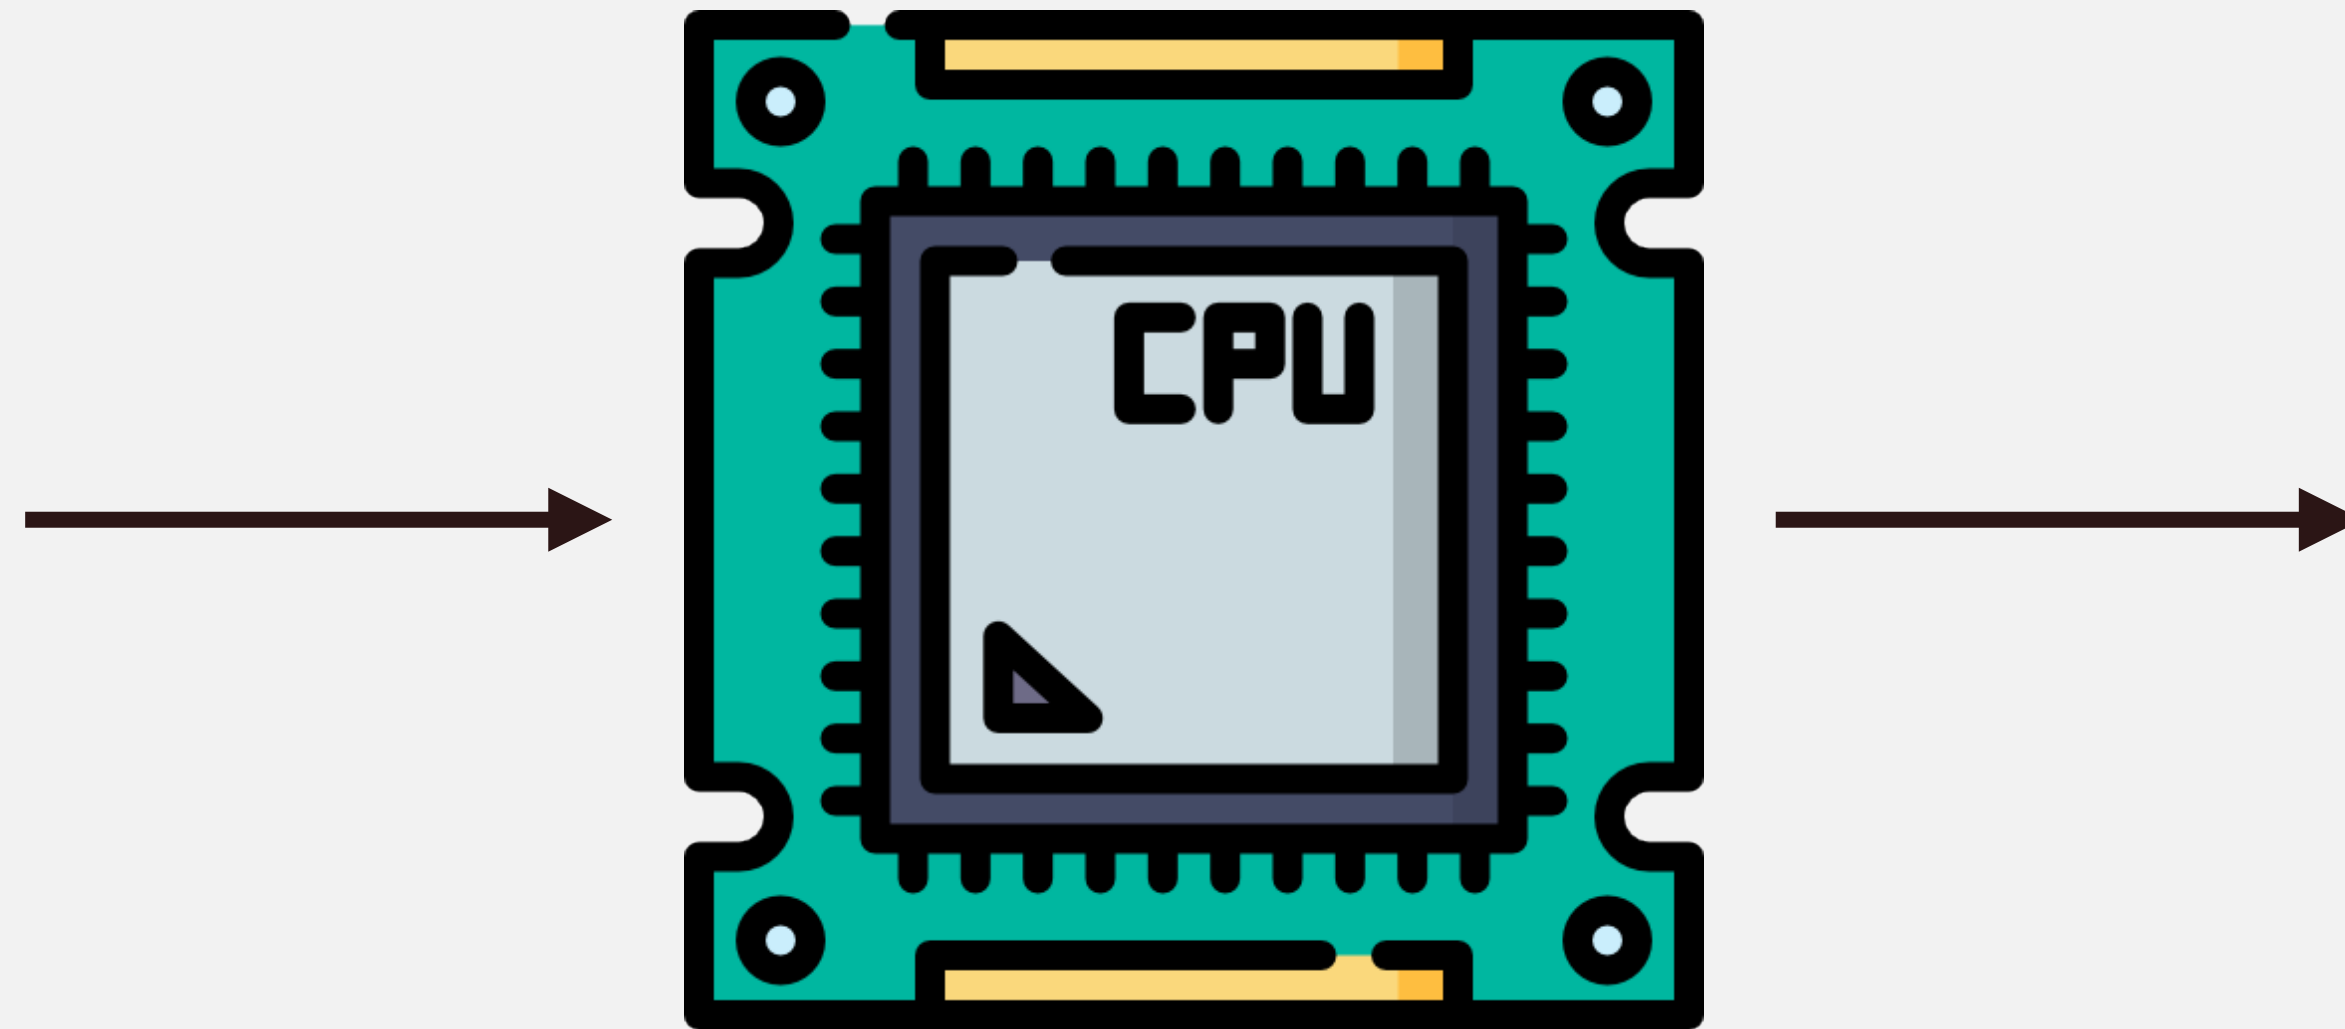
Branch Prediction

```
inst_1
inst_2
inst_3
if (cond_1)
    inst_4
else
    inst_5
inst_6
...
inst_N
```

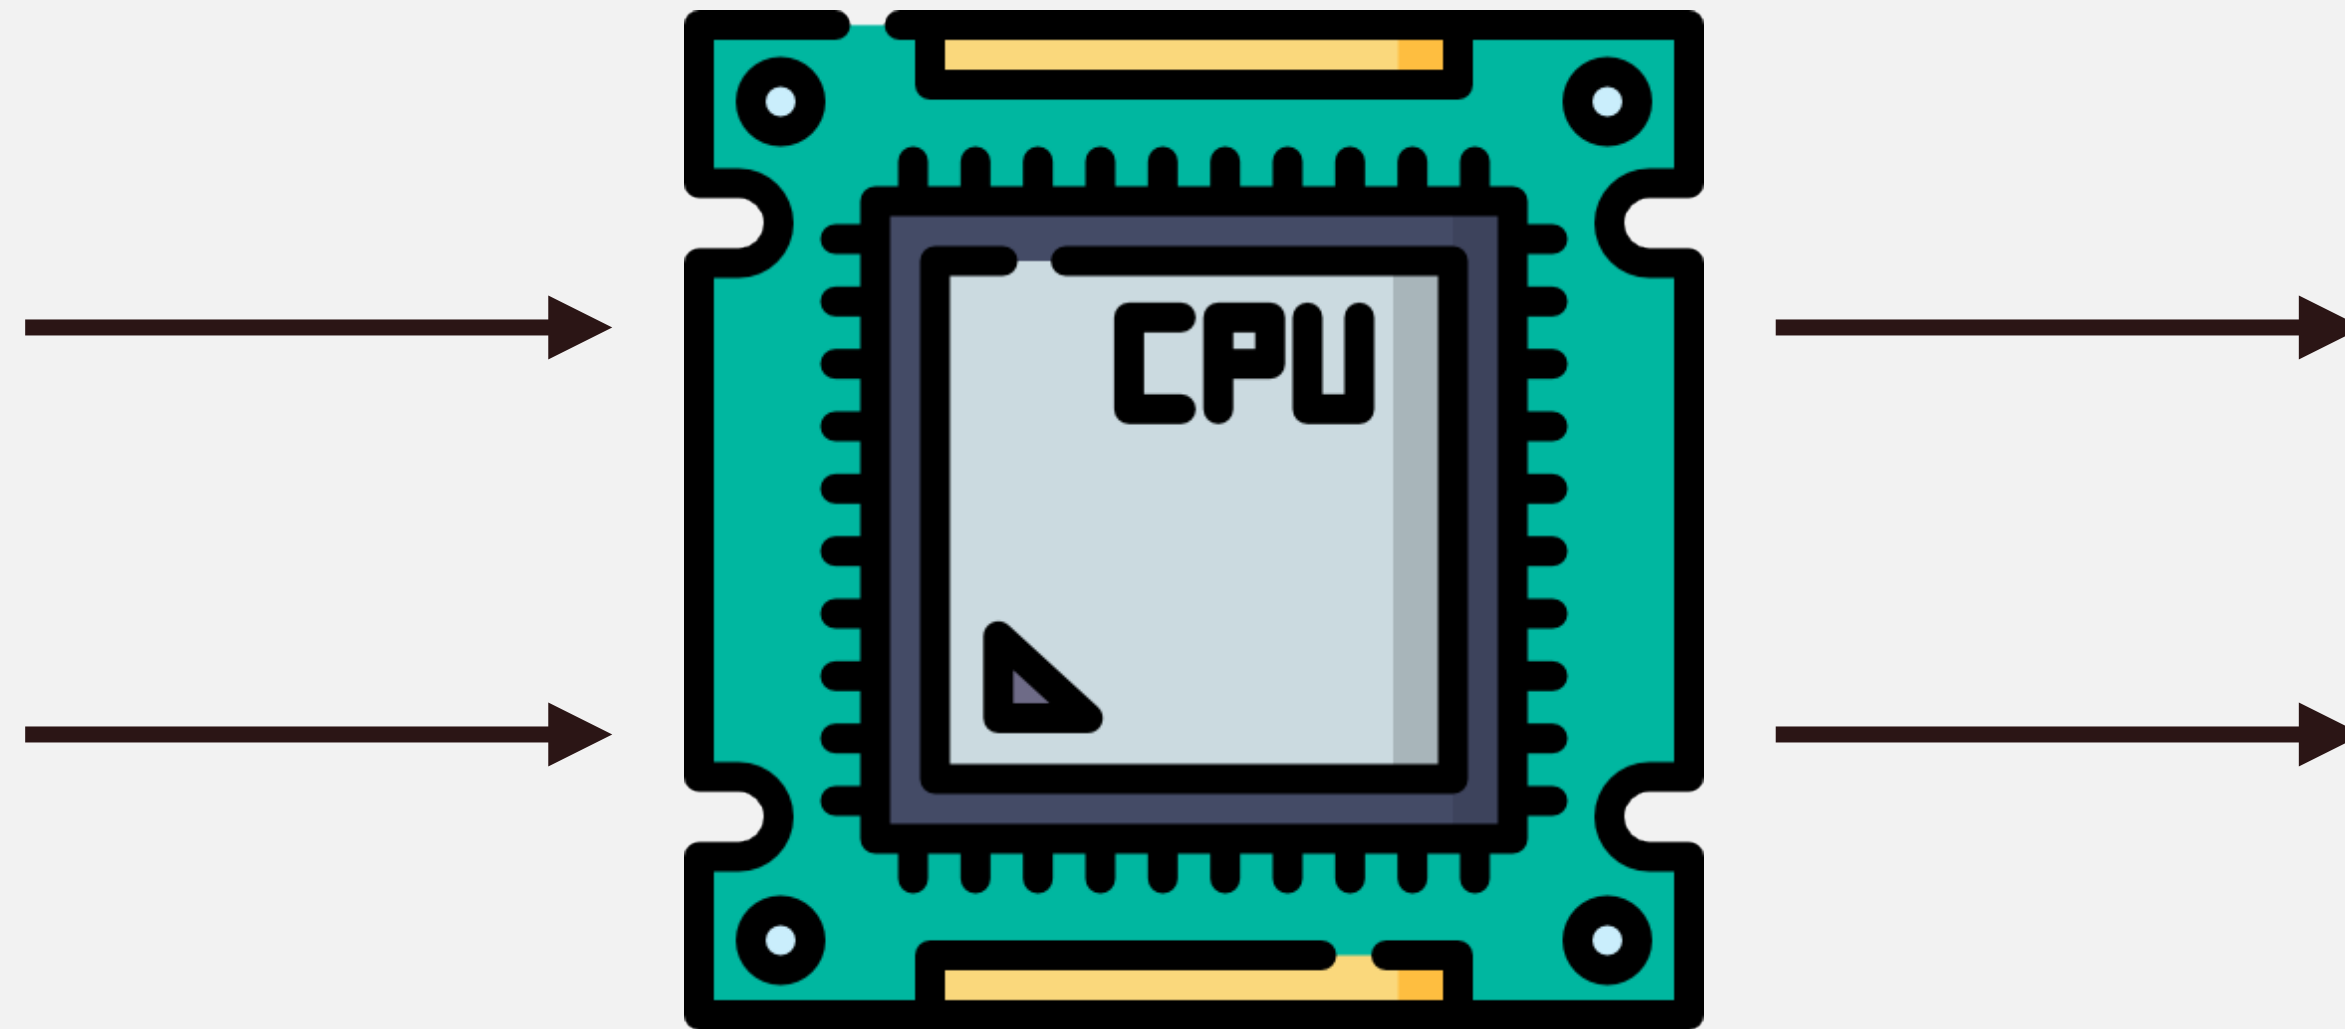


A Branch Misprediction costs ~10-20 CPU cycles

CPU Parallelism

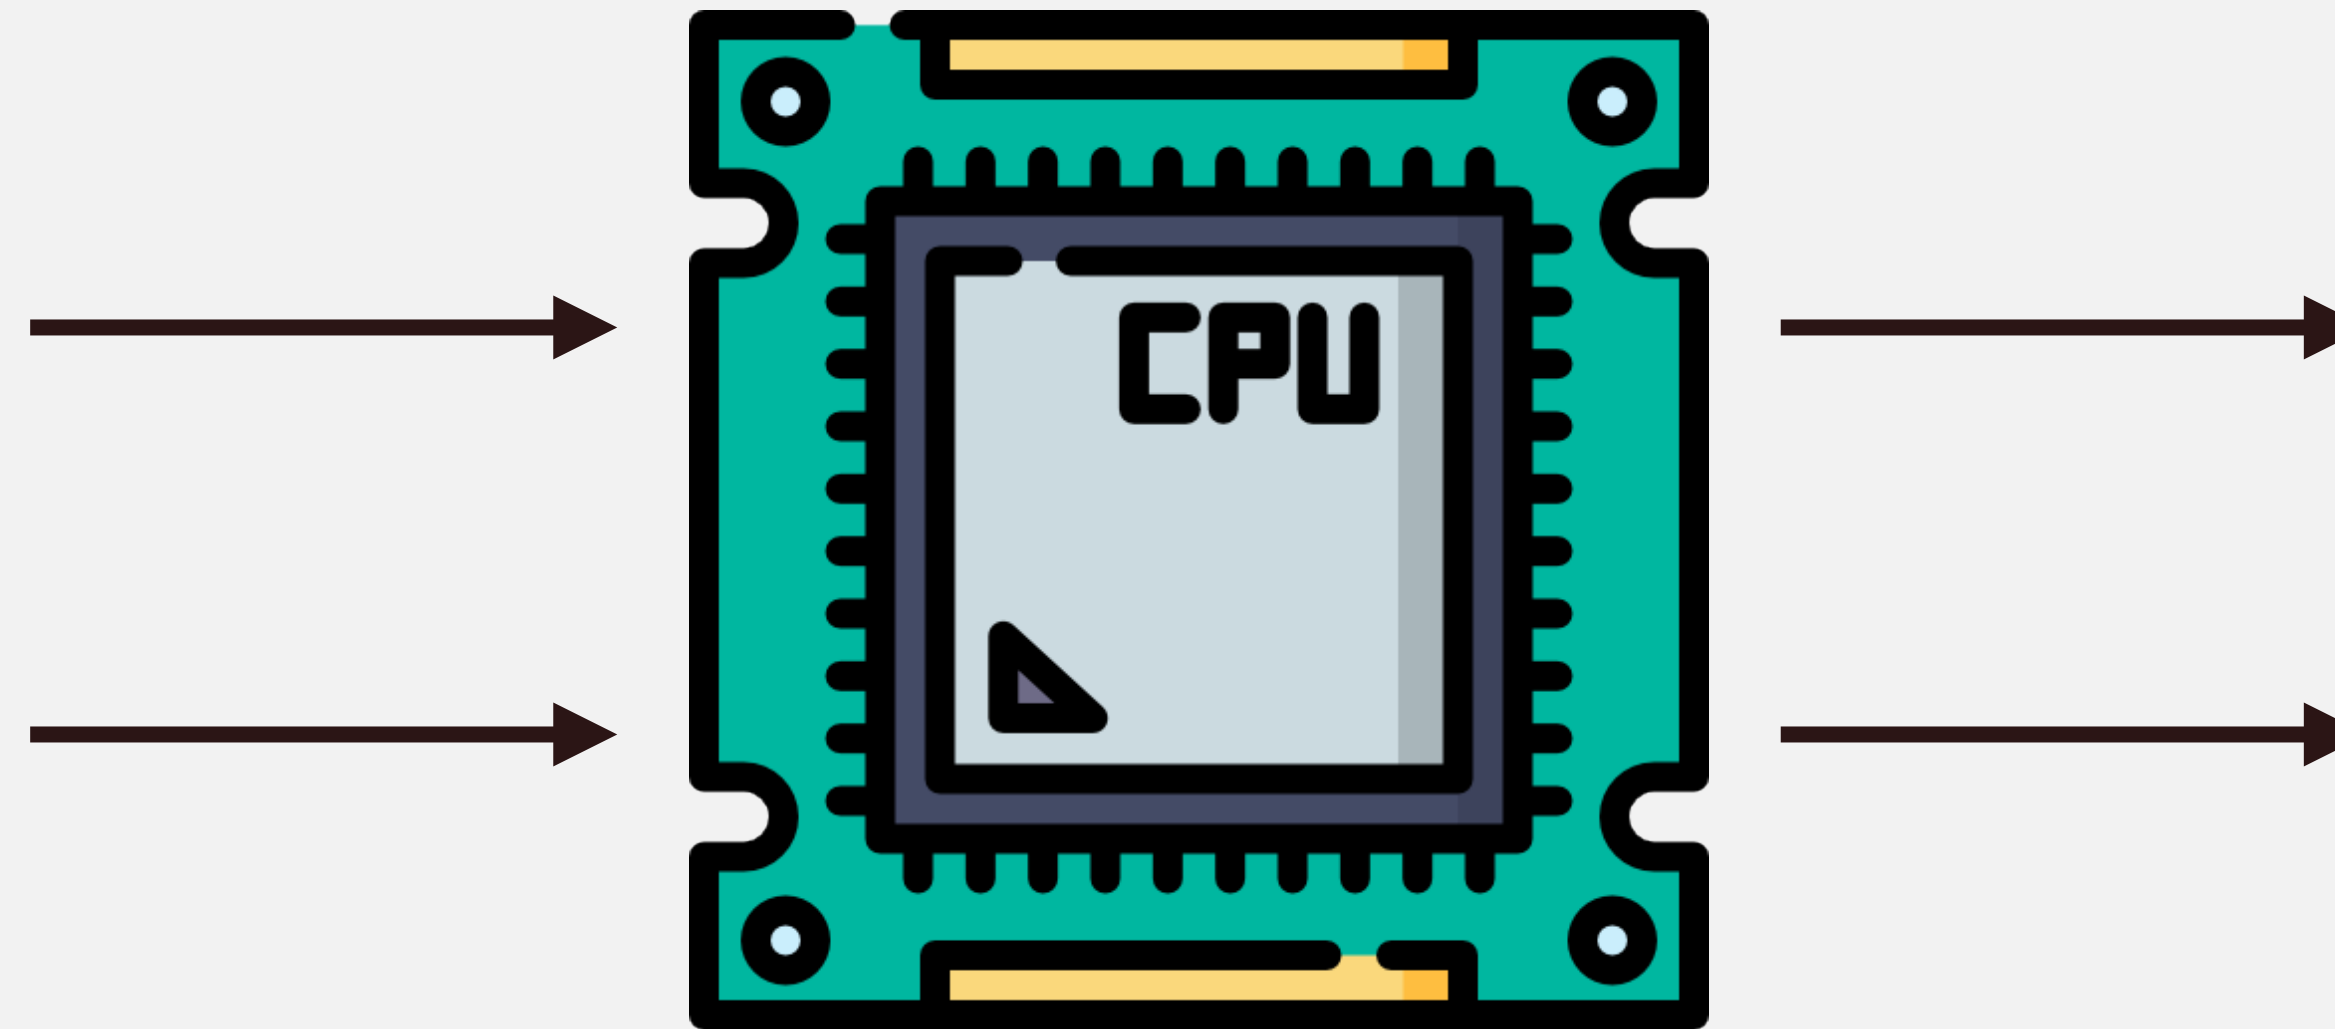


CPU Parallelism



Superscalar

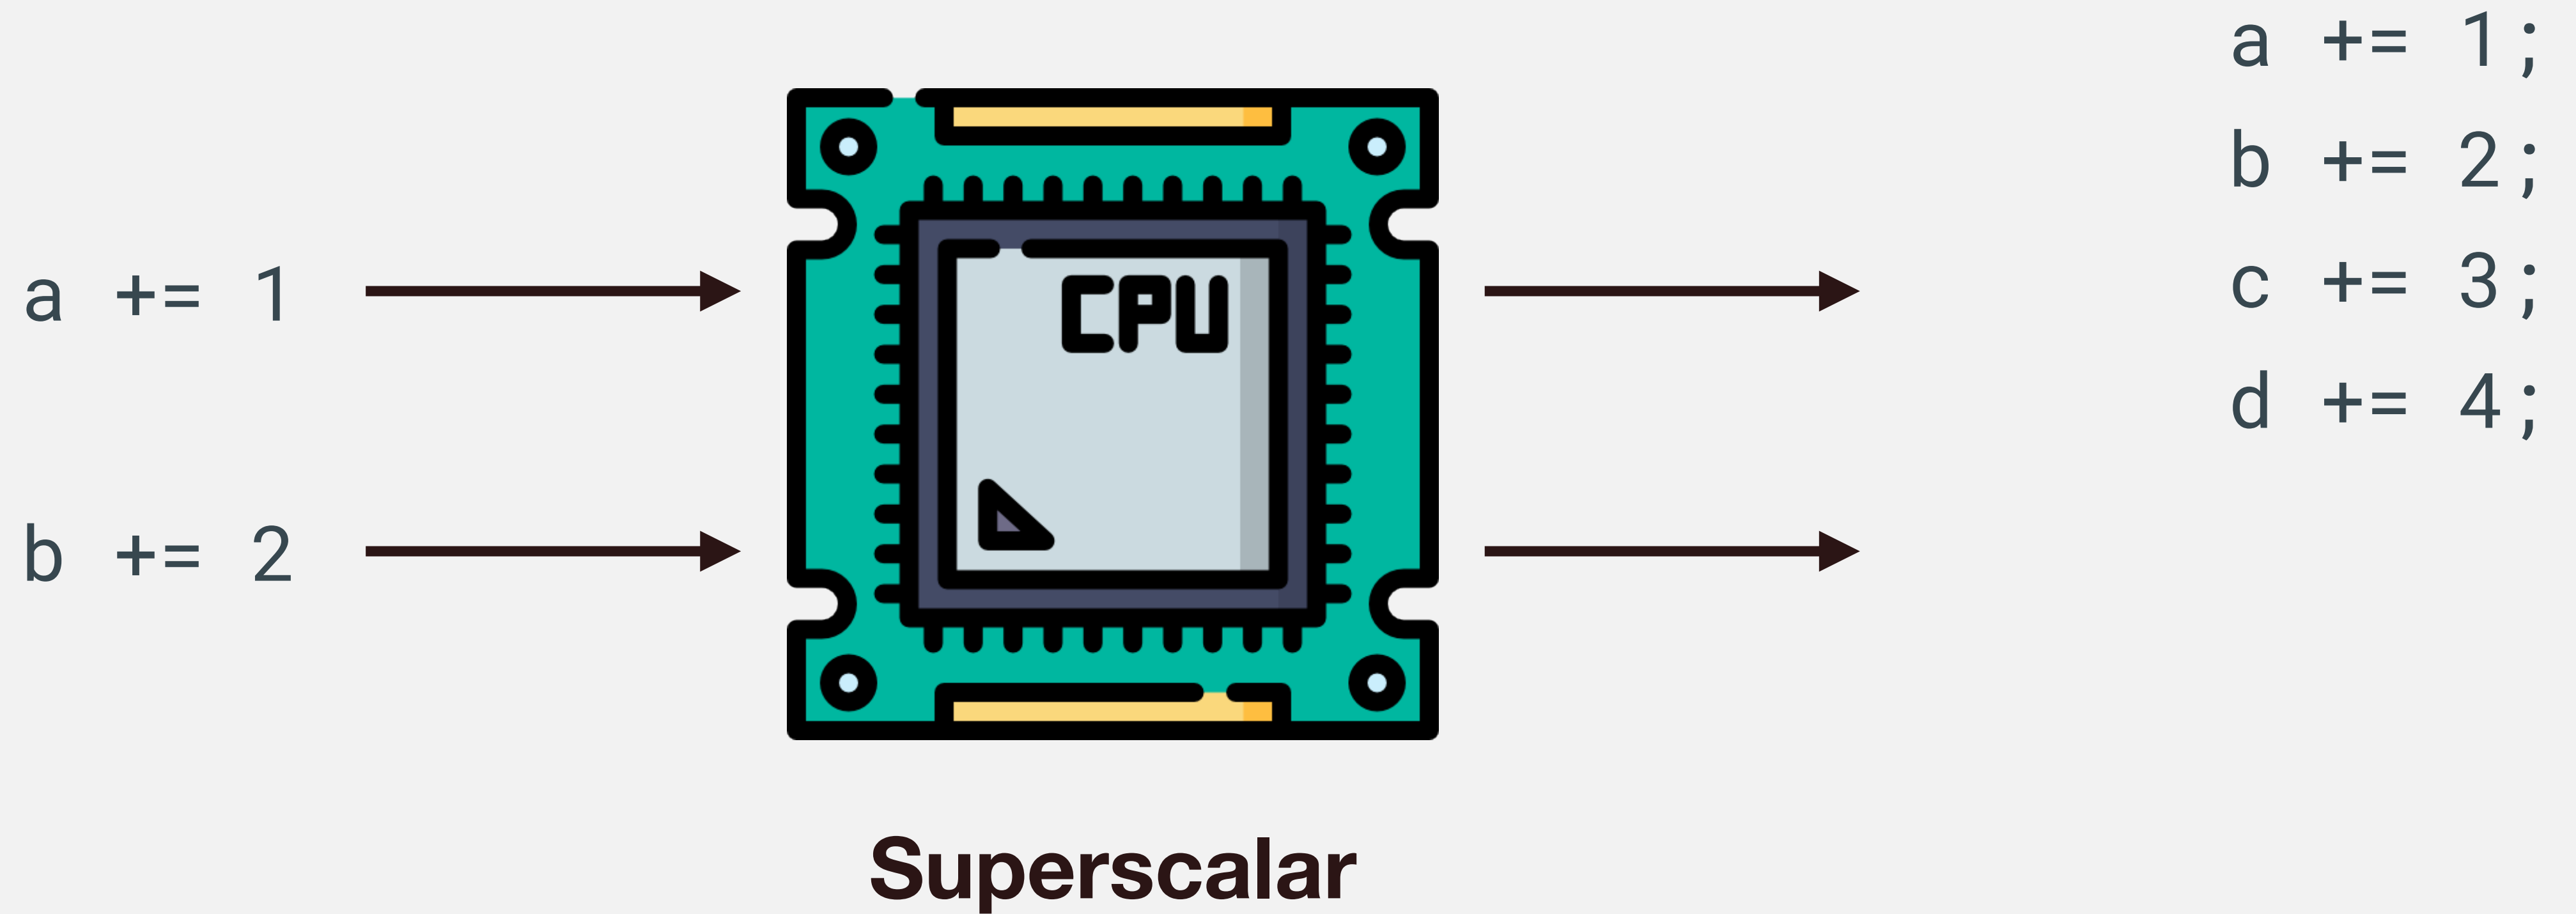
CPU Parallelism



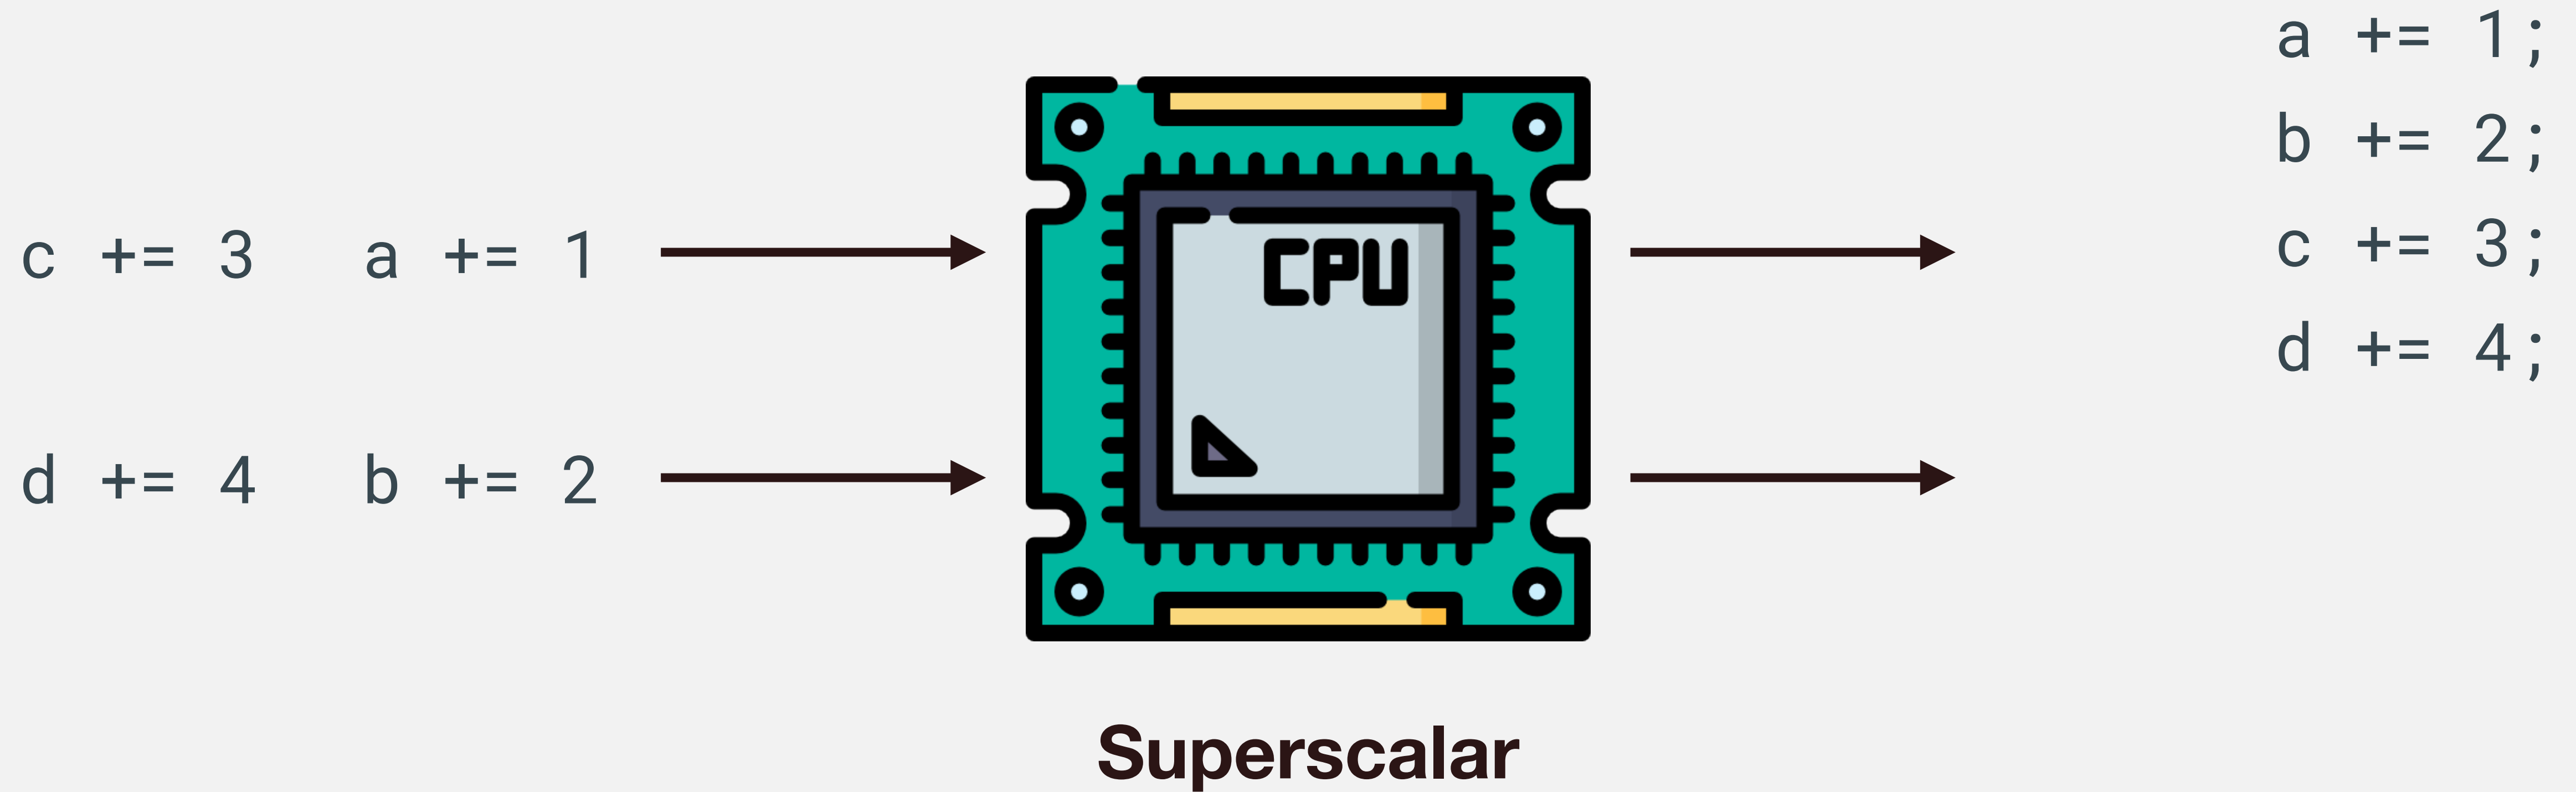
Superscalar

```
a += 1;  
b += 2;  
c += 3;  
d += 4;
```

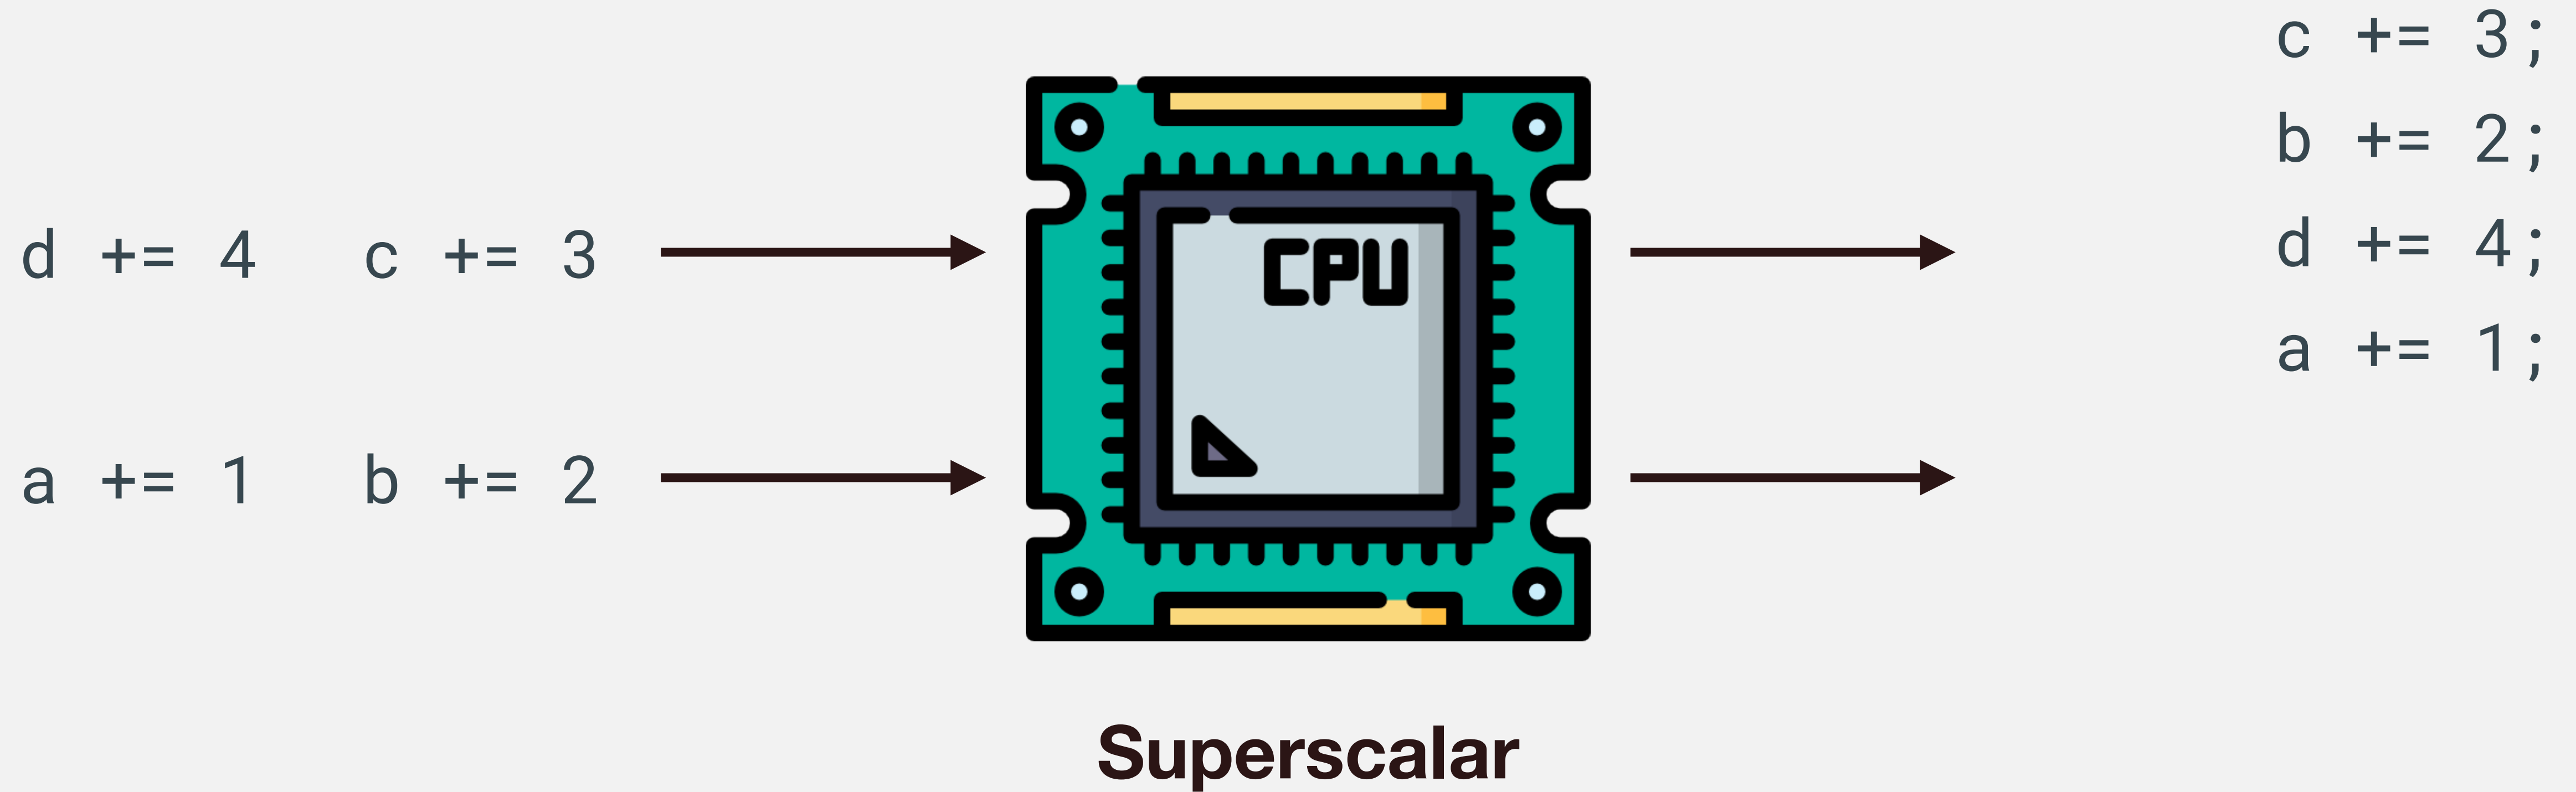
CPU Parallelism



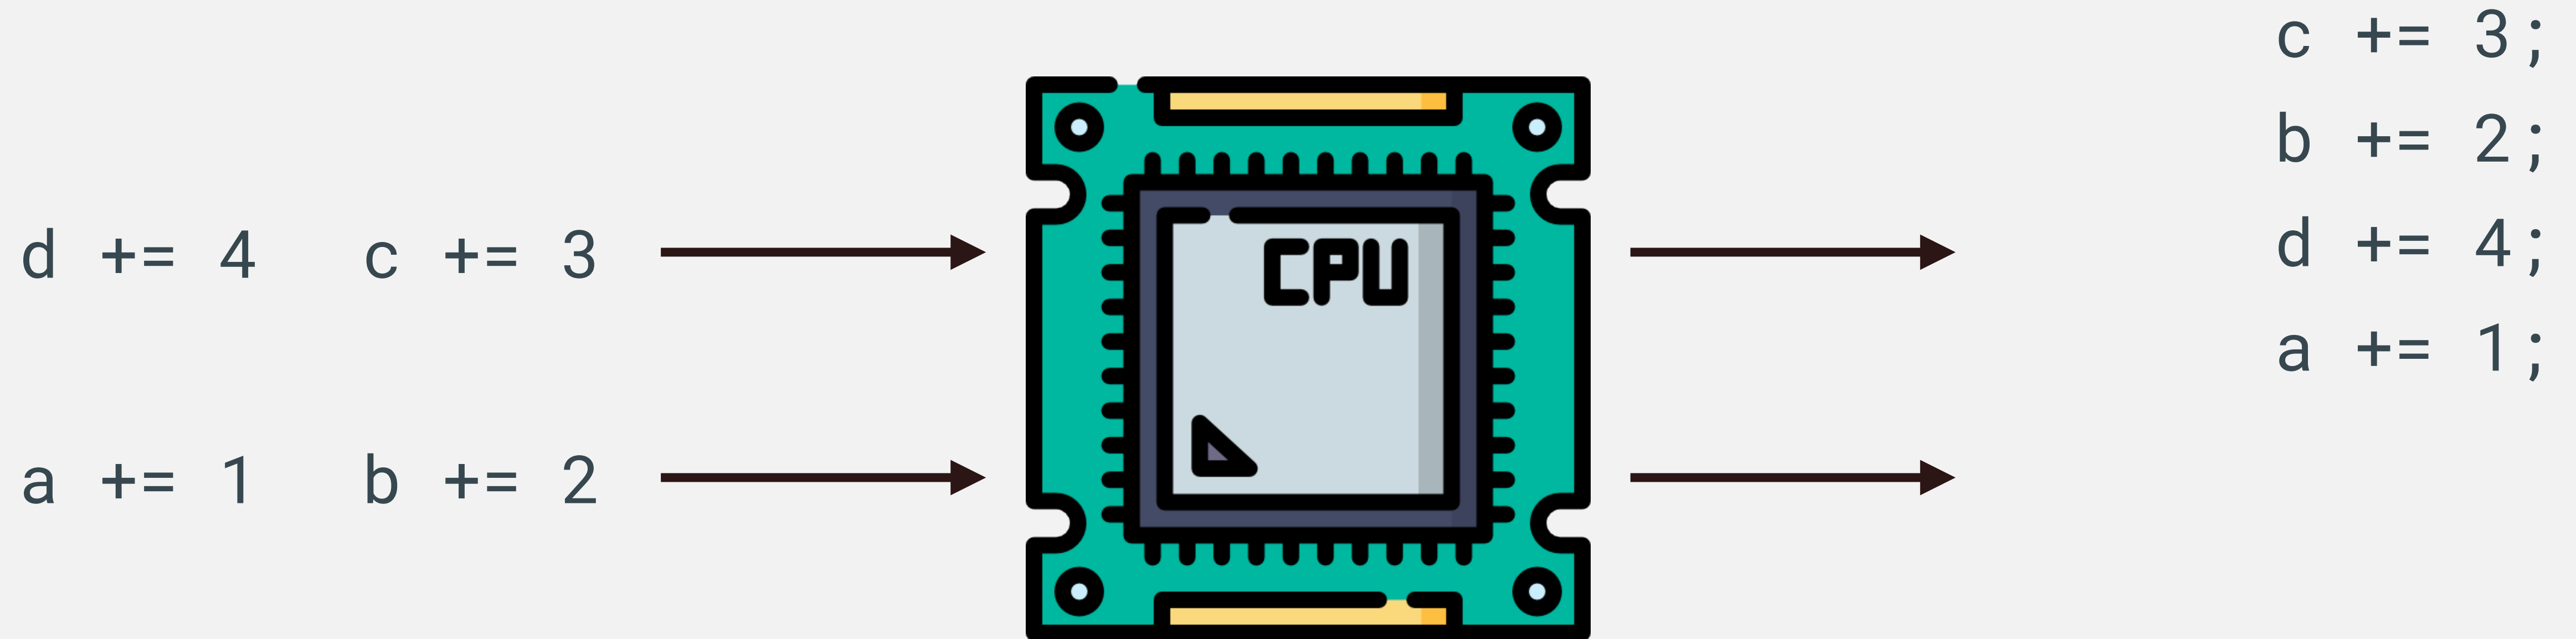
CPU Parallelism



CPU Parallelism



CPU Parallelism



Superscalar

Out-of-Order

Data Dependency

a

1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---

b

2	2	2	2	2	2	2	2
---	---	---	---	---	---	---	---

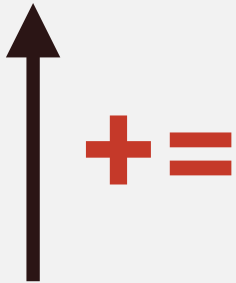
c

3	3	3	3	3	3	3	3
---	---	---	---	---	---	---	---

Data Dependency

a

1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---



b

2	2	2	2	2	2	2	2
---	---	---	---	---	---	---	---

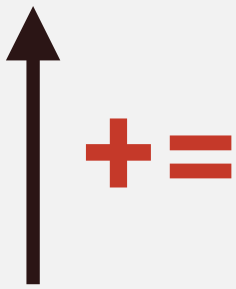
c

3	3	3	3	3	3	3	3
---	---	---	---	---	---	---	---

Data Dependency

a

3	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---



b

2	2	2	2	2	2	2	2
---	---	---	---	---	---	---	---

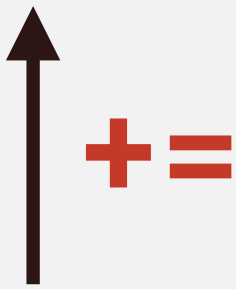
c

3	3	3	3	3	3	3	3
---	---	---	---	---	---	---	---

Data Dependency

a

3	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---



b

2	2	2	2	2	2	2	2
---	---	---	---	---	---	---	---



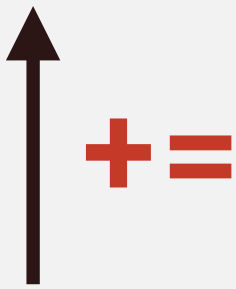
c

3	3	3	3	3	3	3	3
---	---	---	---	---	---	---	---

Data Dependency

a

3	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---



b

2	5	2	2	2	2	2	2
---	---	---	---	---	---	---	---



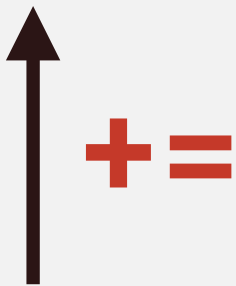
c

3	3	3	3	3	3	3	3
---	---	---	---	---	---	---	---

Data Dependency

a

3	6	1	1	1	1	1	1
---	---	---	---	---	---	---	---



b

2	5	2	2	2	2	2	2
---	---	---	---	---	---	---	---



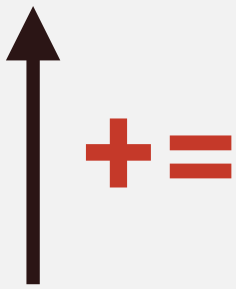
c

3	3	3	3	3	3	3	3
---	---	---	---	---	---	---	---

Data Dependency

a

3	6	1	1	1	1	1	1
---	---	---	---	---	---	---	---



b

2	5	5	2	2	2	2	2
---	---	---	---	---	---	---	---



c

3	3	3	3	3	3	3	3
---	---	---	---	---	---	---	---

Data Dependency

a

3	6	1	1	1	1	1	1
---	---	---	---	---	---	---	---



+=



...

b

2	5	5	2	2	2	2	2
---	---	---	---	---	---	---	---



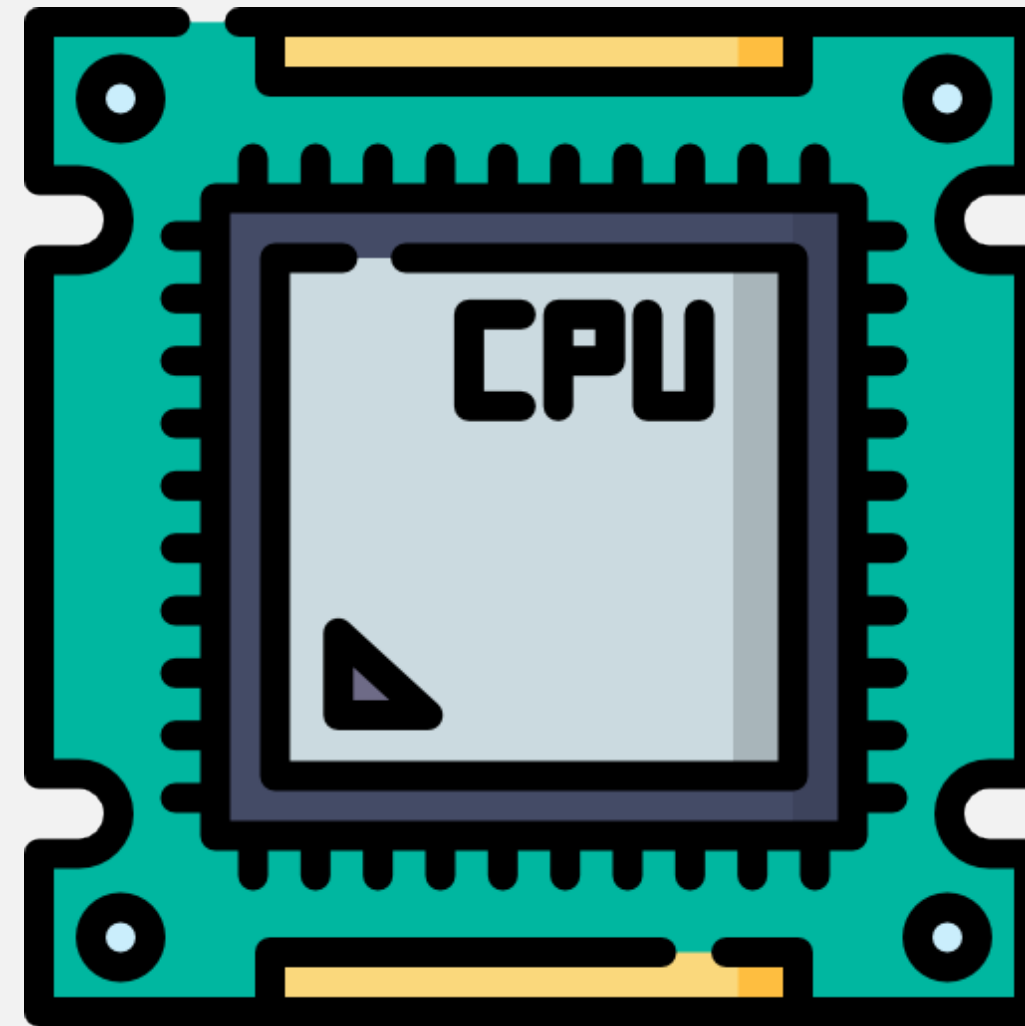
...

c

3	3	3	3	3	3	3	3
---	---	---	---	---	---	---	---

CPU Parallelism

```
a += 3    a += 1  
b += 4    b += 2
```



```
a += 1;  
a += 3;  
b += 2;  
b += 4;
```

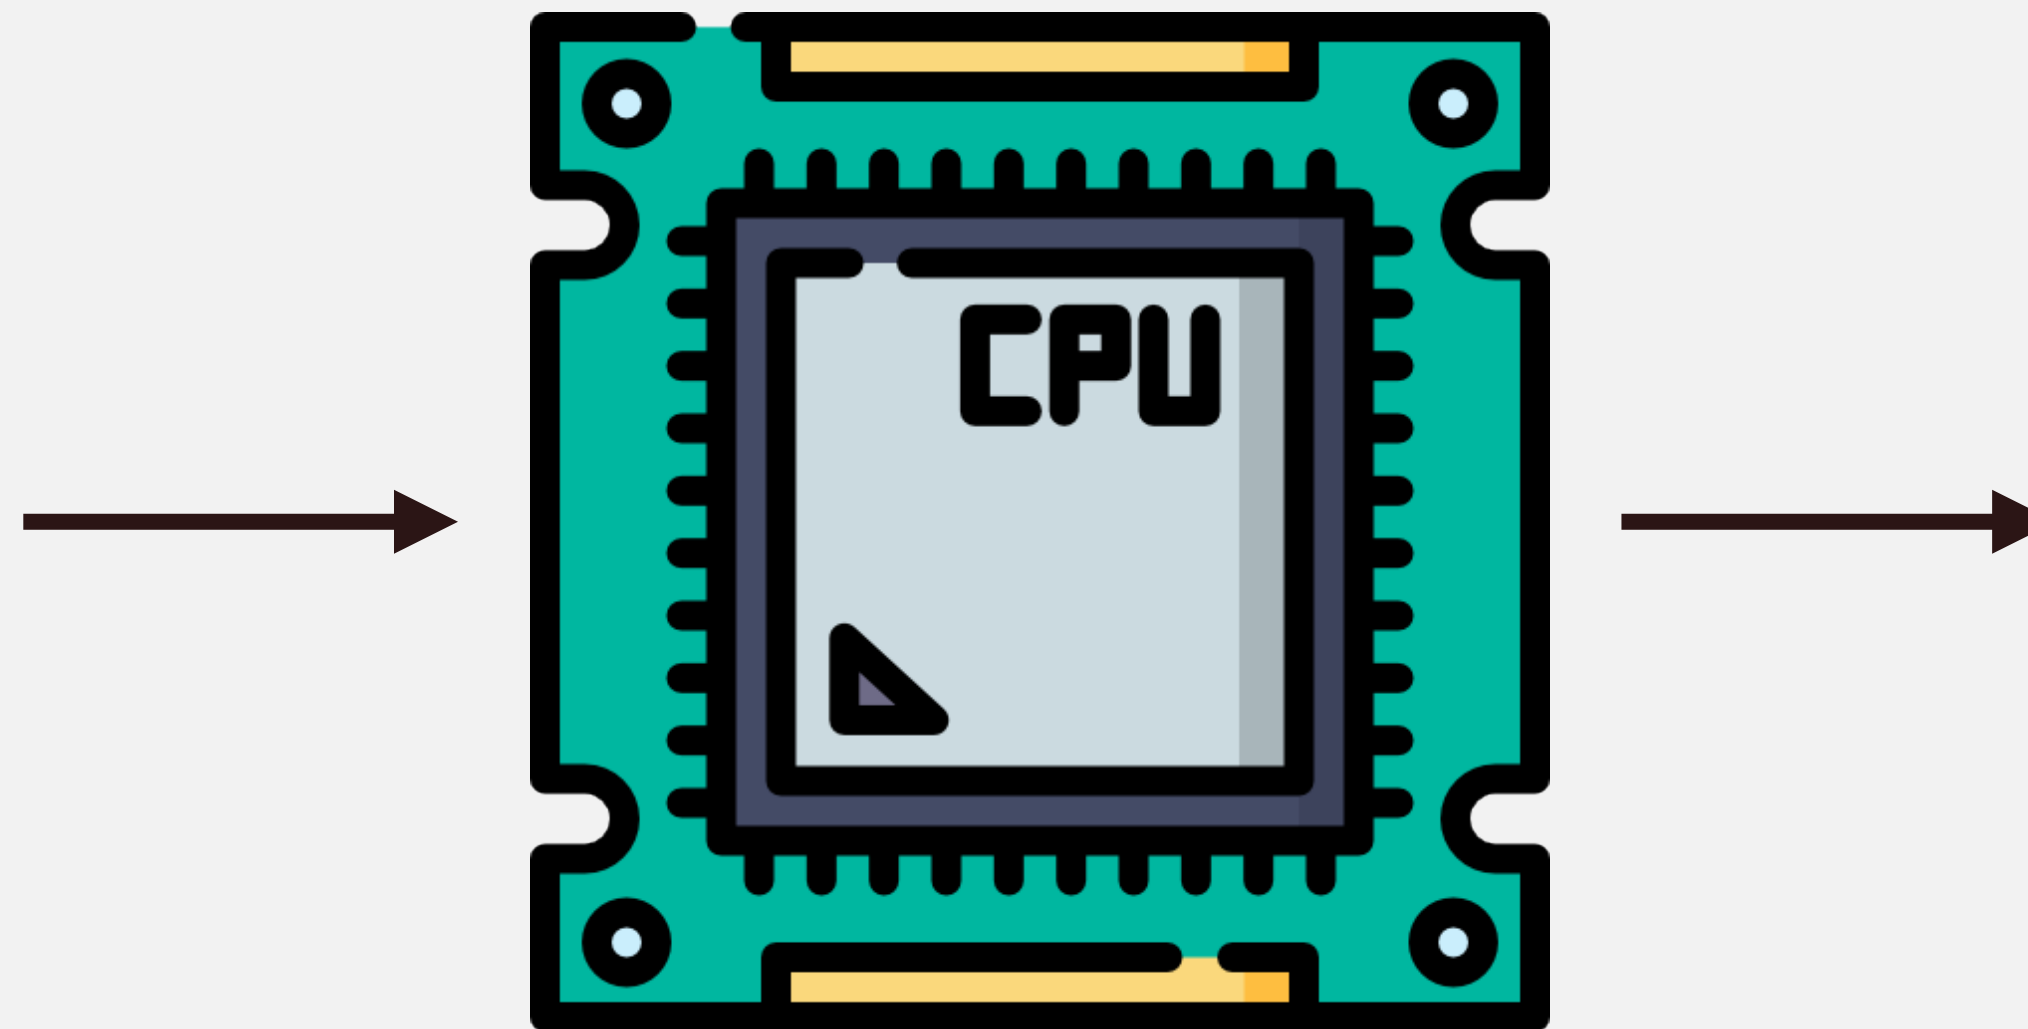
Loop Unrolling

- Reduces # of loop test
- Help compiler improve parallelism

Superscalar

Out-of-Order

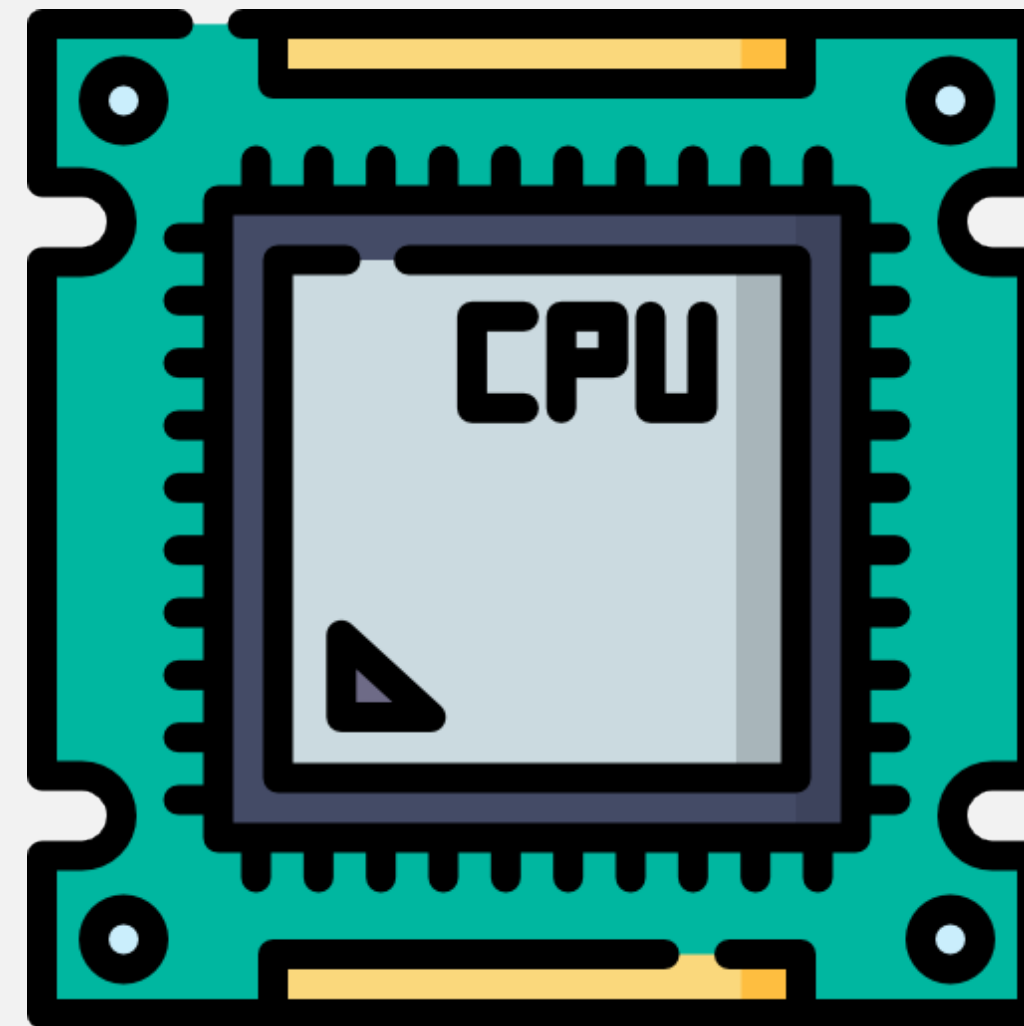
Data Parallelism (SIMD)



```
a += 1;  
b += 2;  
c += 3;  
d += 4;
```

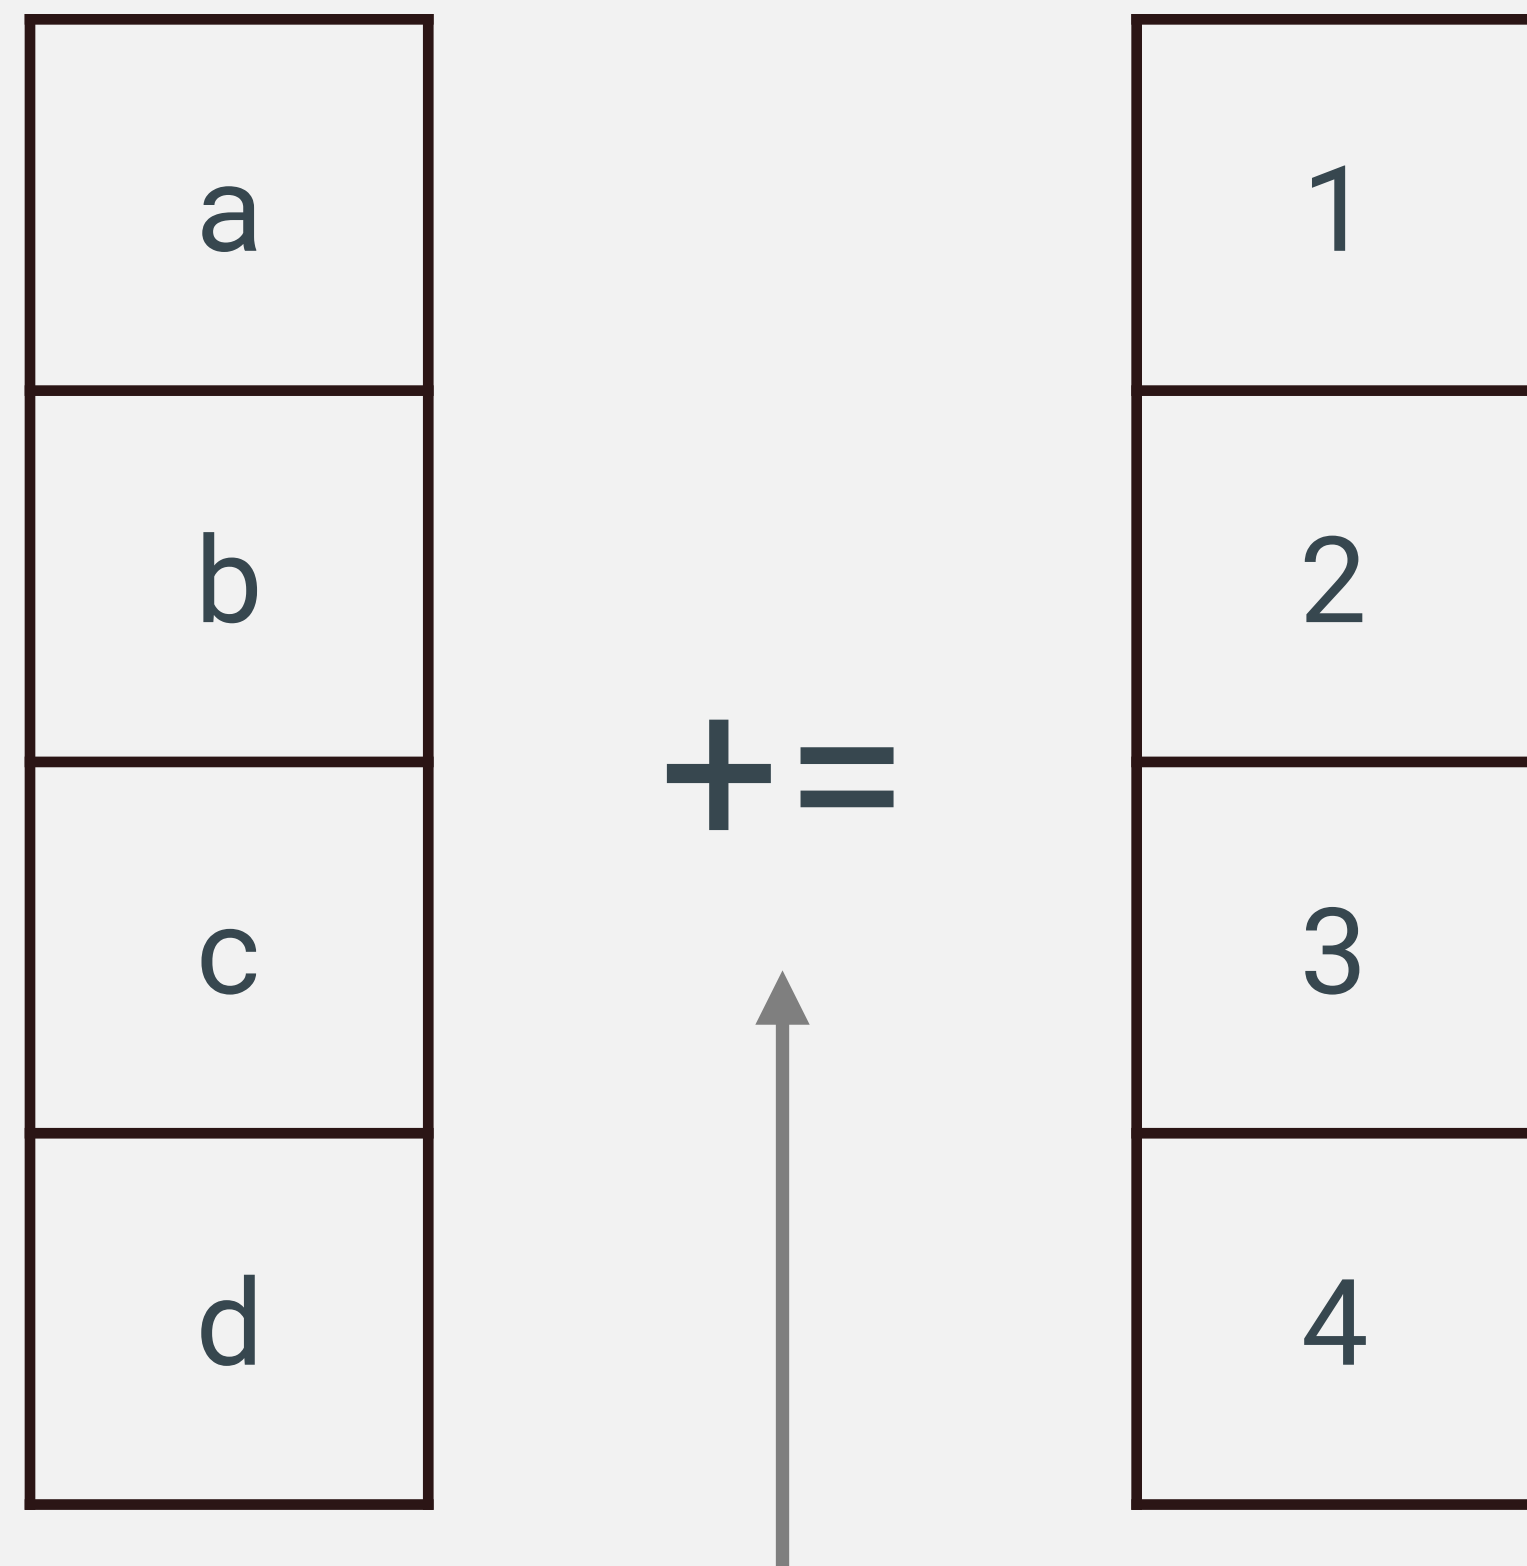
Data Parallelism (SIMD)

```
d+=4; c+=3; b+=2; a+=1;
```

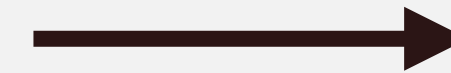
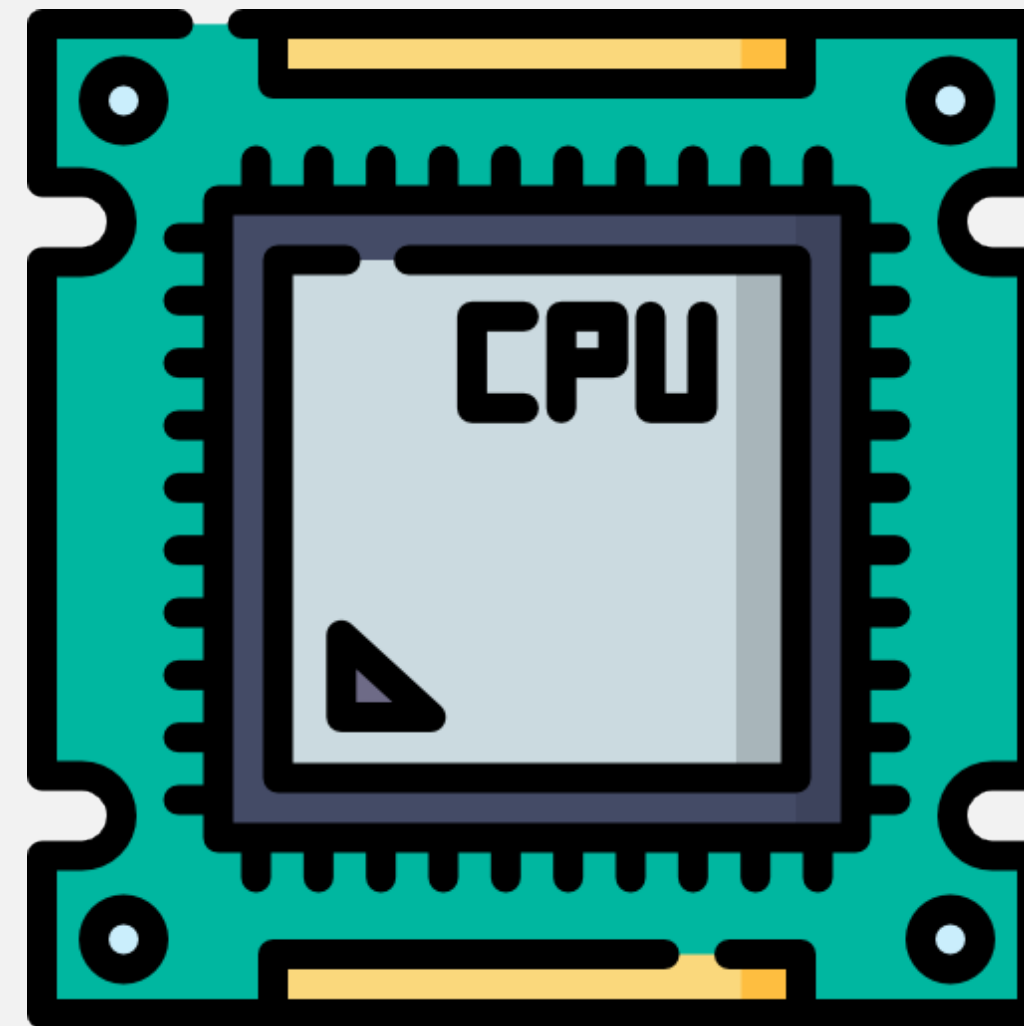


```
a += 1;  
b += 2;  
c += 3;  
d += 4;
```

Data Parallelism (SIMD)

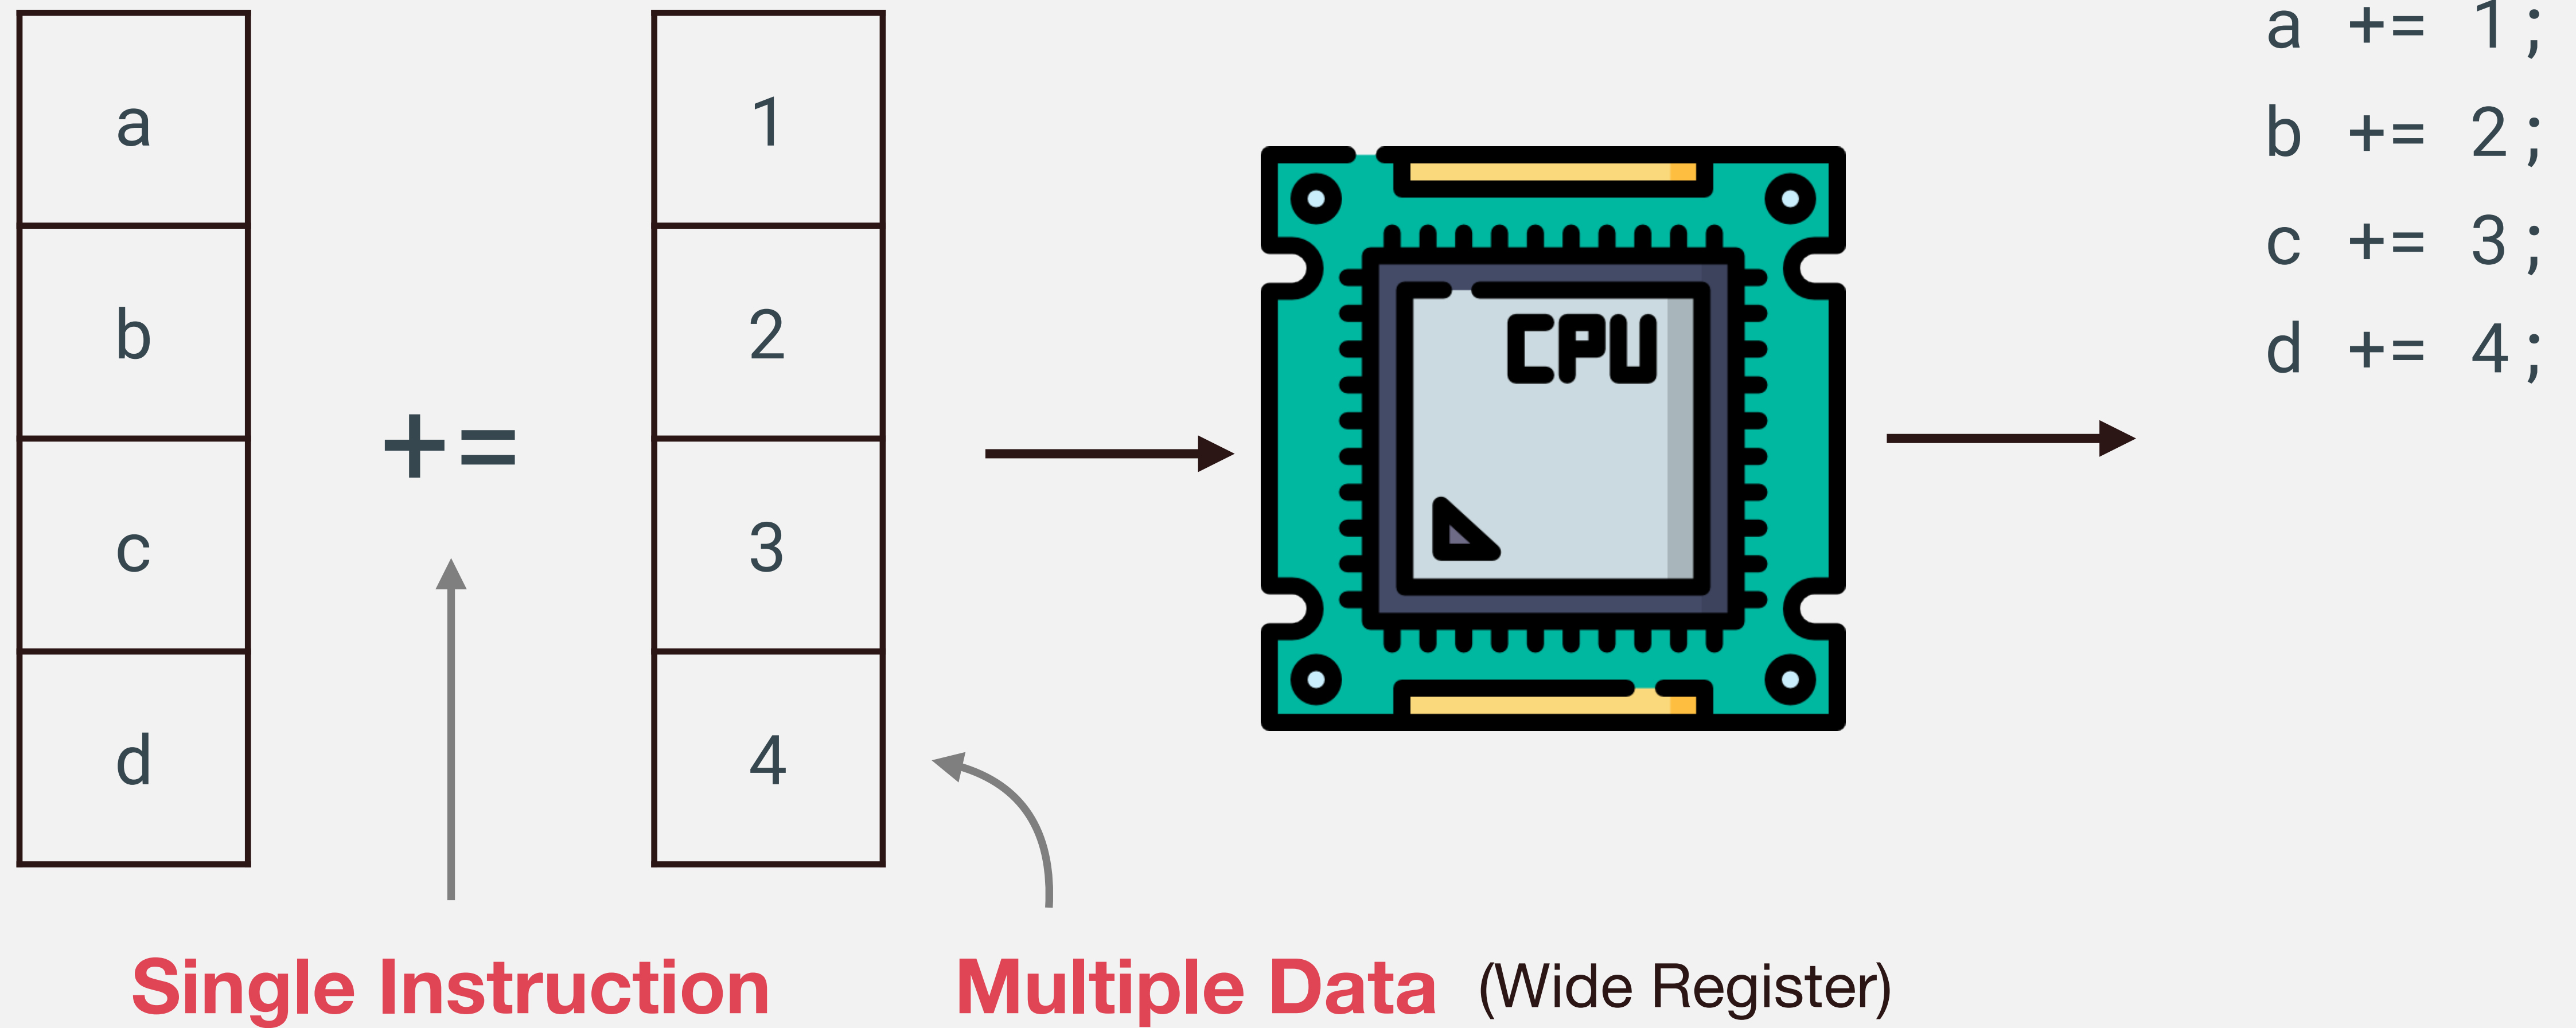


Single Instruction



```
a += 1;  
b += 2;  
c += 3;  
d += 4;
```

Data Parallelism (SIMD)

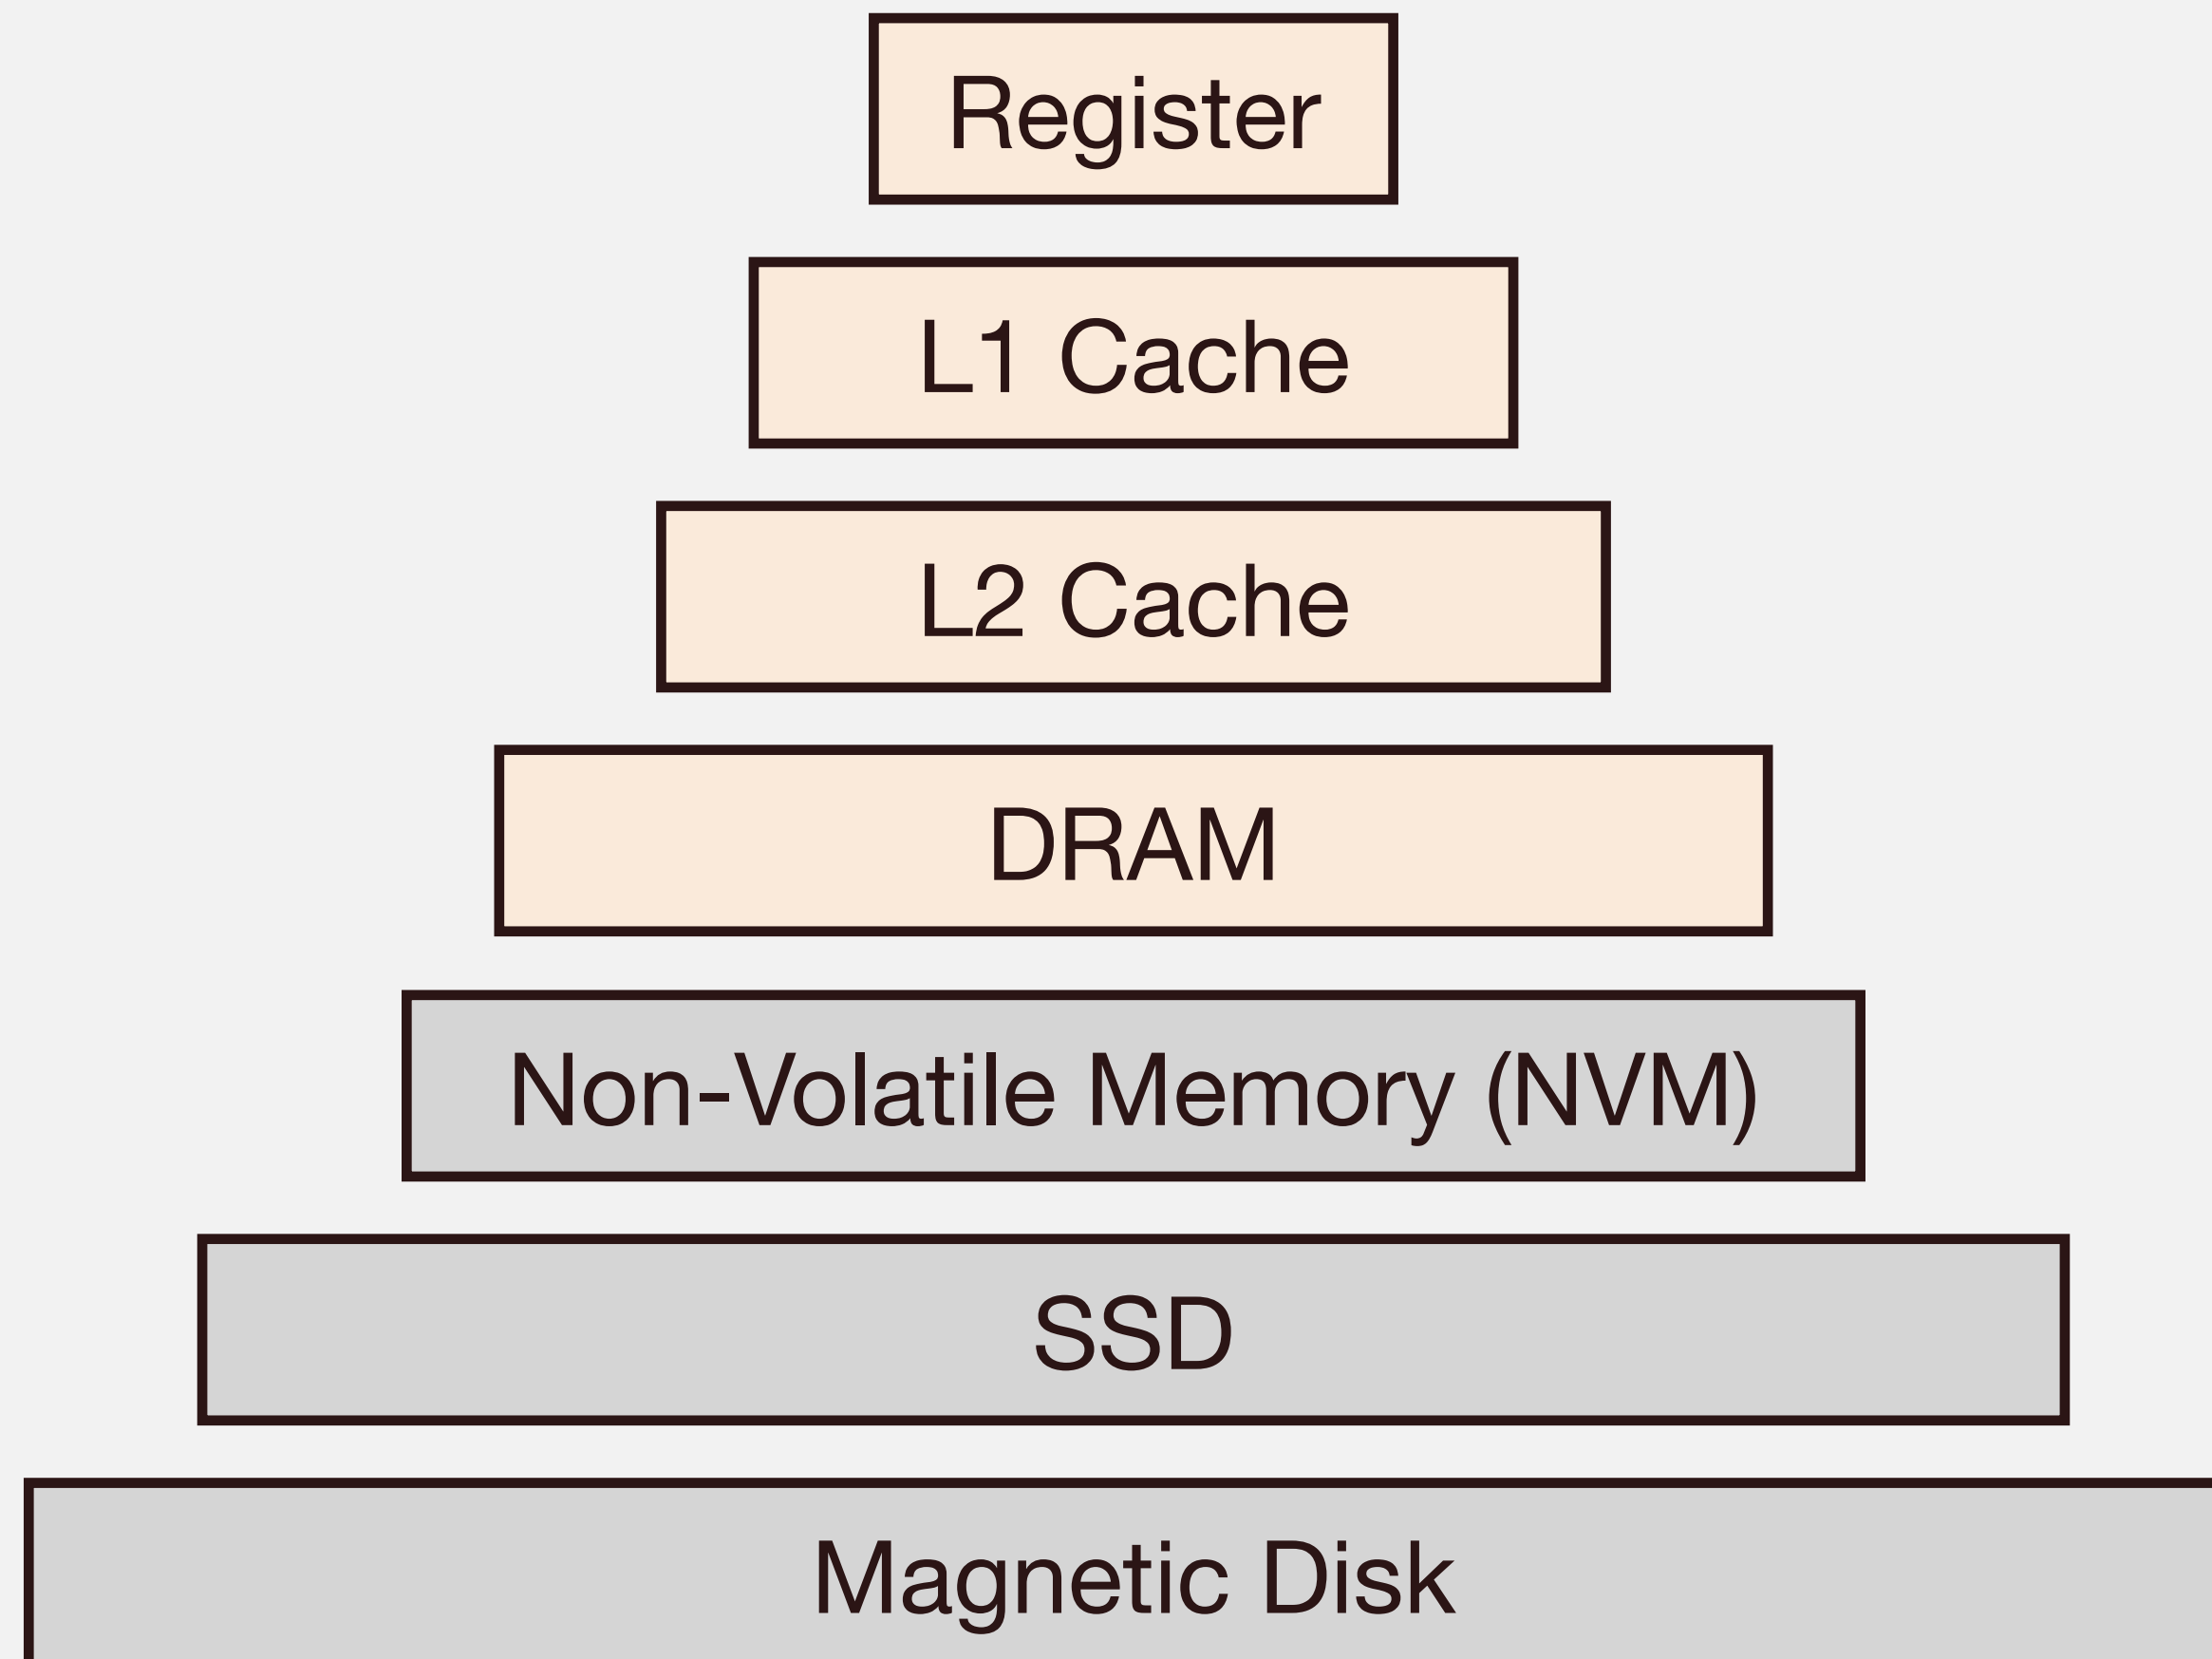


Low-Level Optimization Techniques

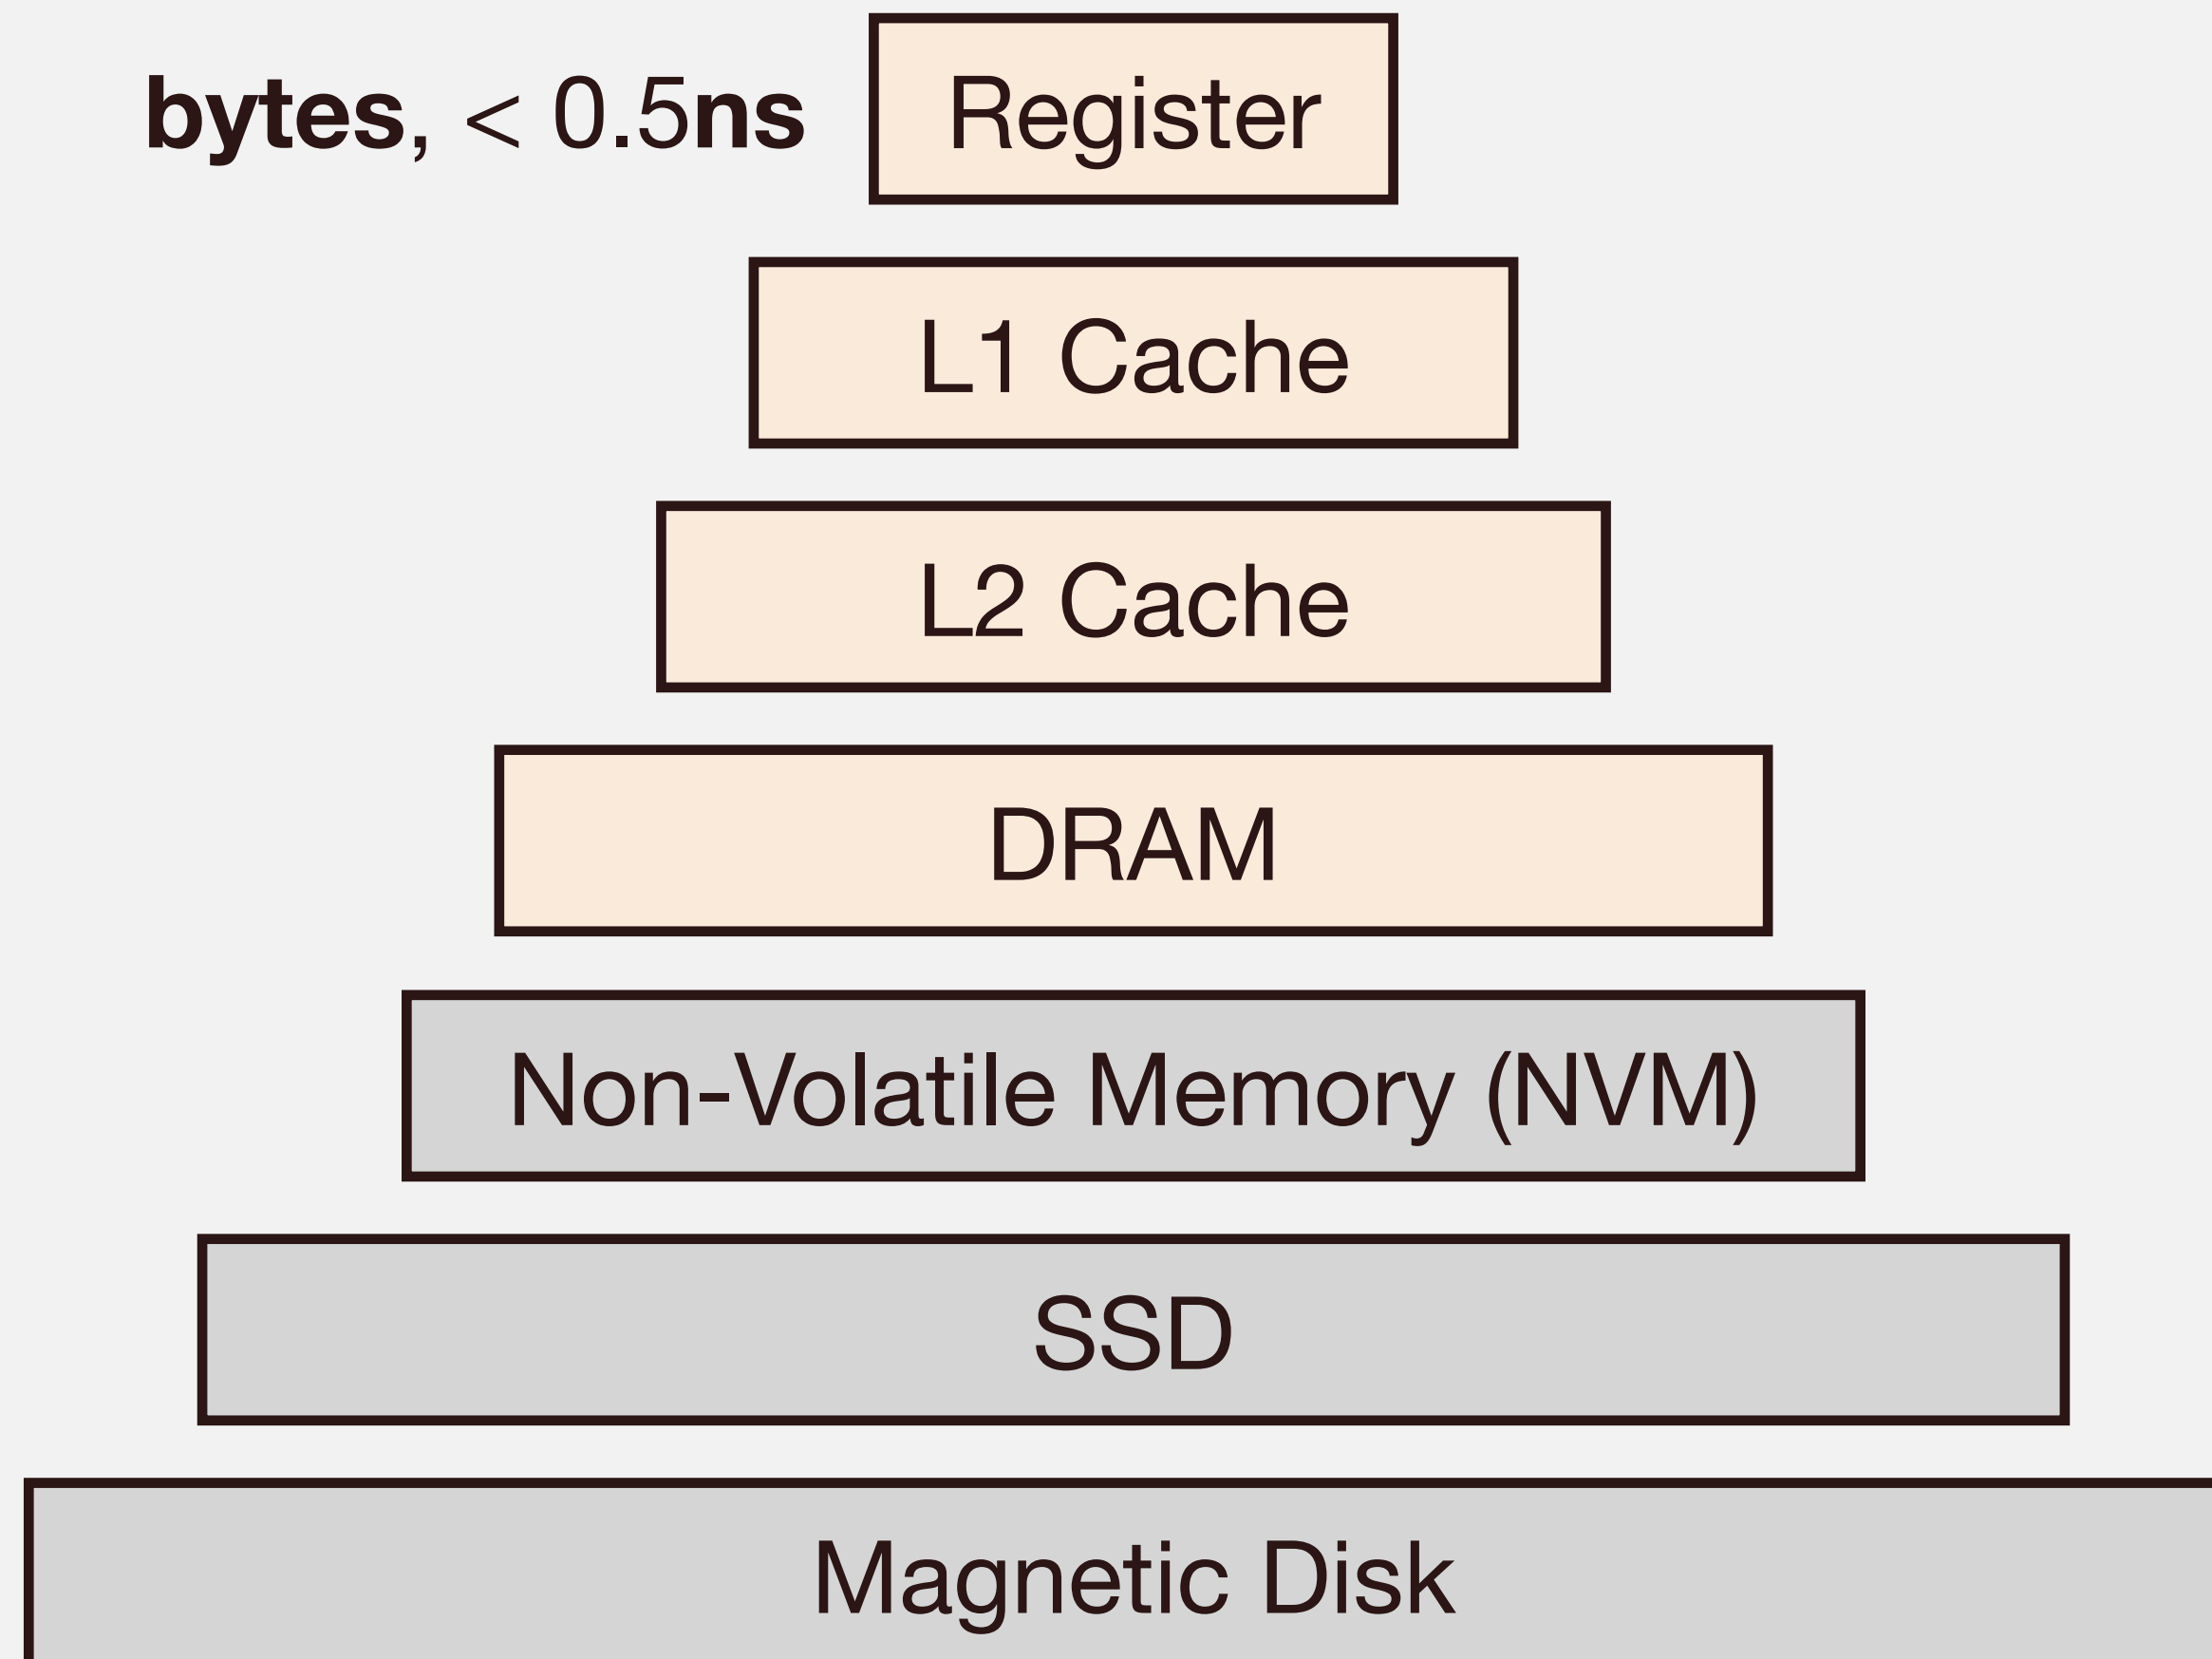
CPU

- Reduce Operator Strength
- Reduce Branch Misprediction
- Reduce Data Dependency
- SIMD

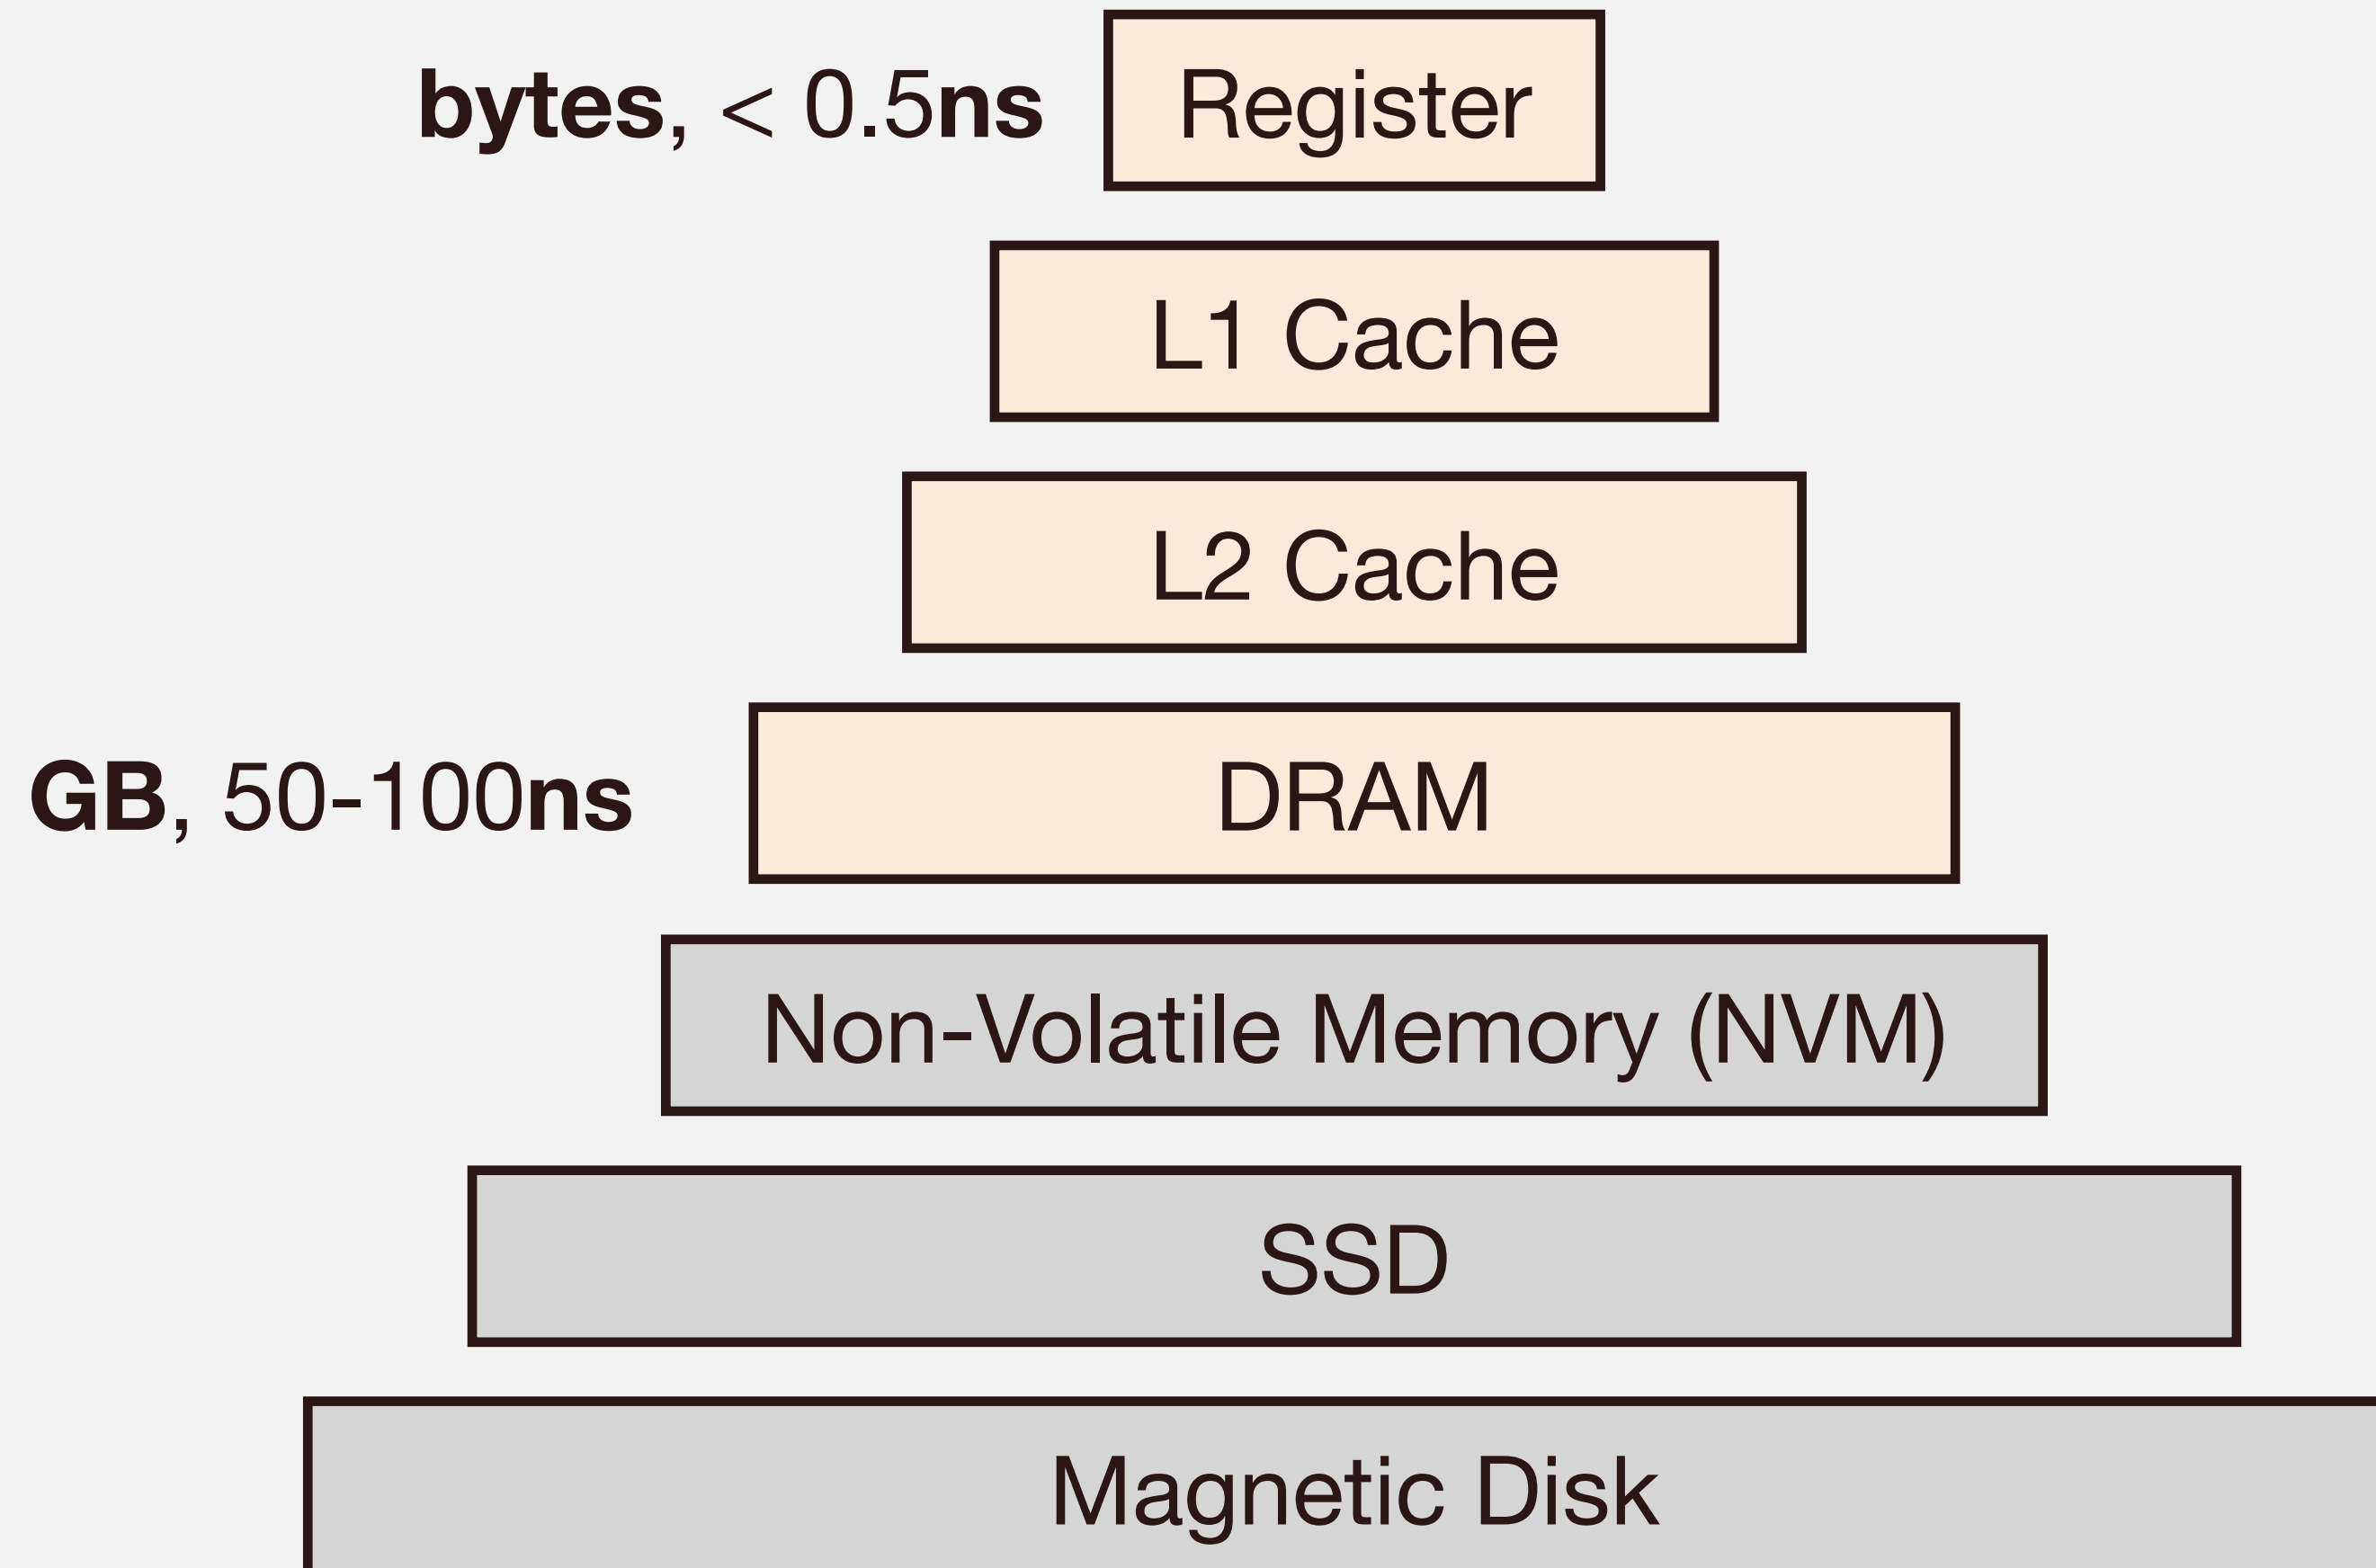
Storage Hierarchy



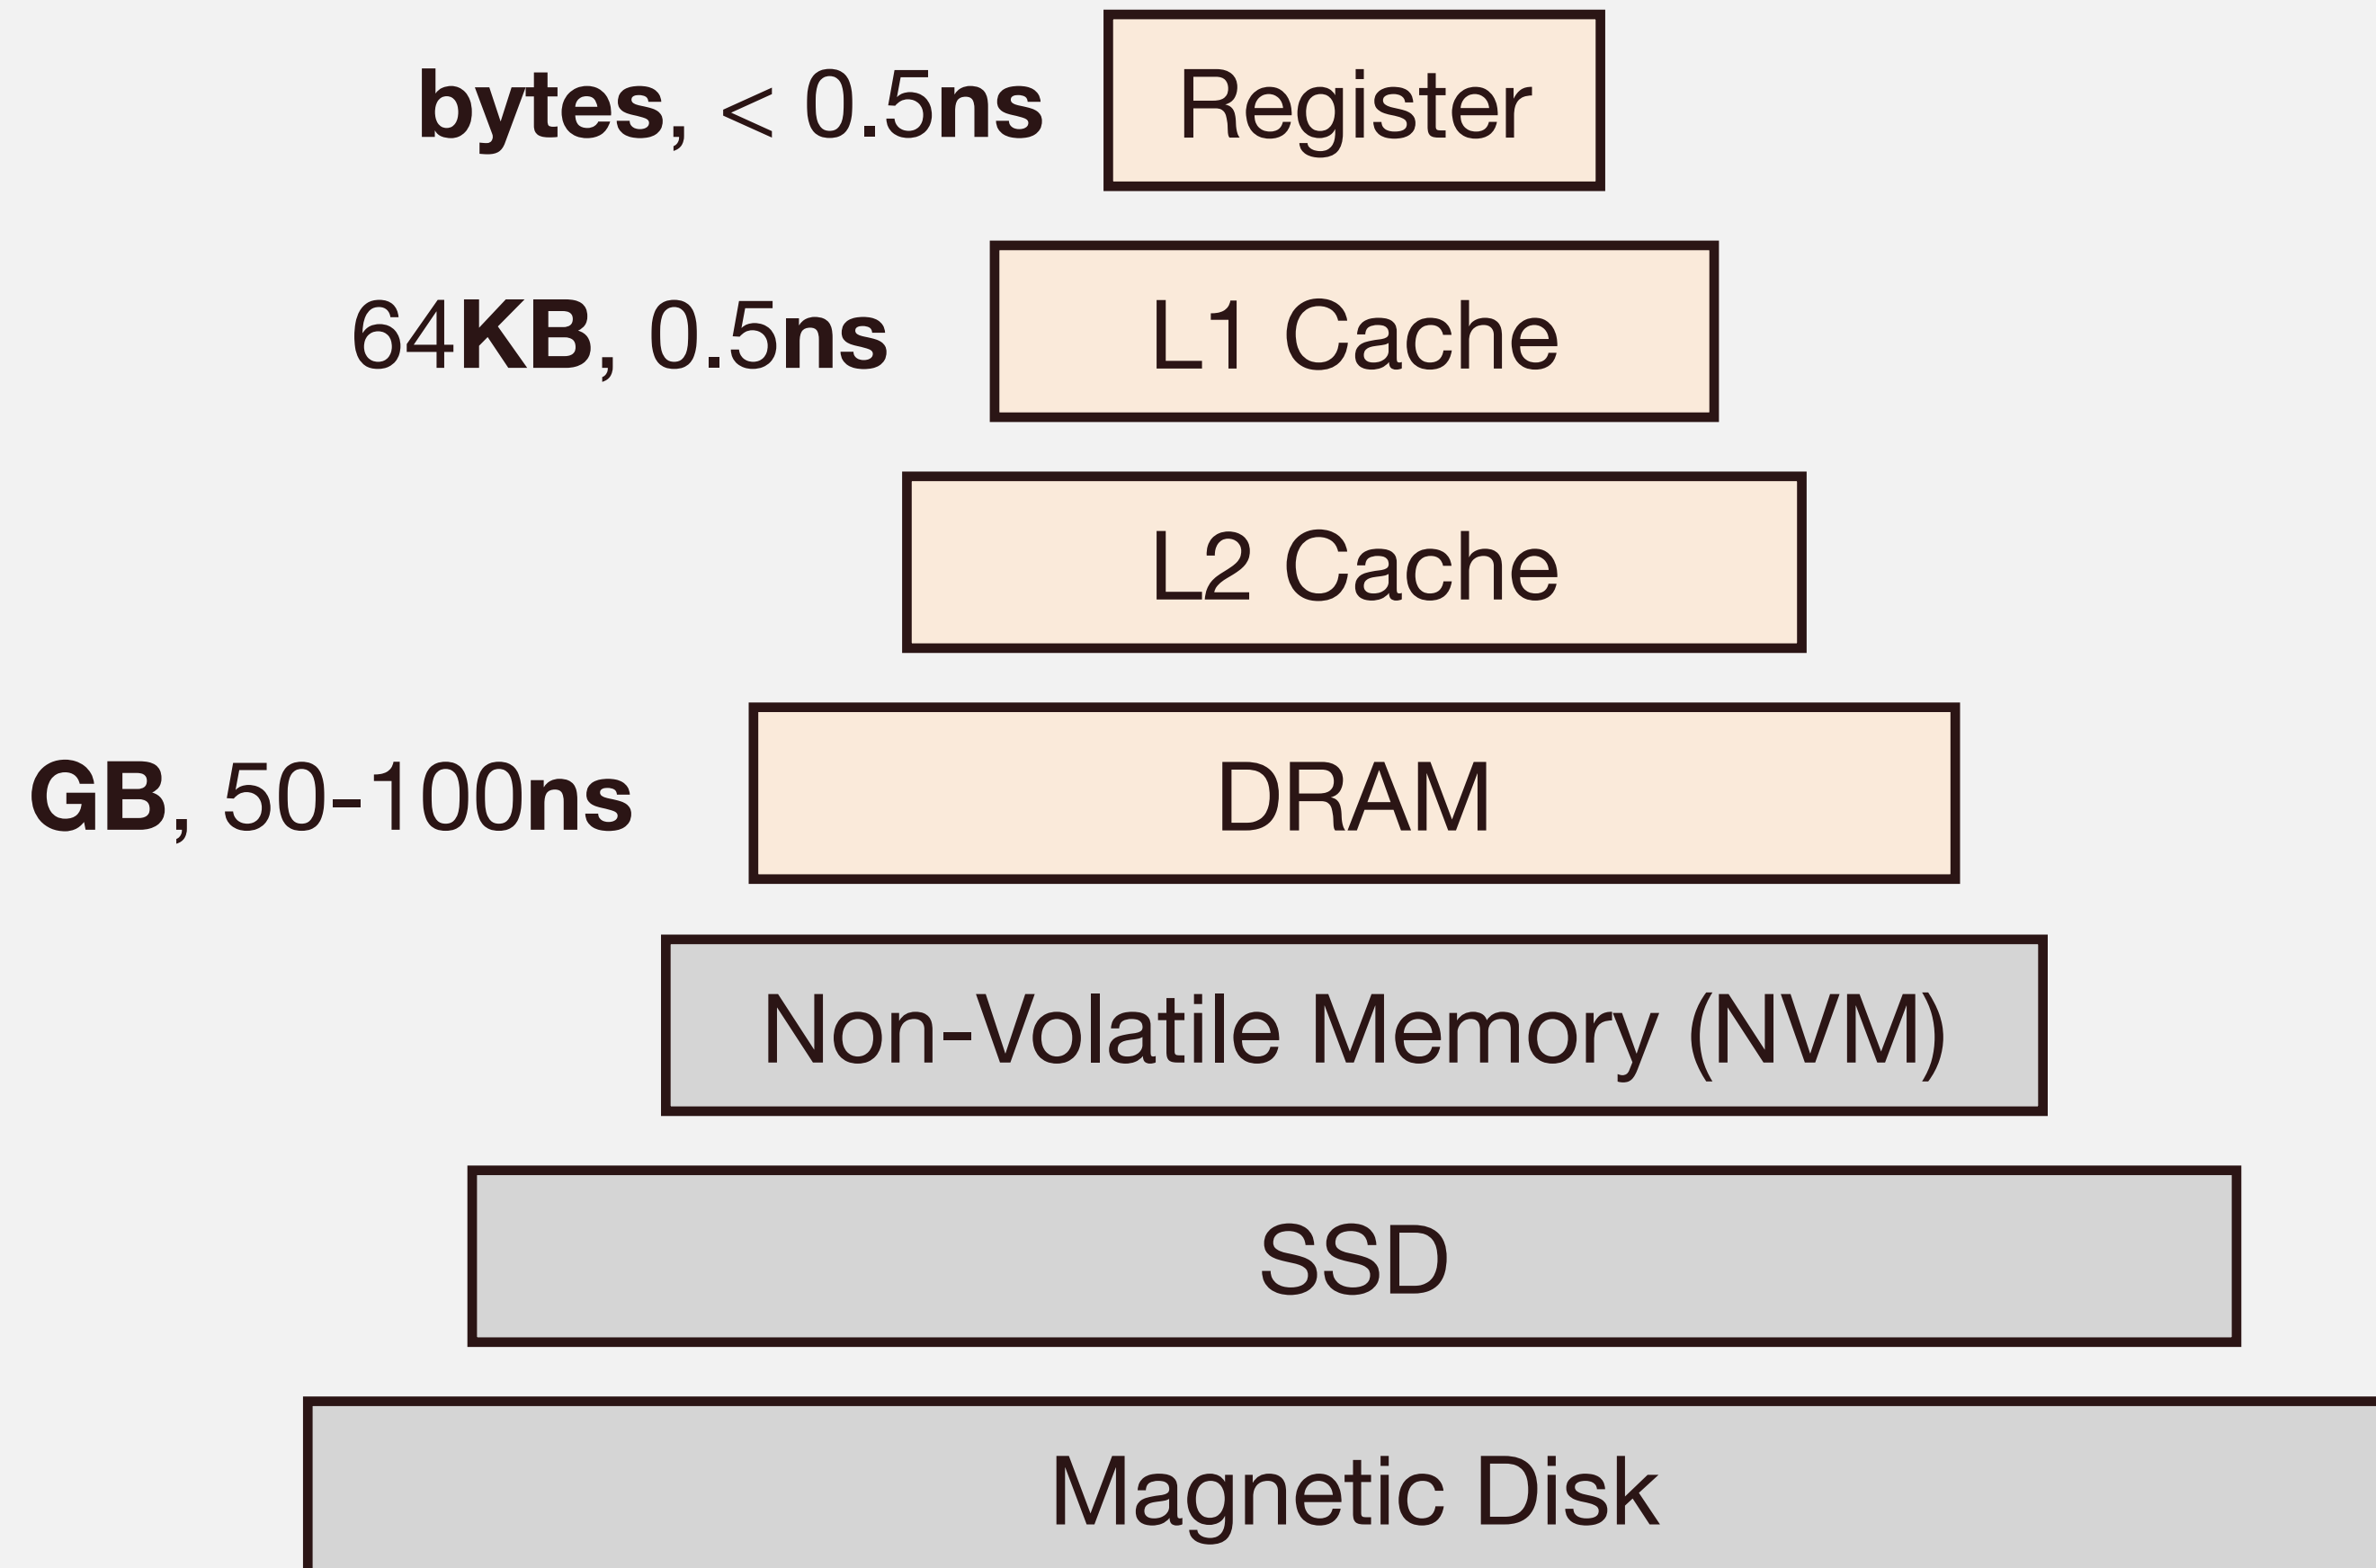
Storage Hierarchy



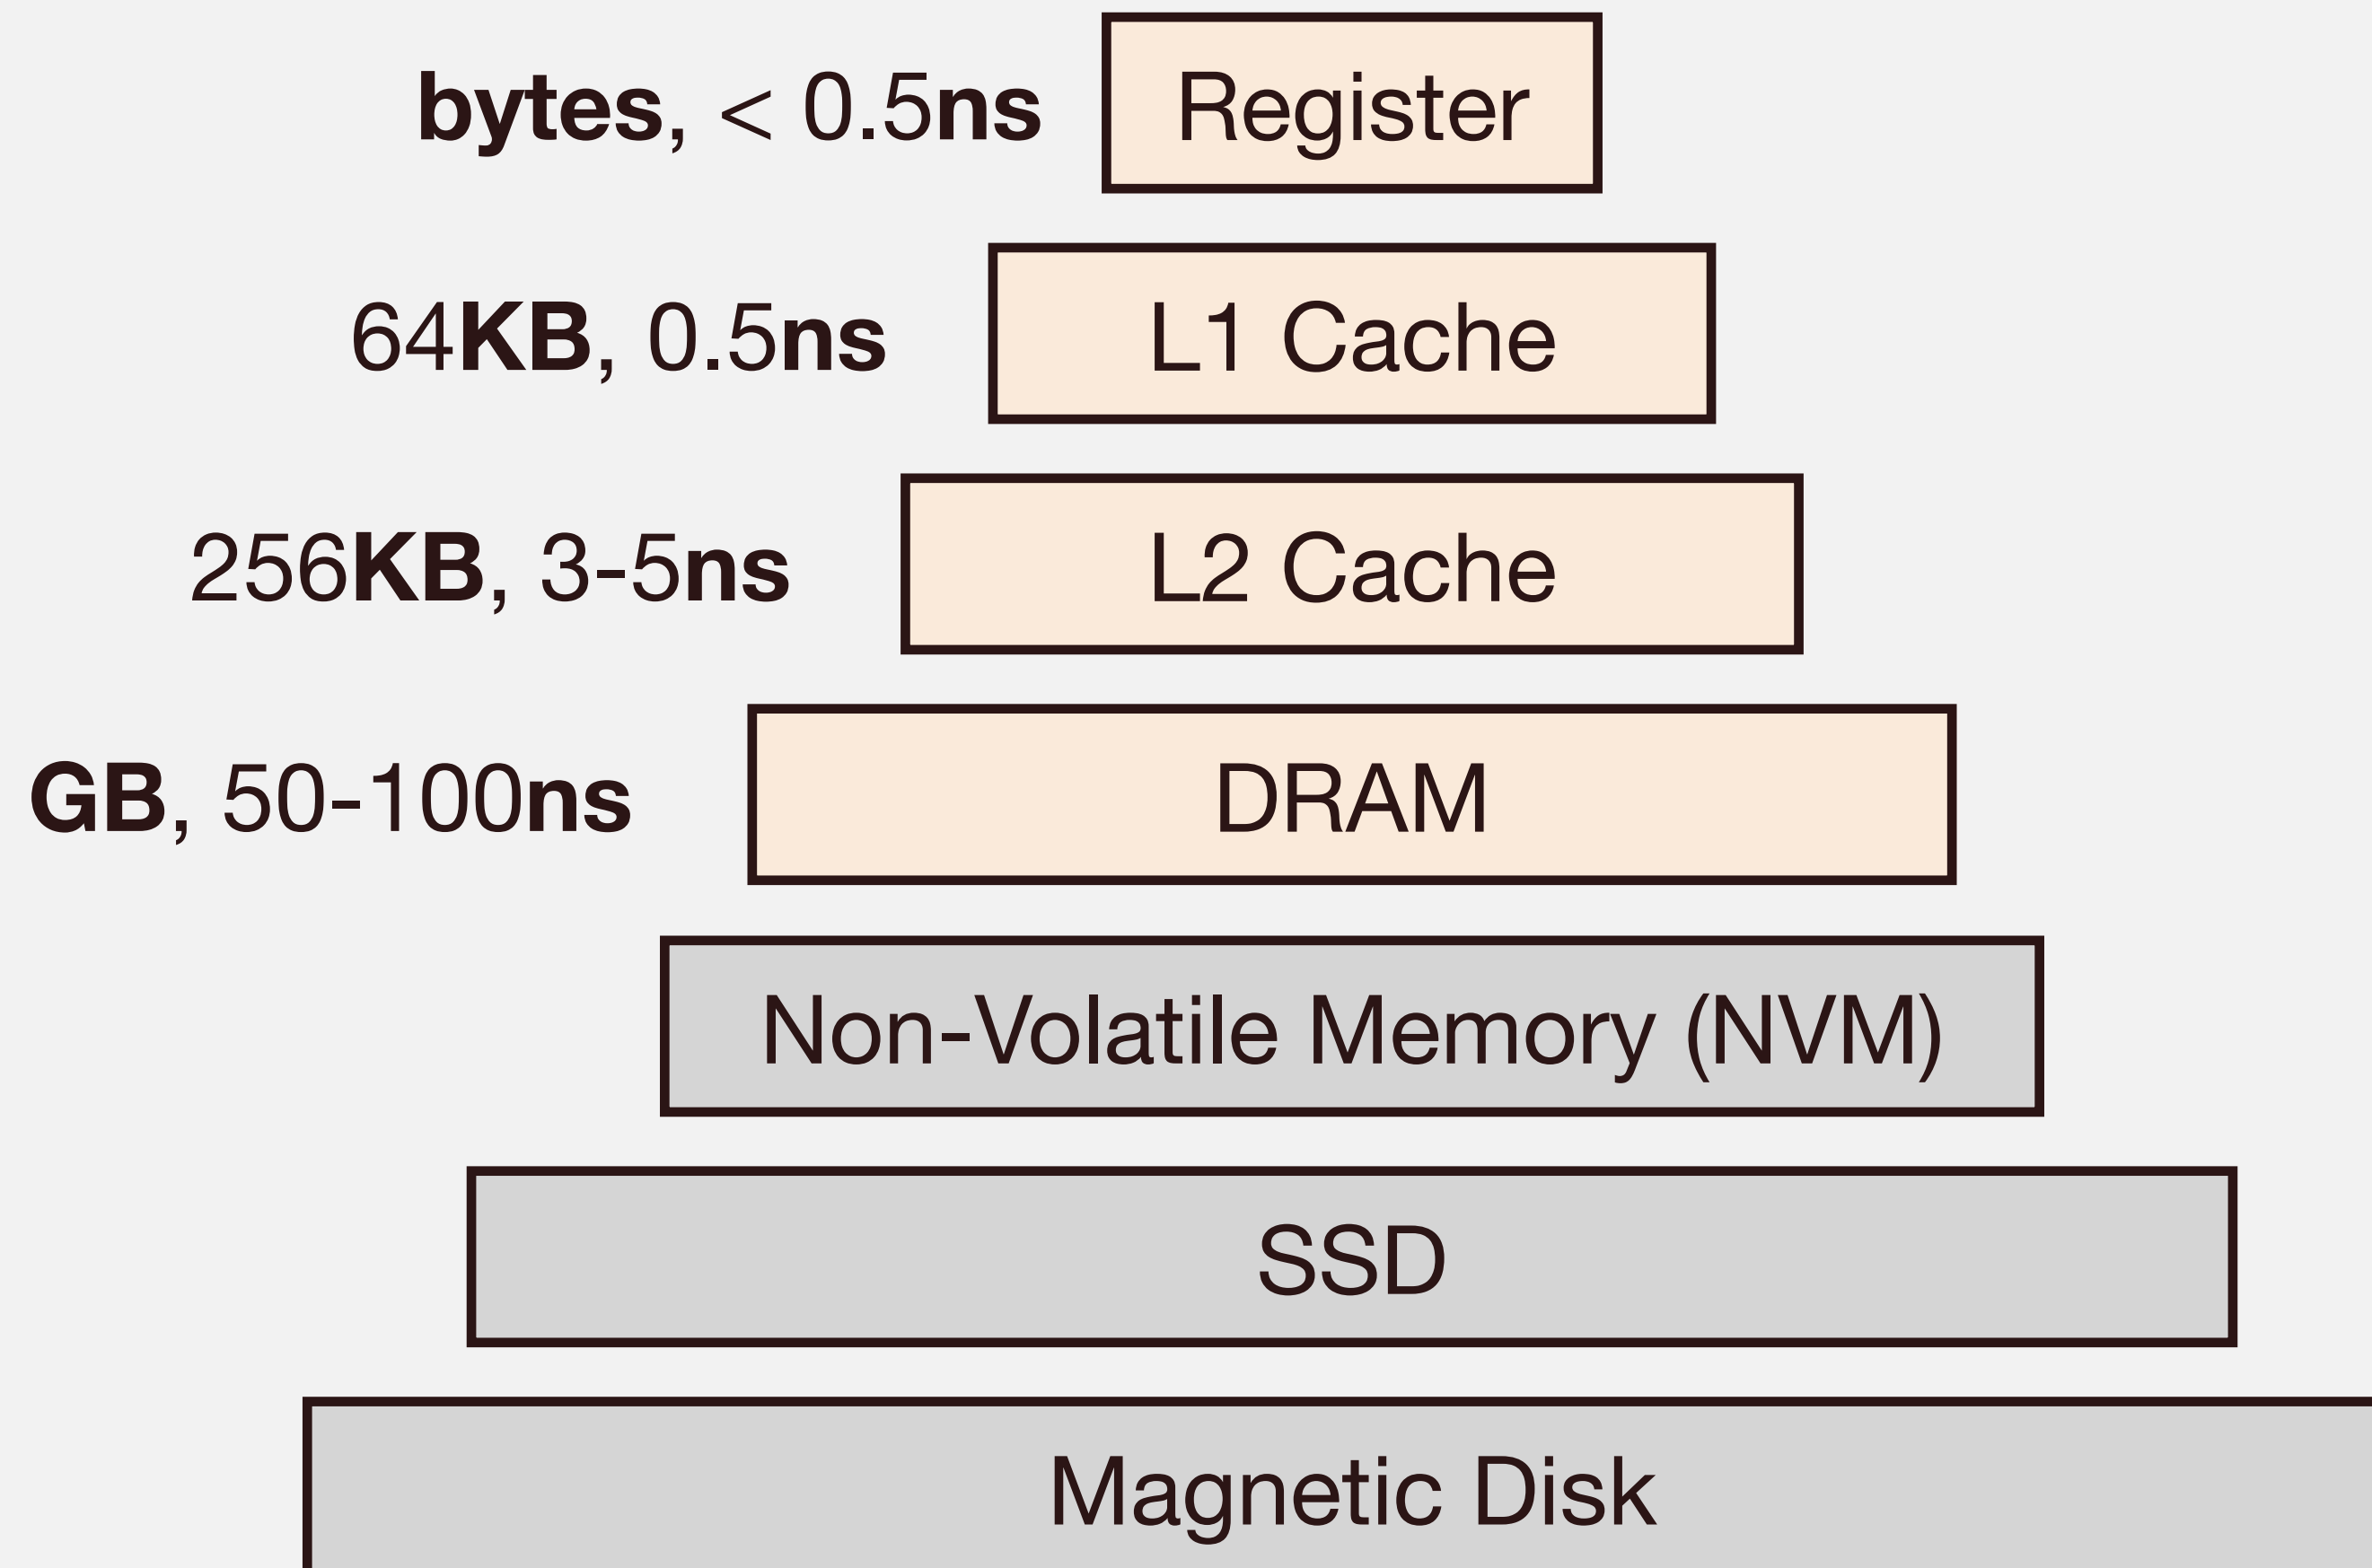
Storage Hierarchy



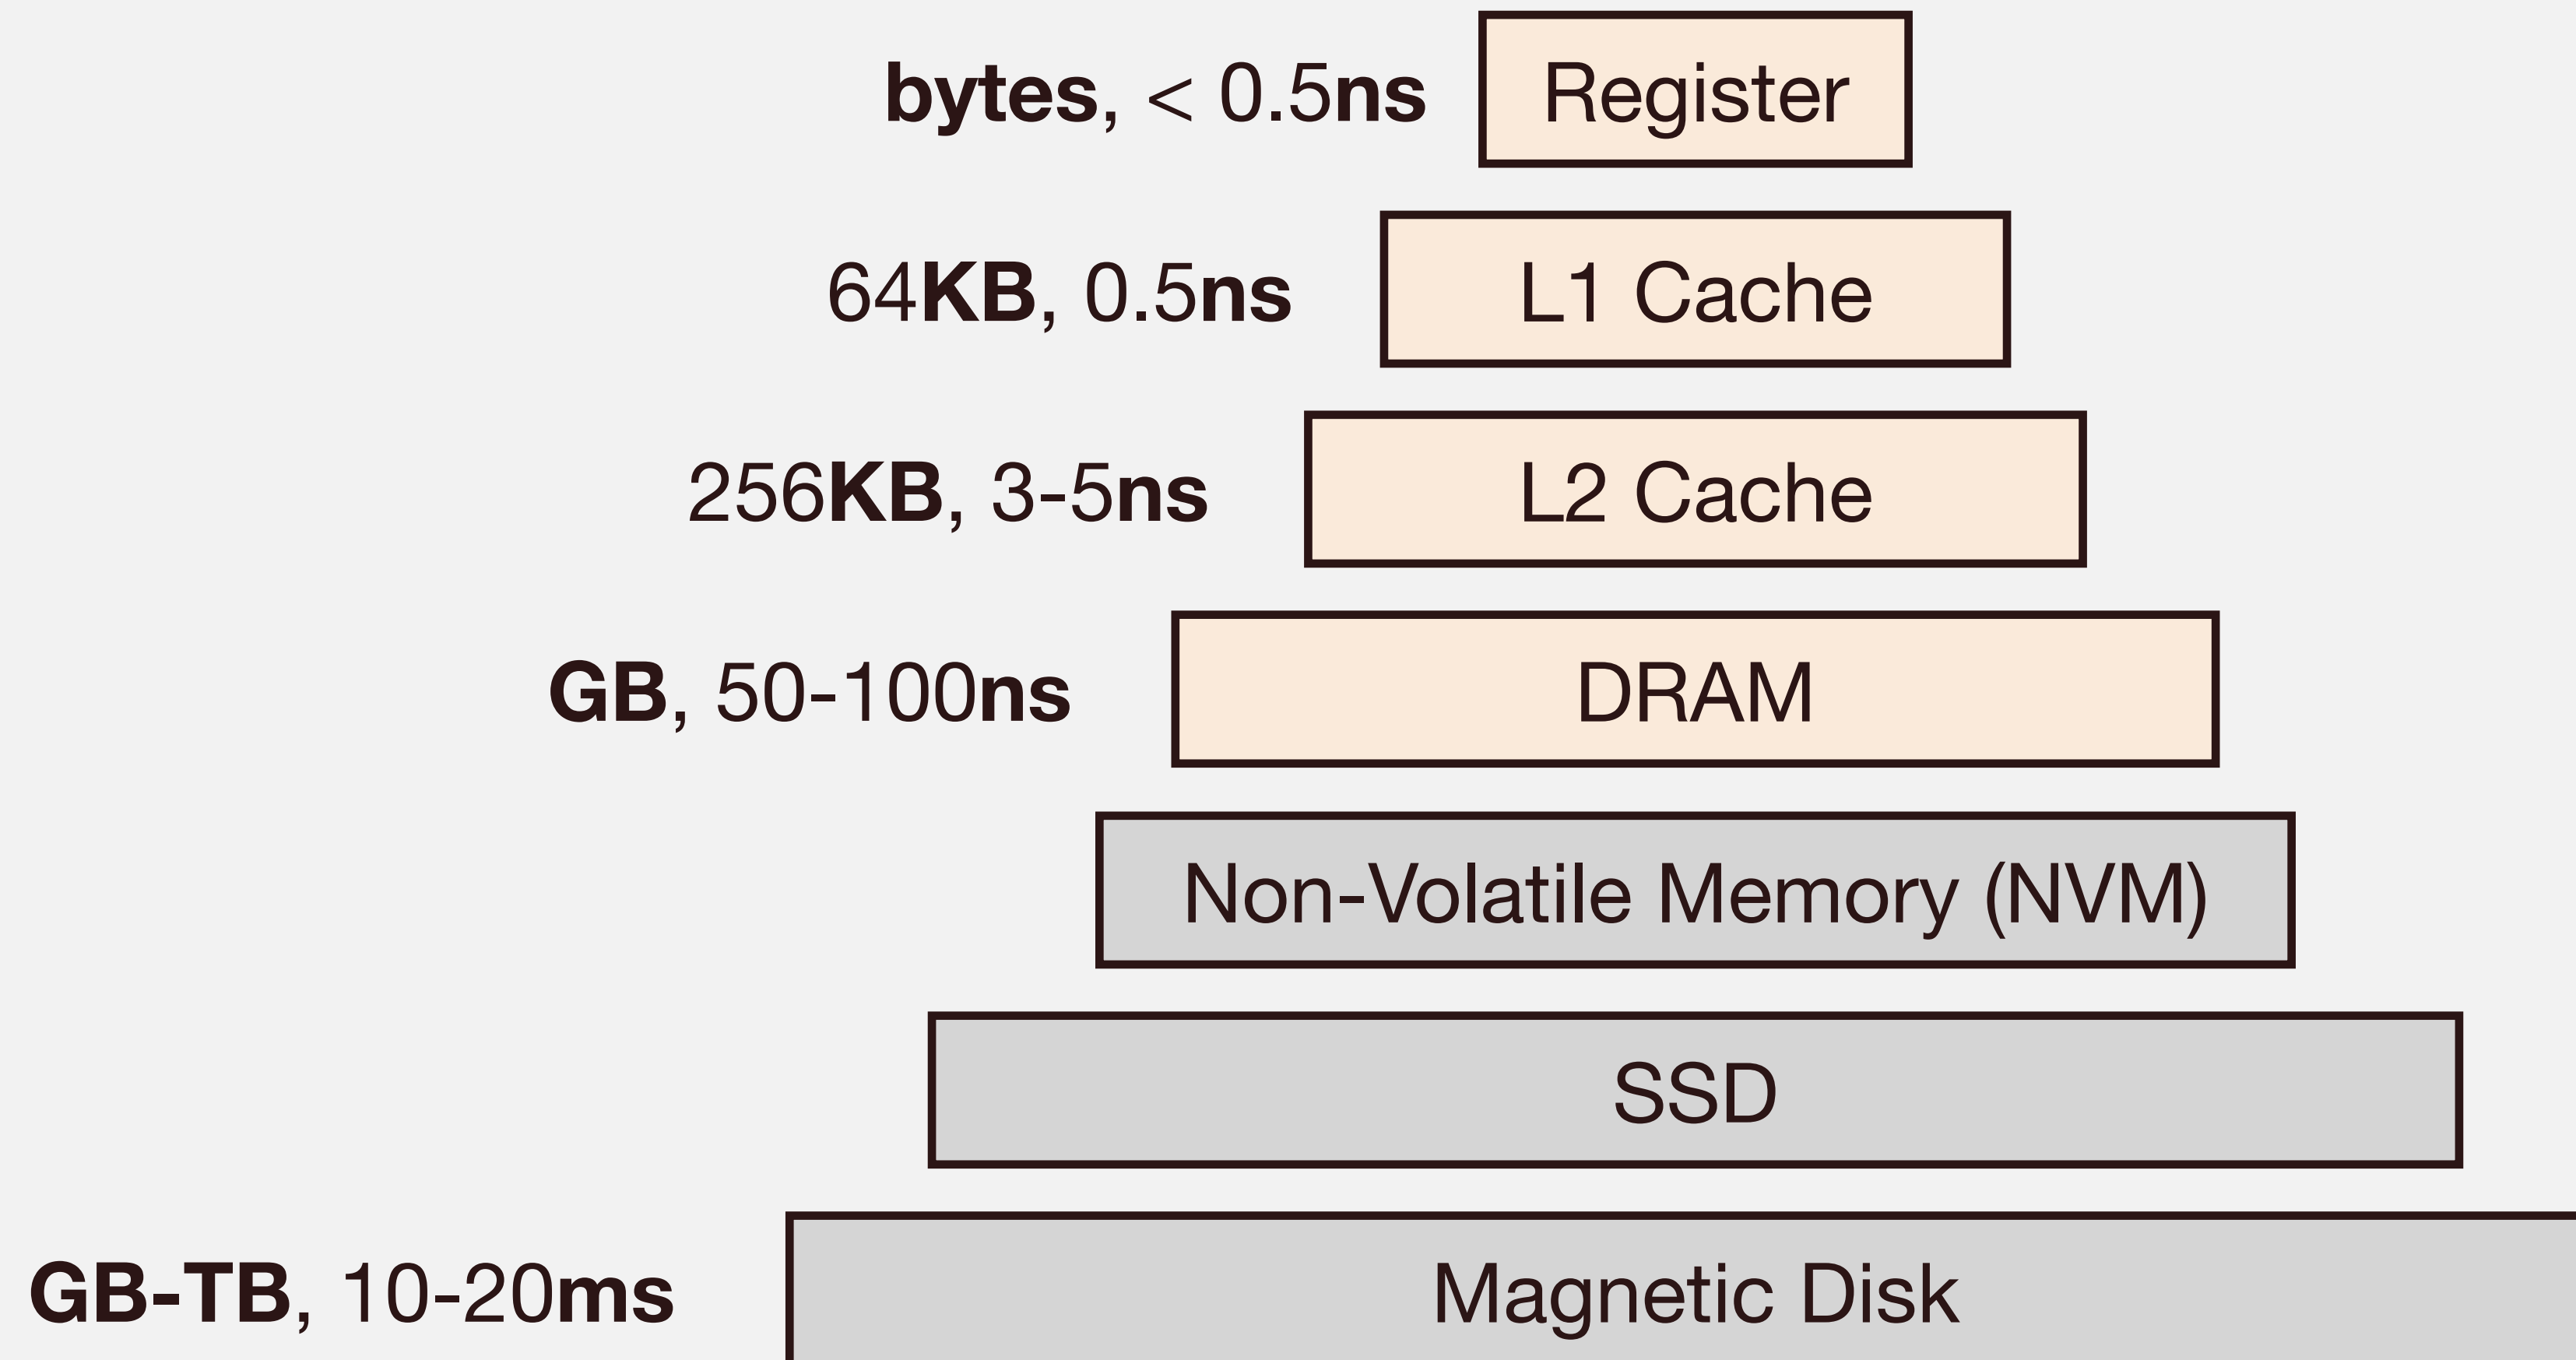
Storage Hierarchy



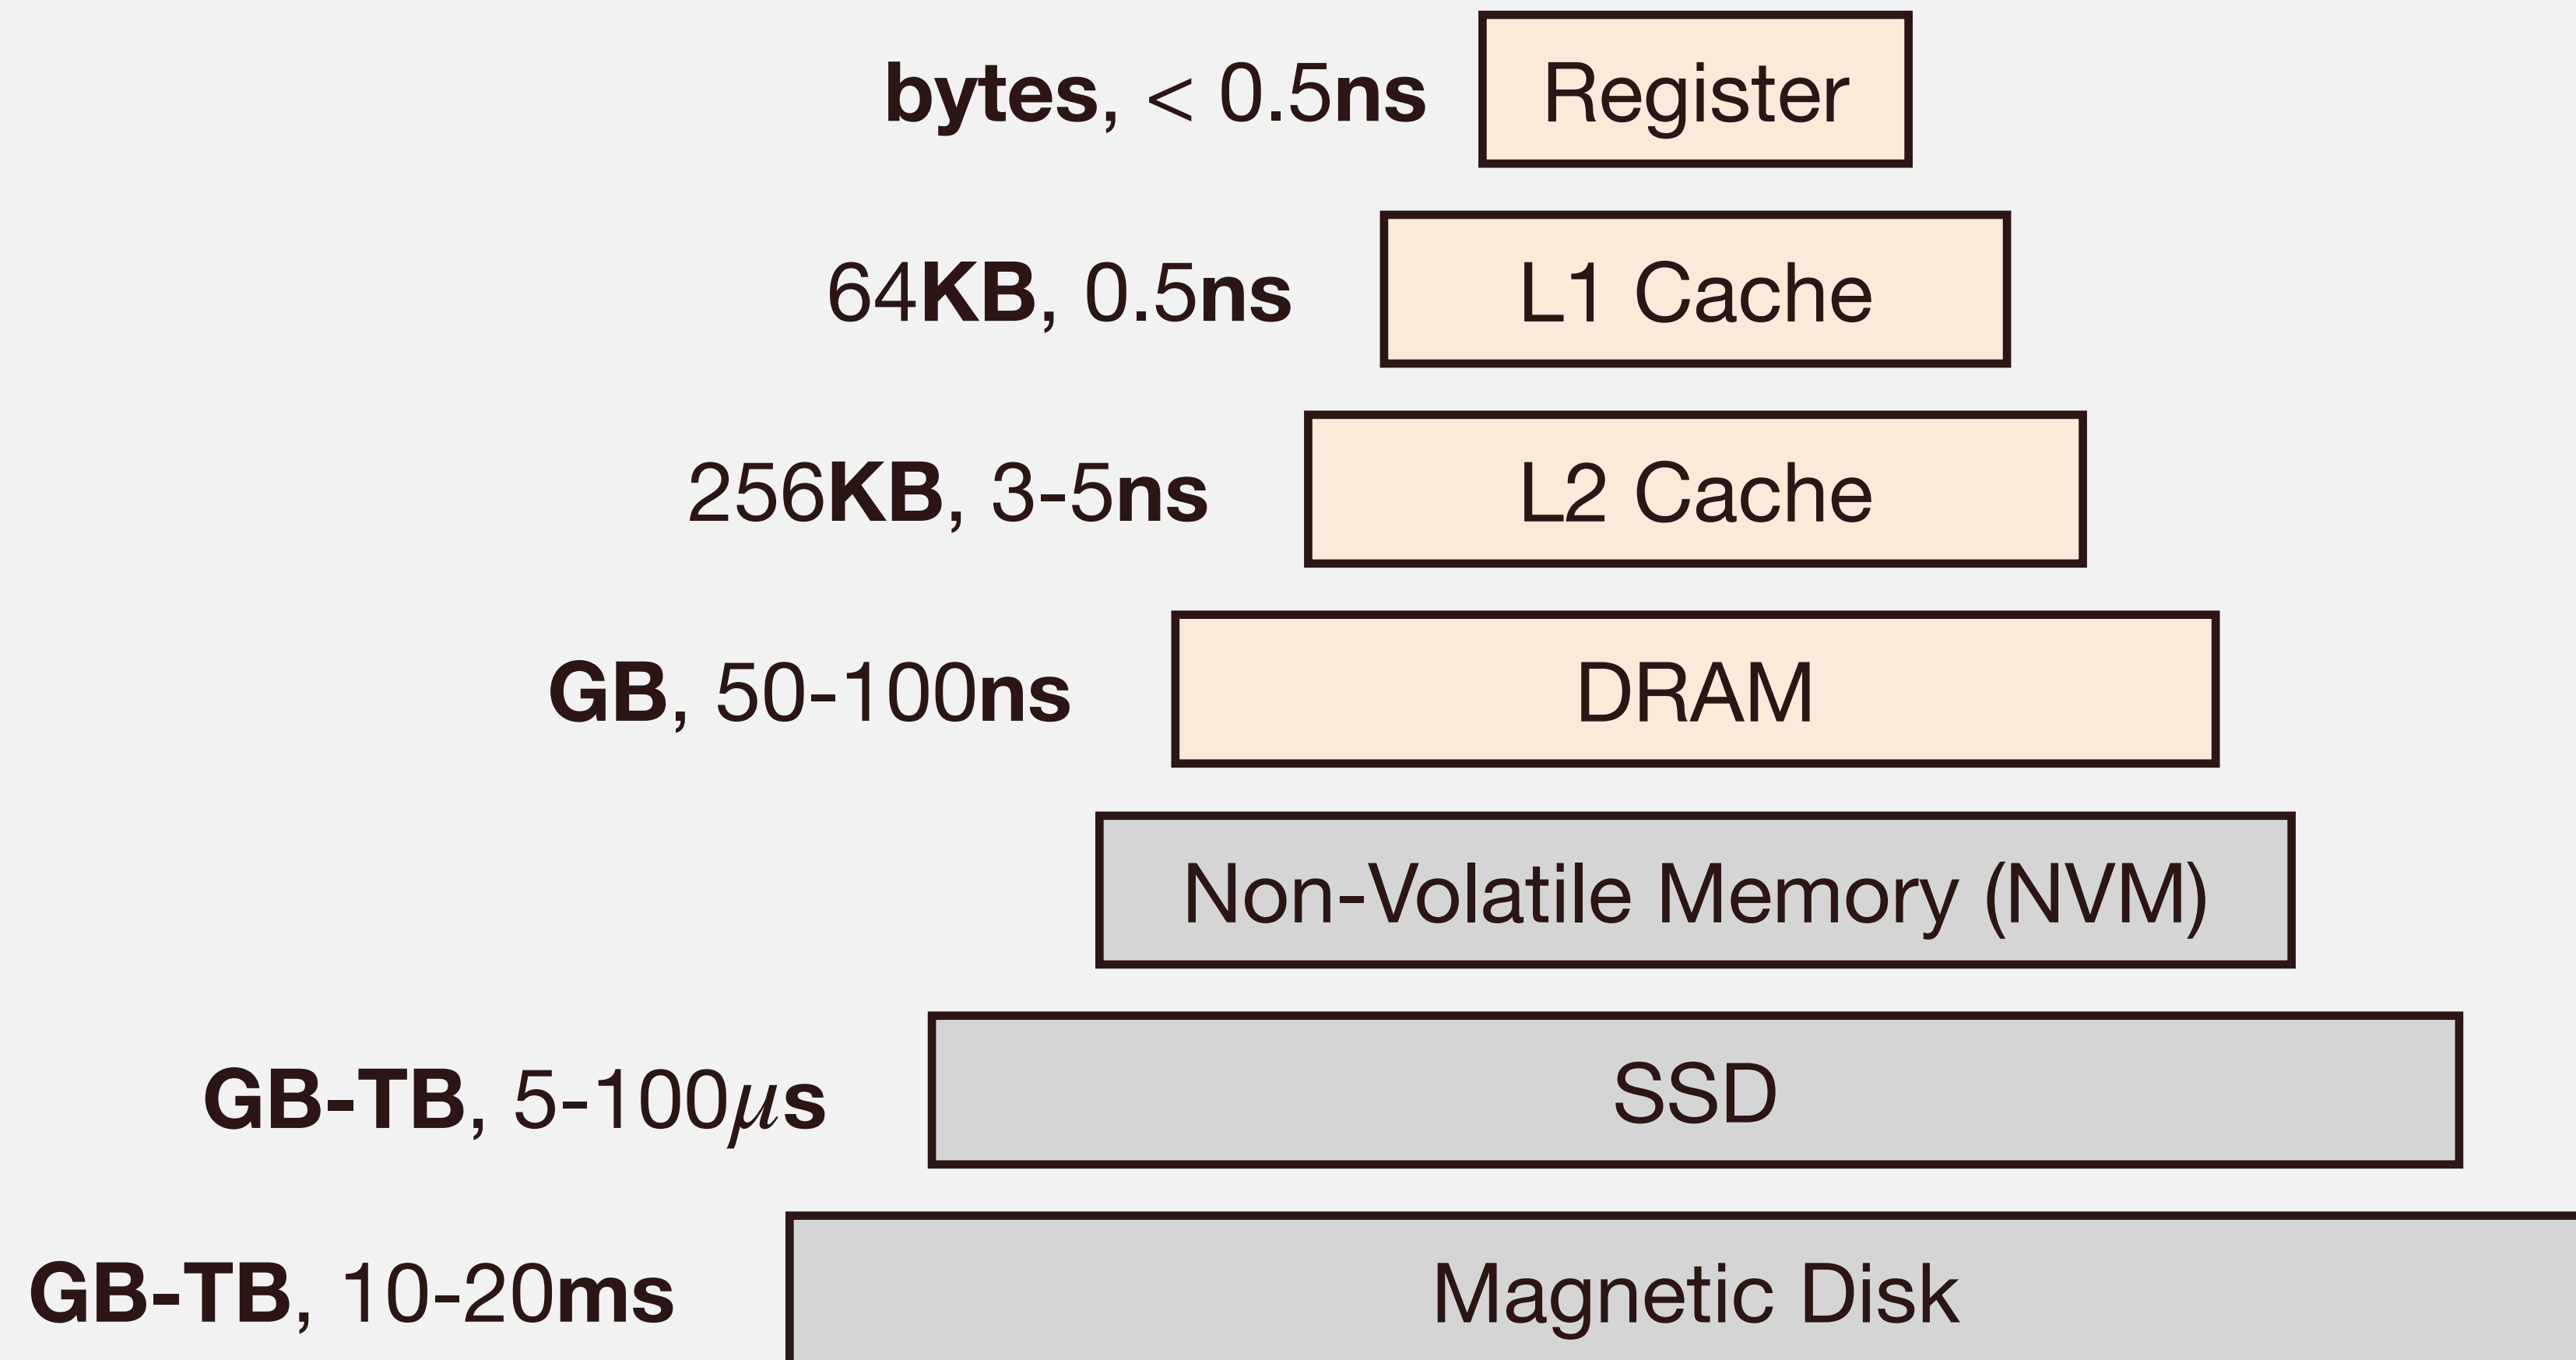
Storage Hierarchy



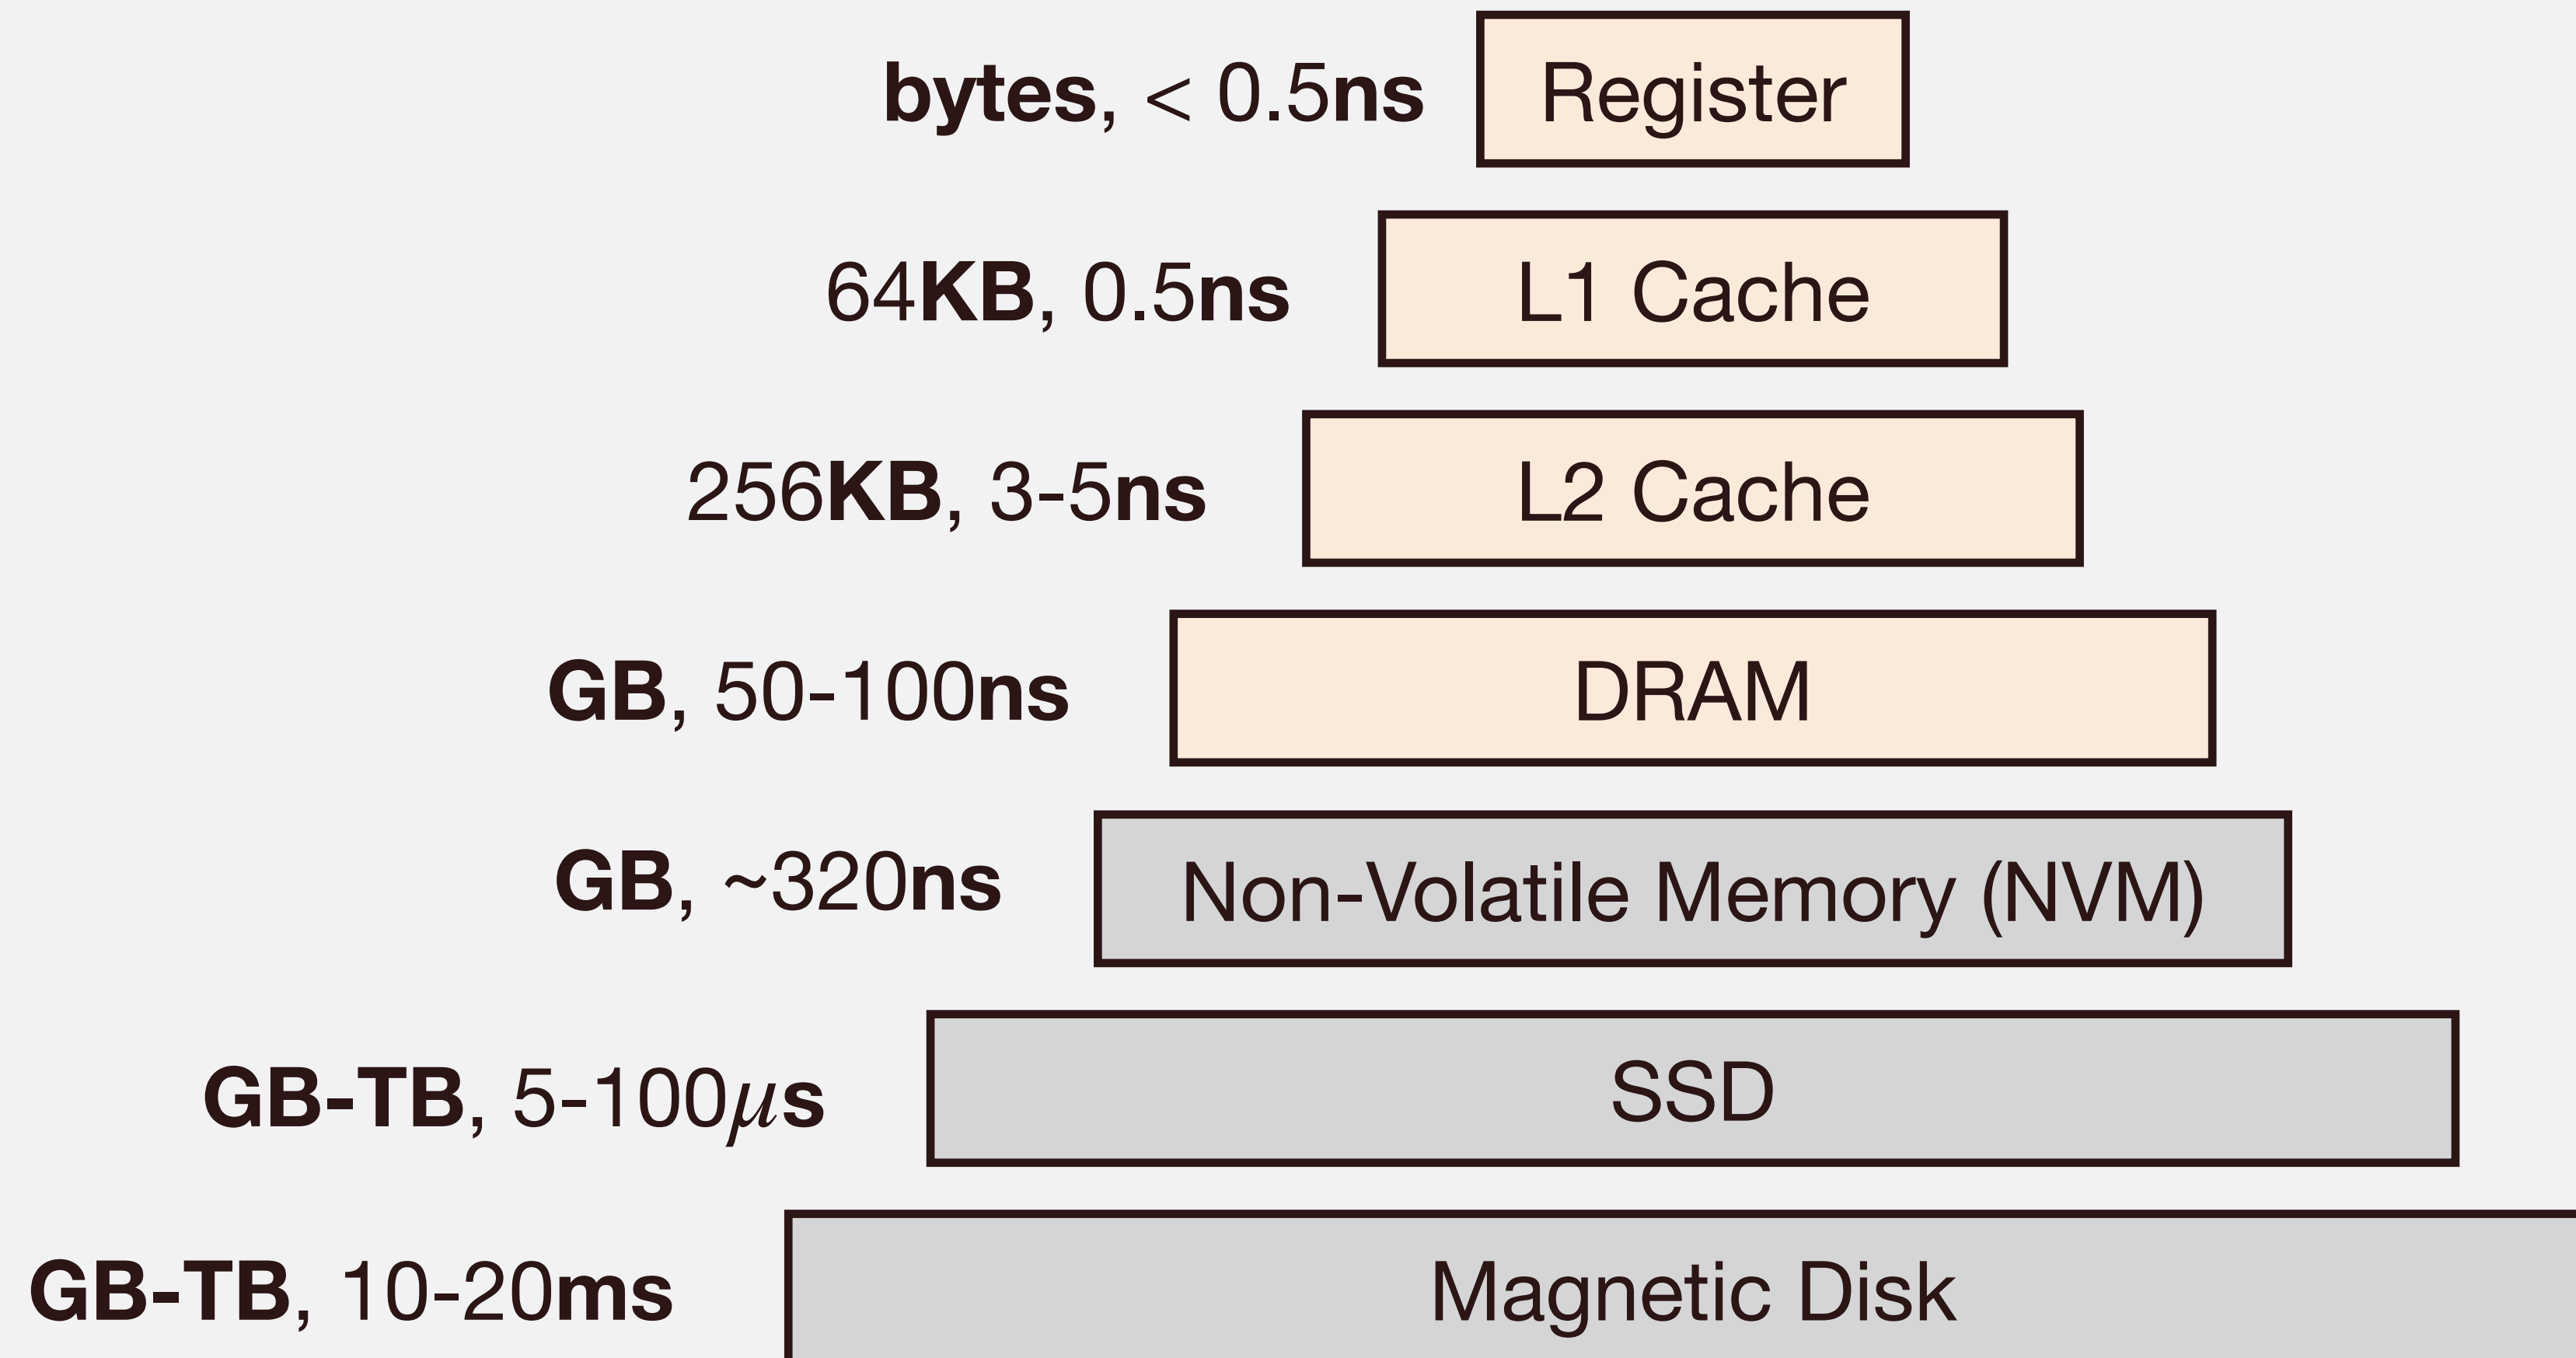
Storage Hierarchy



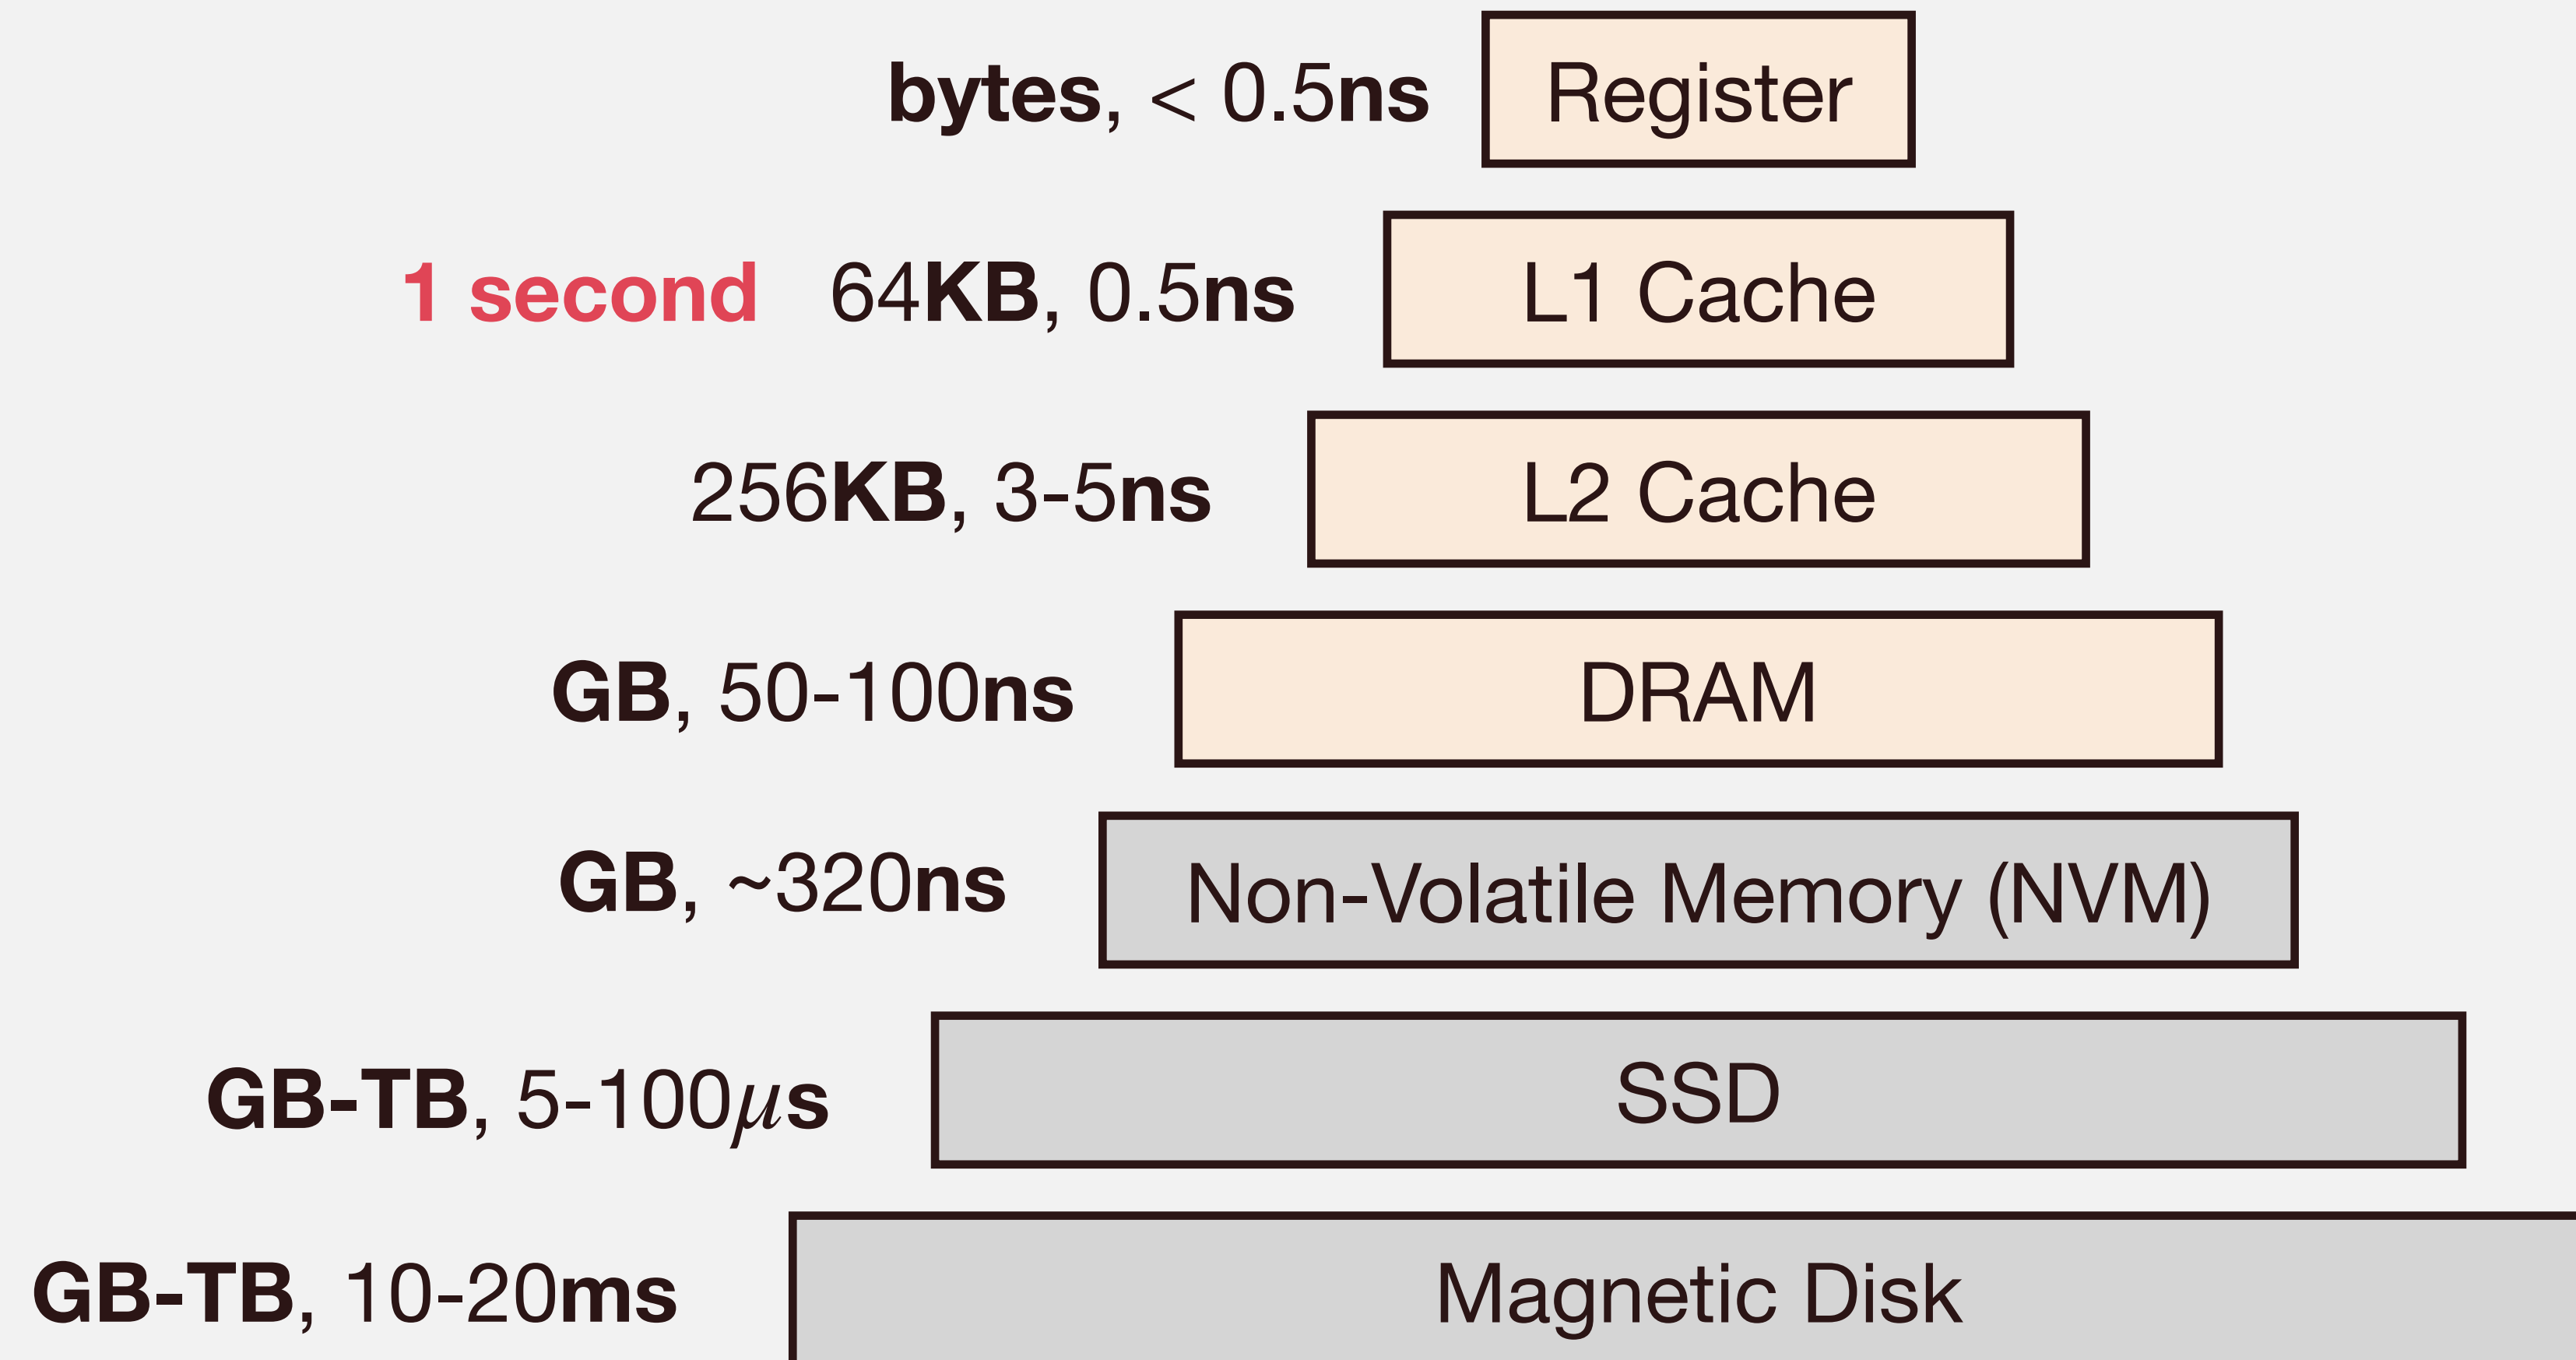
Storage Hierarchy



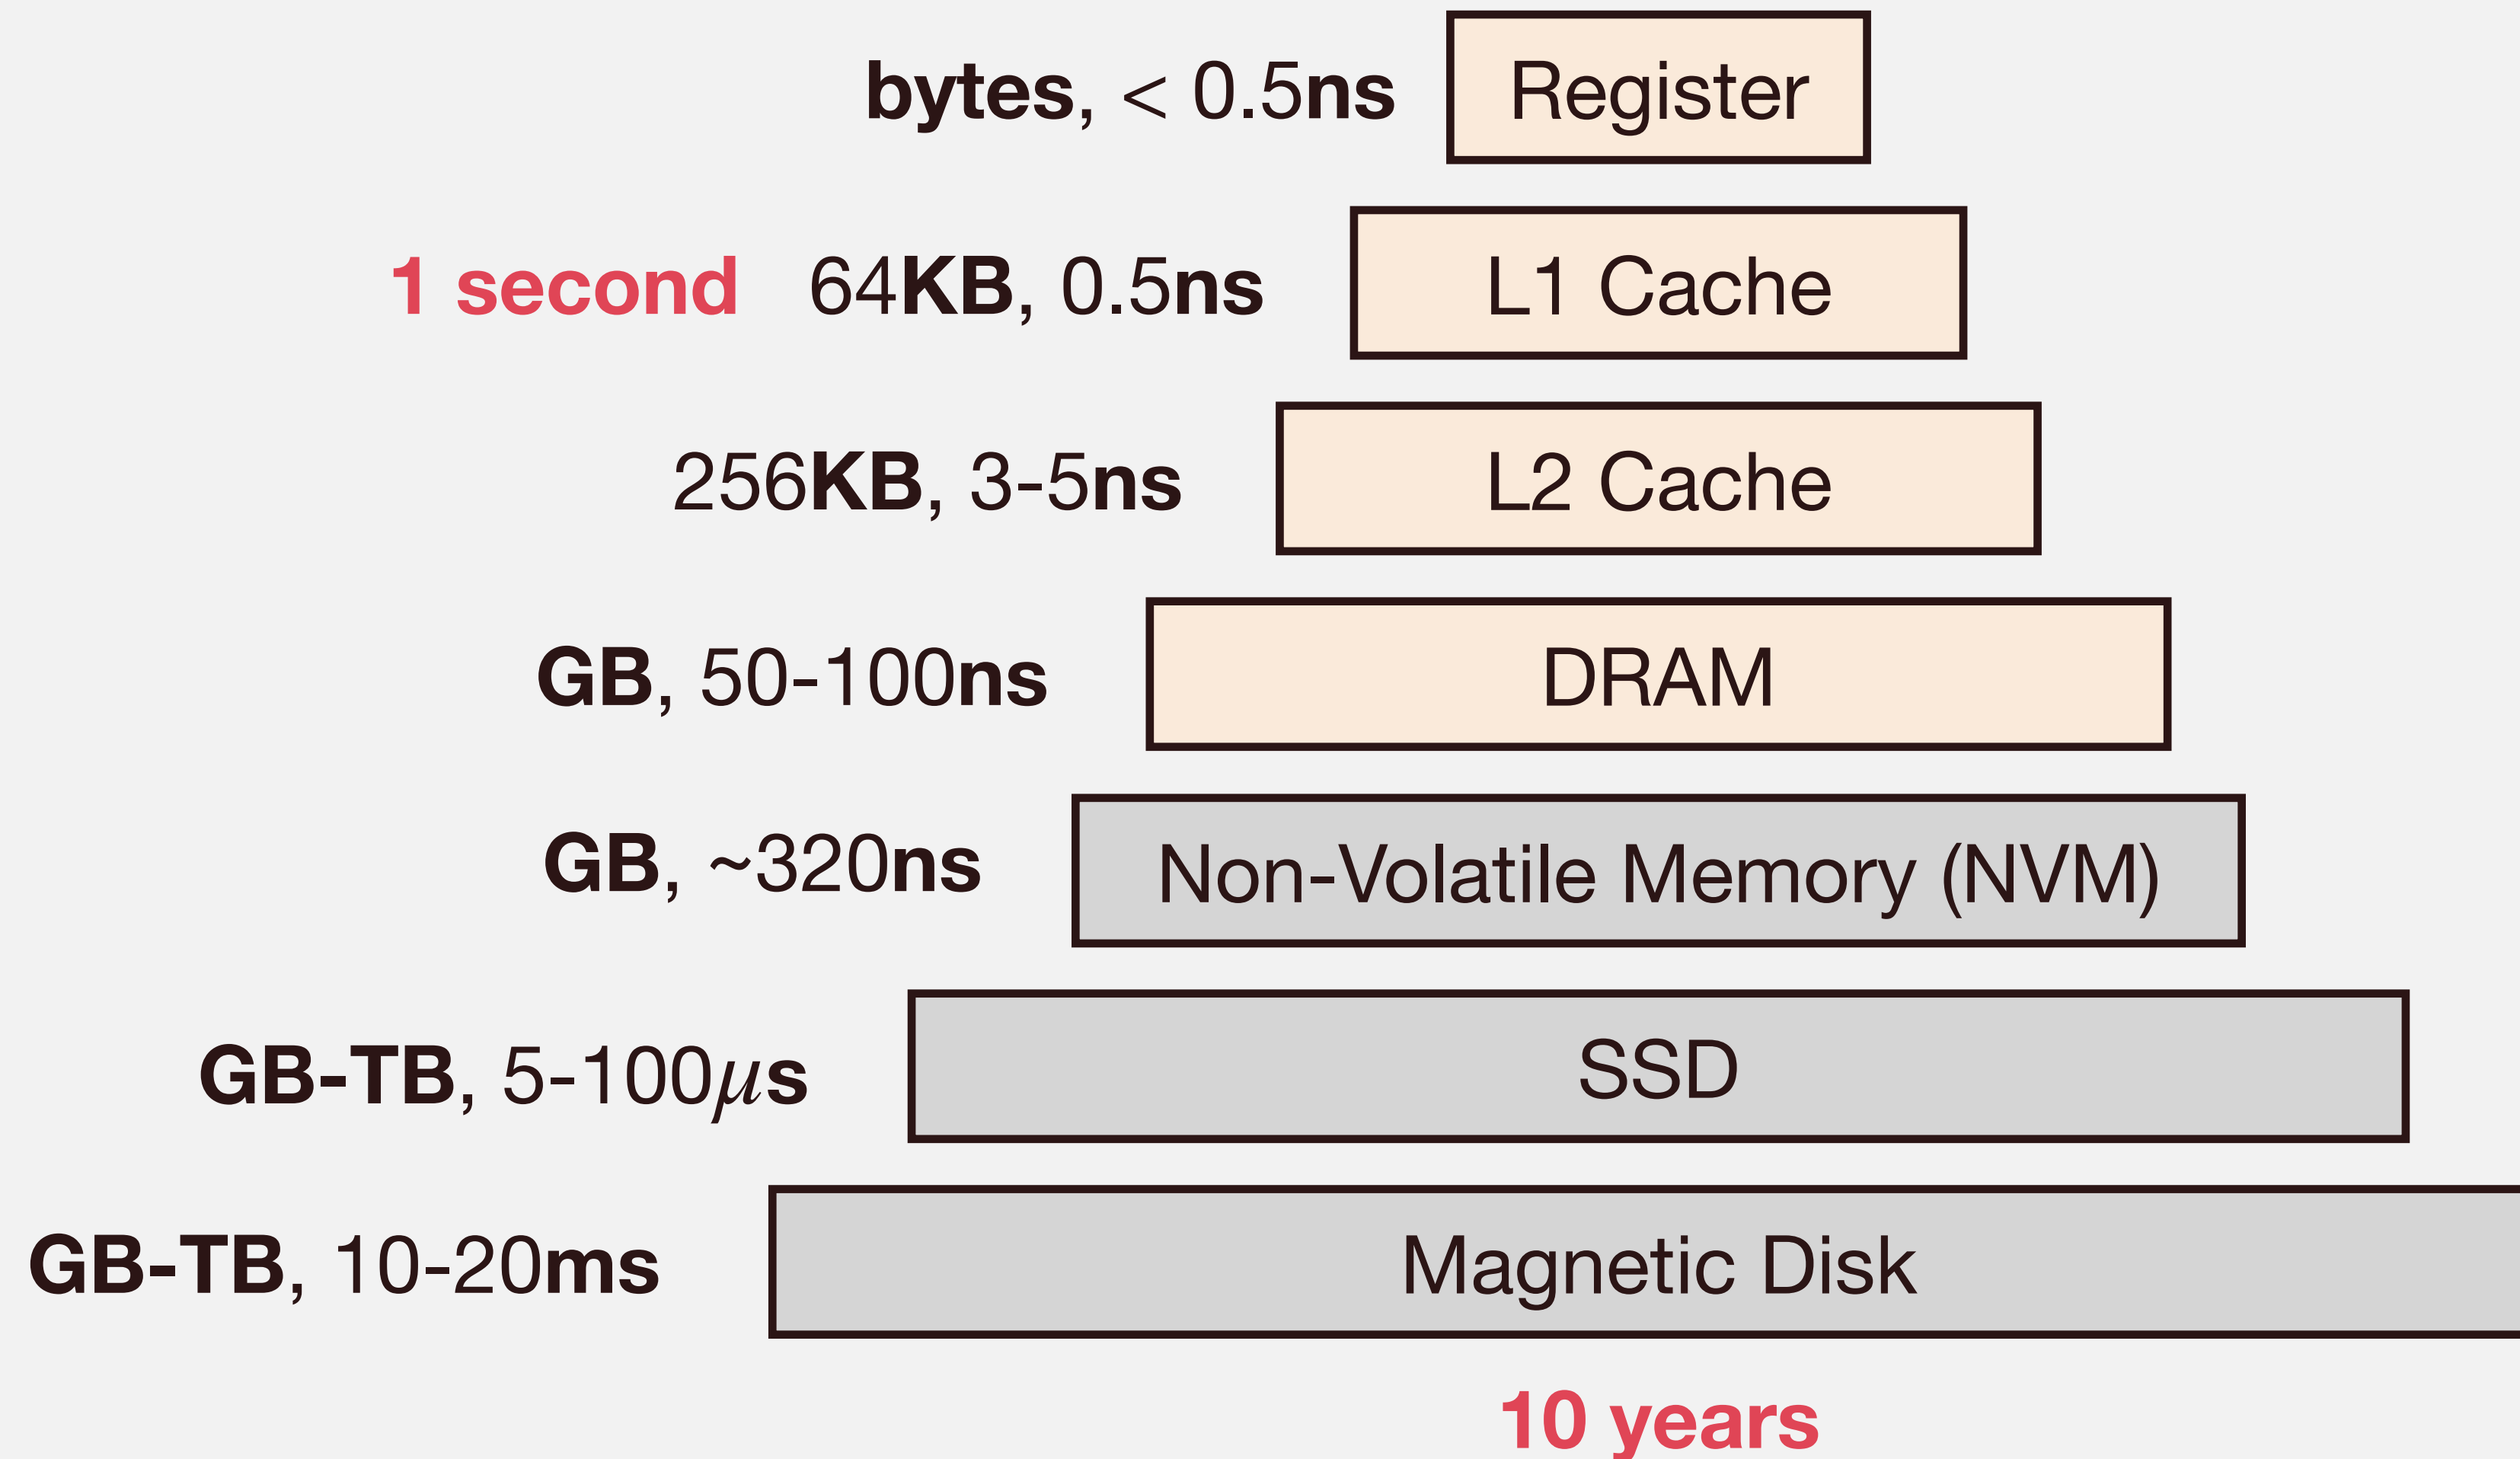
Storage Hierarchy



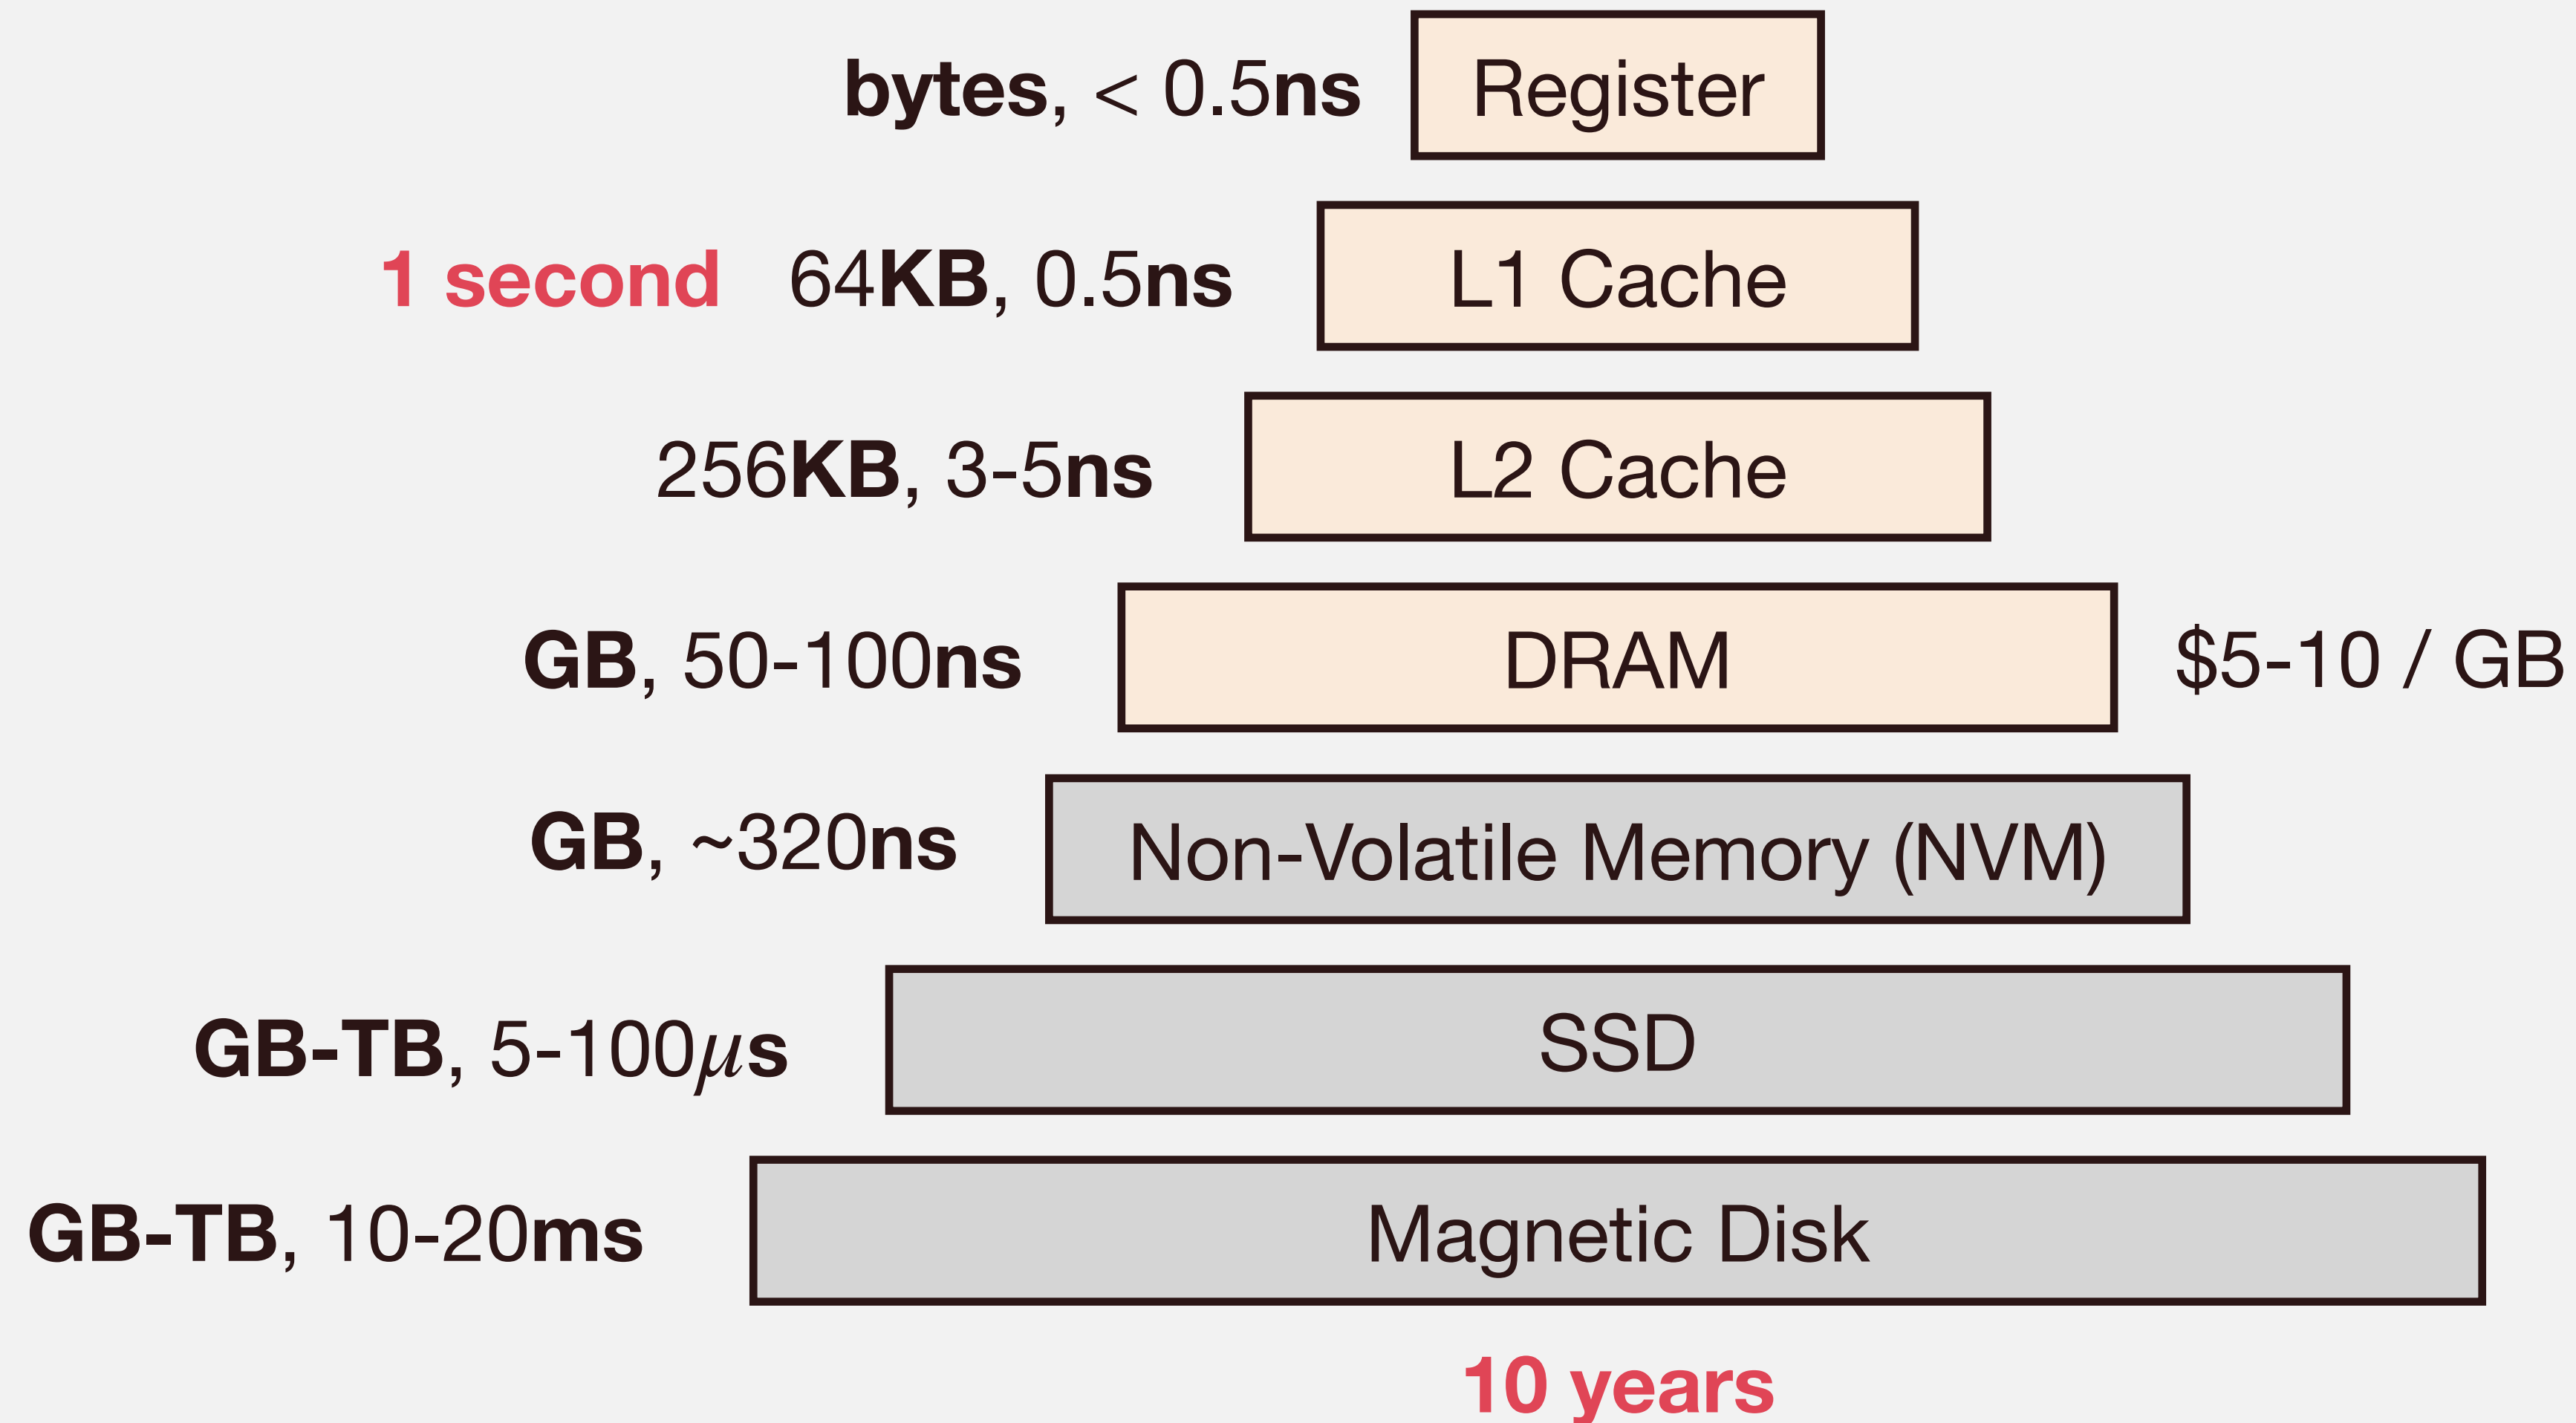
Storage Hierarchy



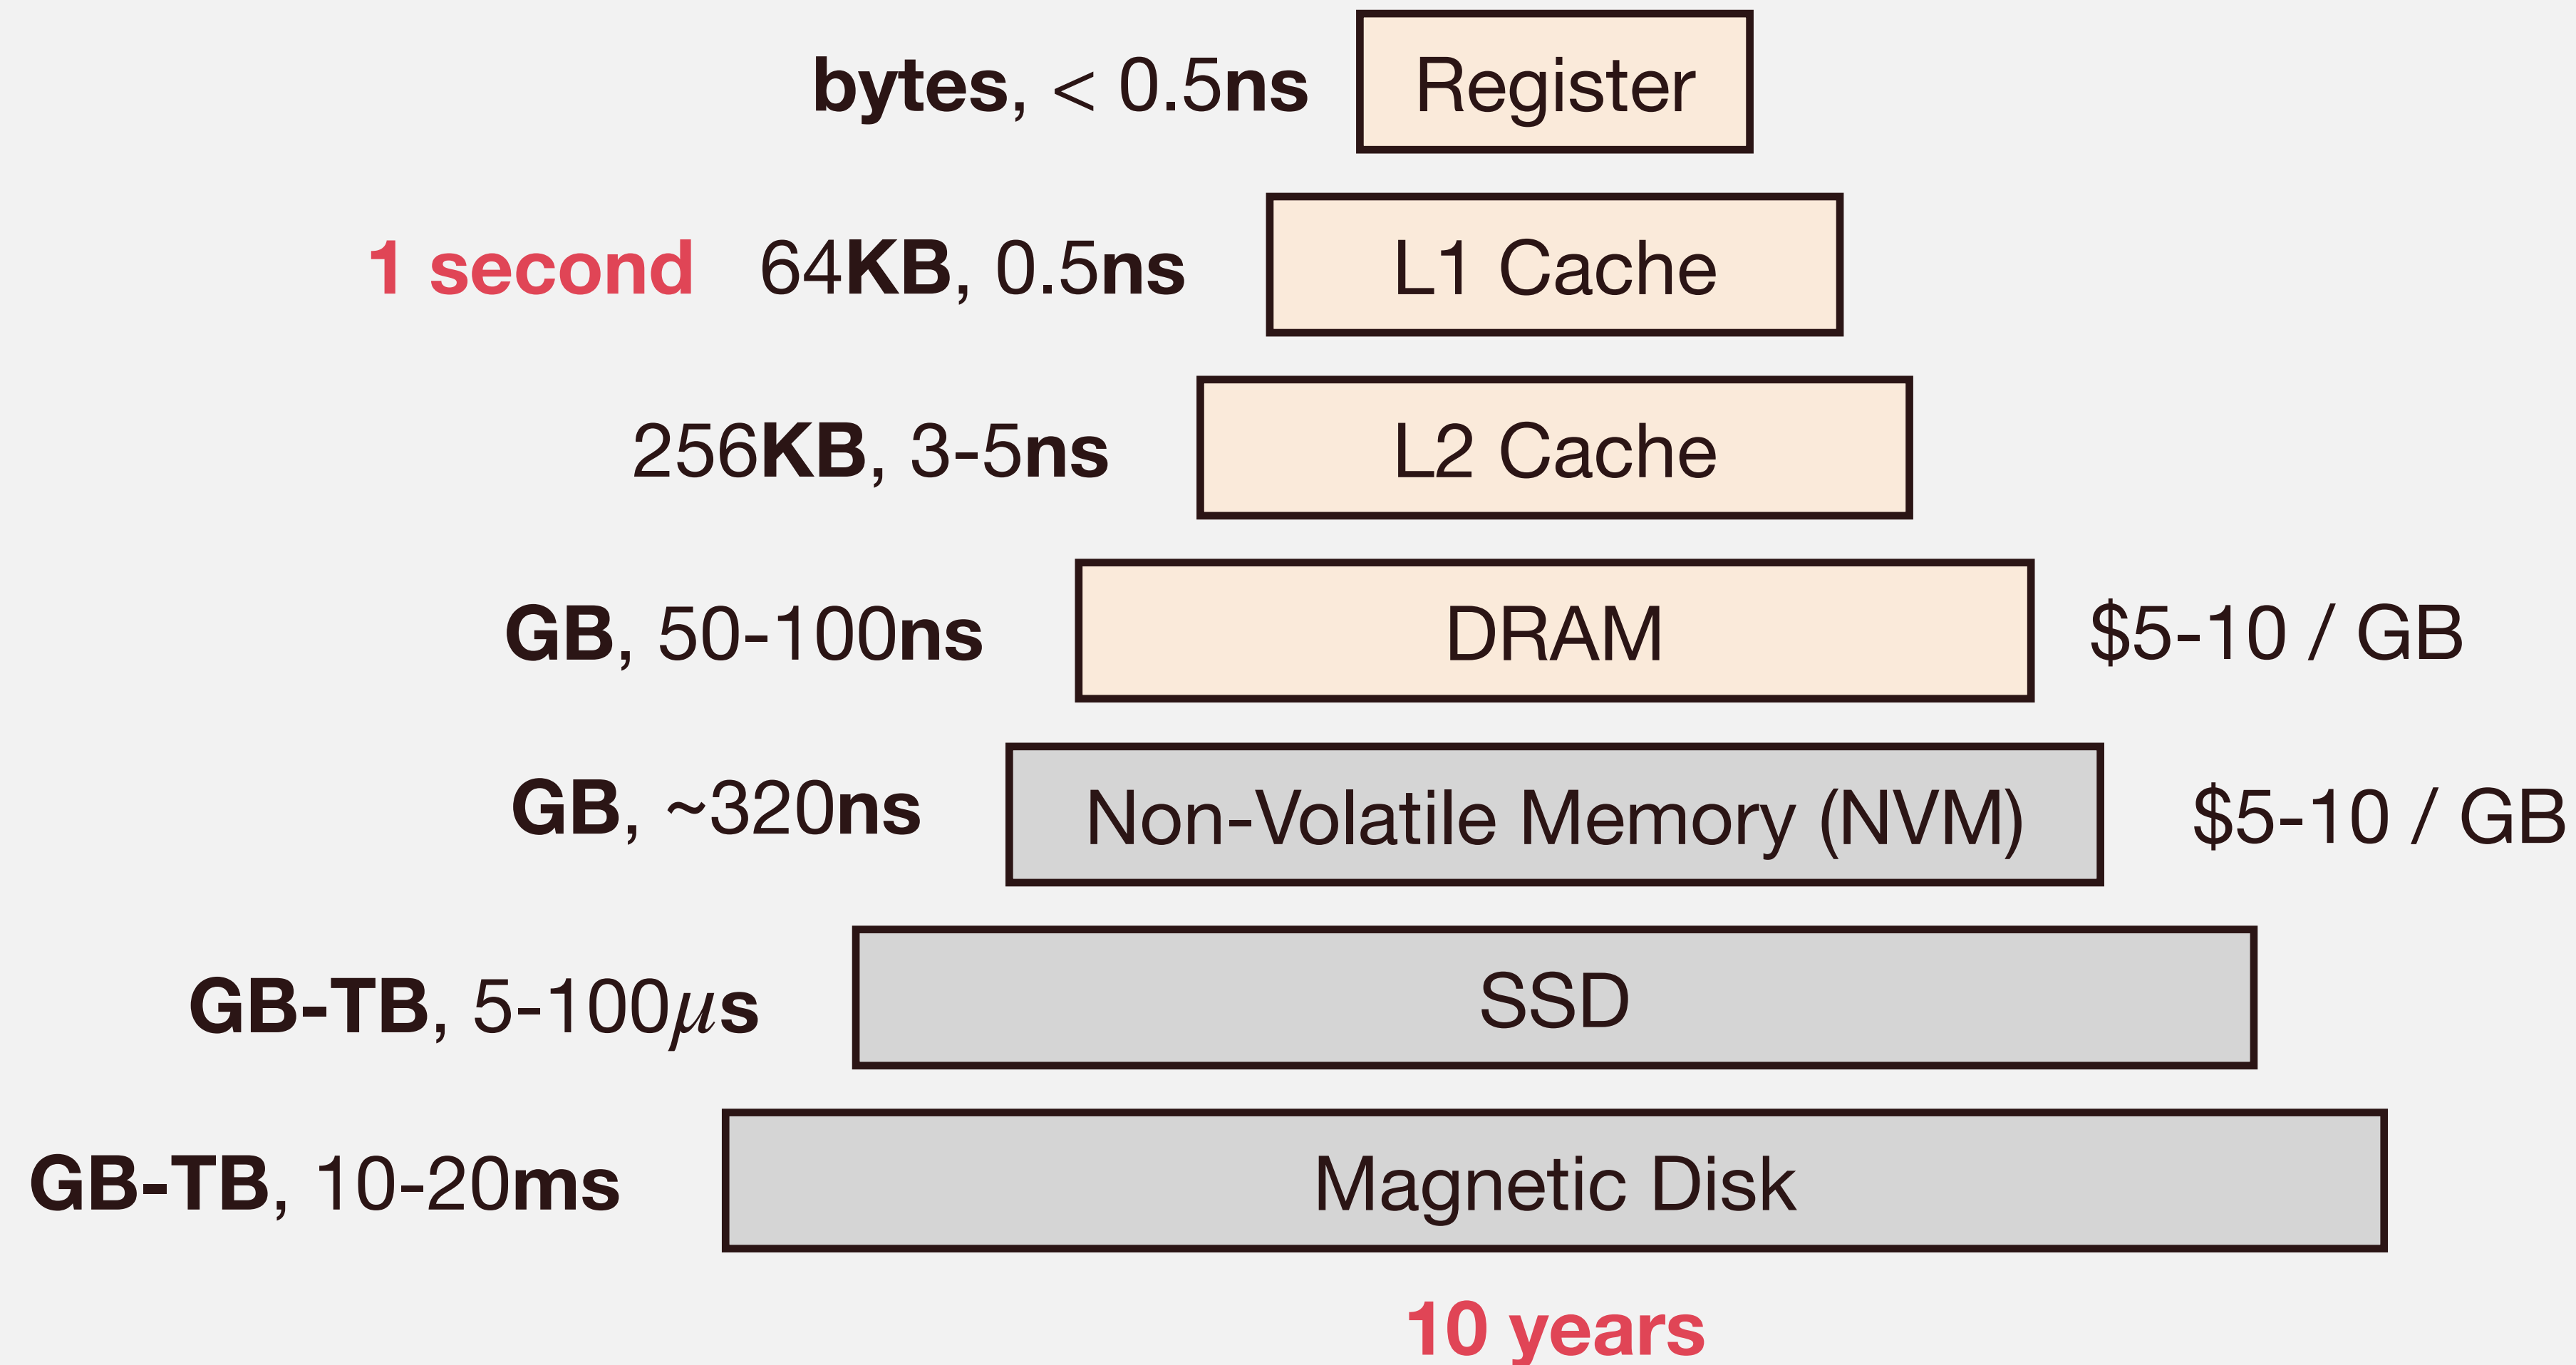
Storage Hierarchy



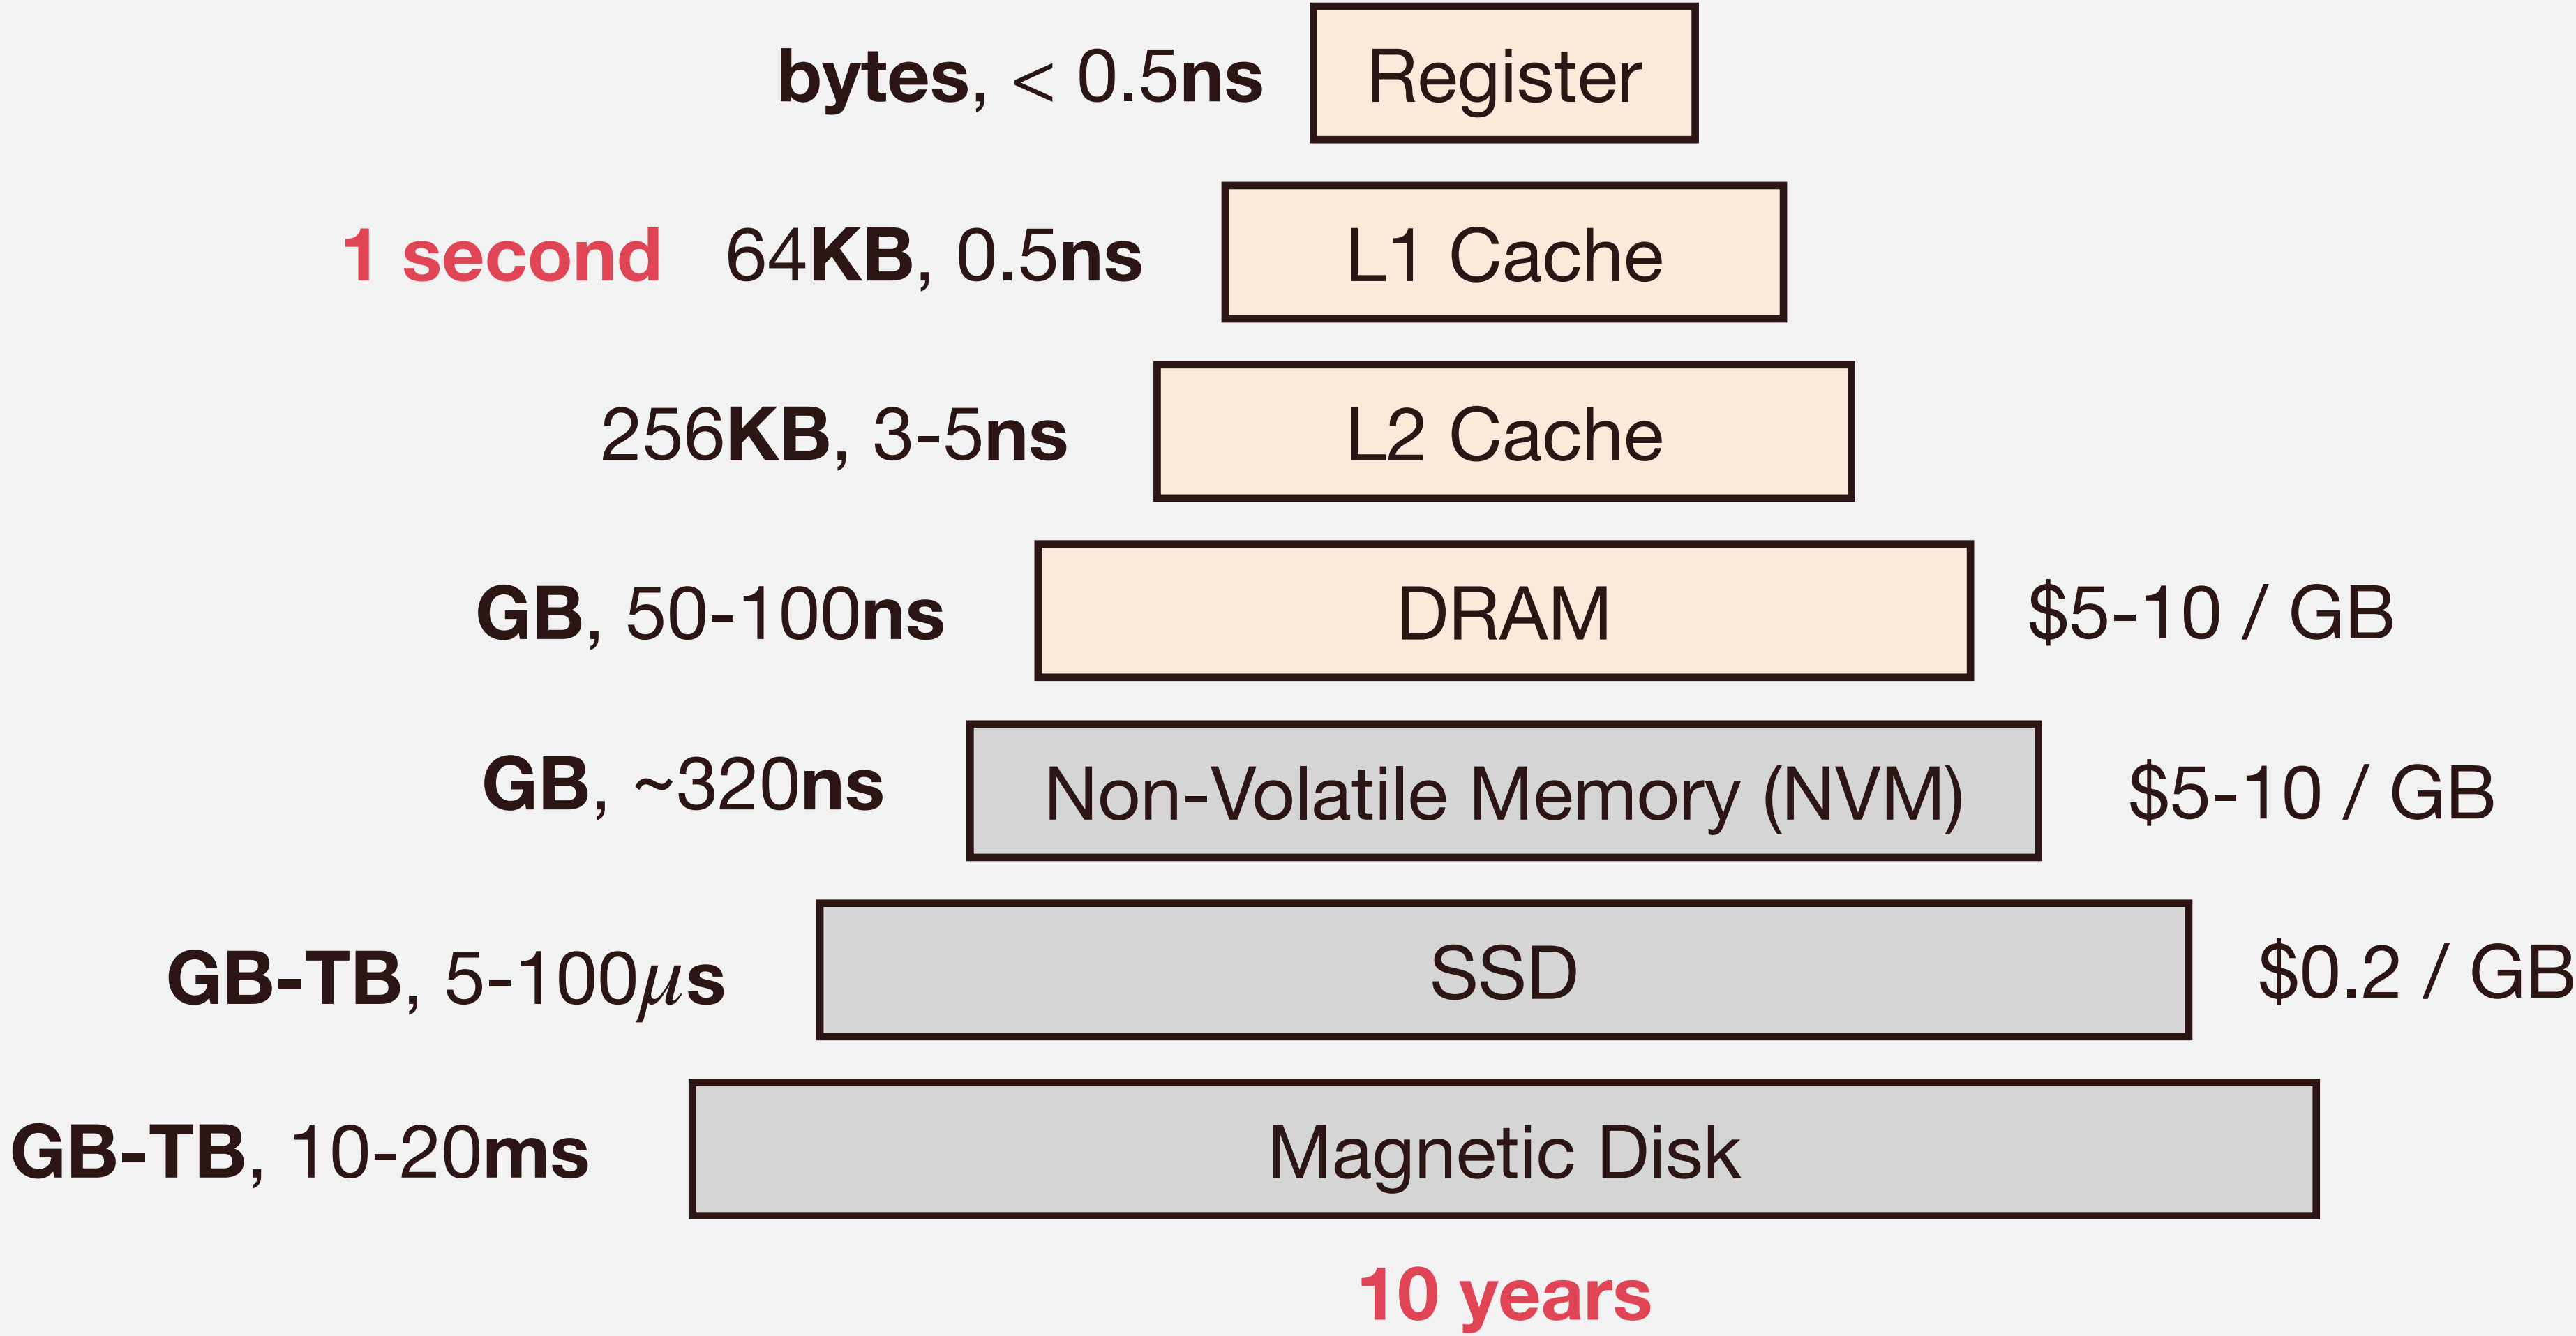
Storage Hierarchy



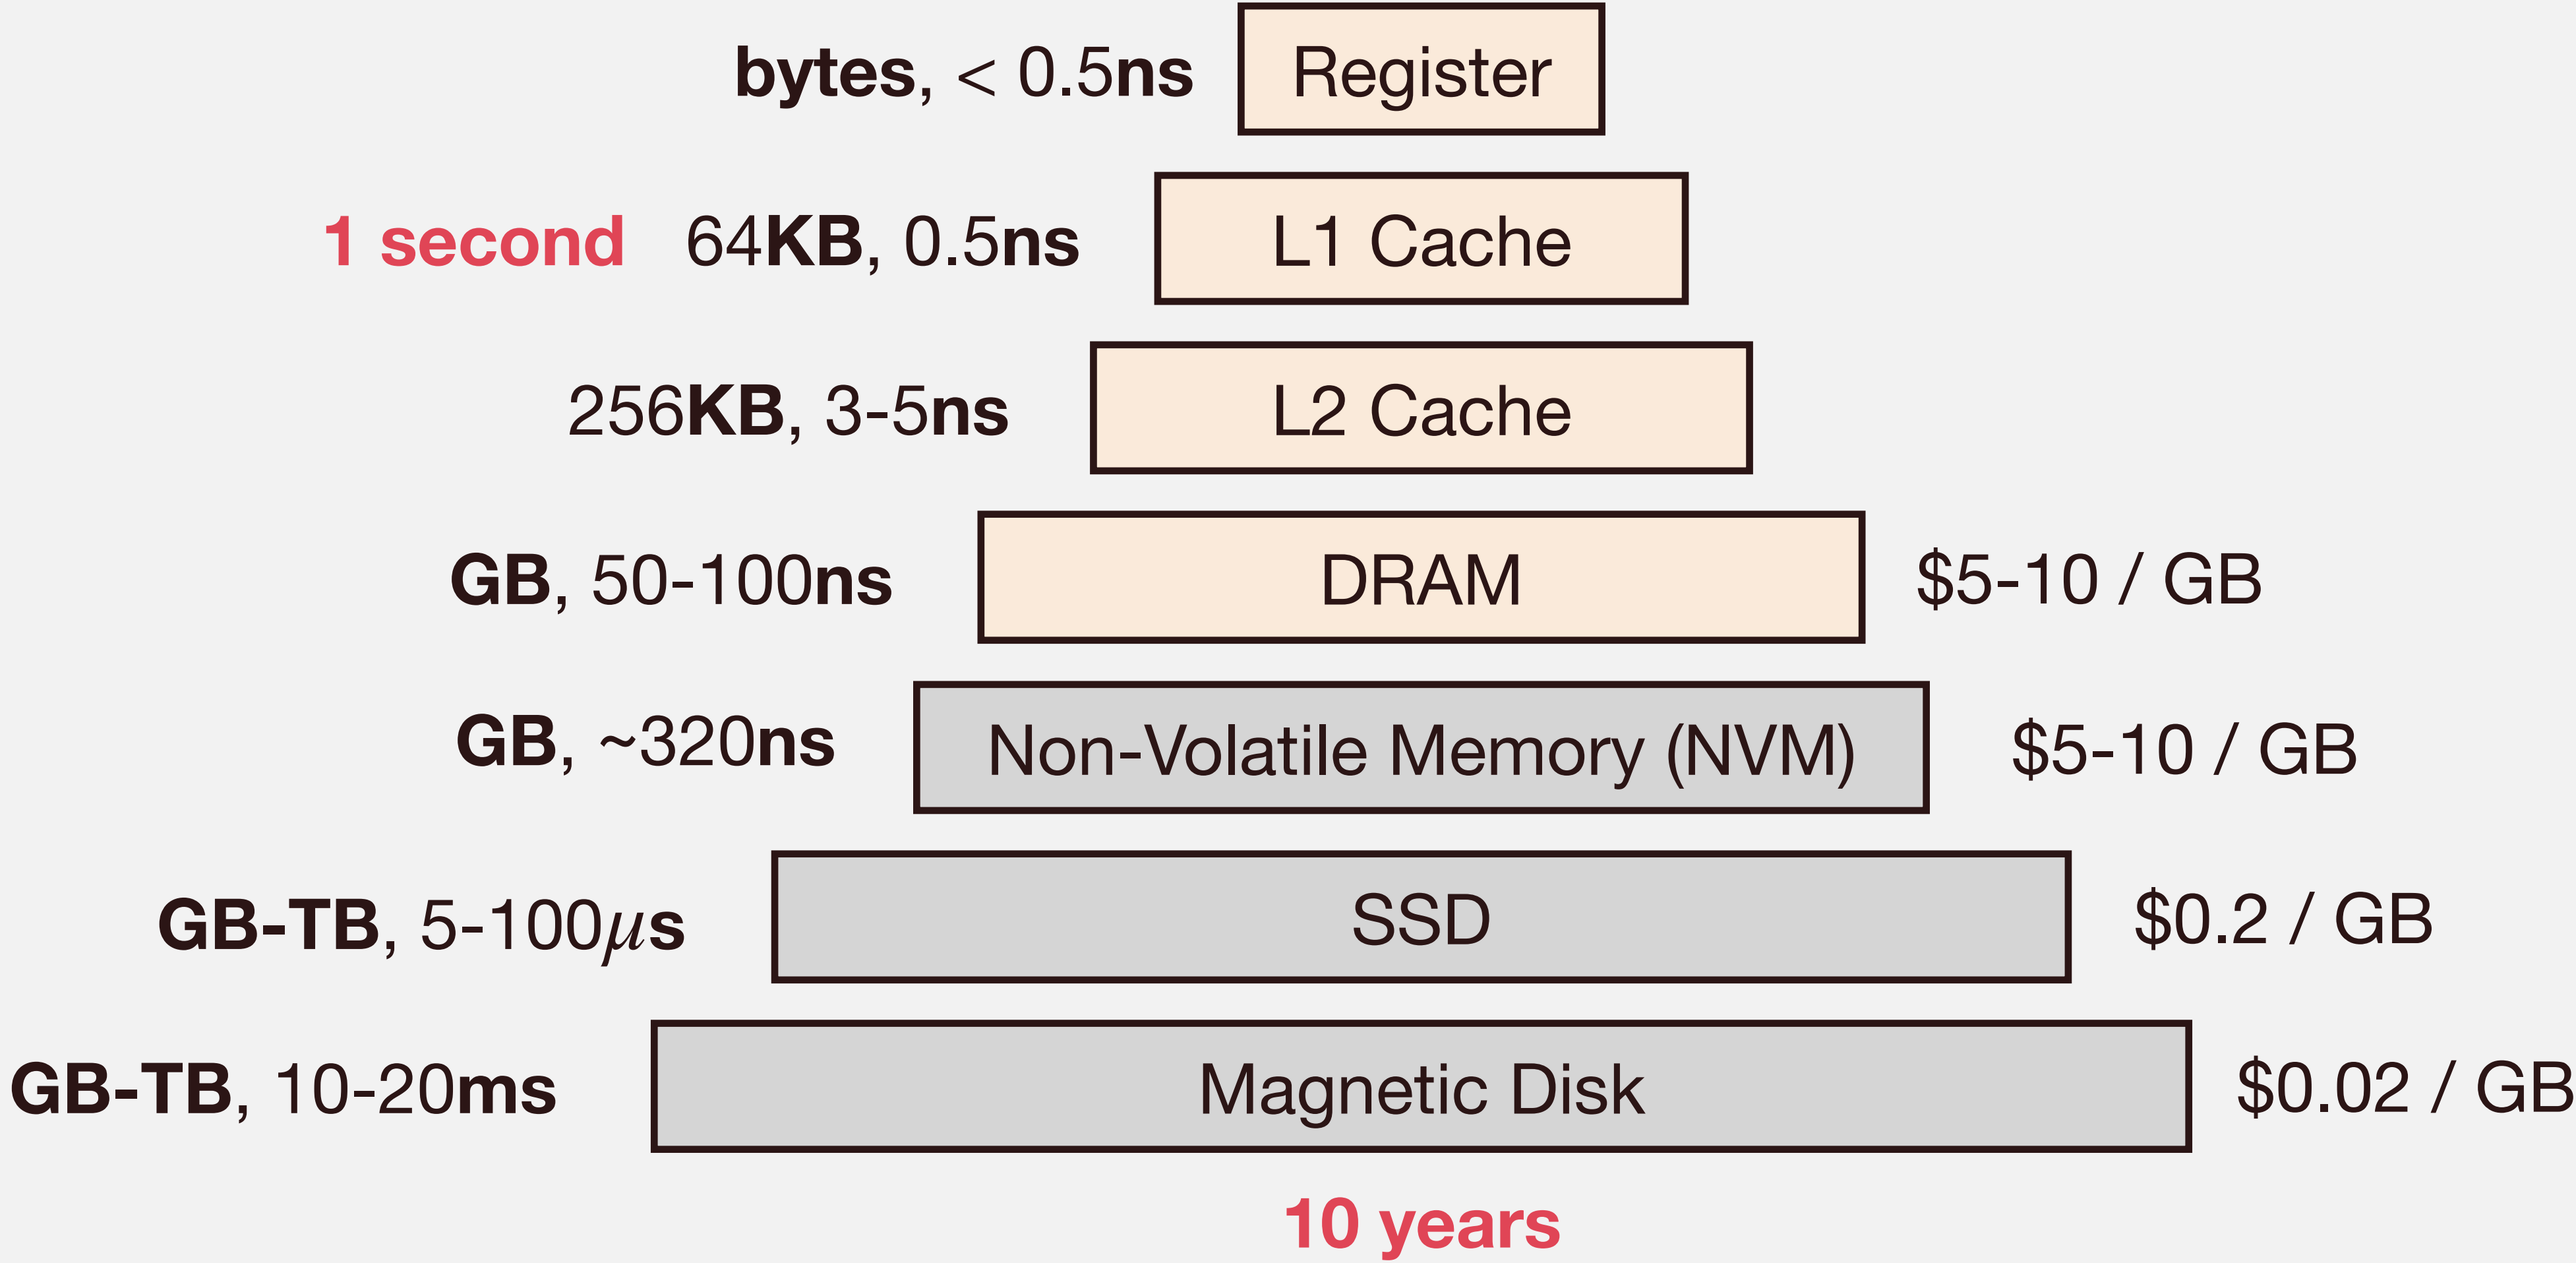
Storage Hierarchy



Storage Hierarchy

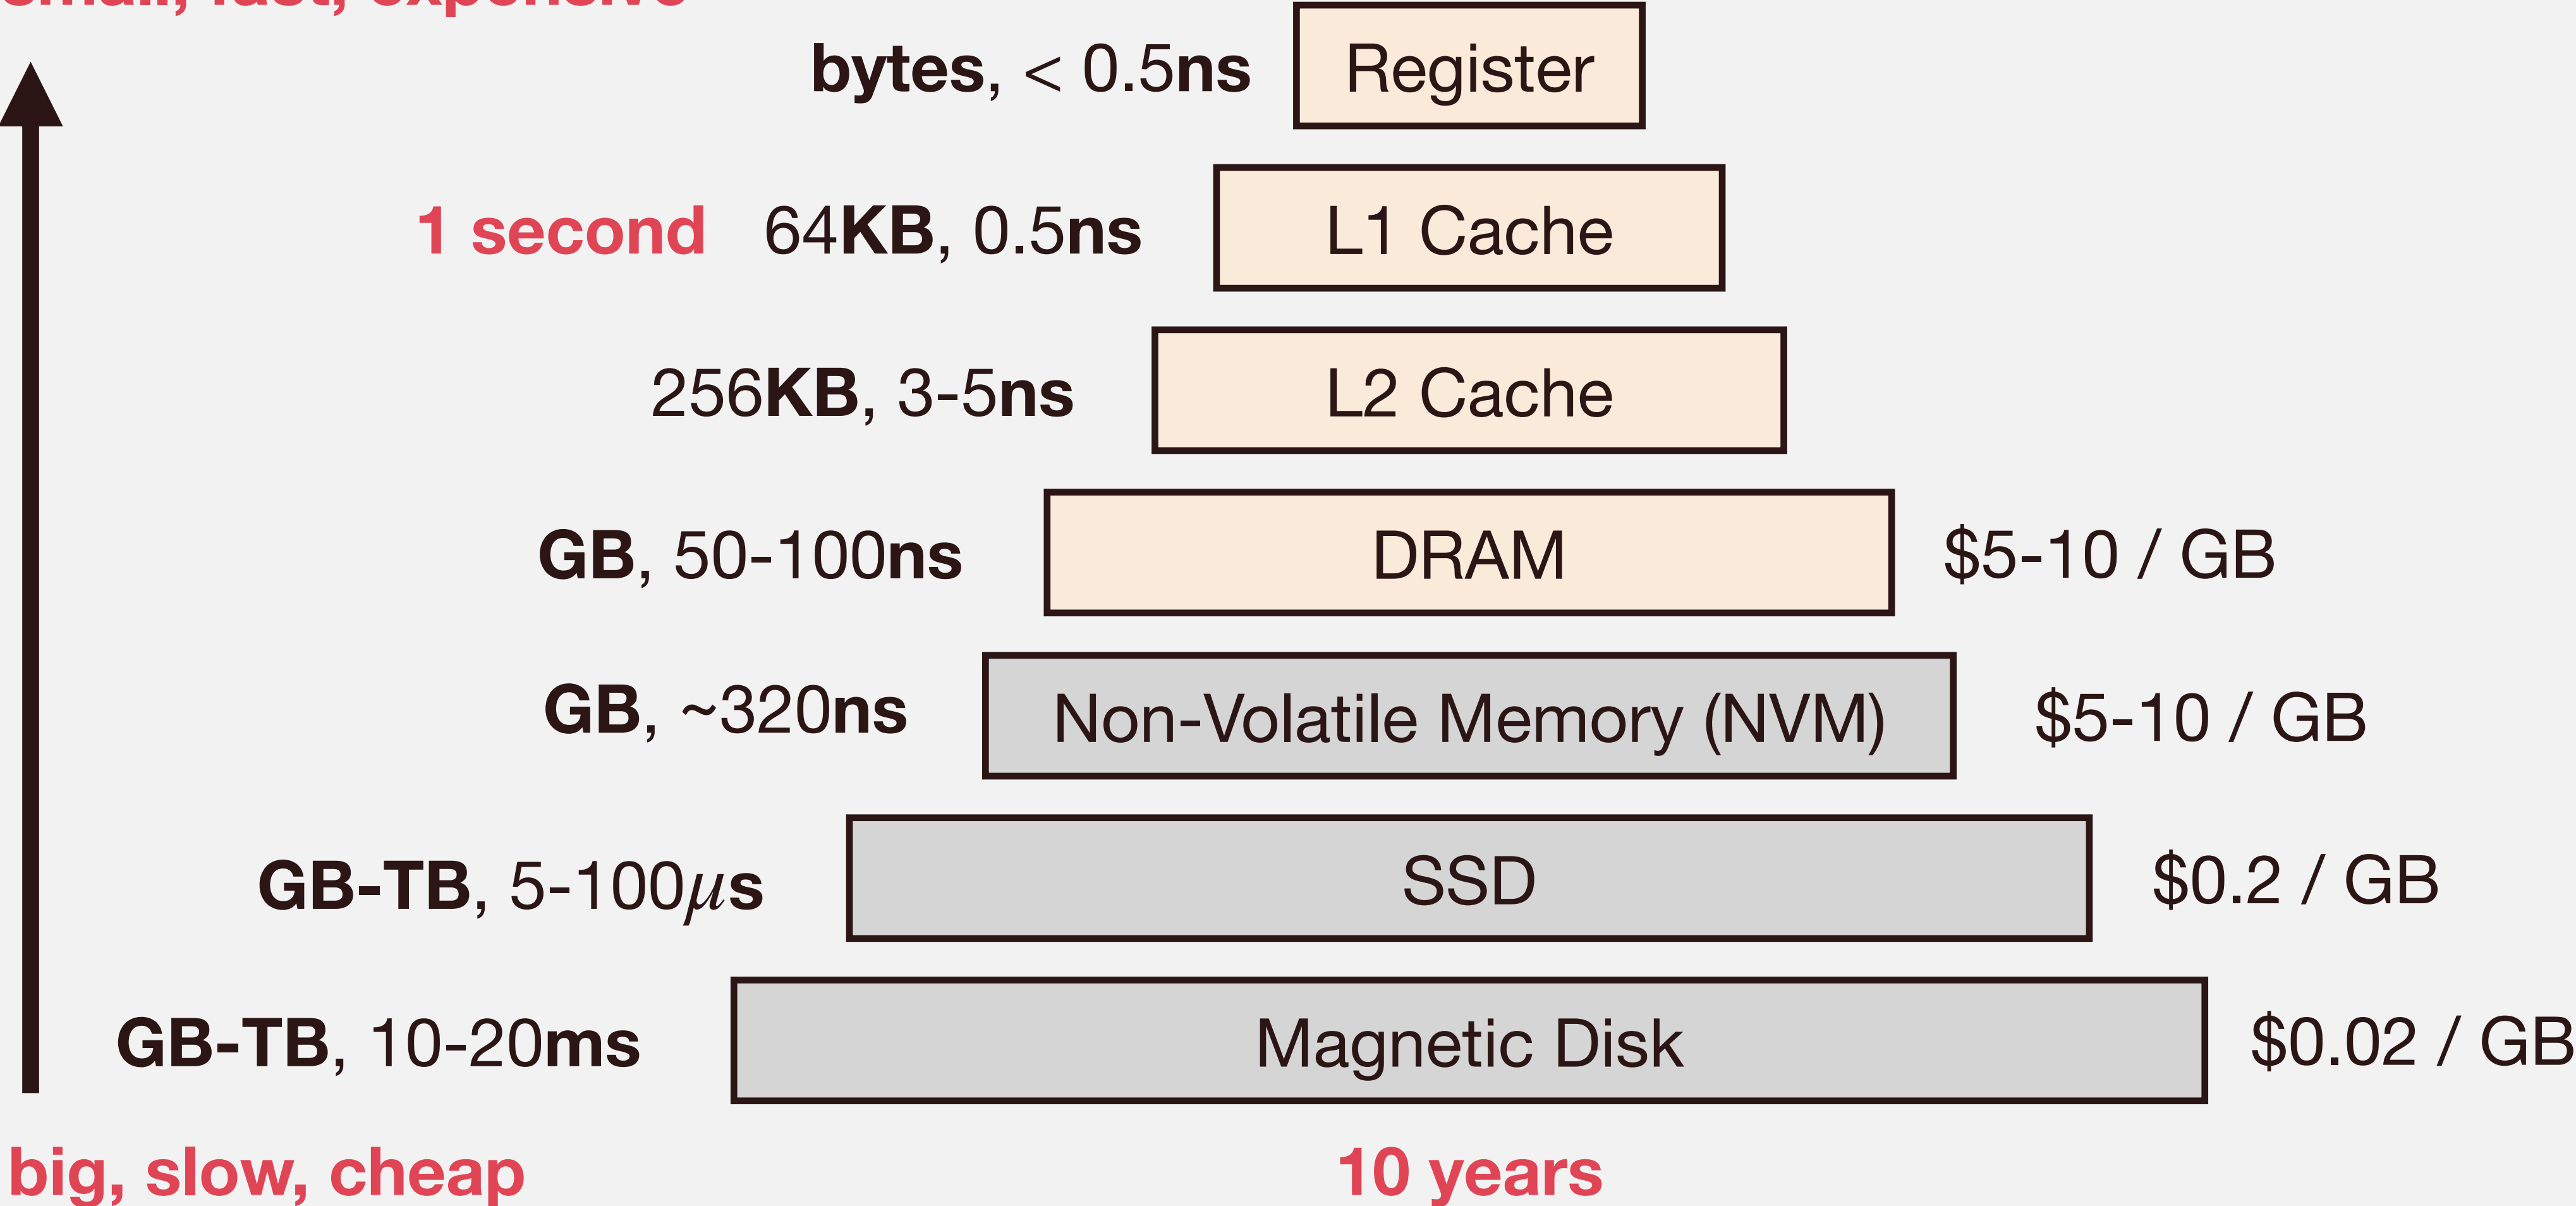


Storage Hierarchy

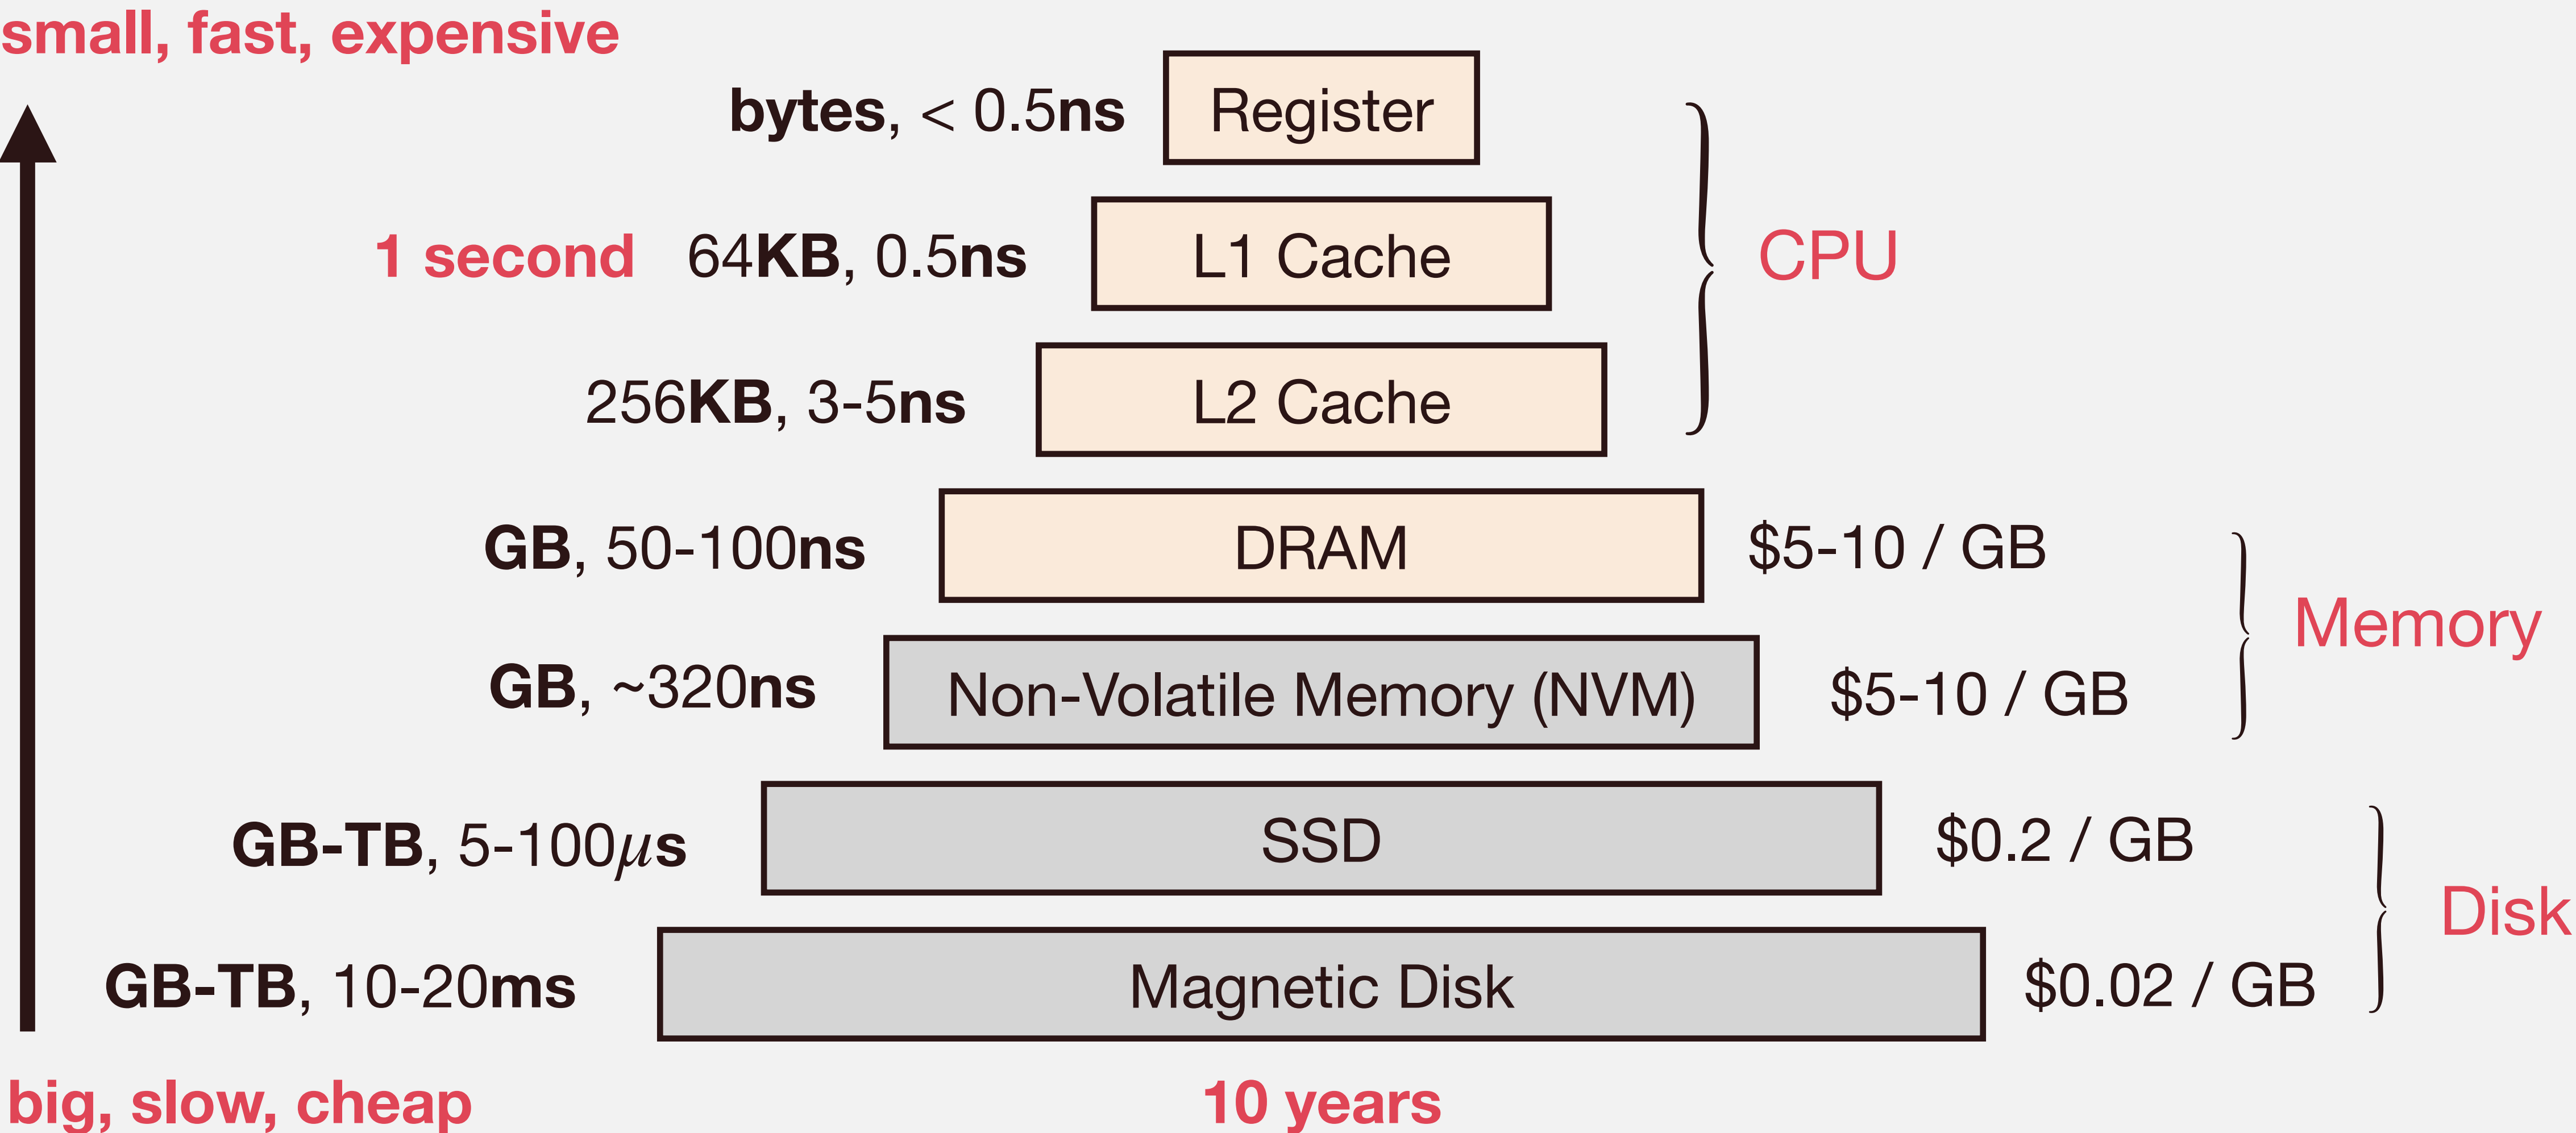


Storage Hierarchy

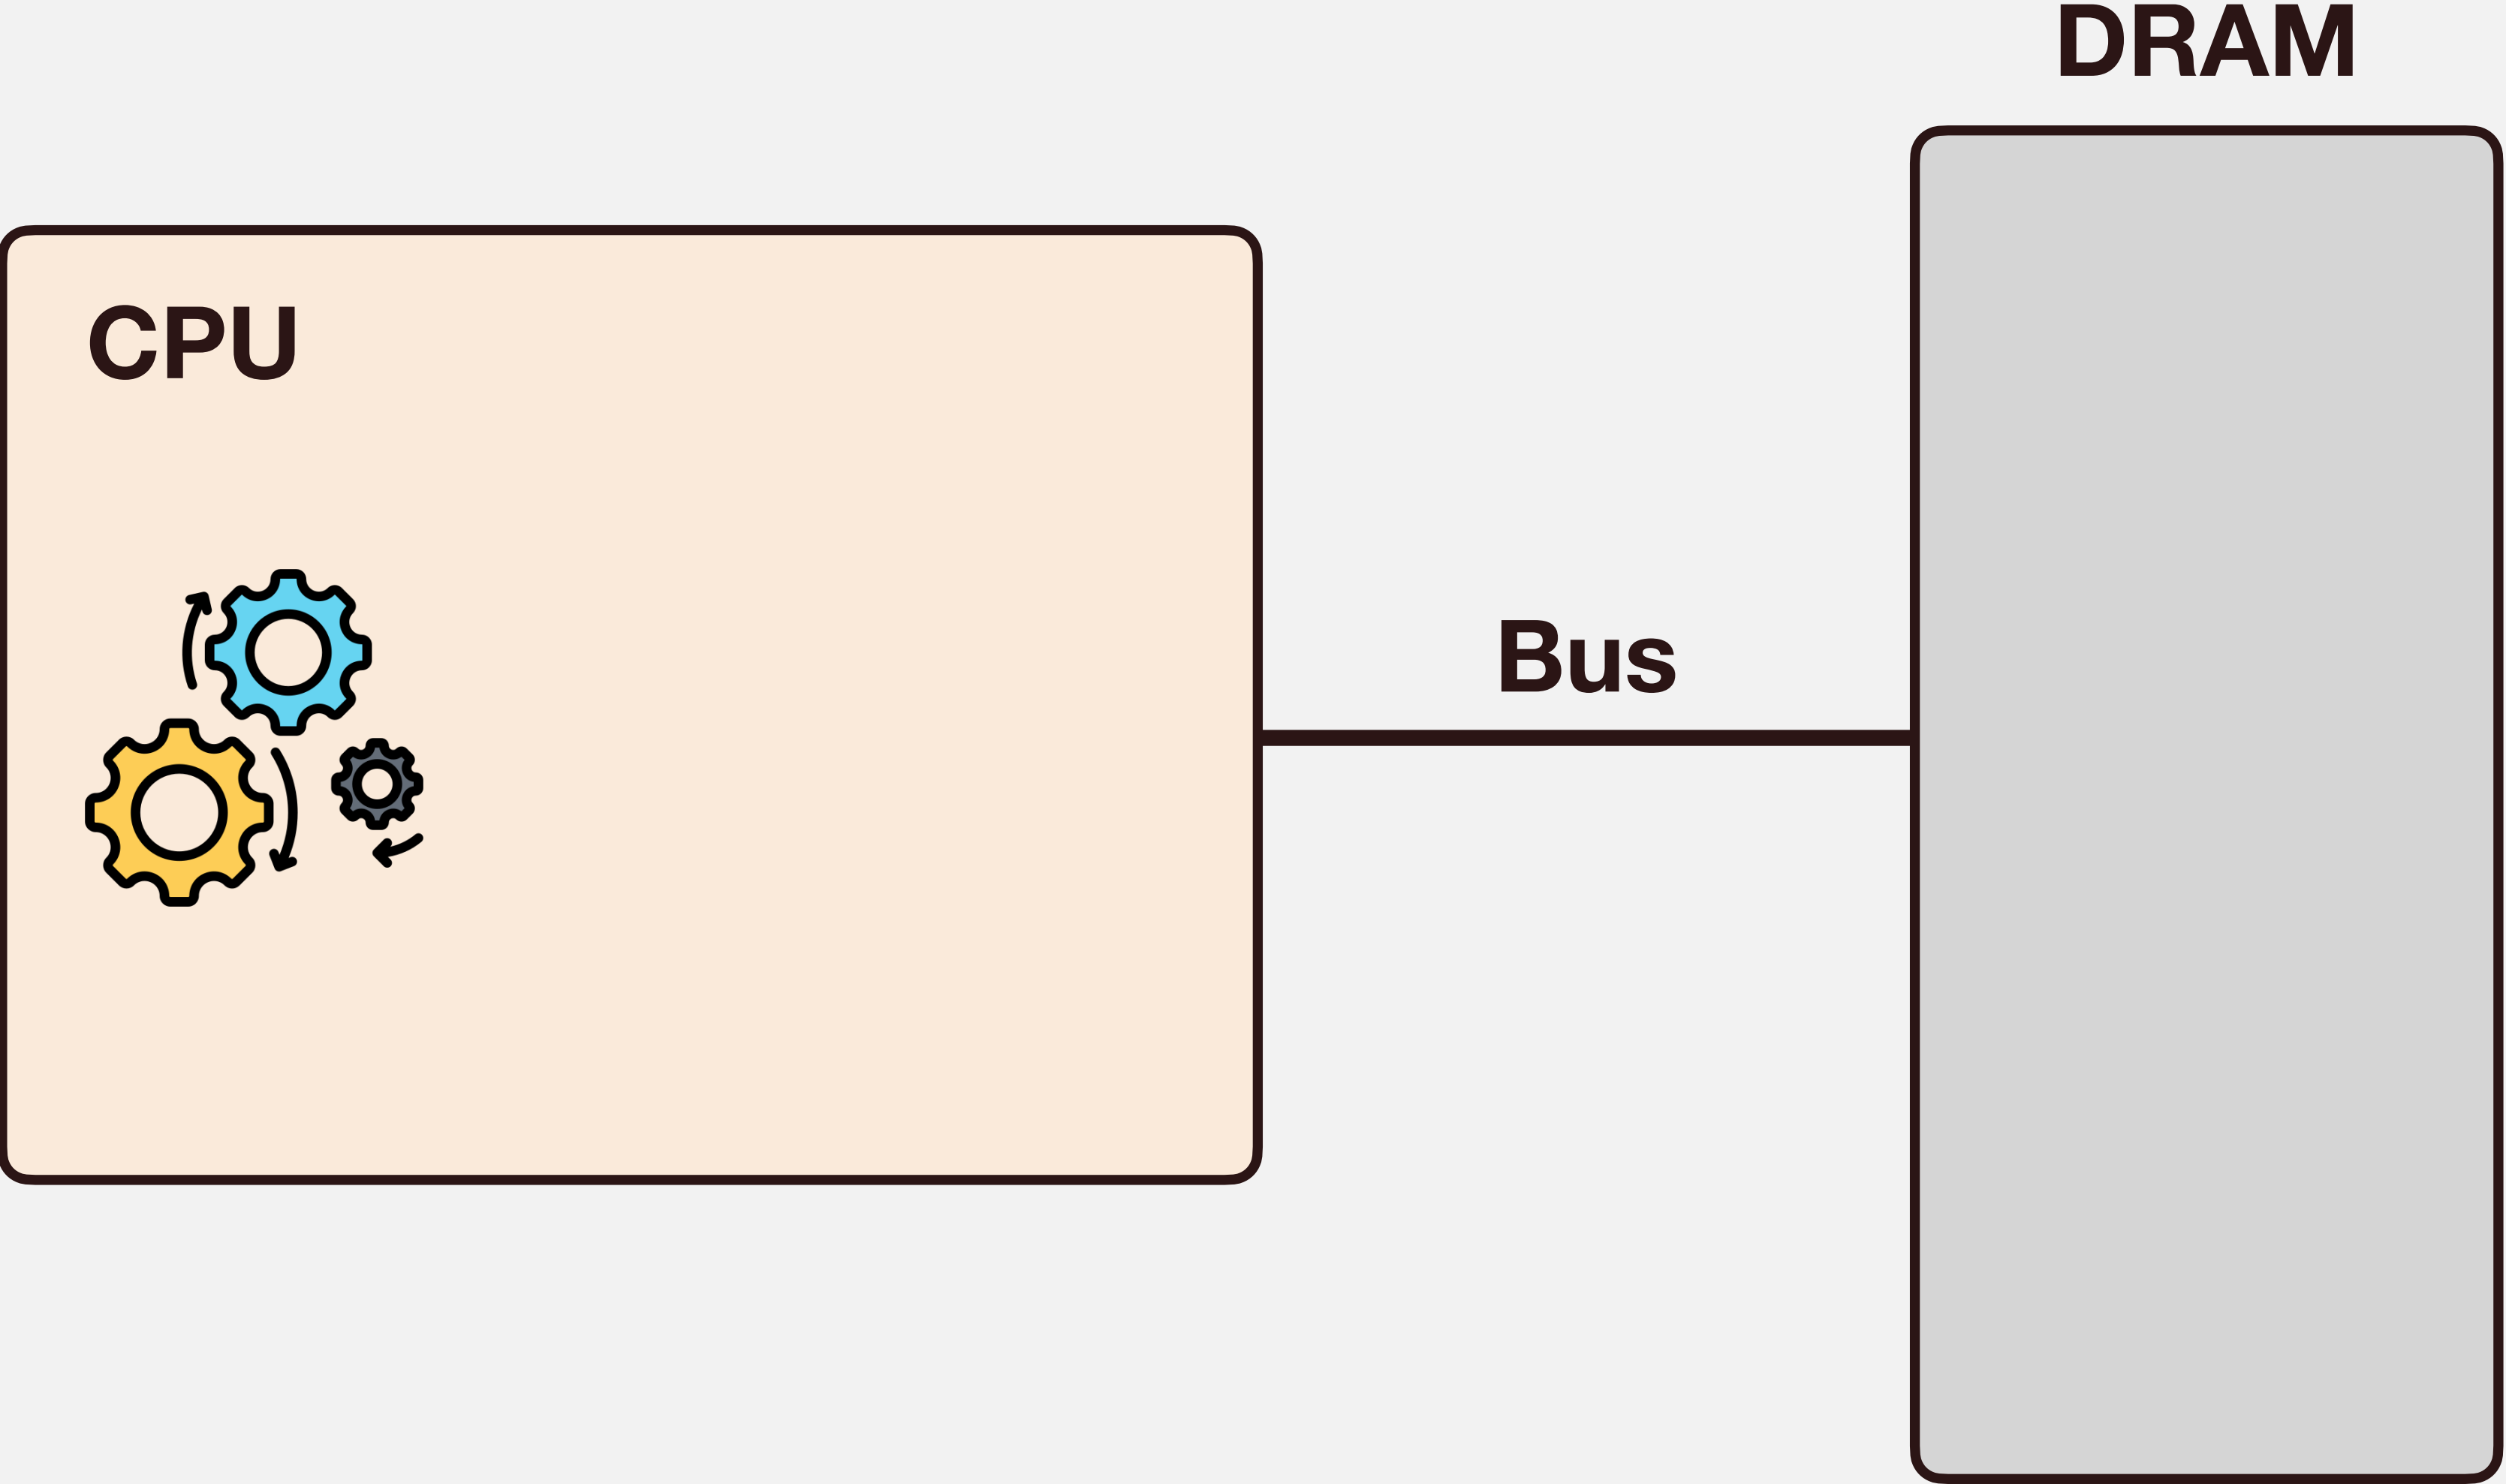
small, fast, expensive



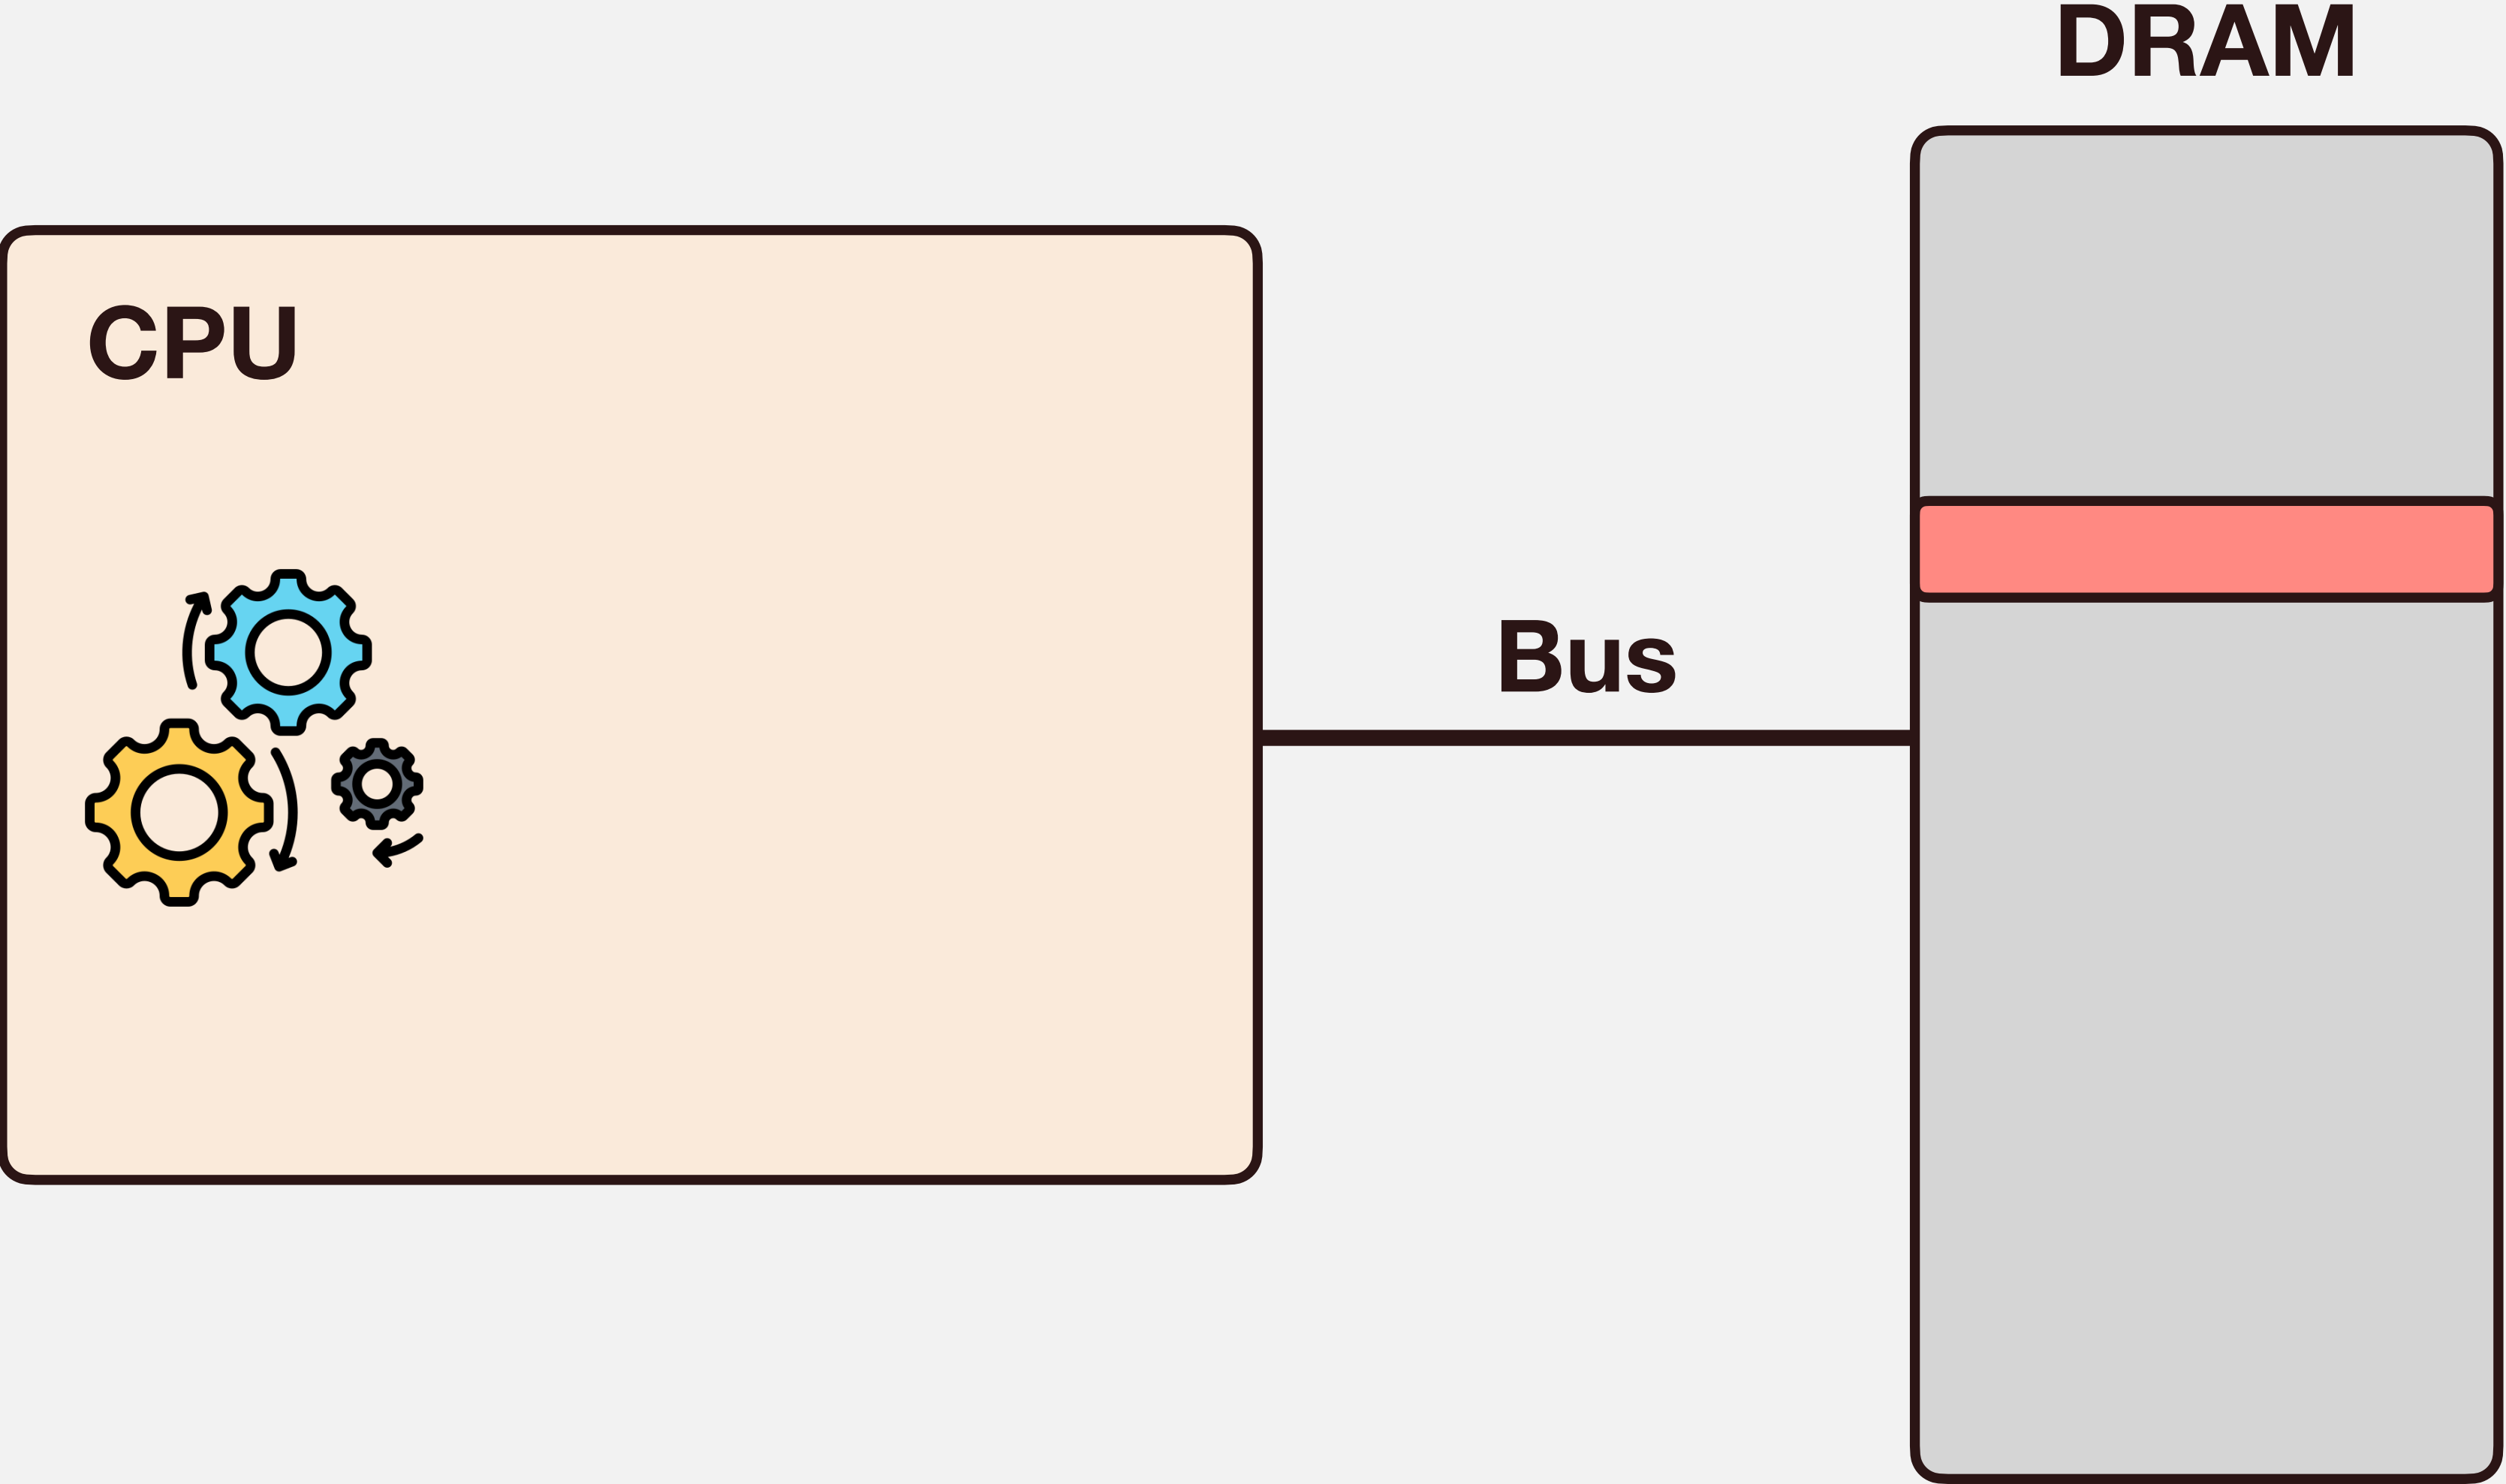
Storage Hierarchy



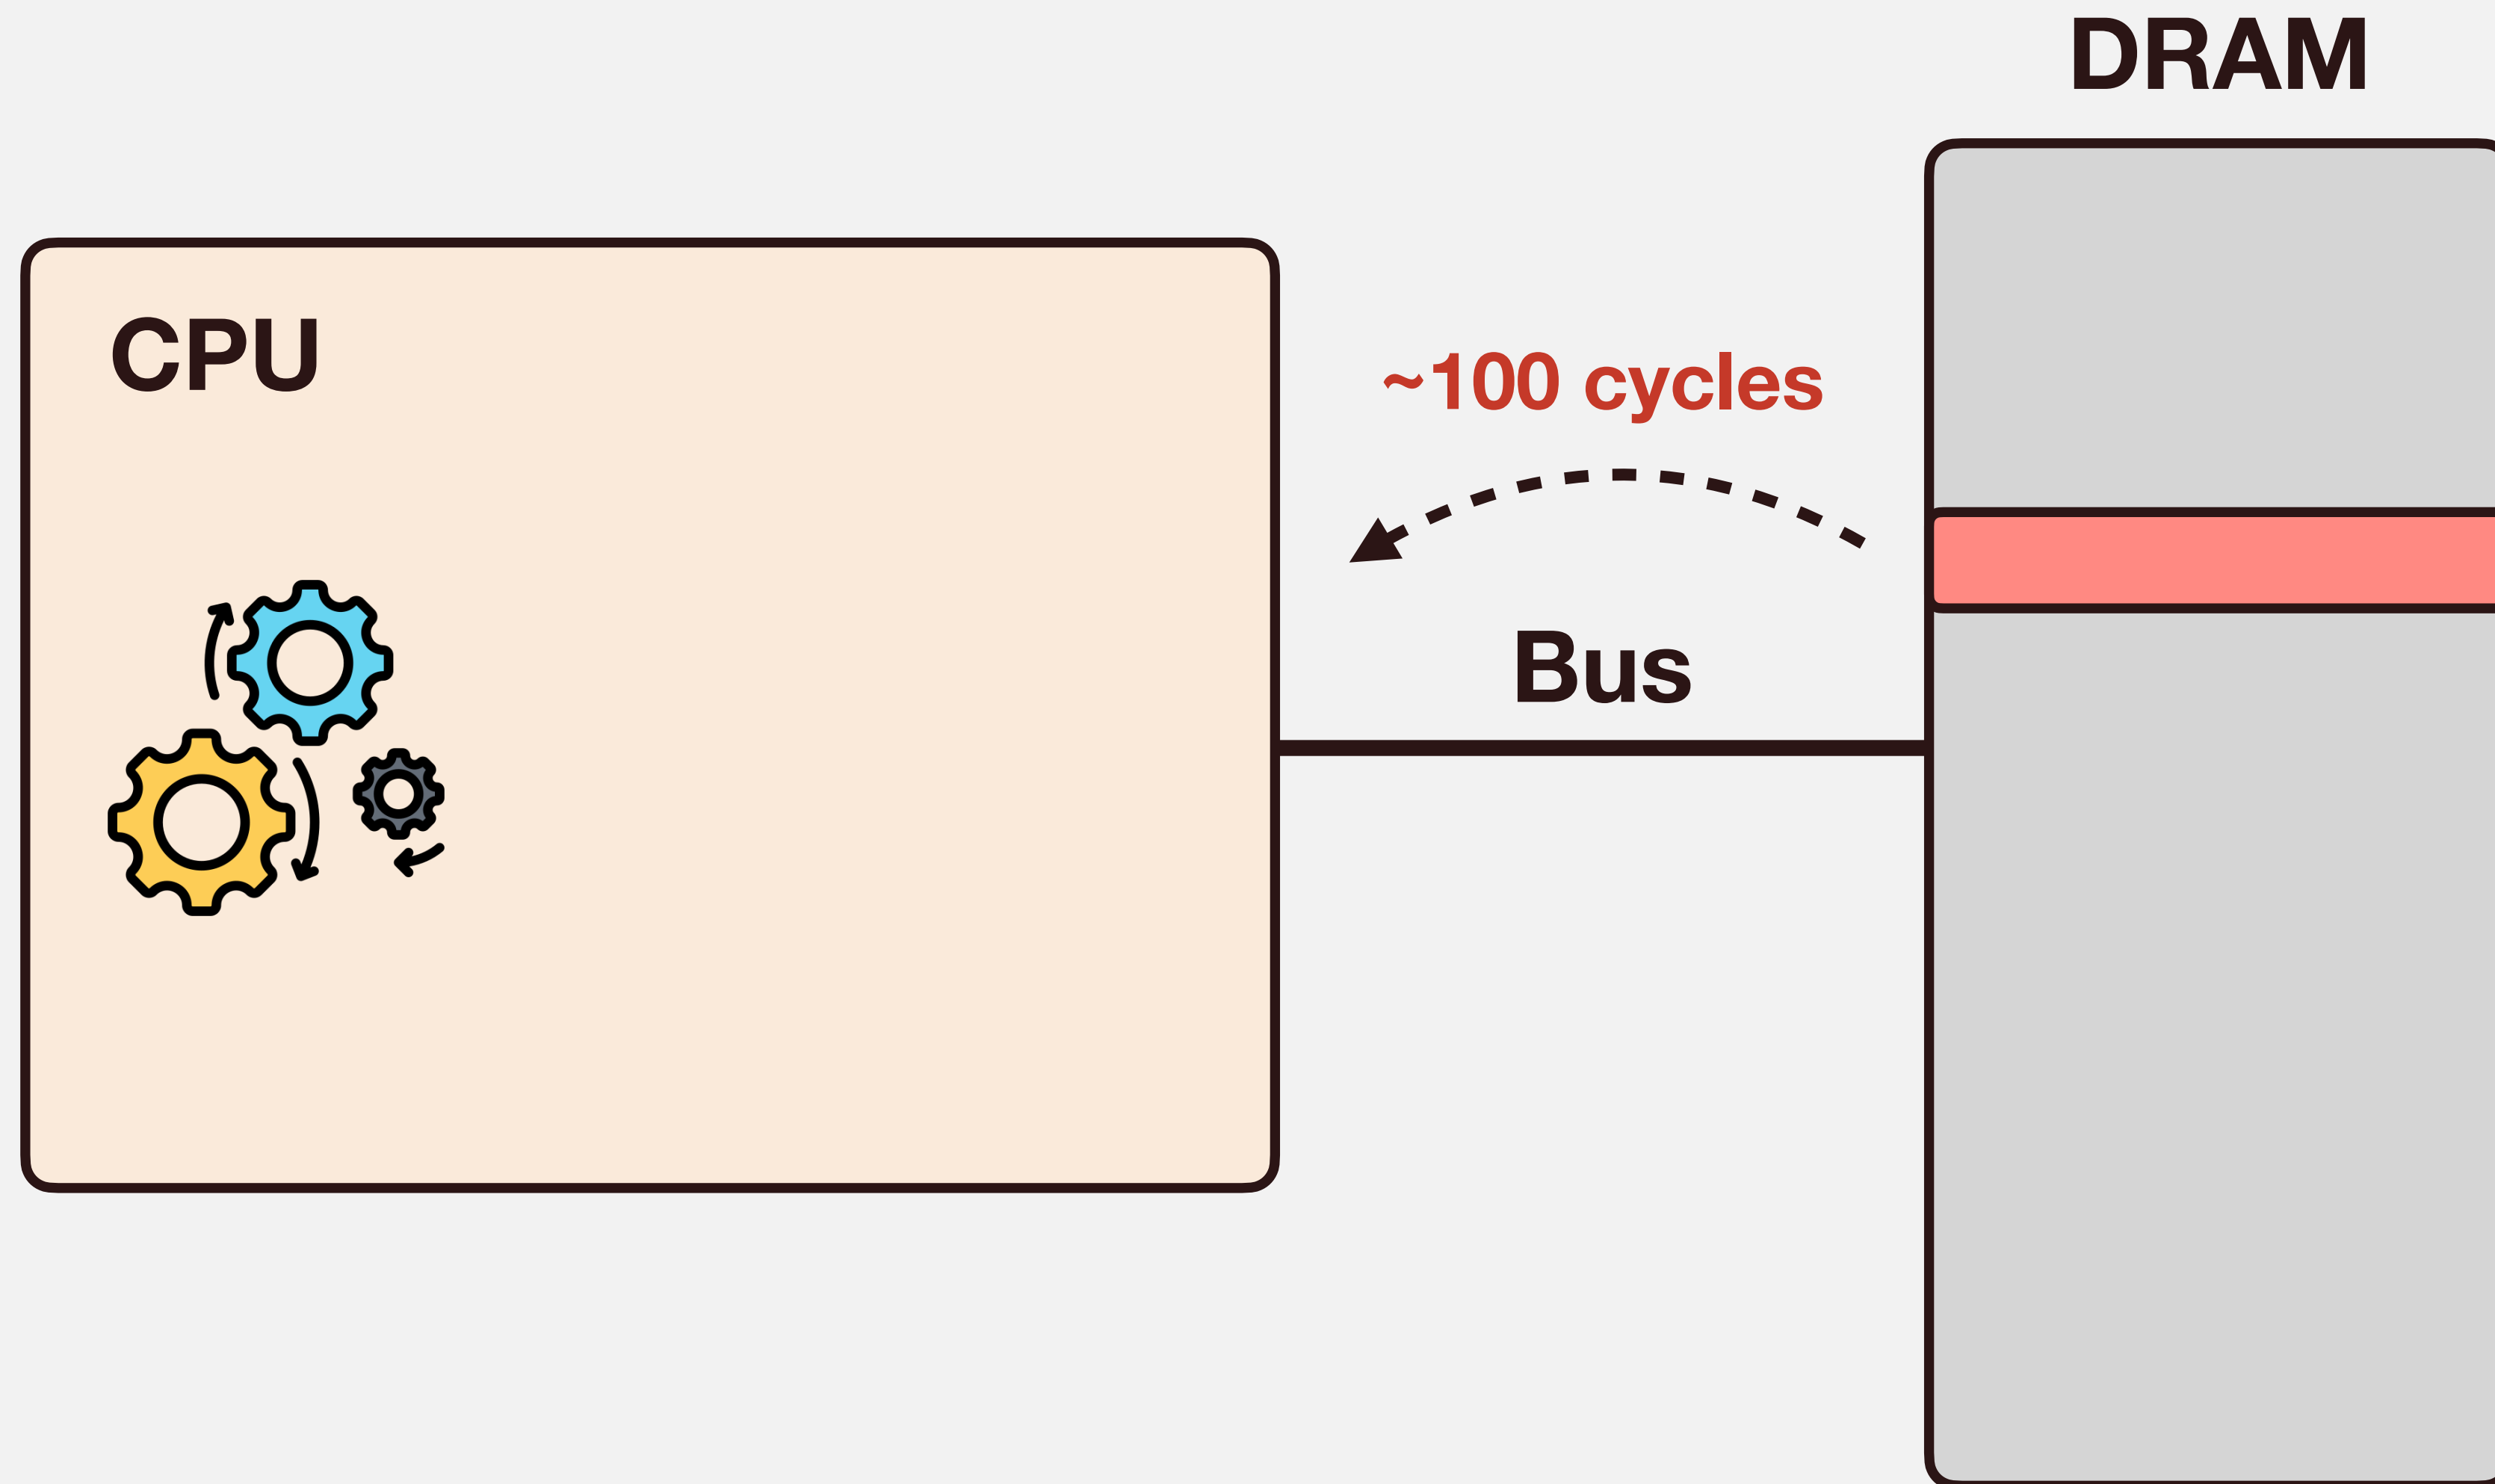
CPU Cache



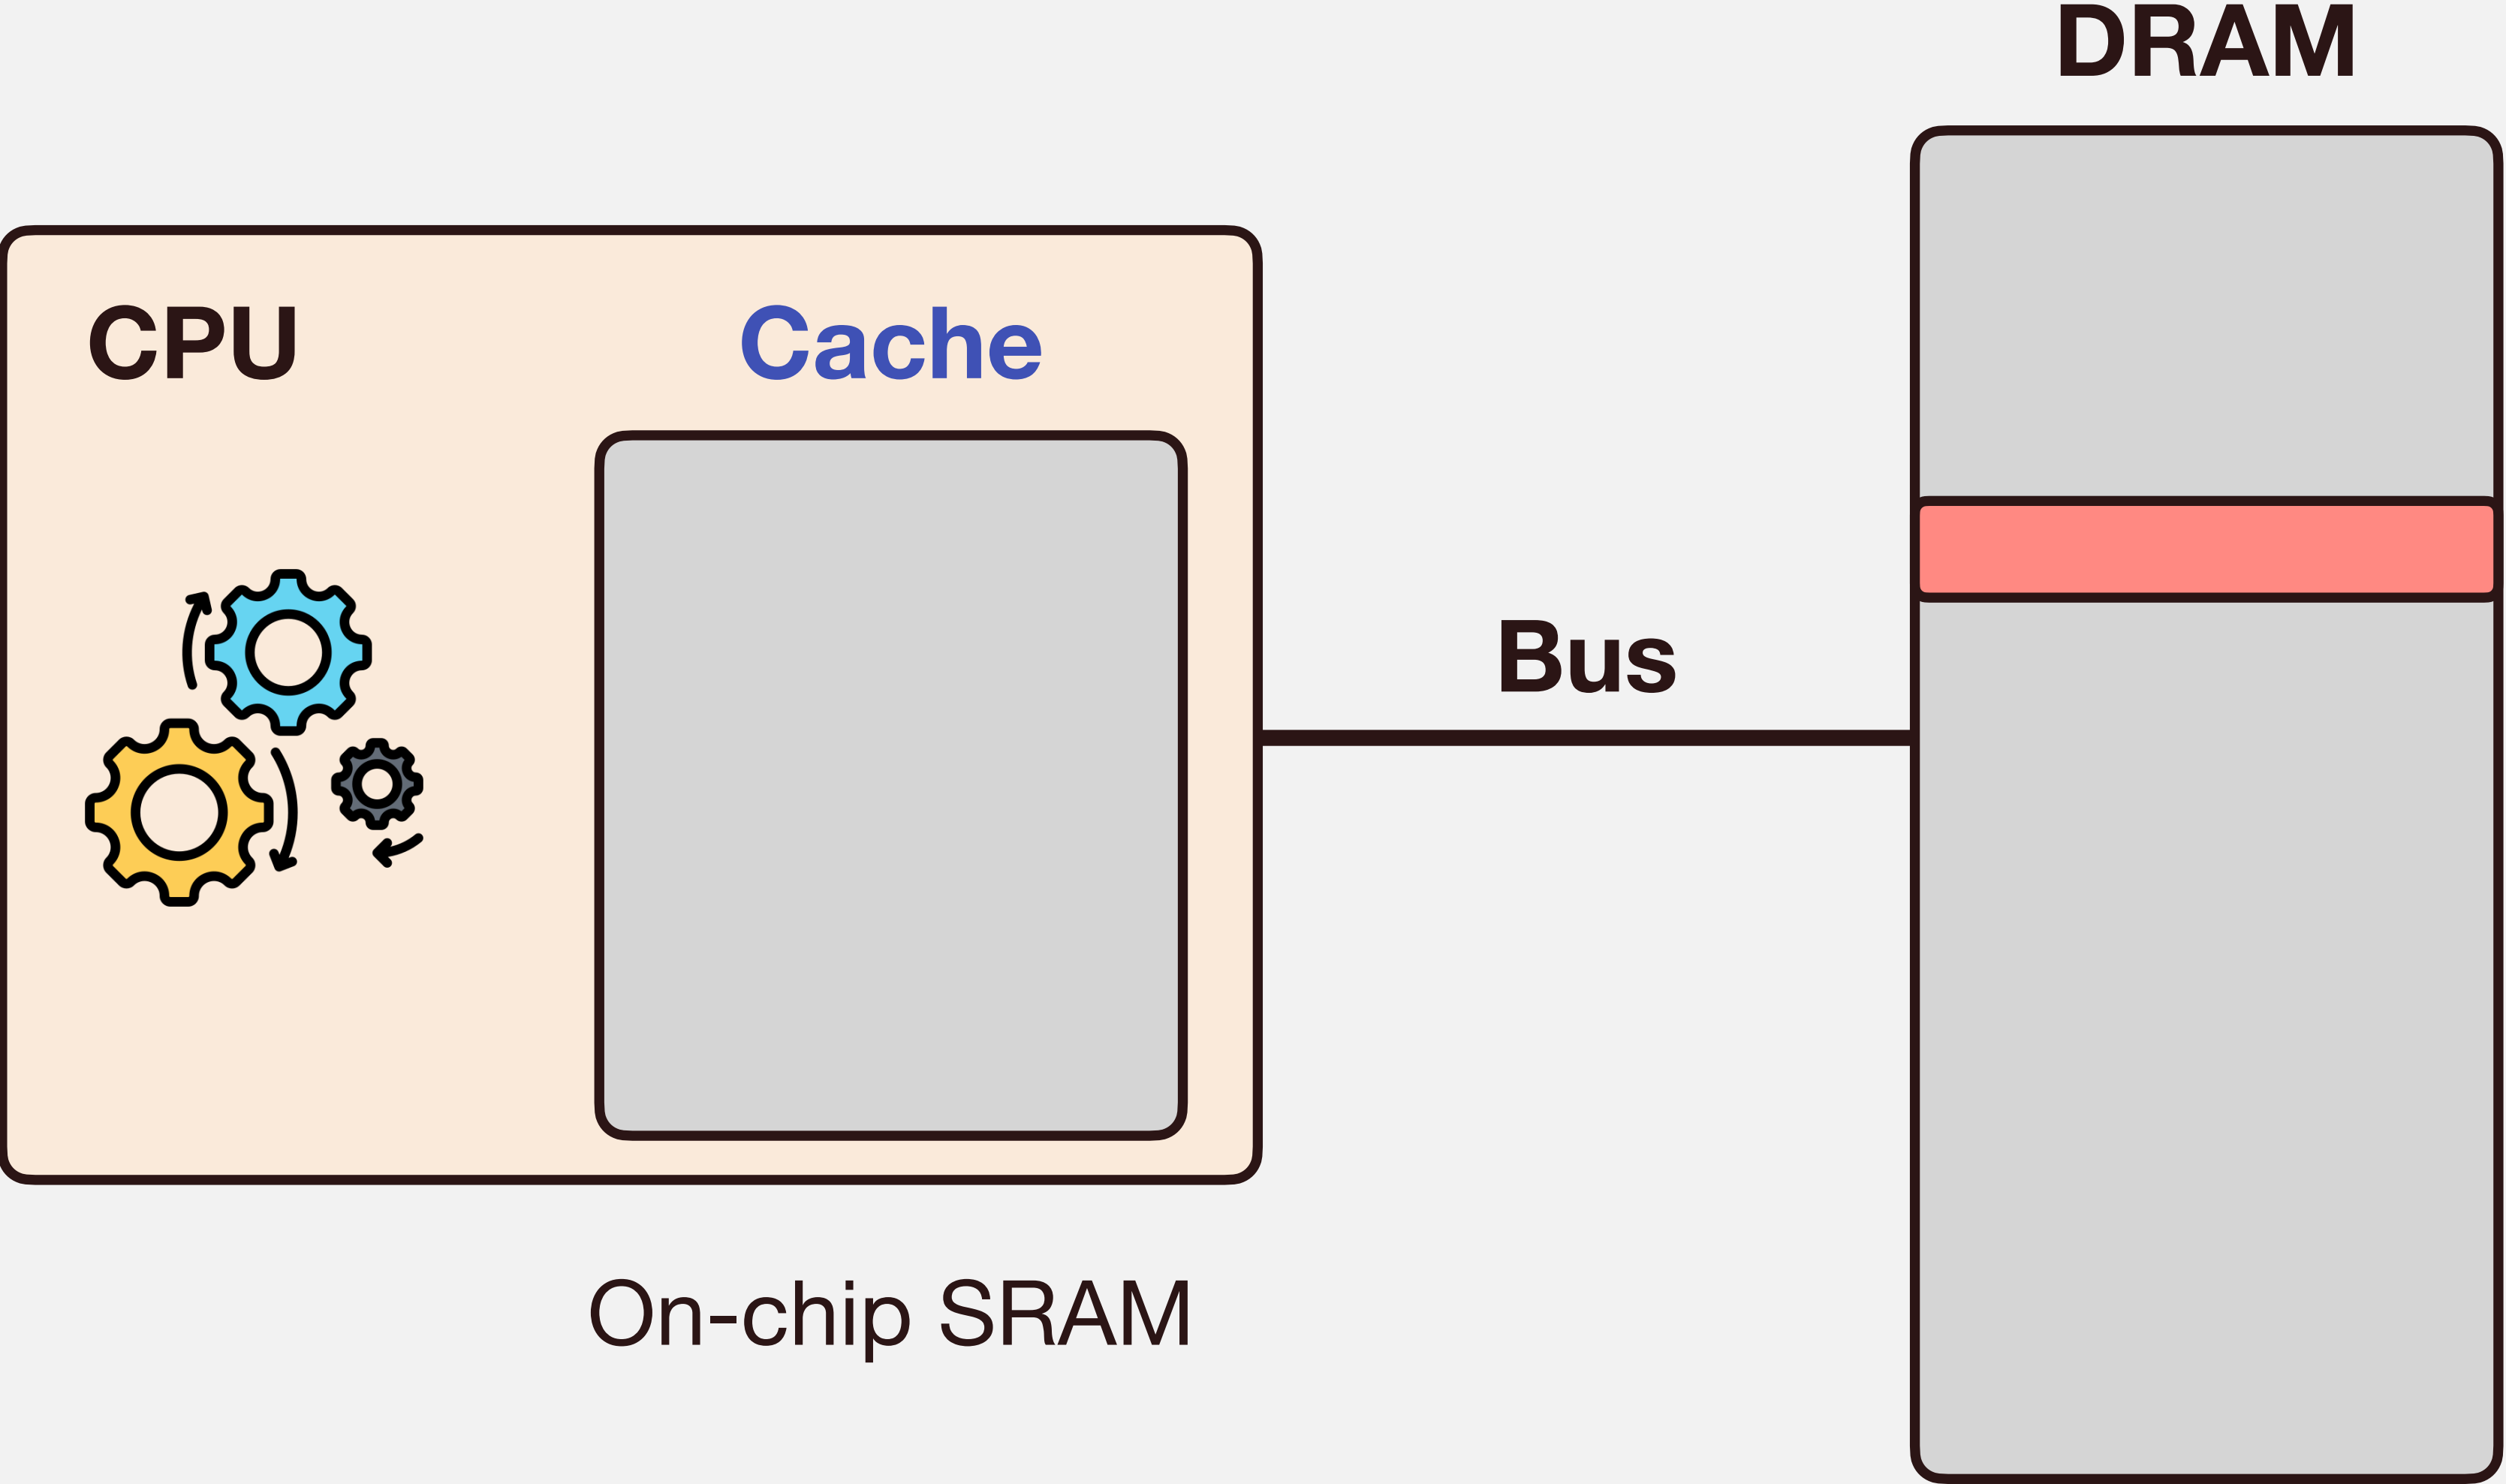
CPU Cache



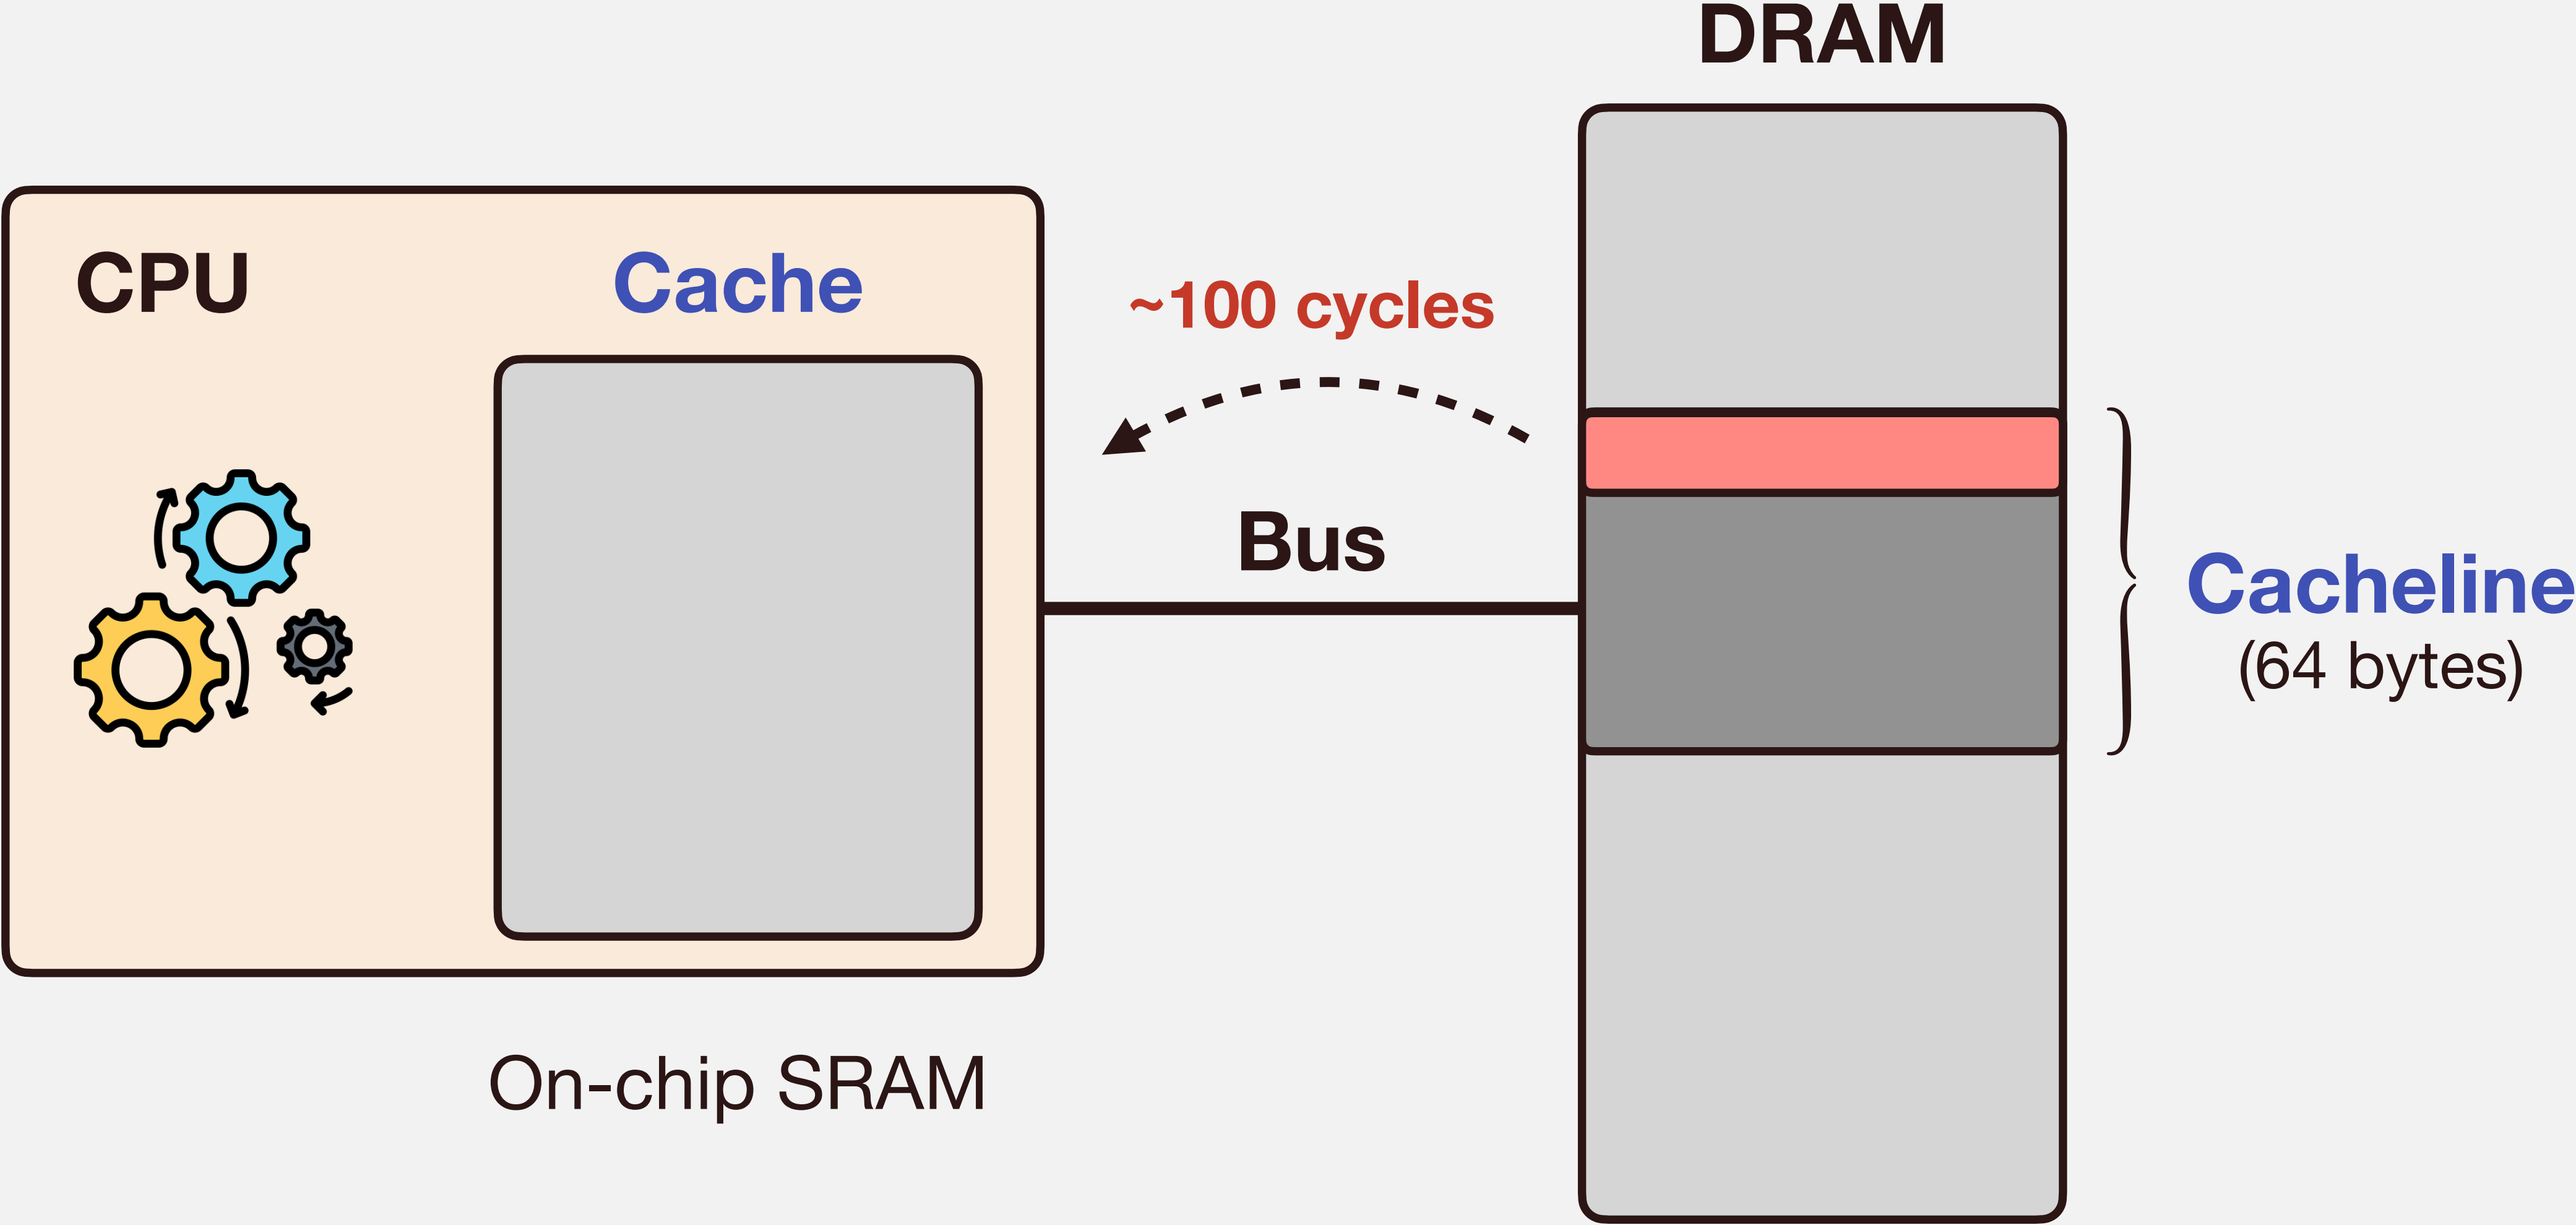
CPU Cache



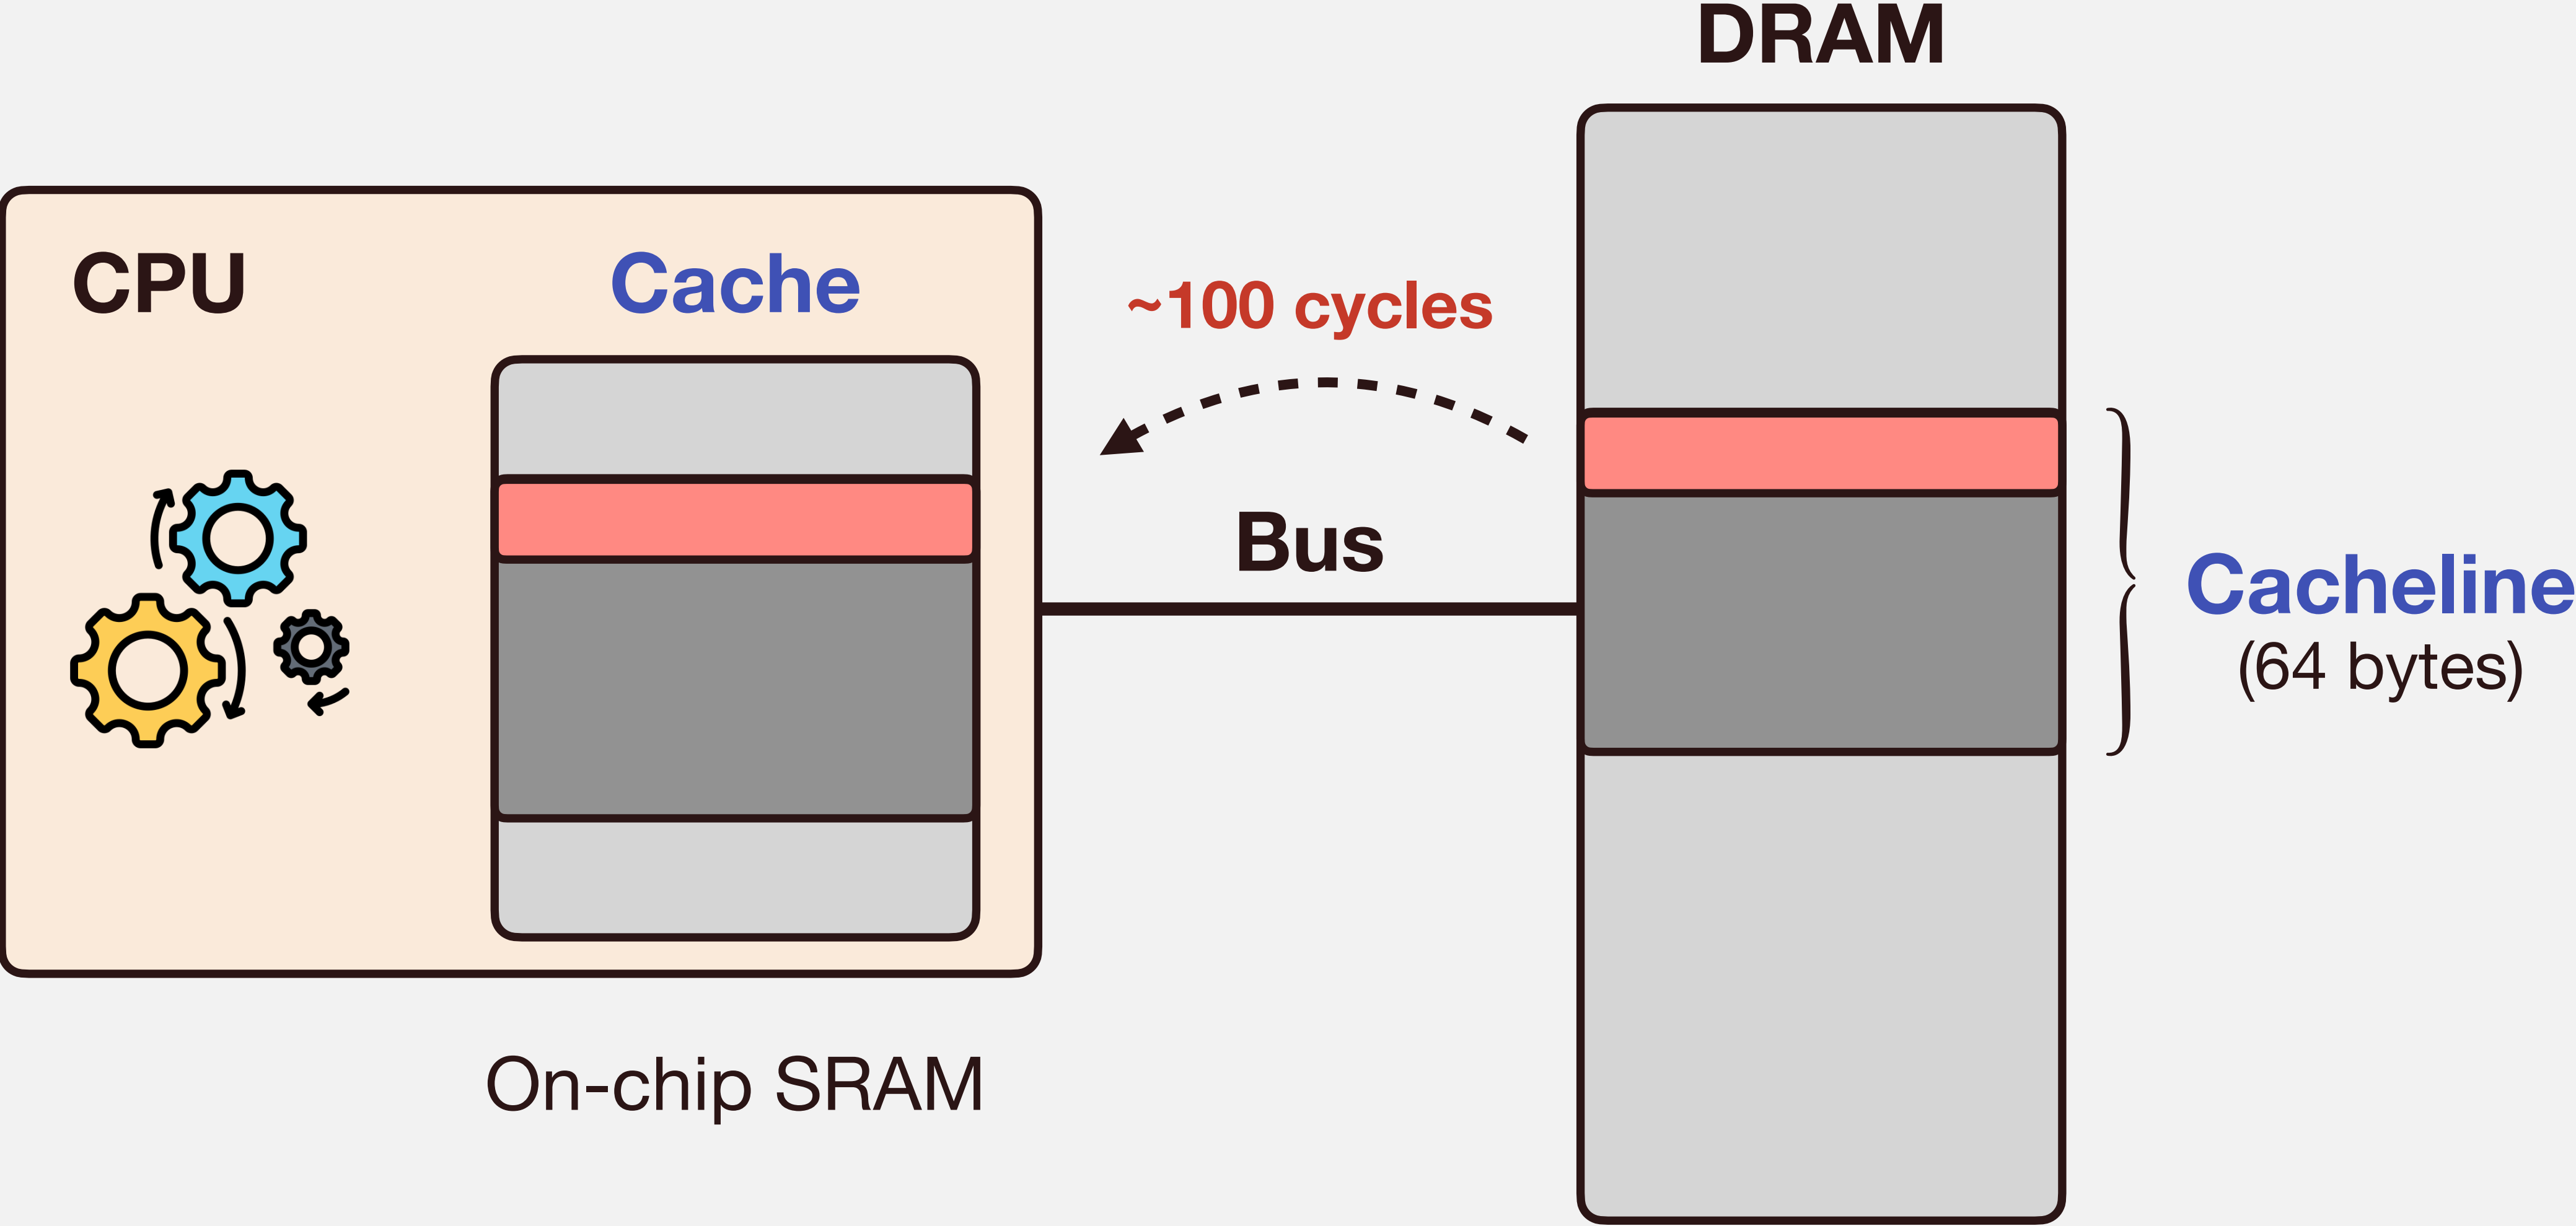
CPU Cache



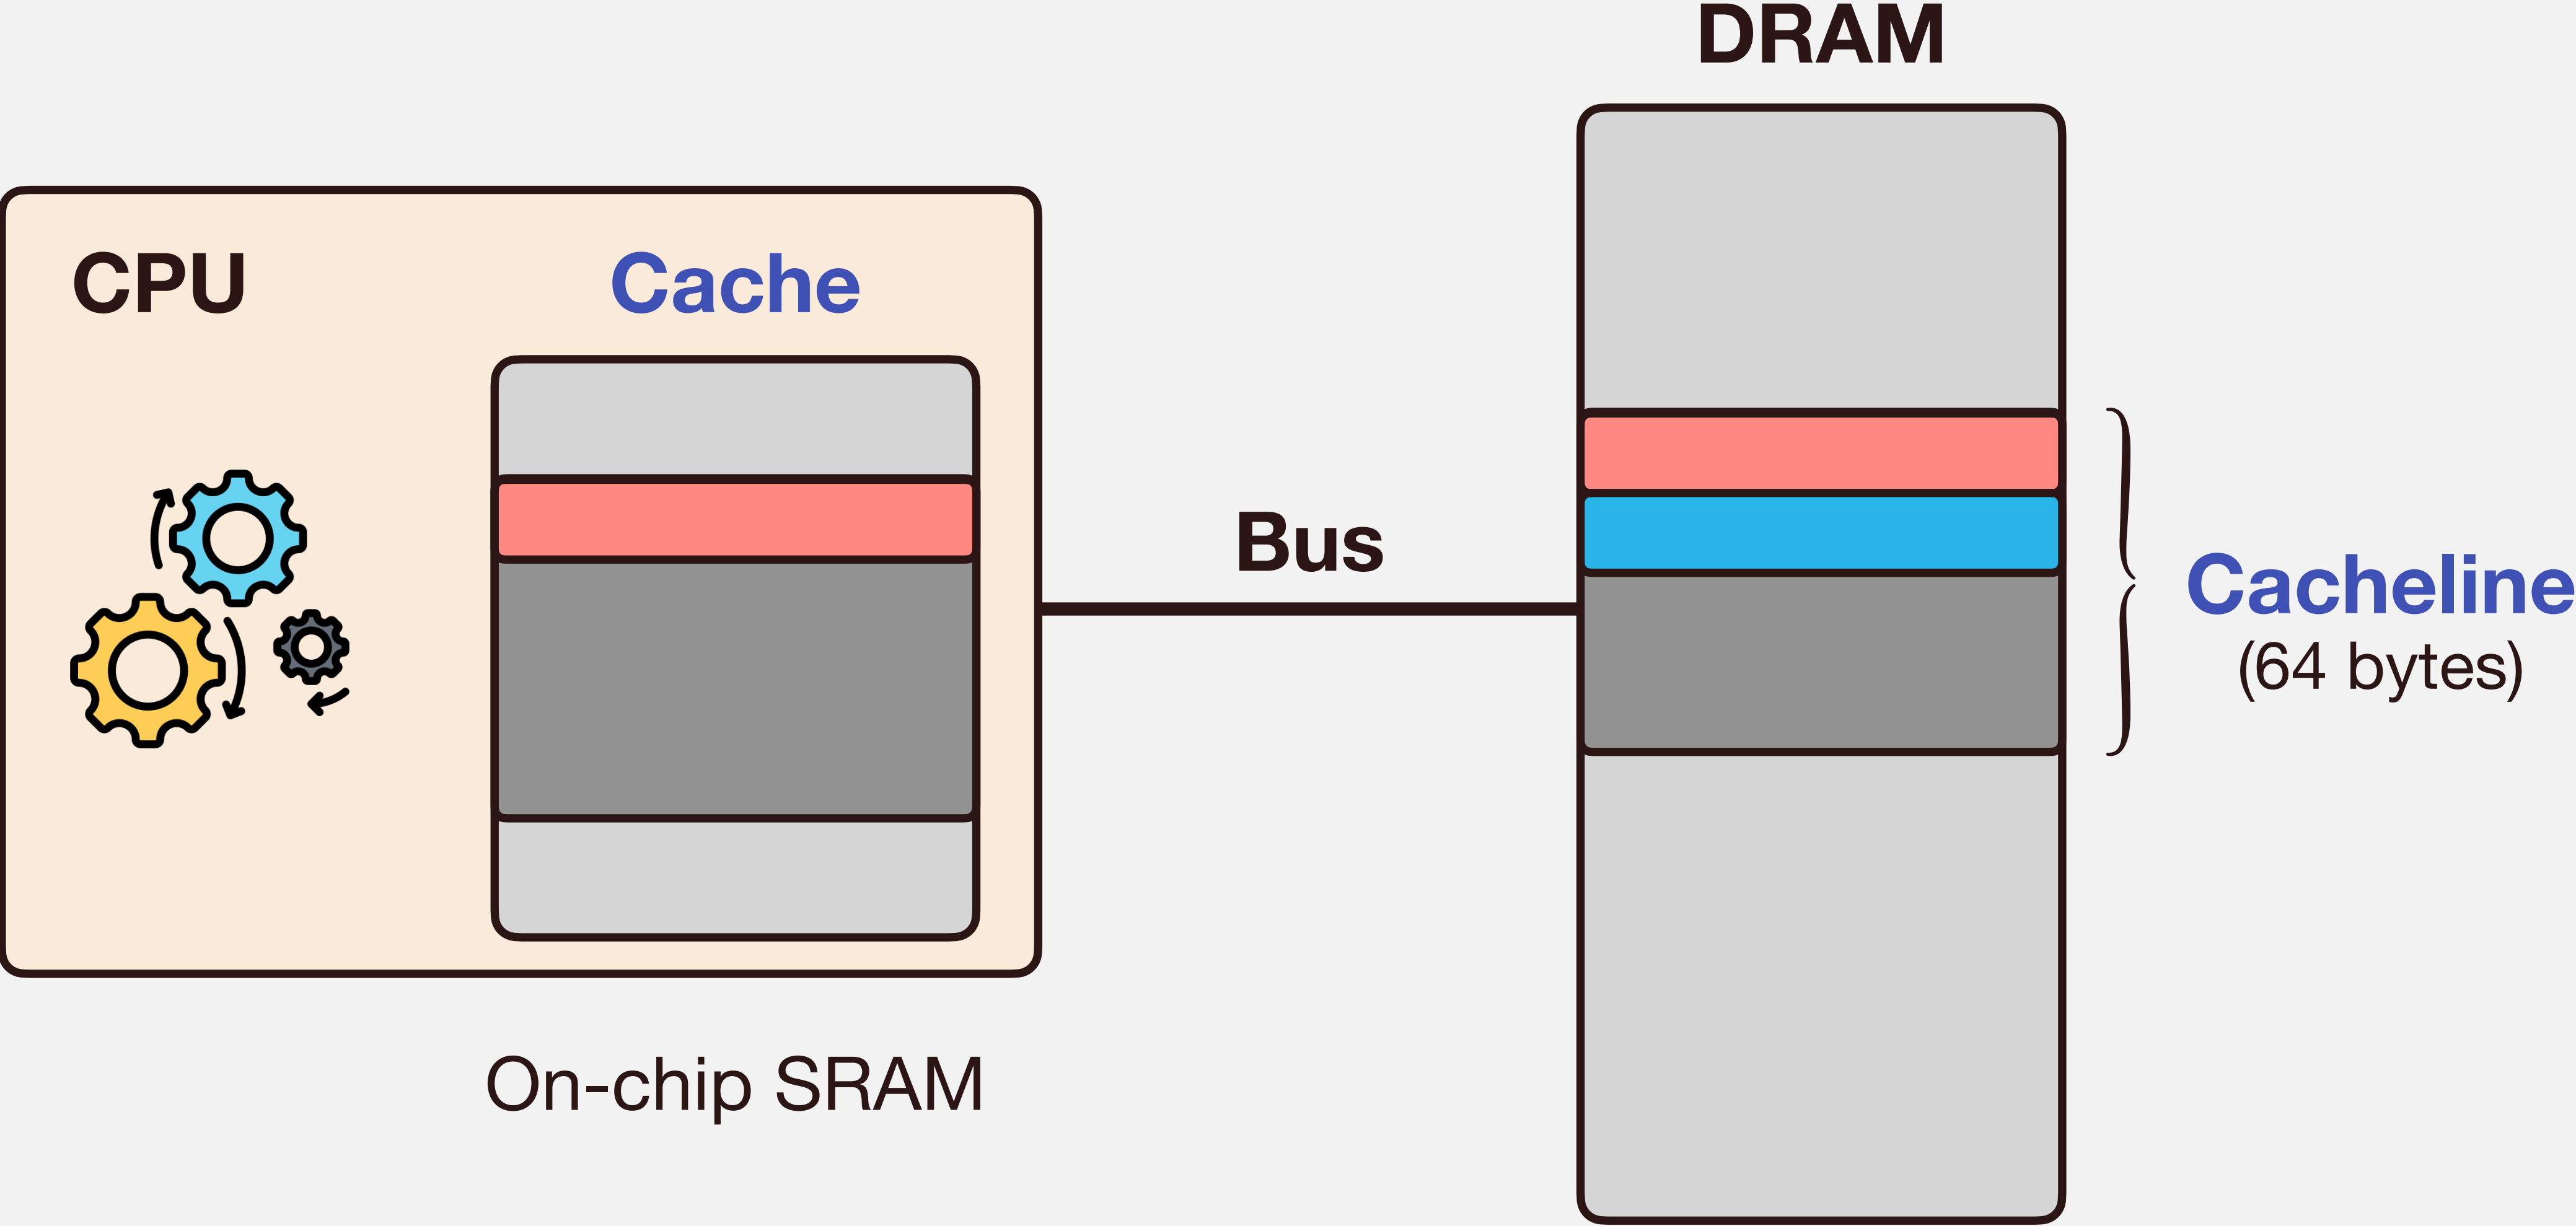
CPU Cache



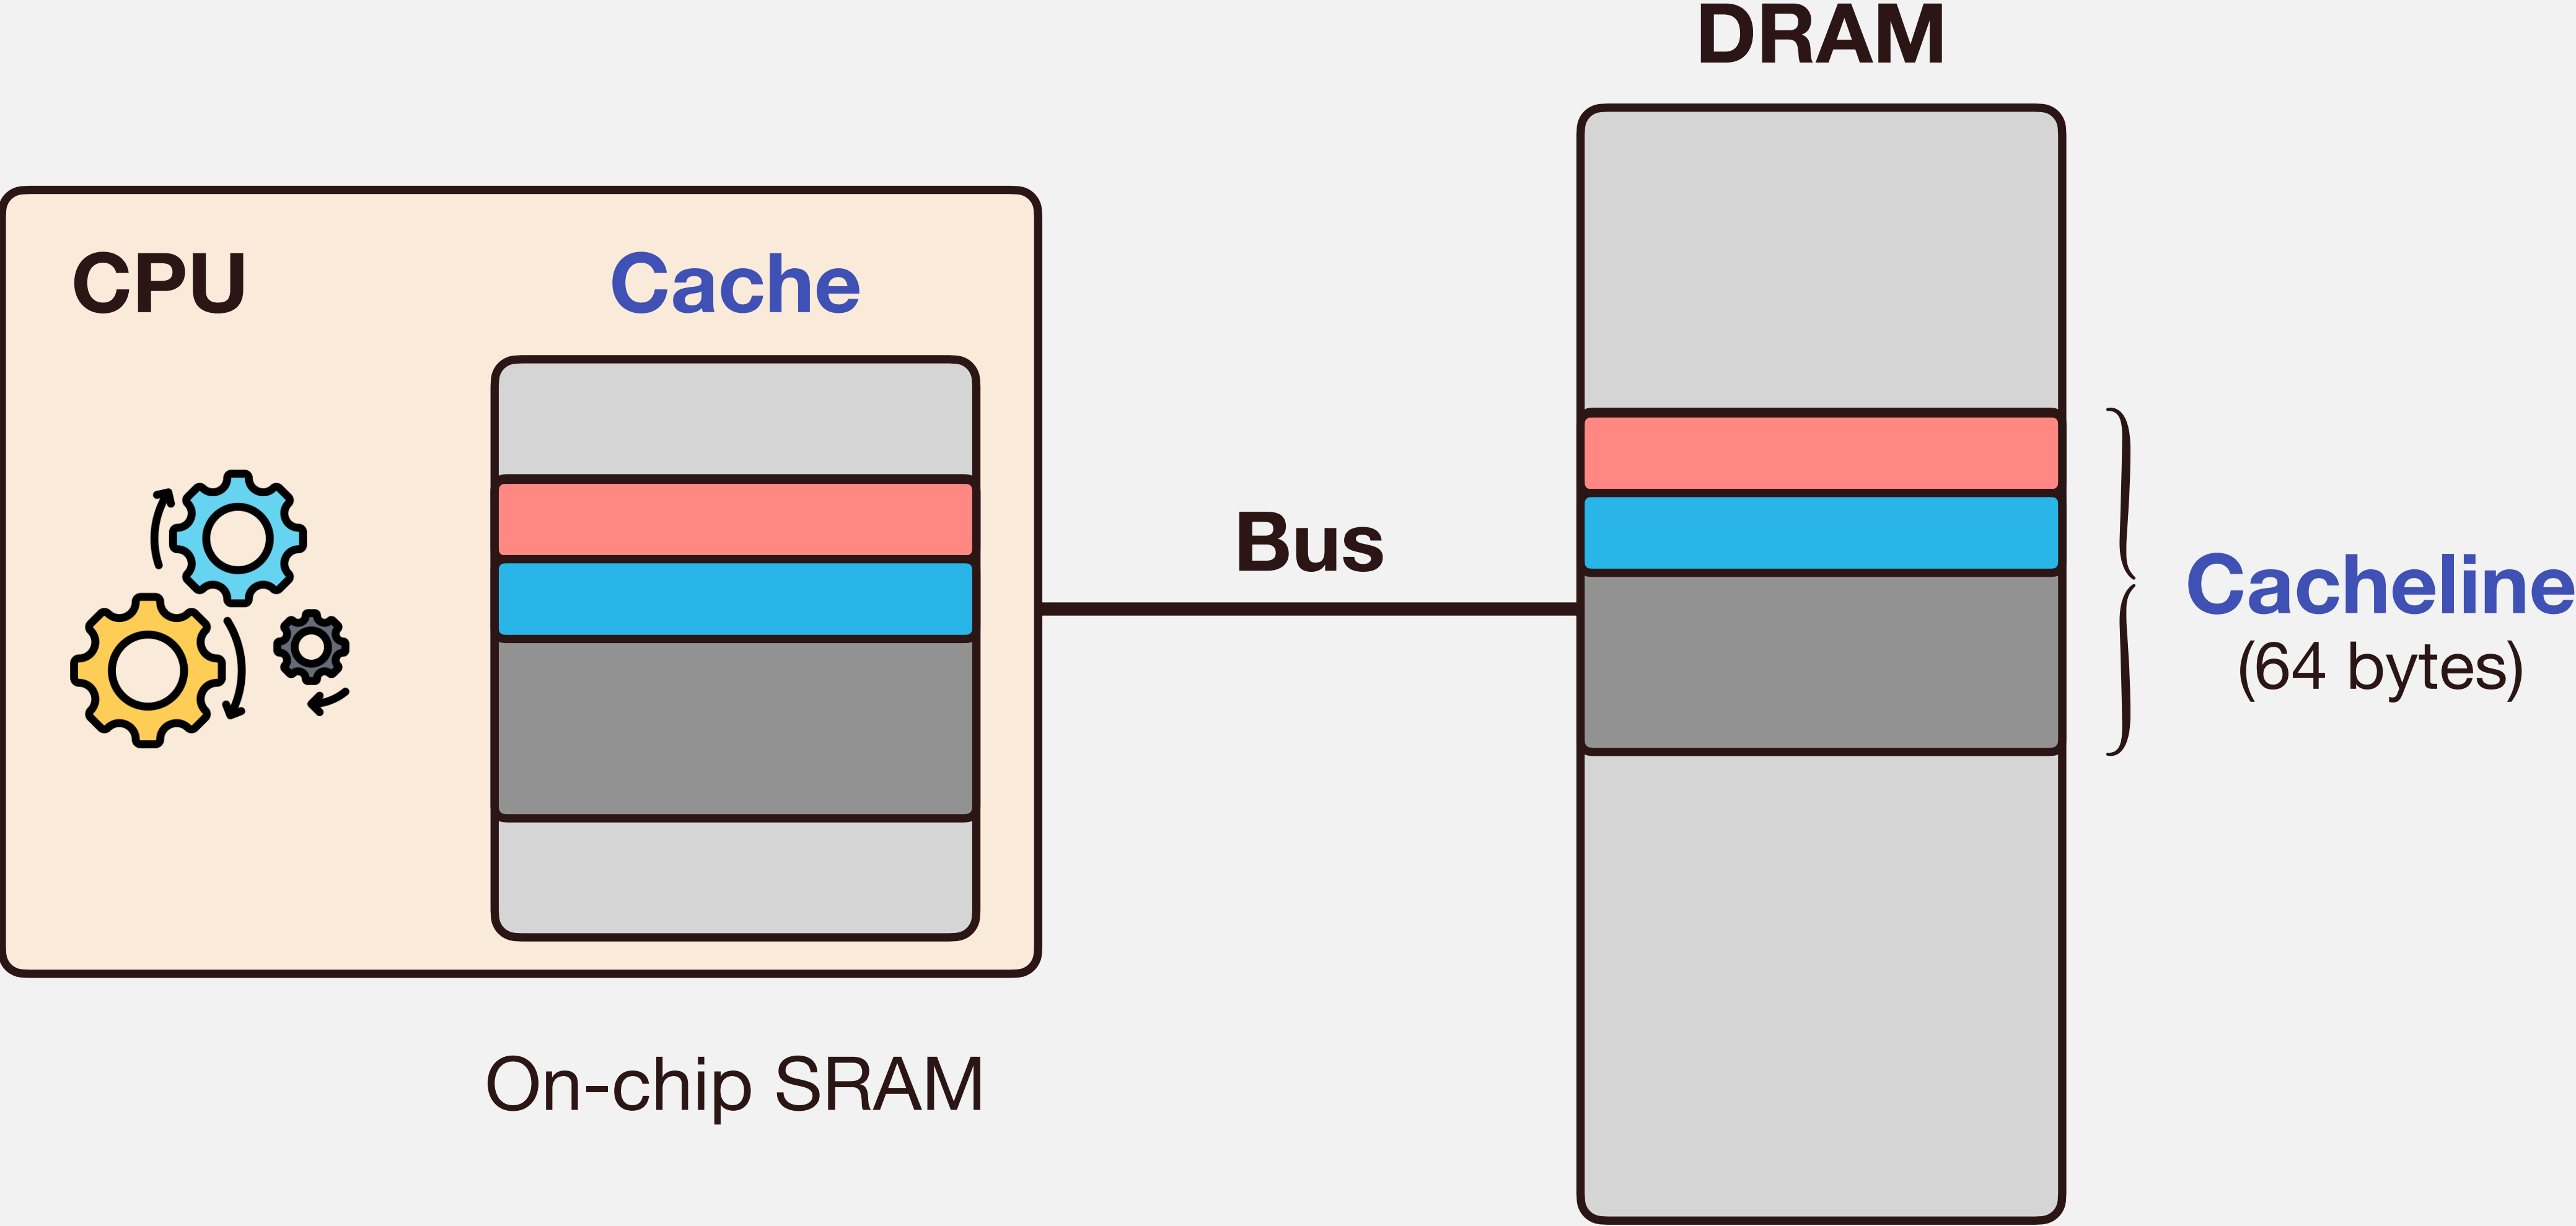
CPU Cache



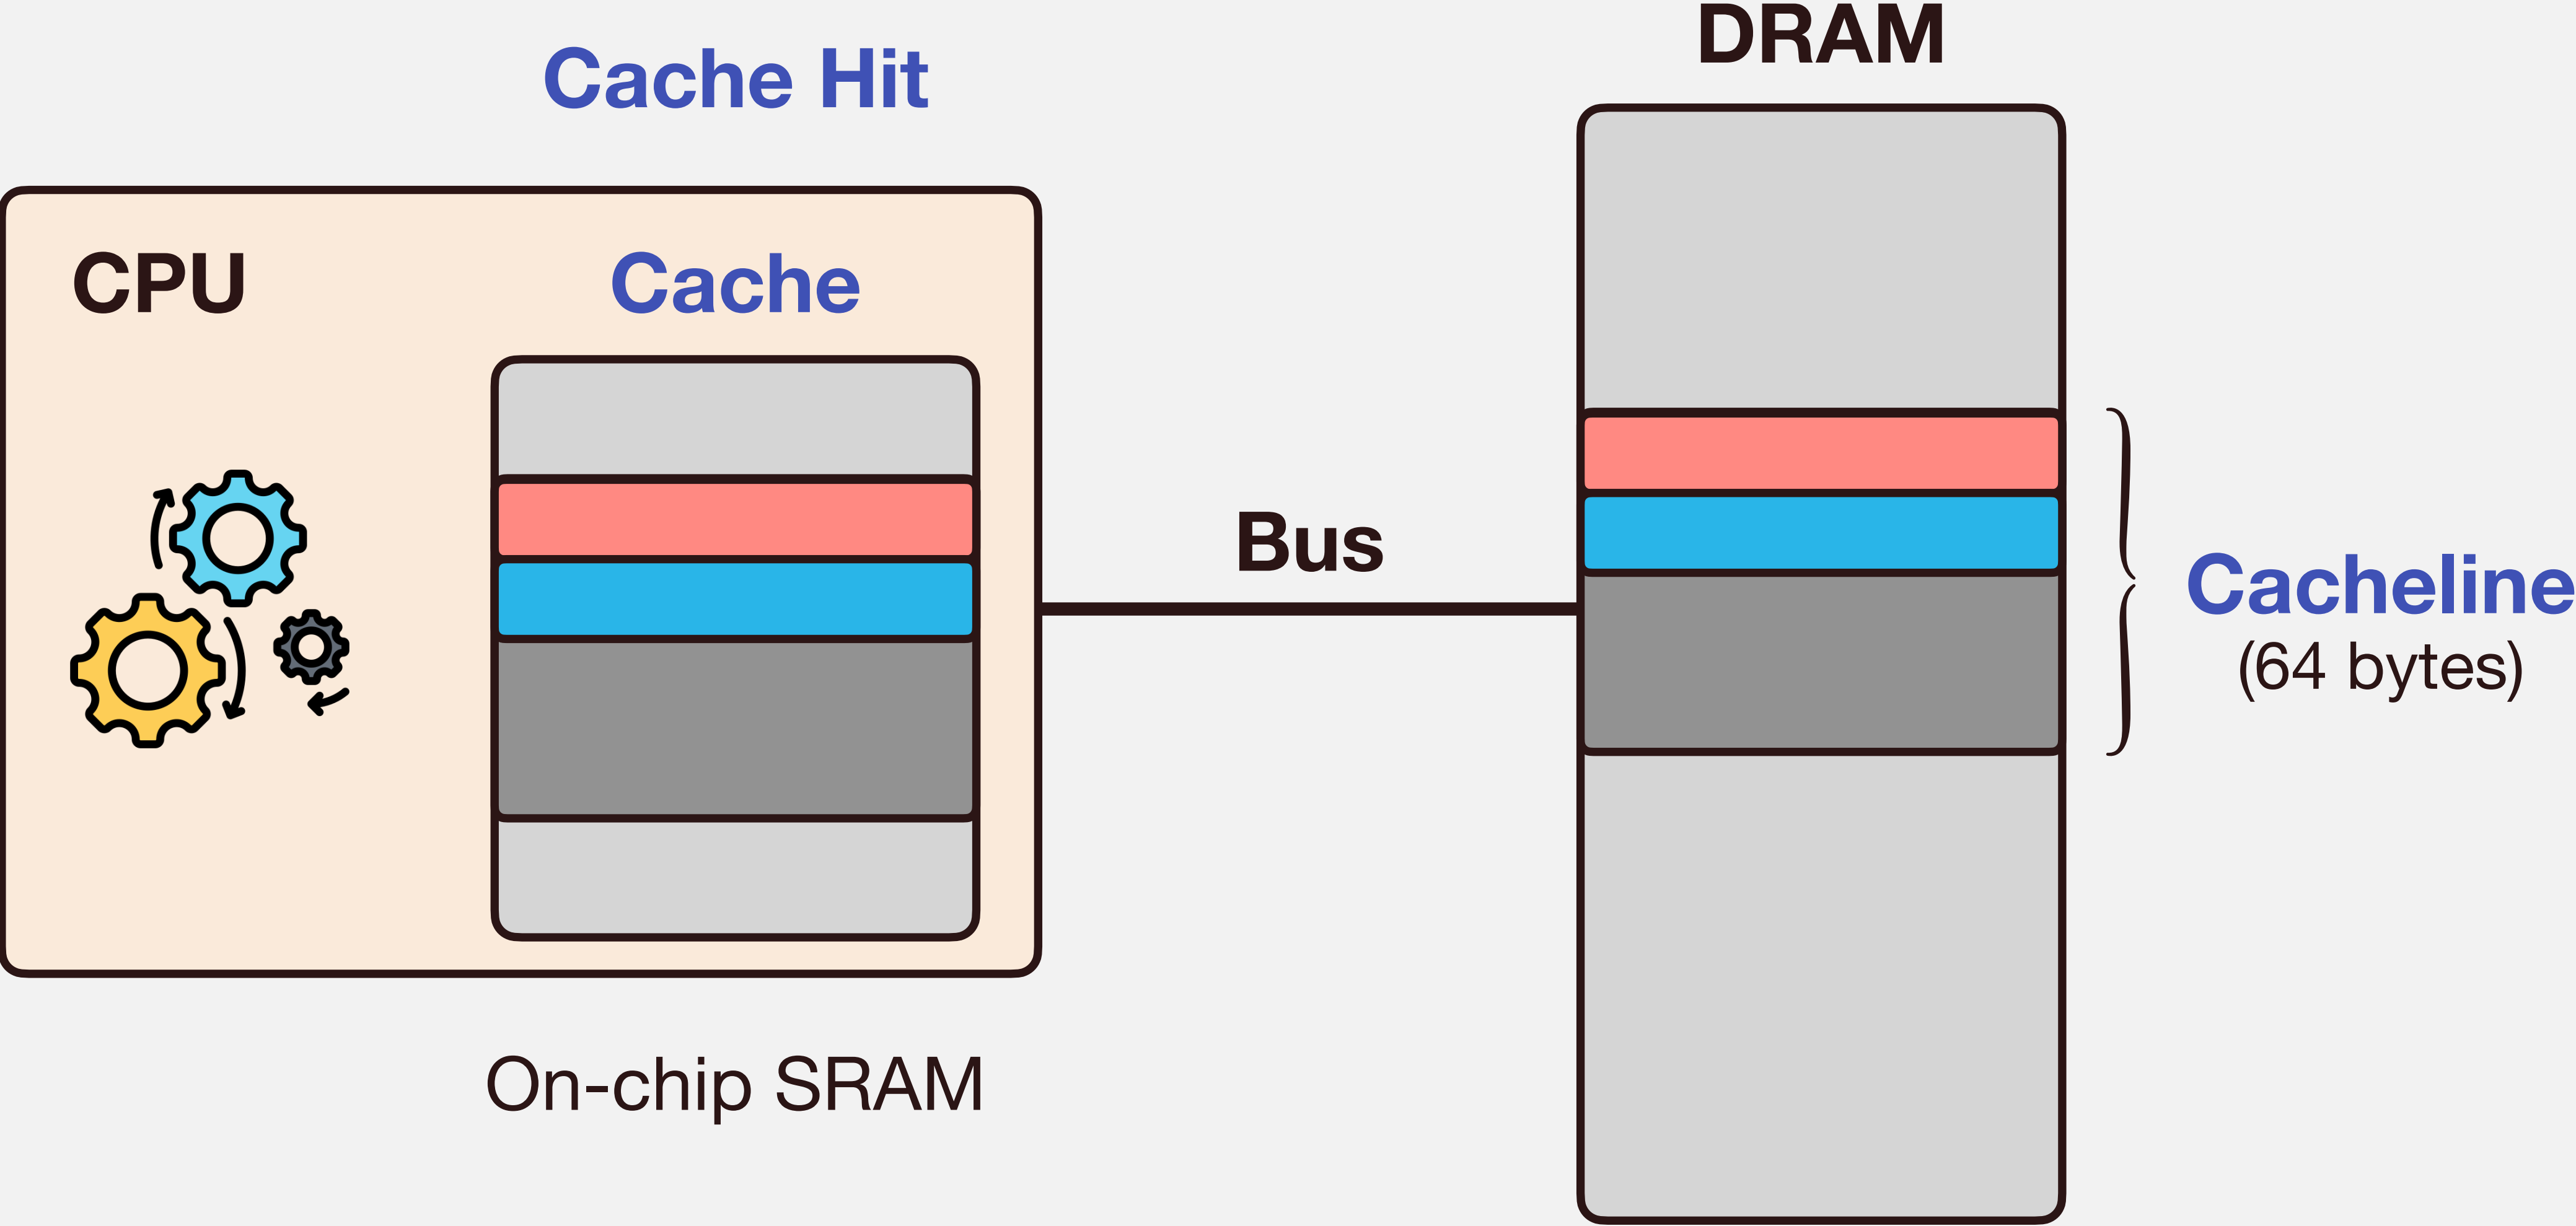
CPU Cache



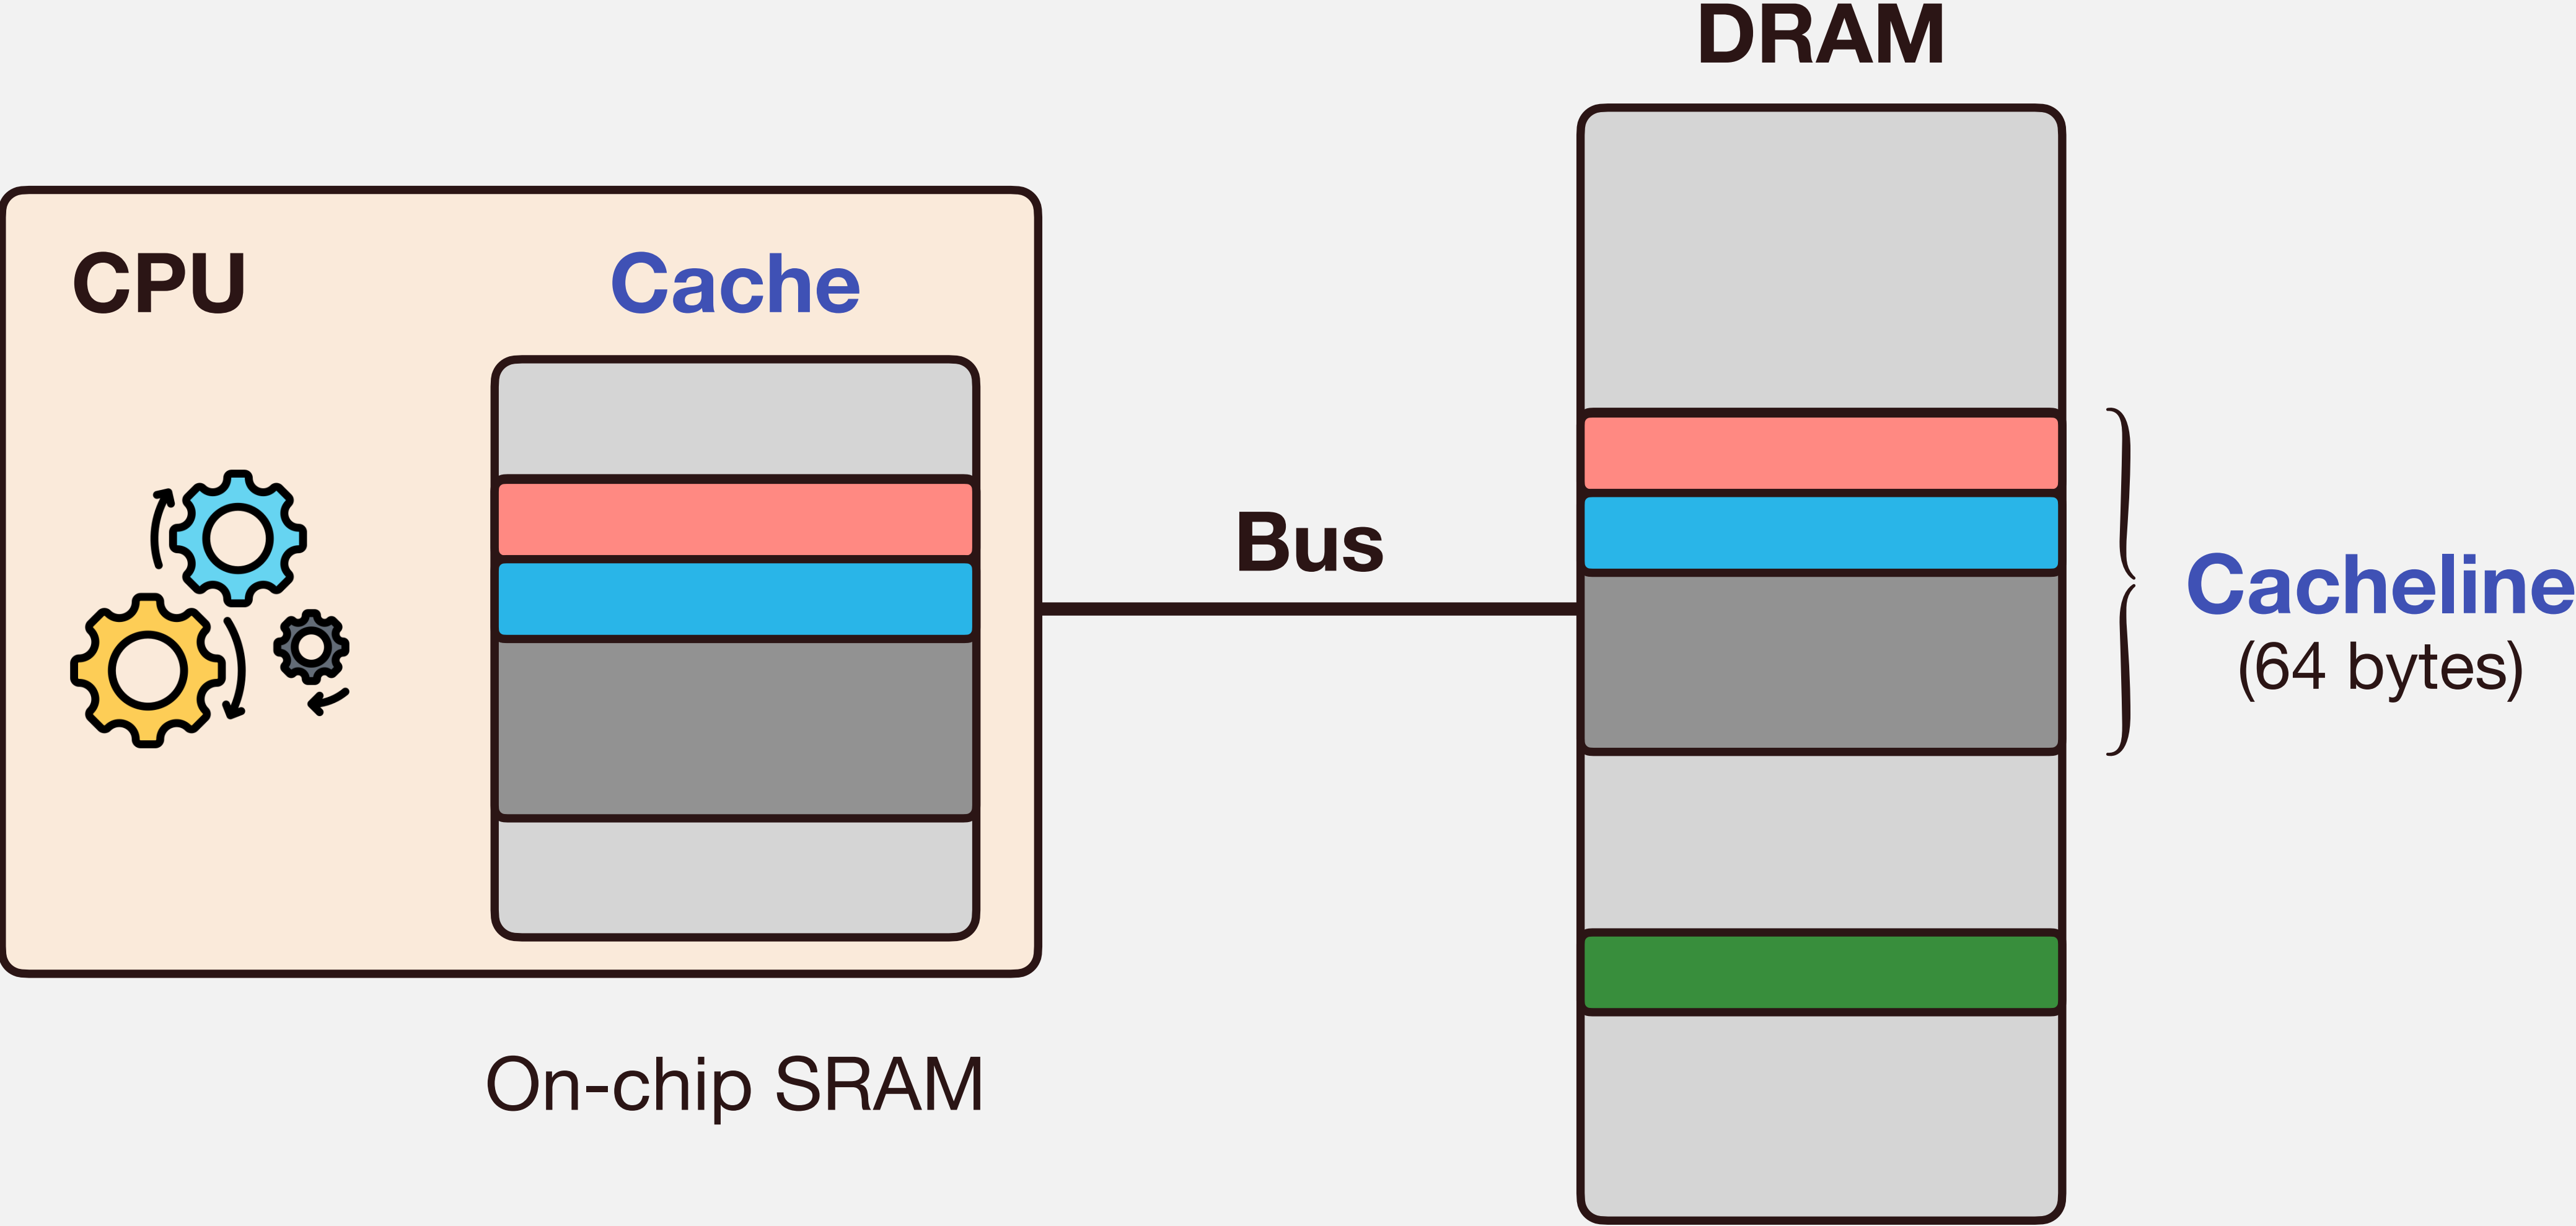
CPU Cache



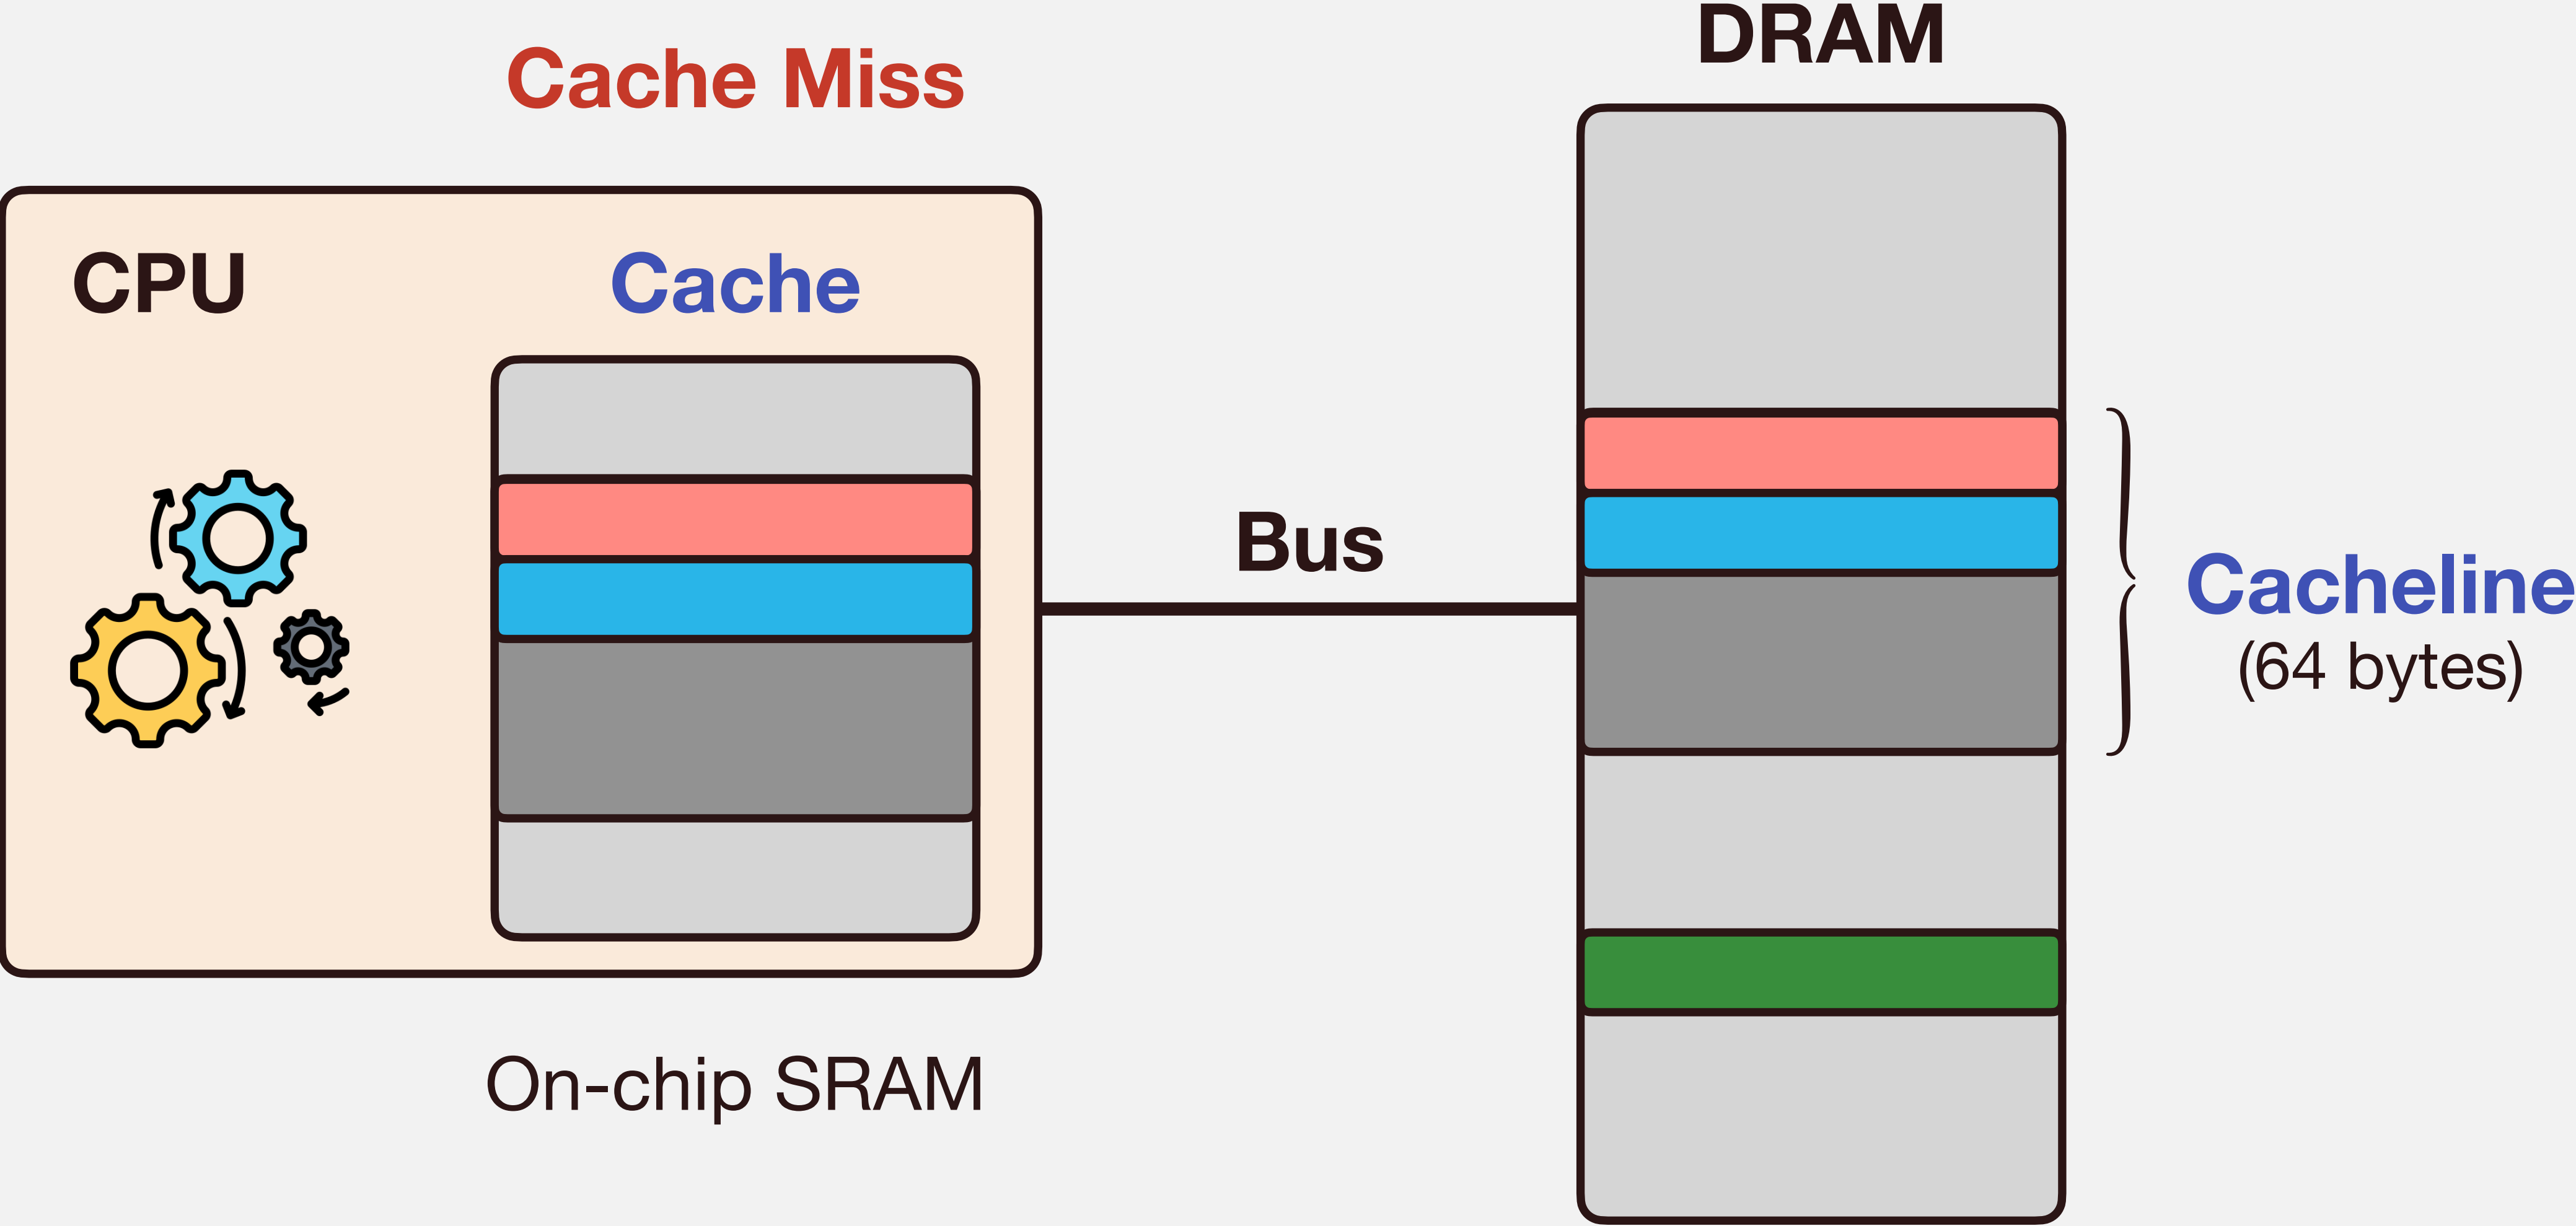
CPU Cache



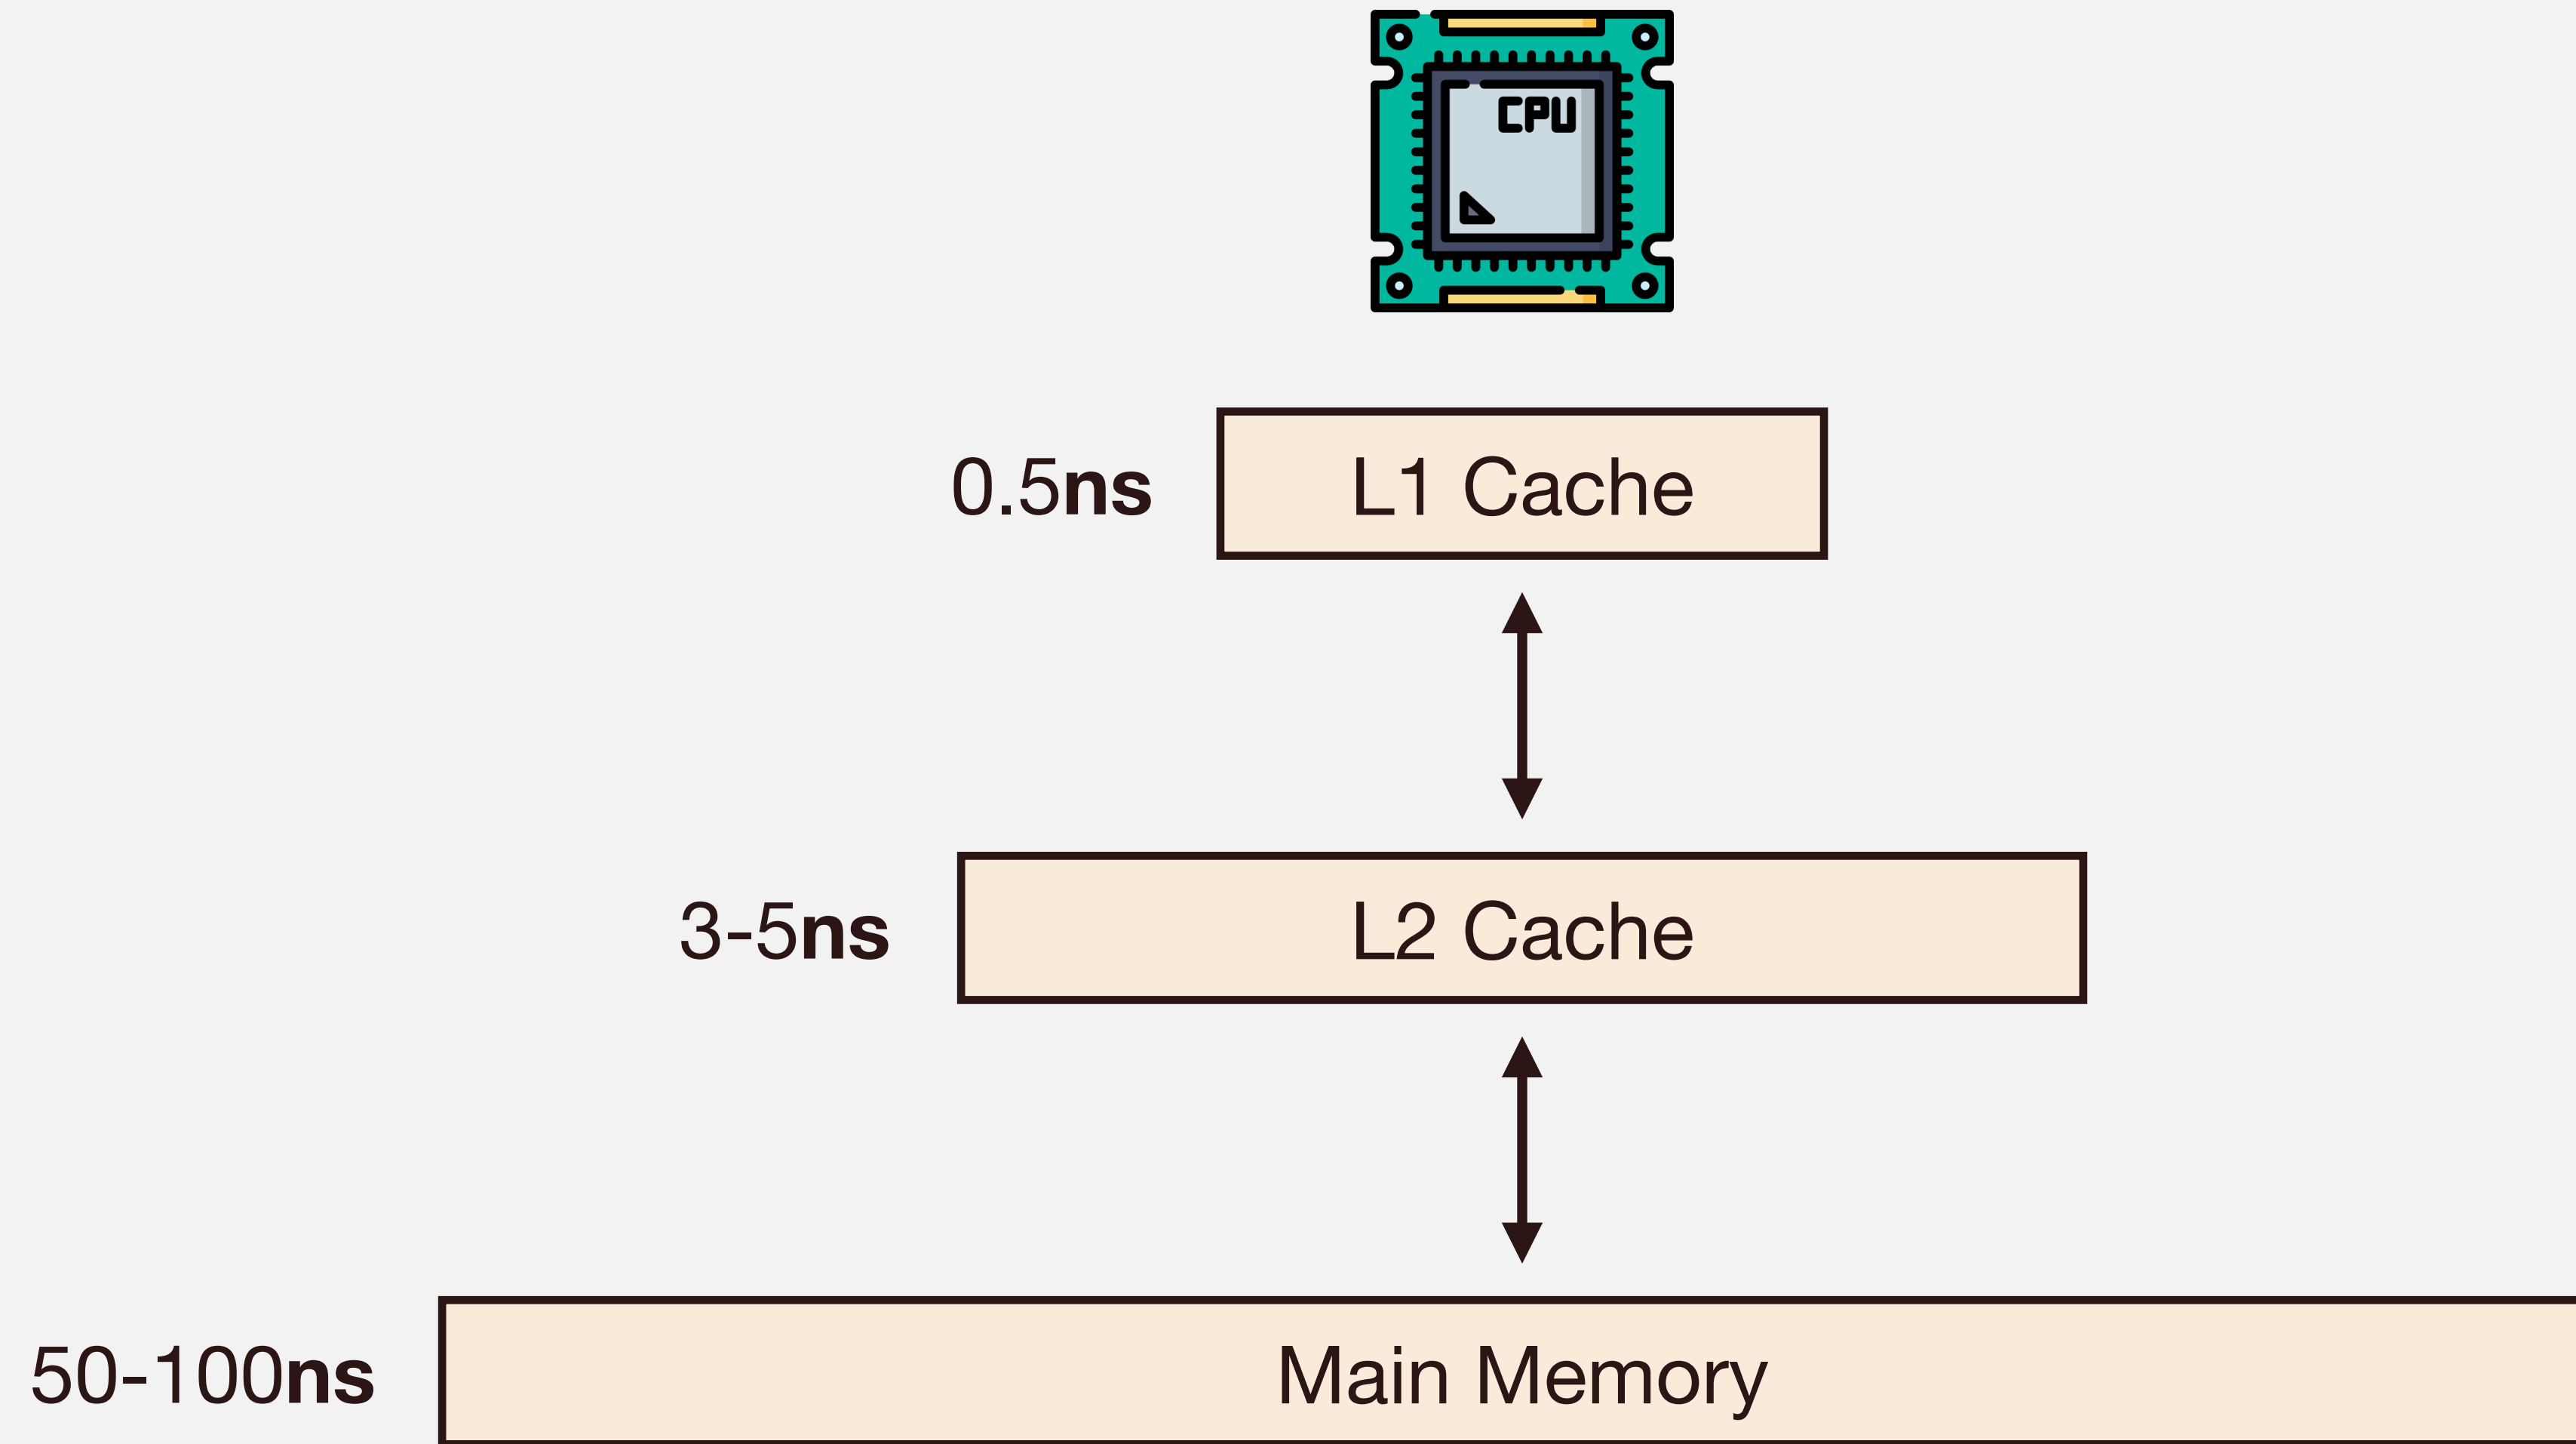
CPU Cache



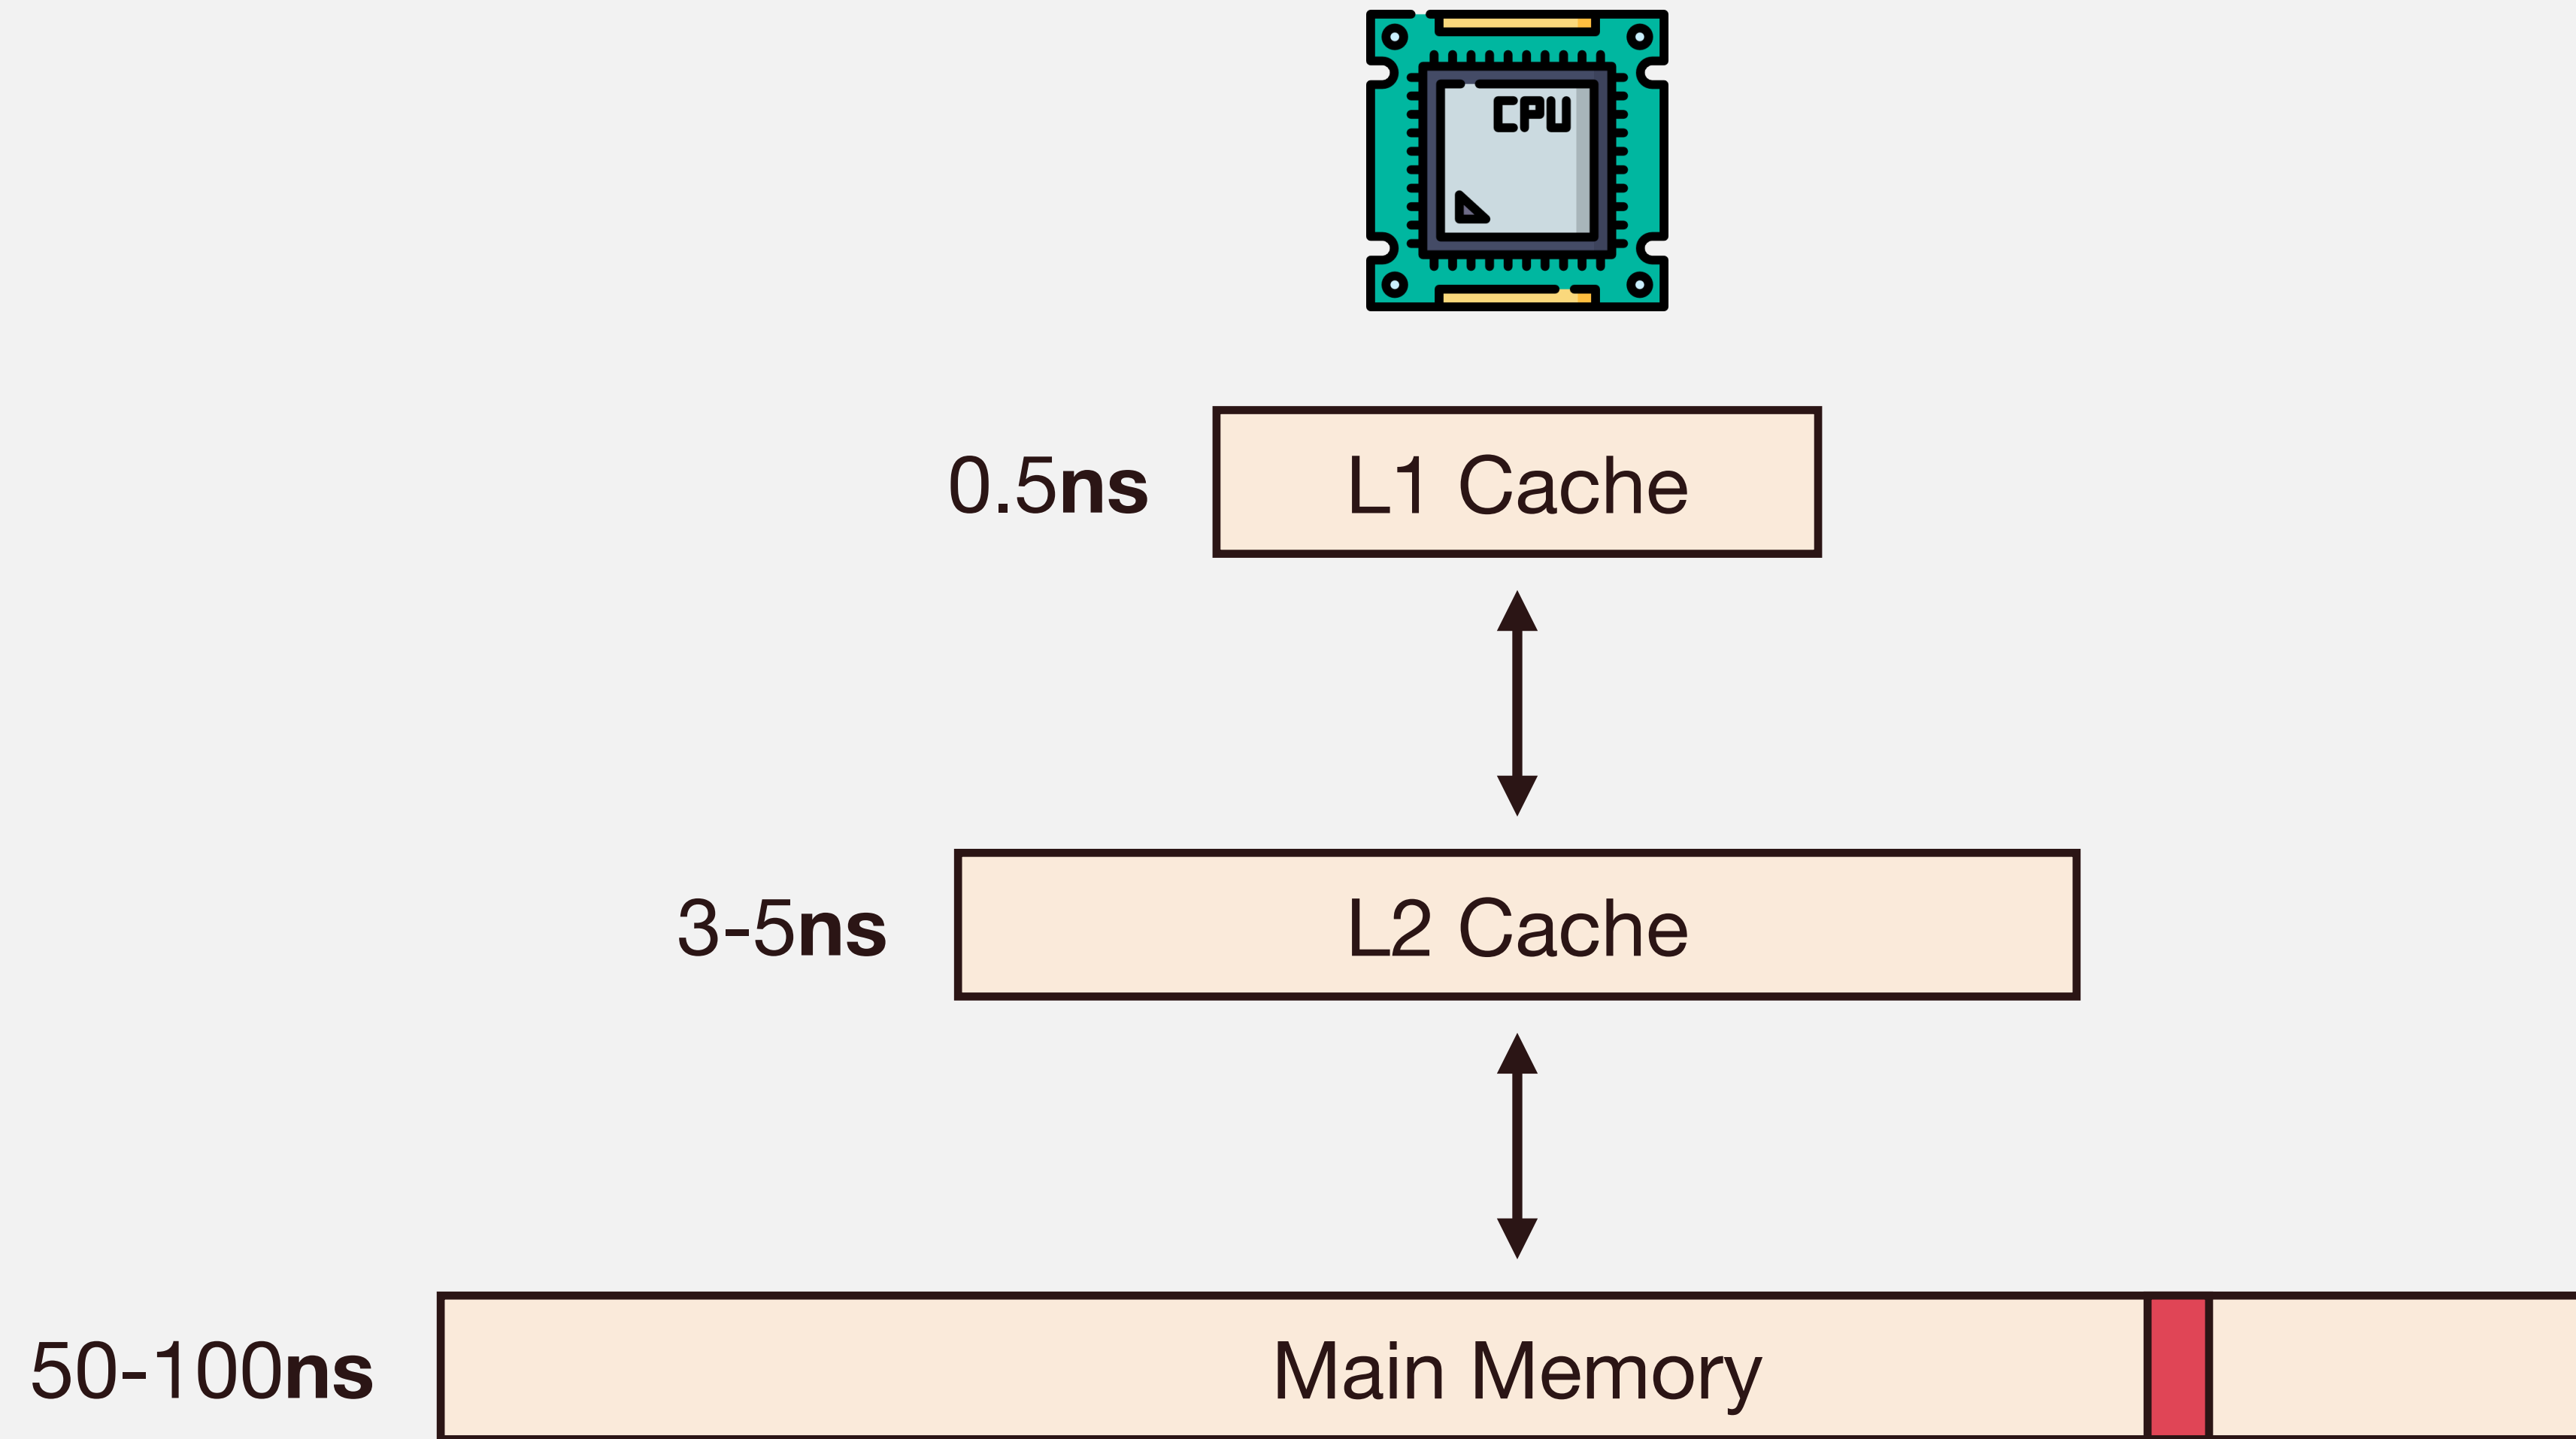
CPU Cache



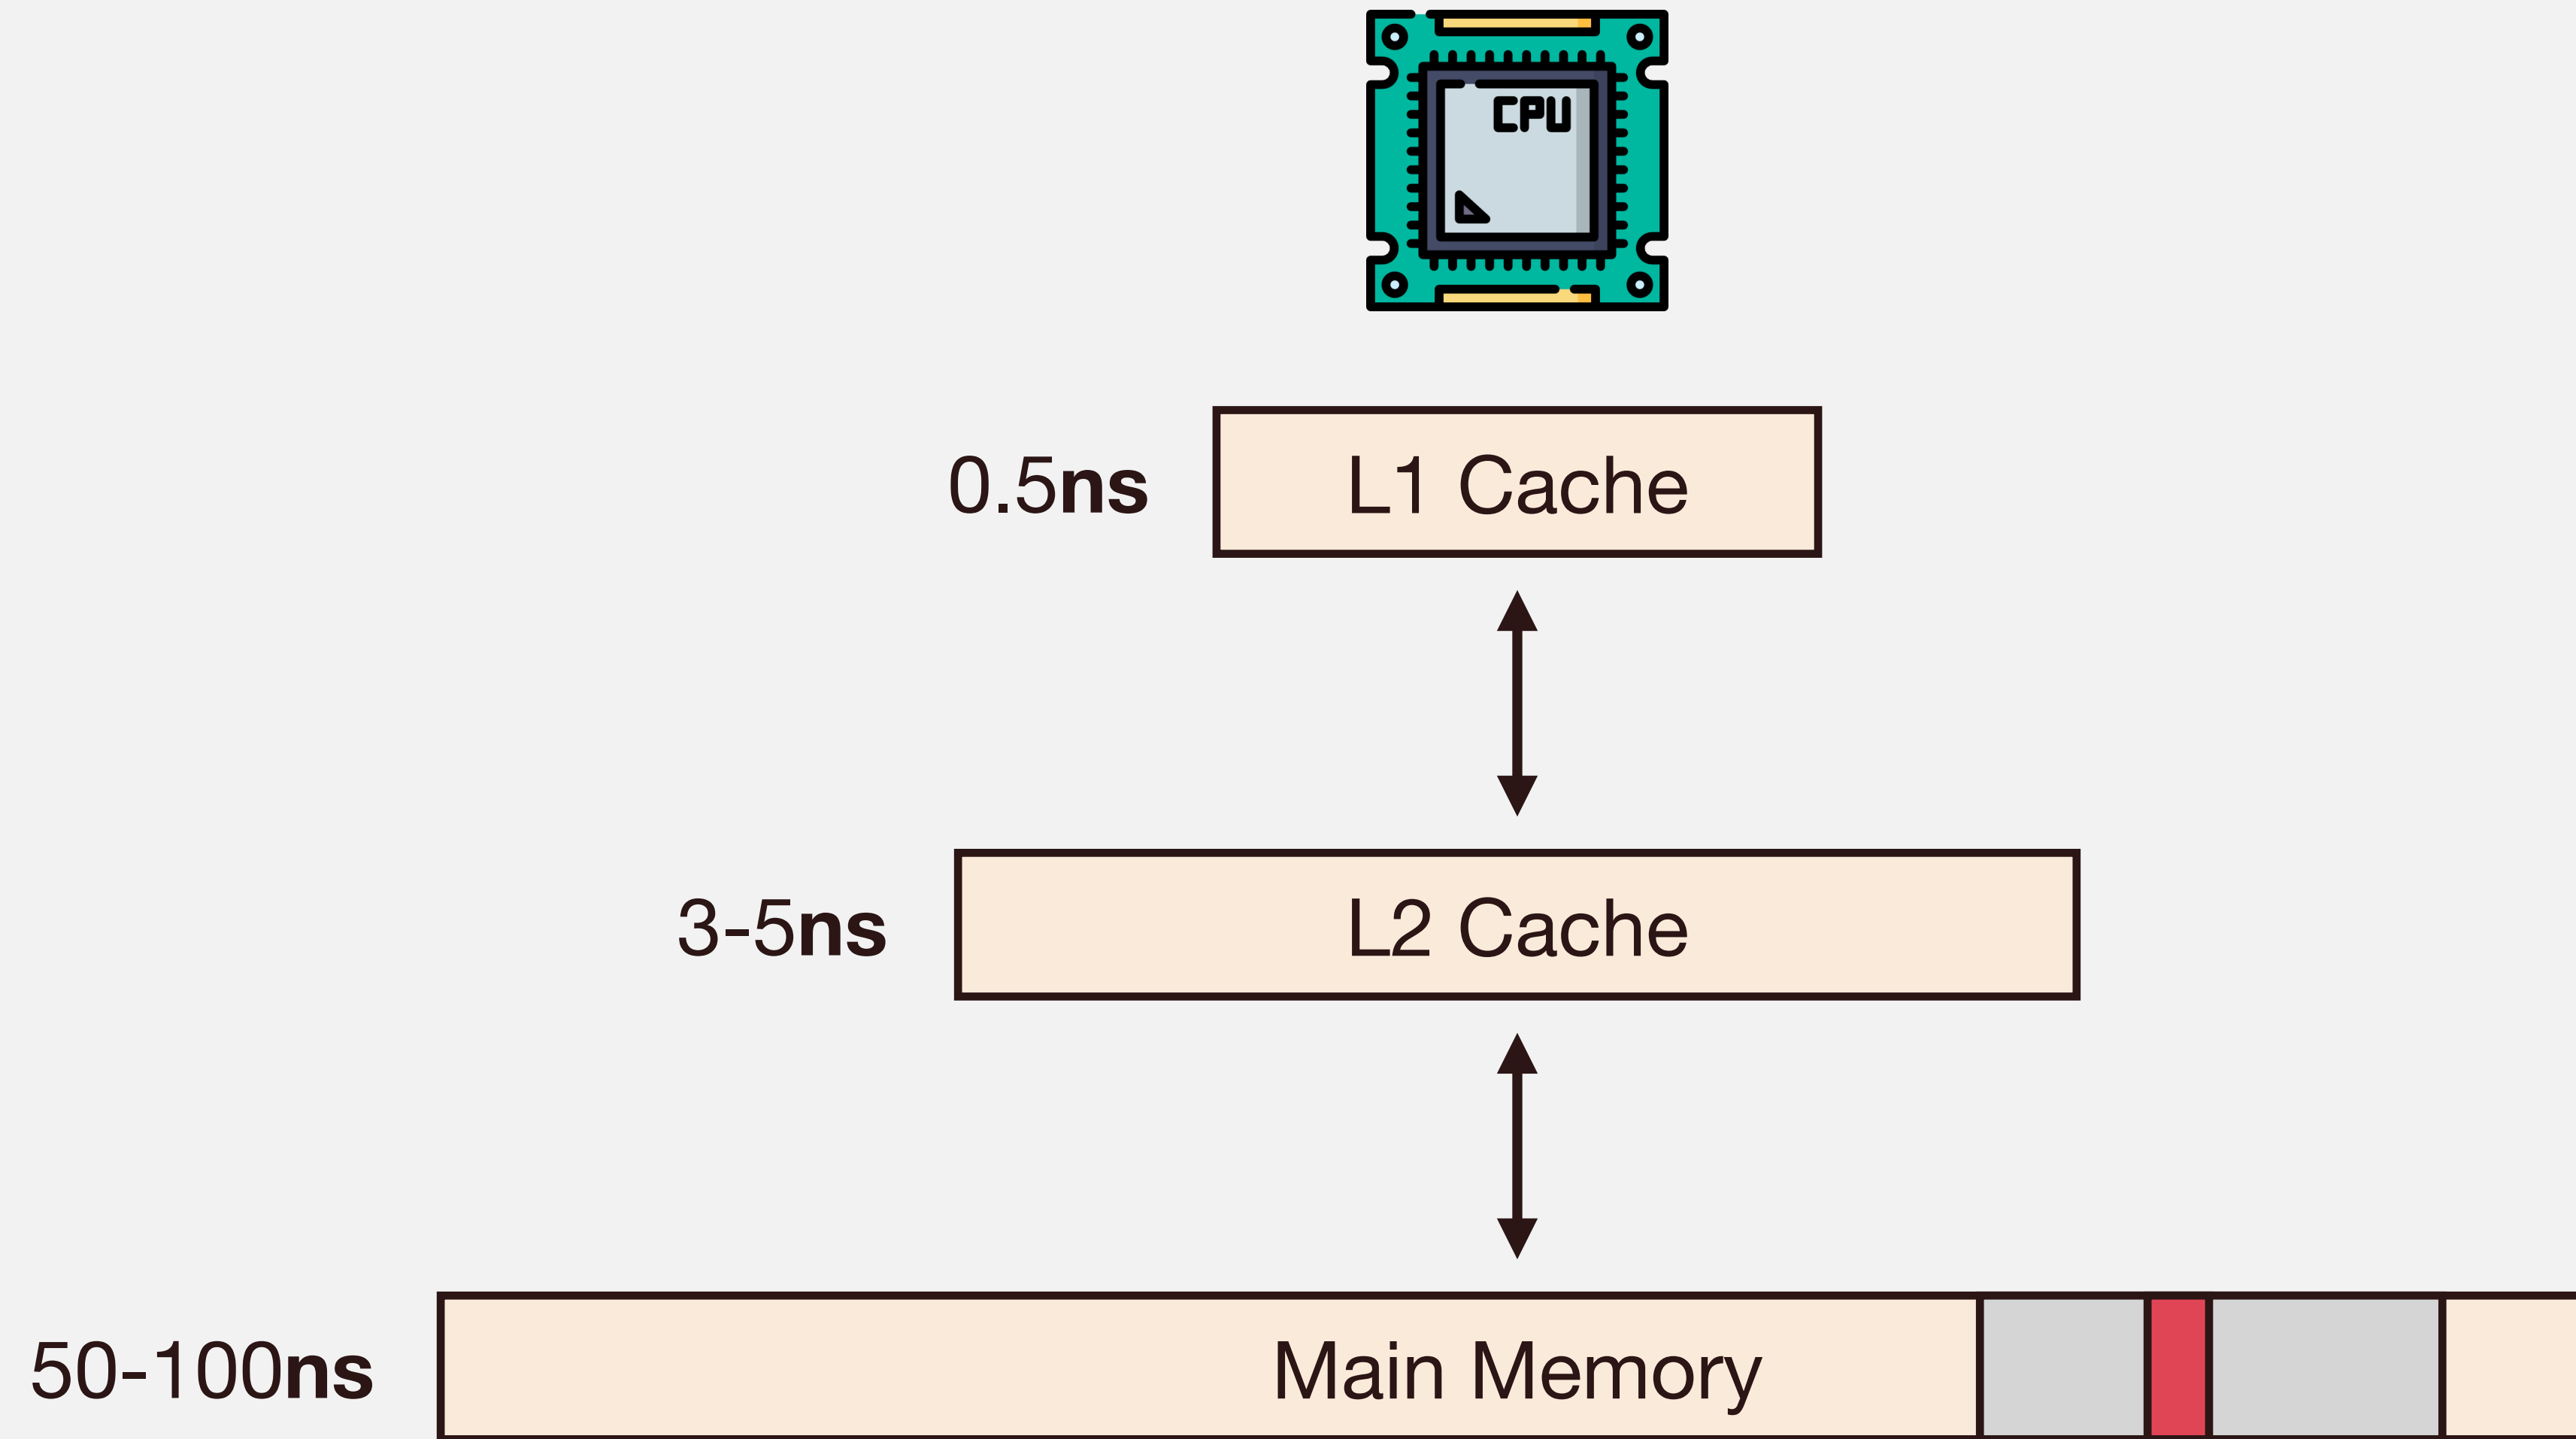
Cacheline



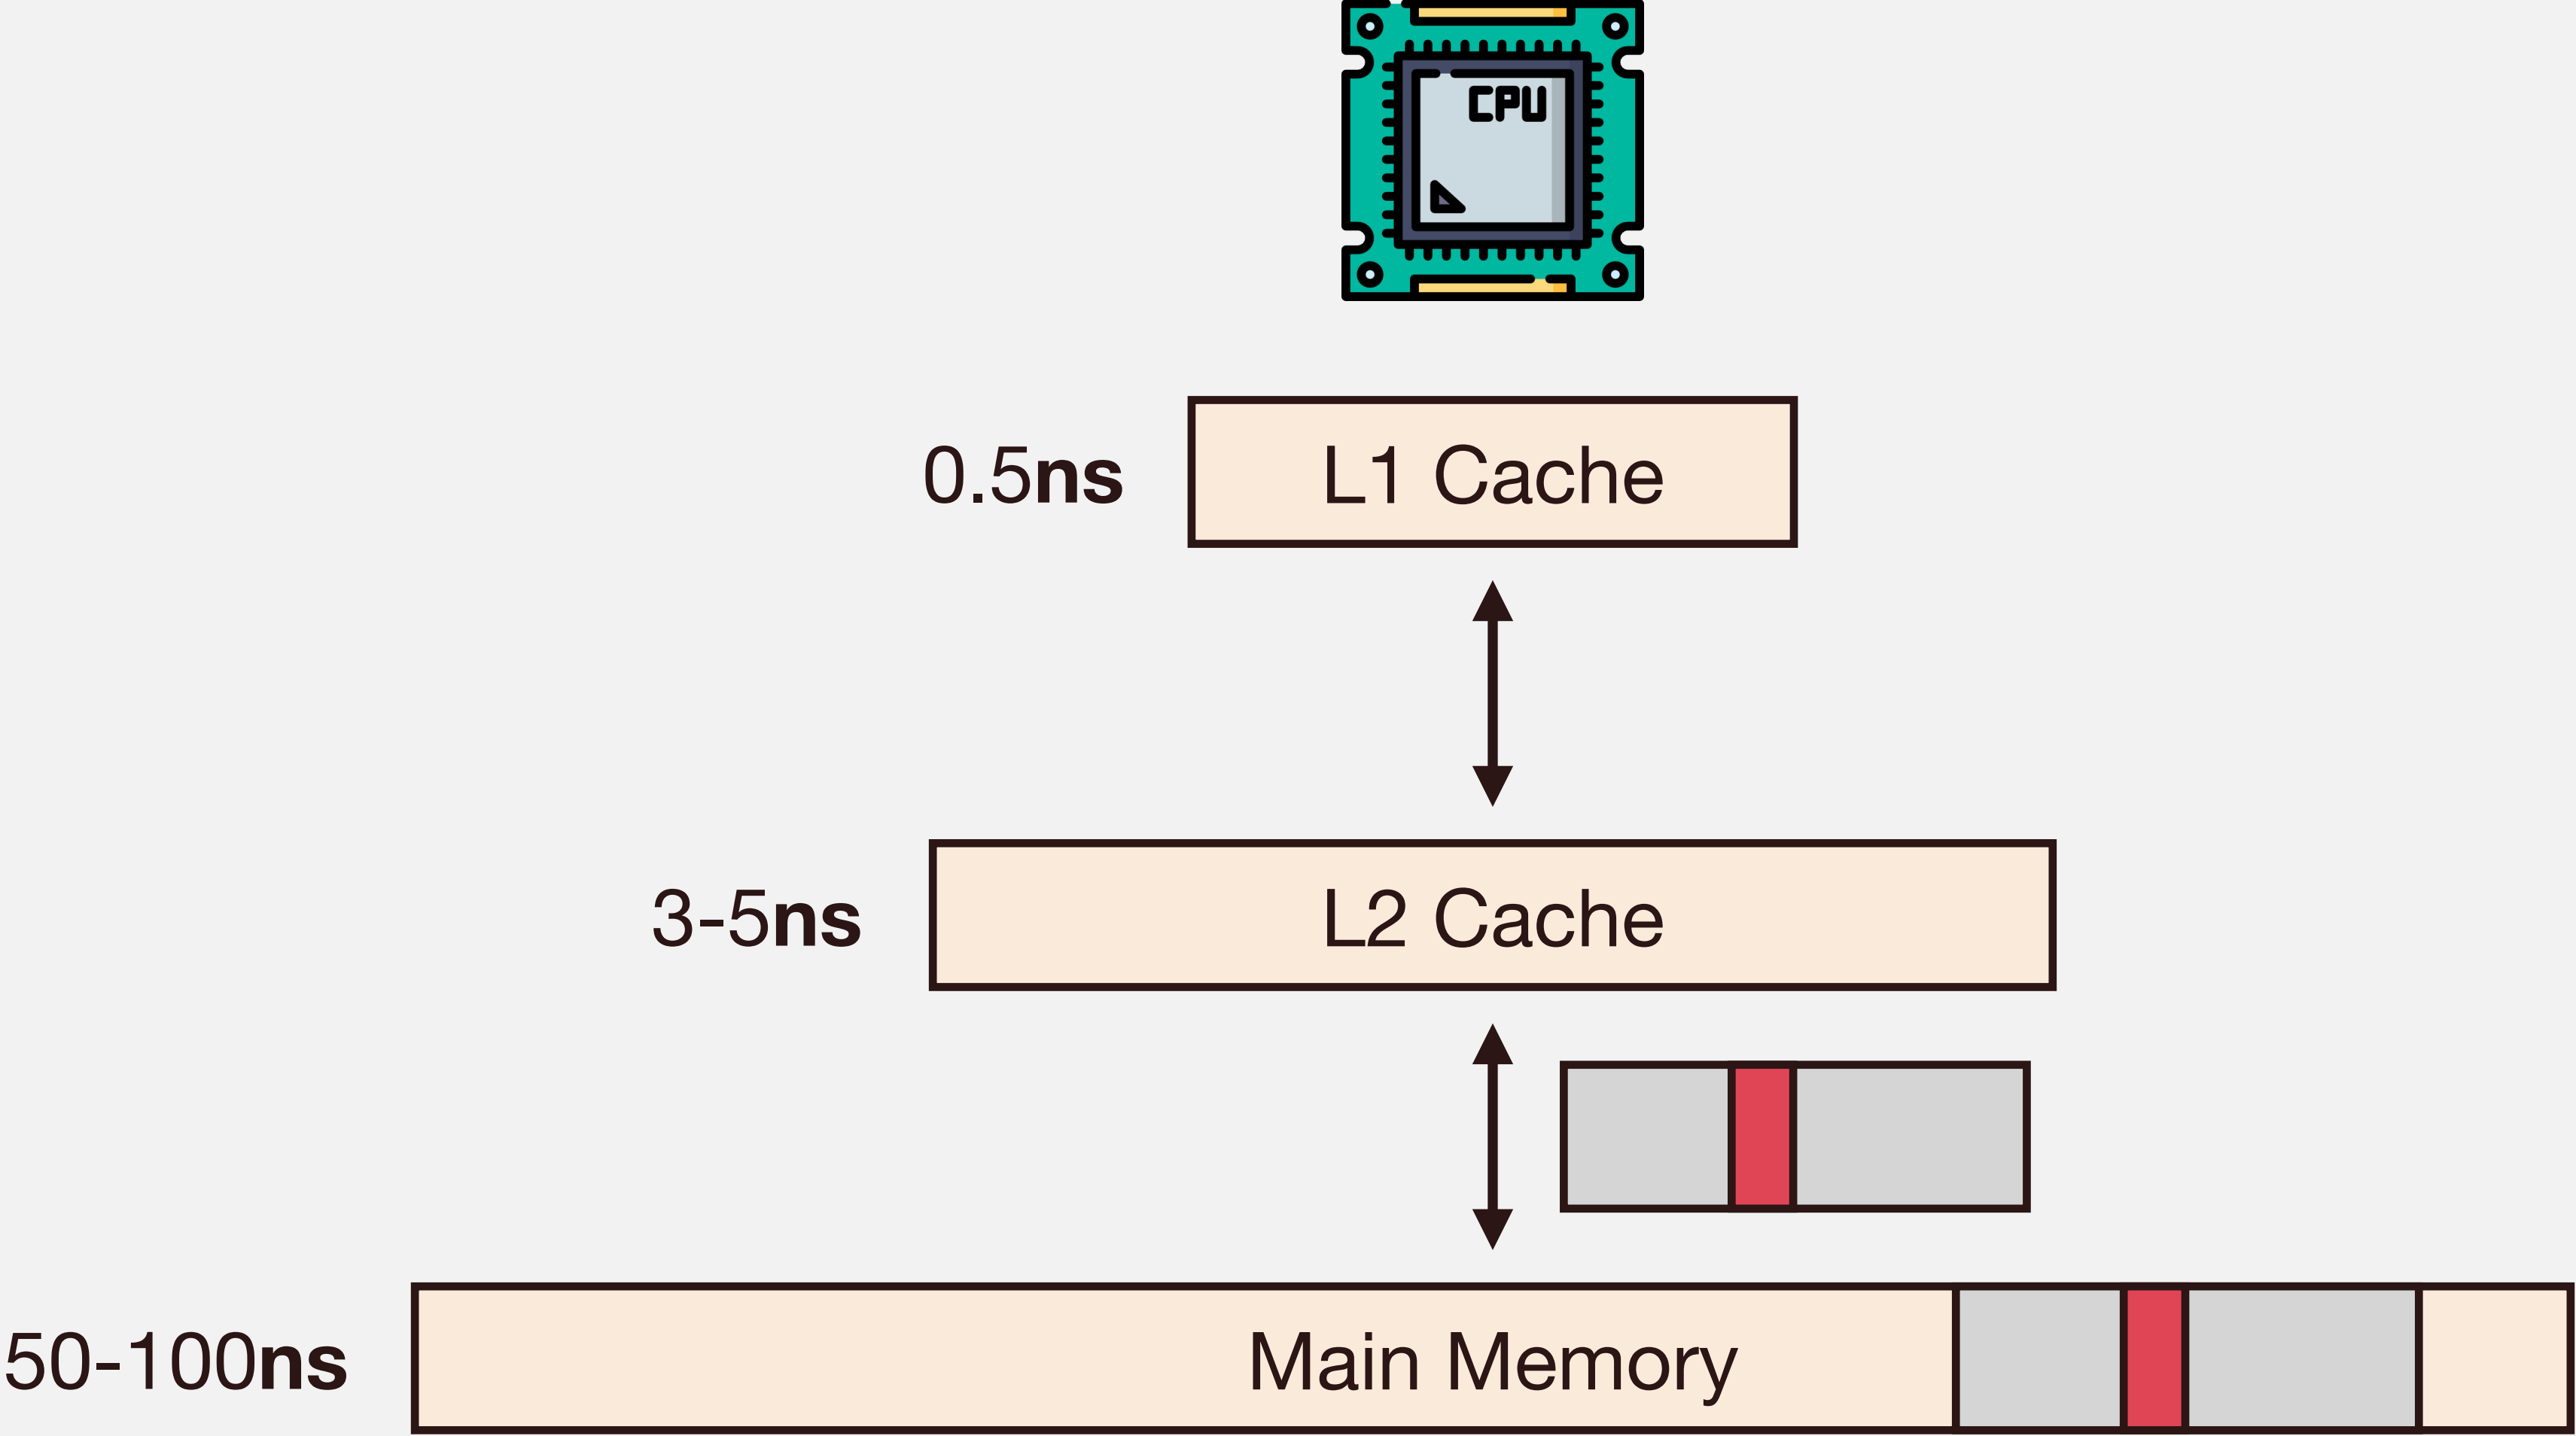
Cacheline



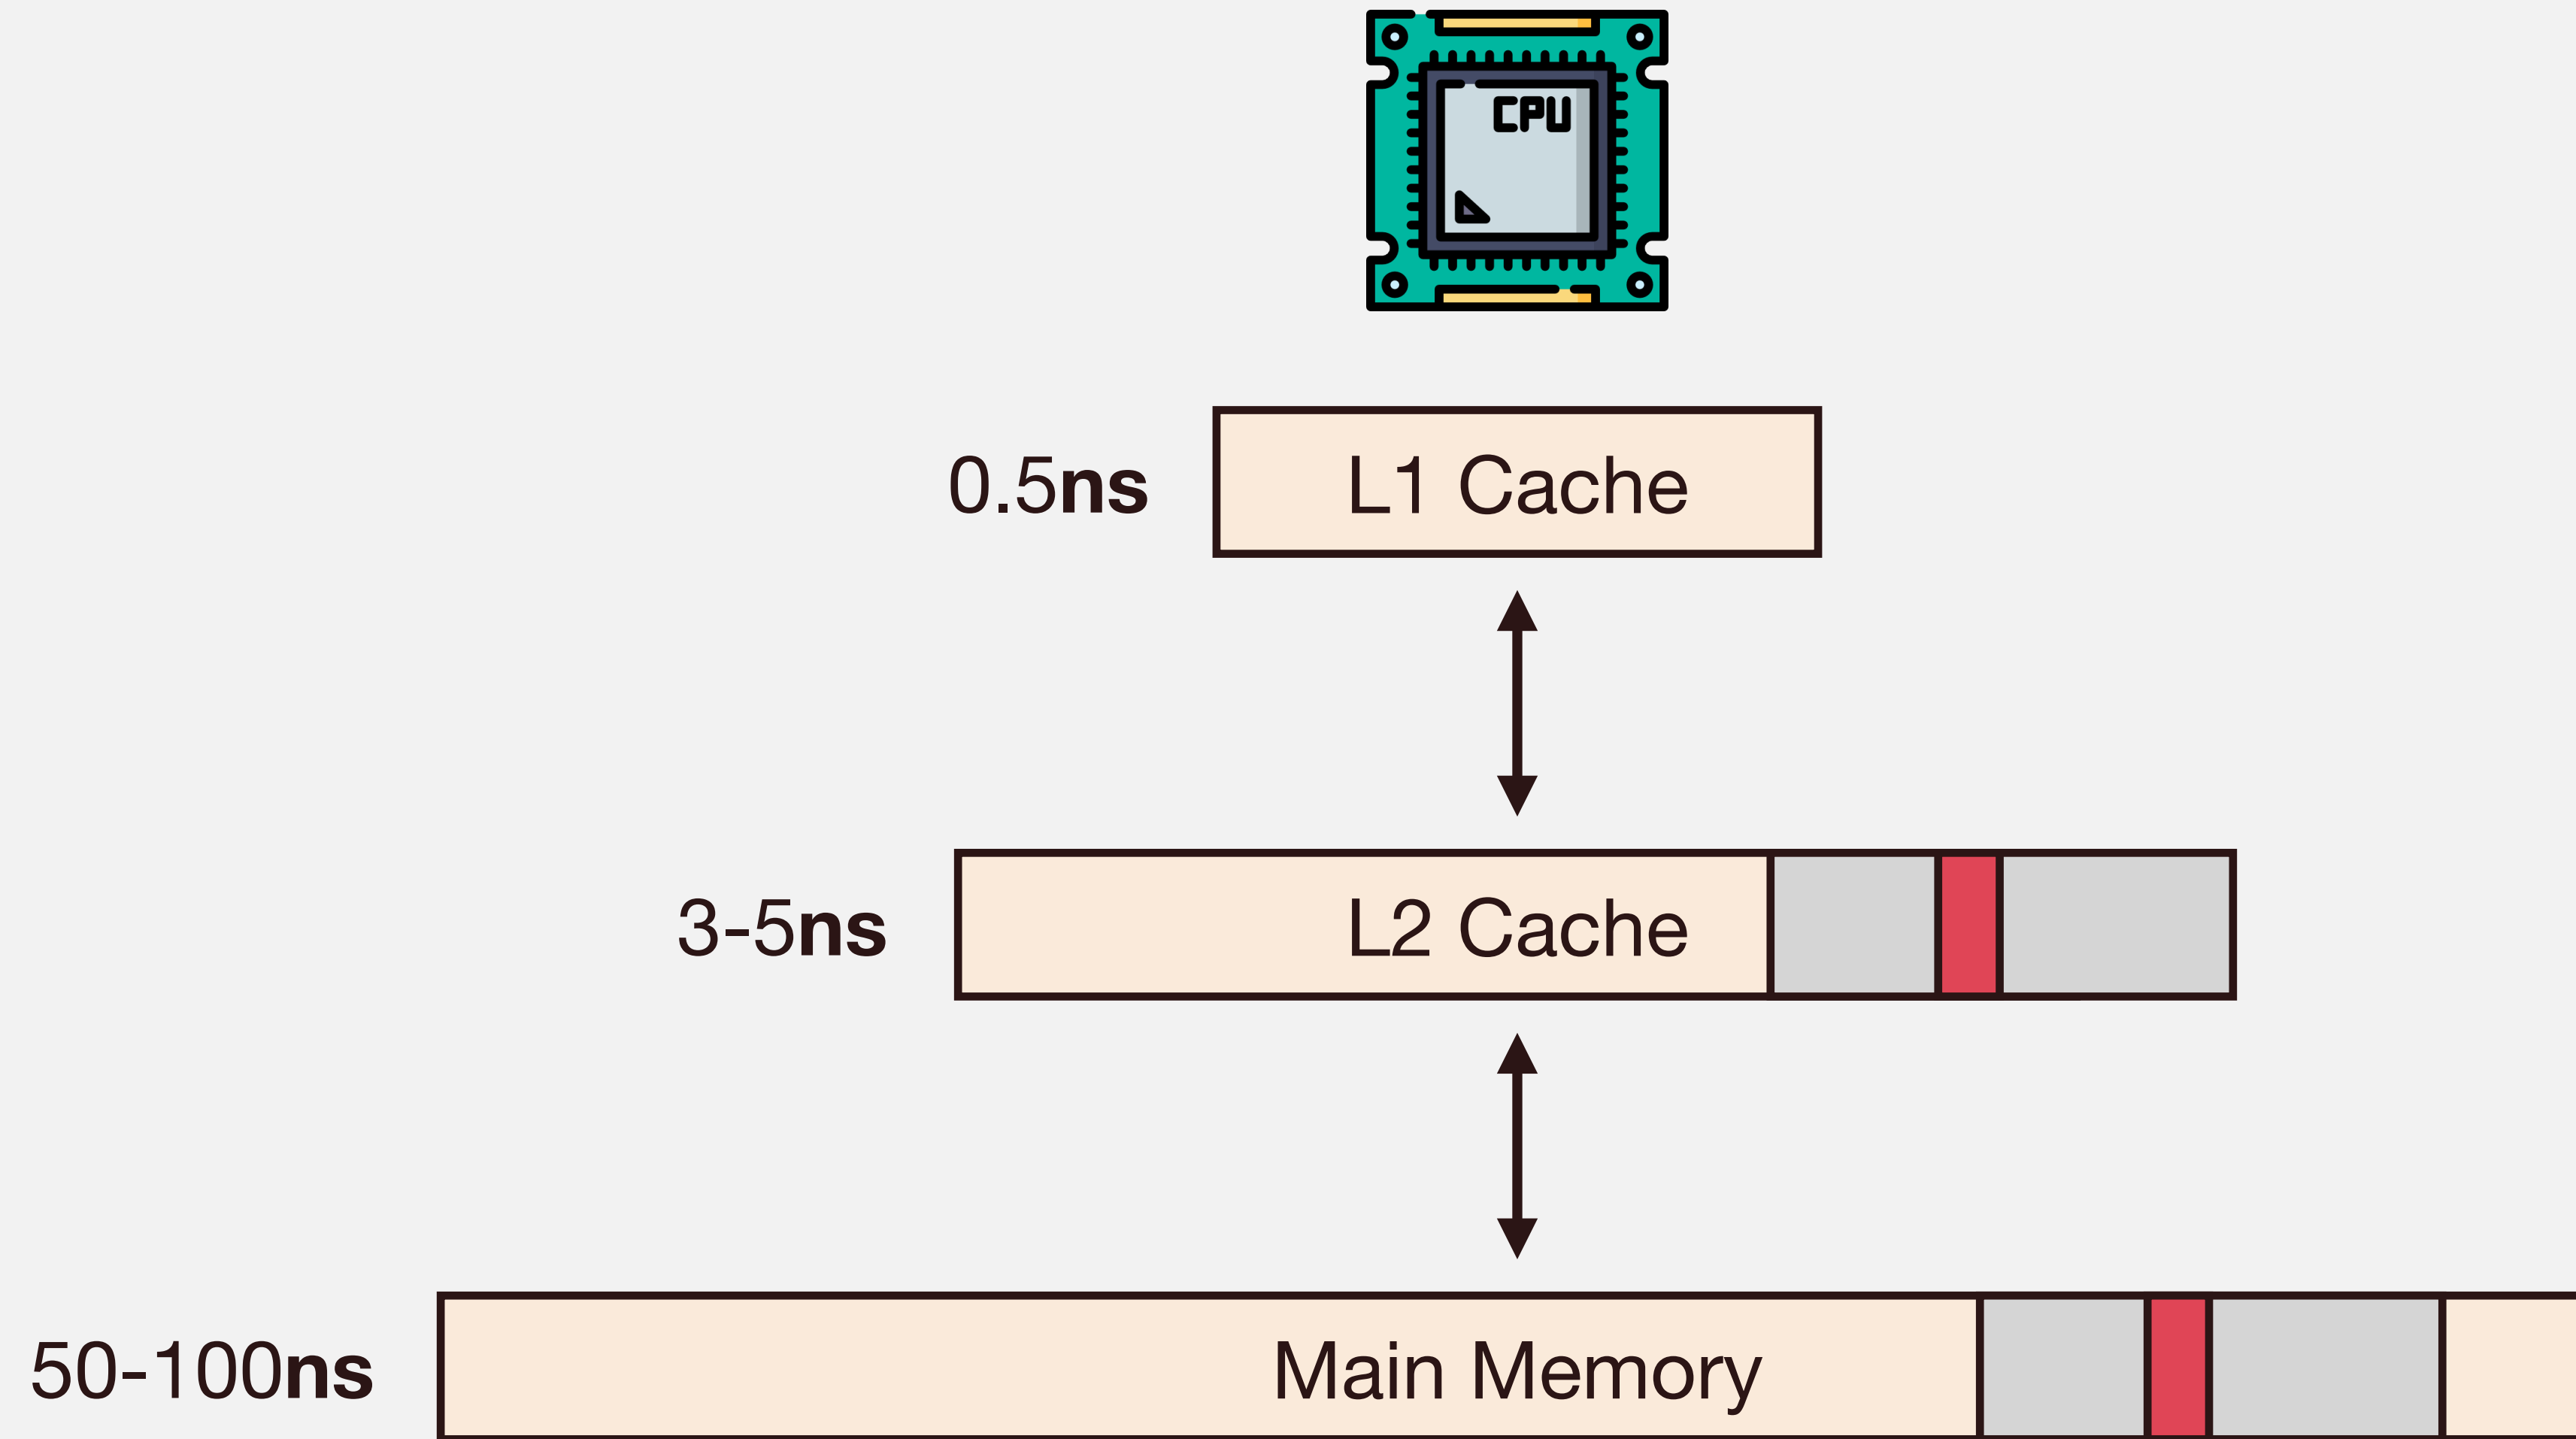
Cacheline



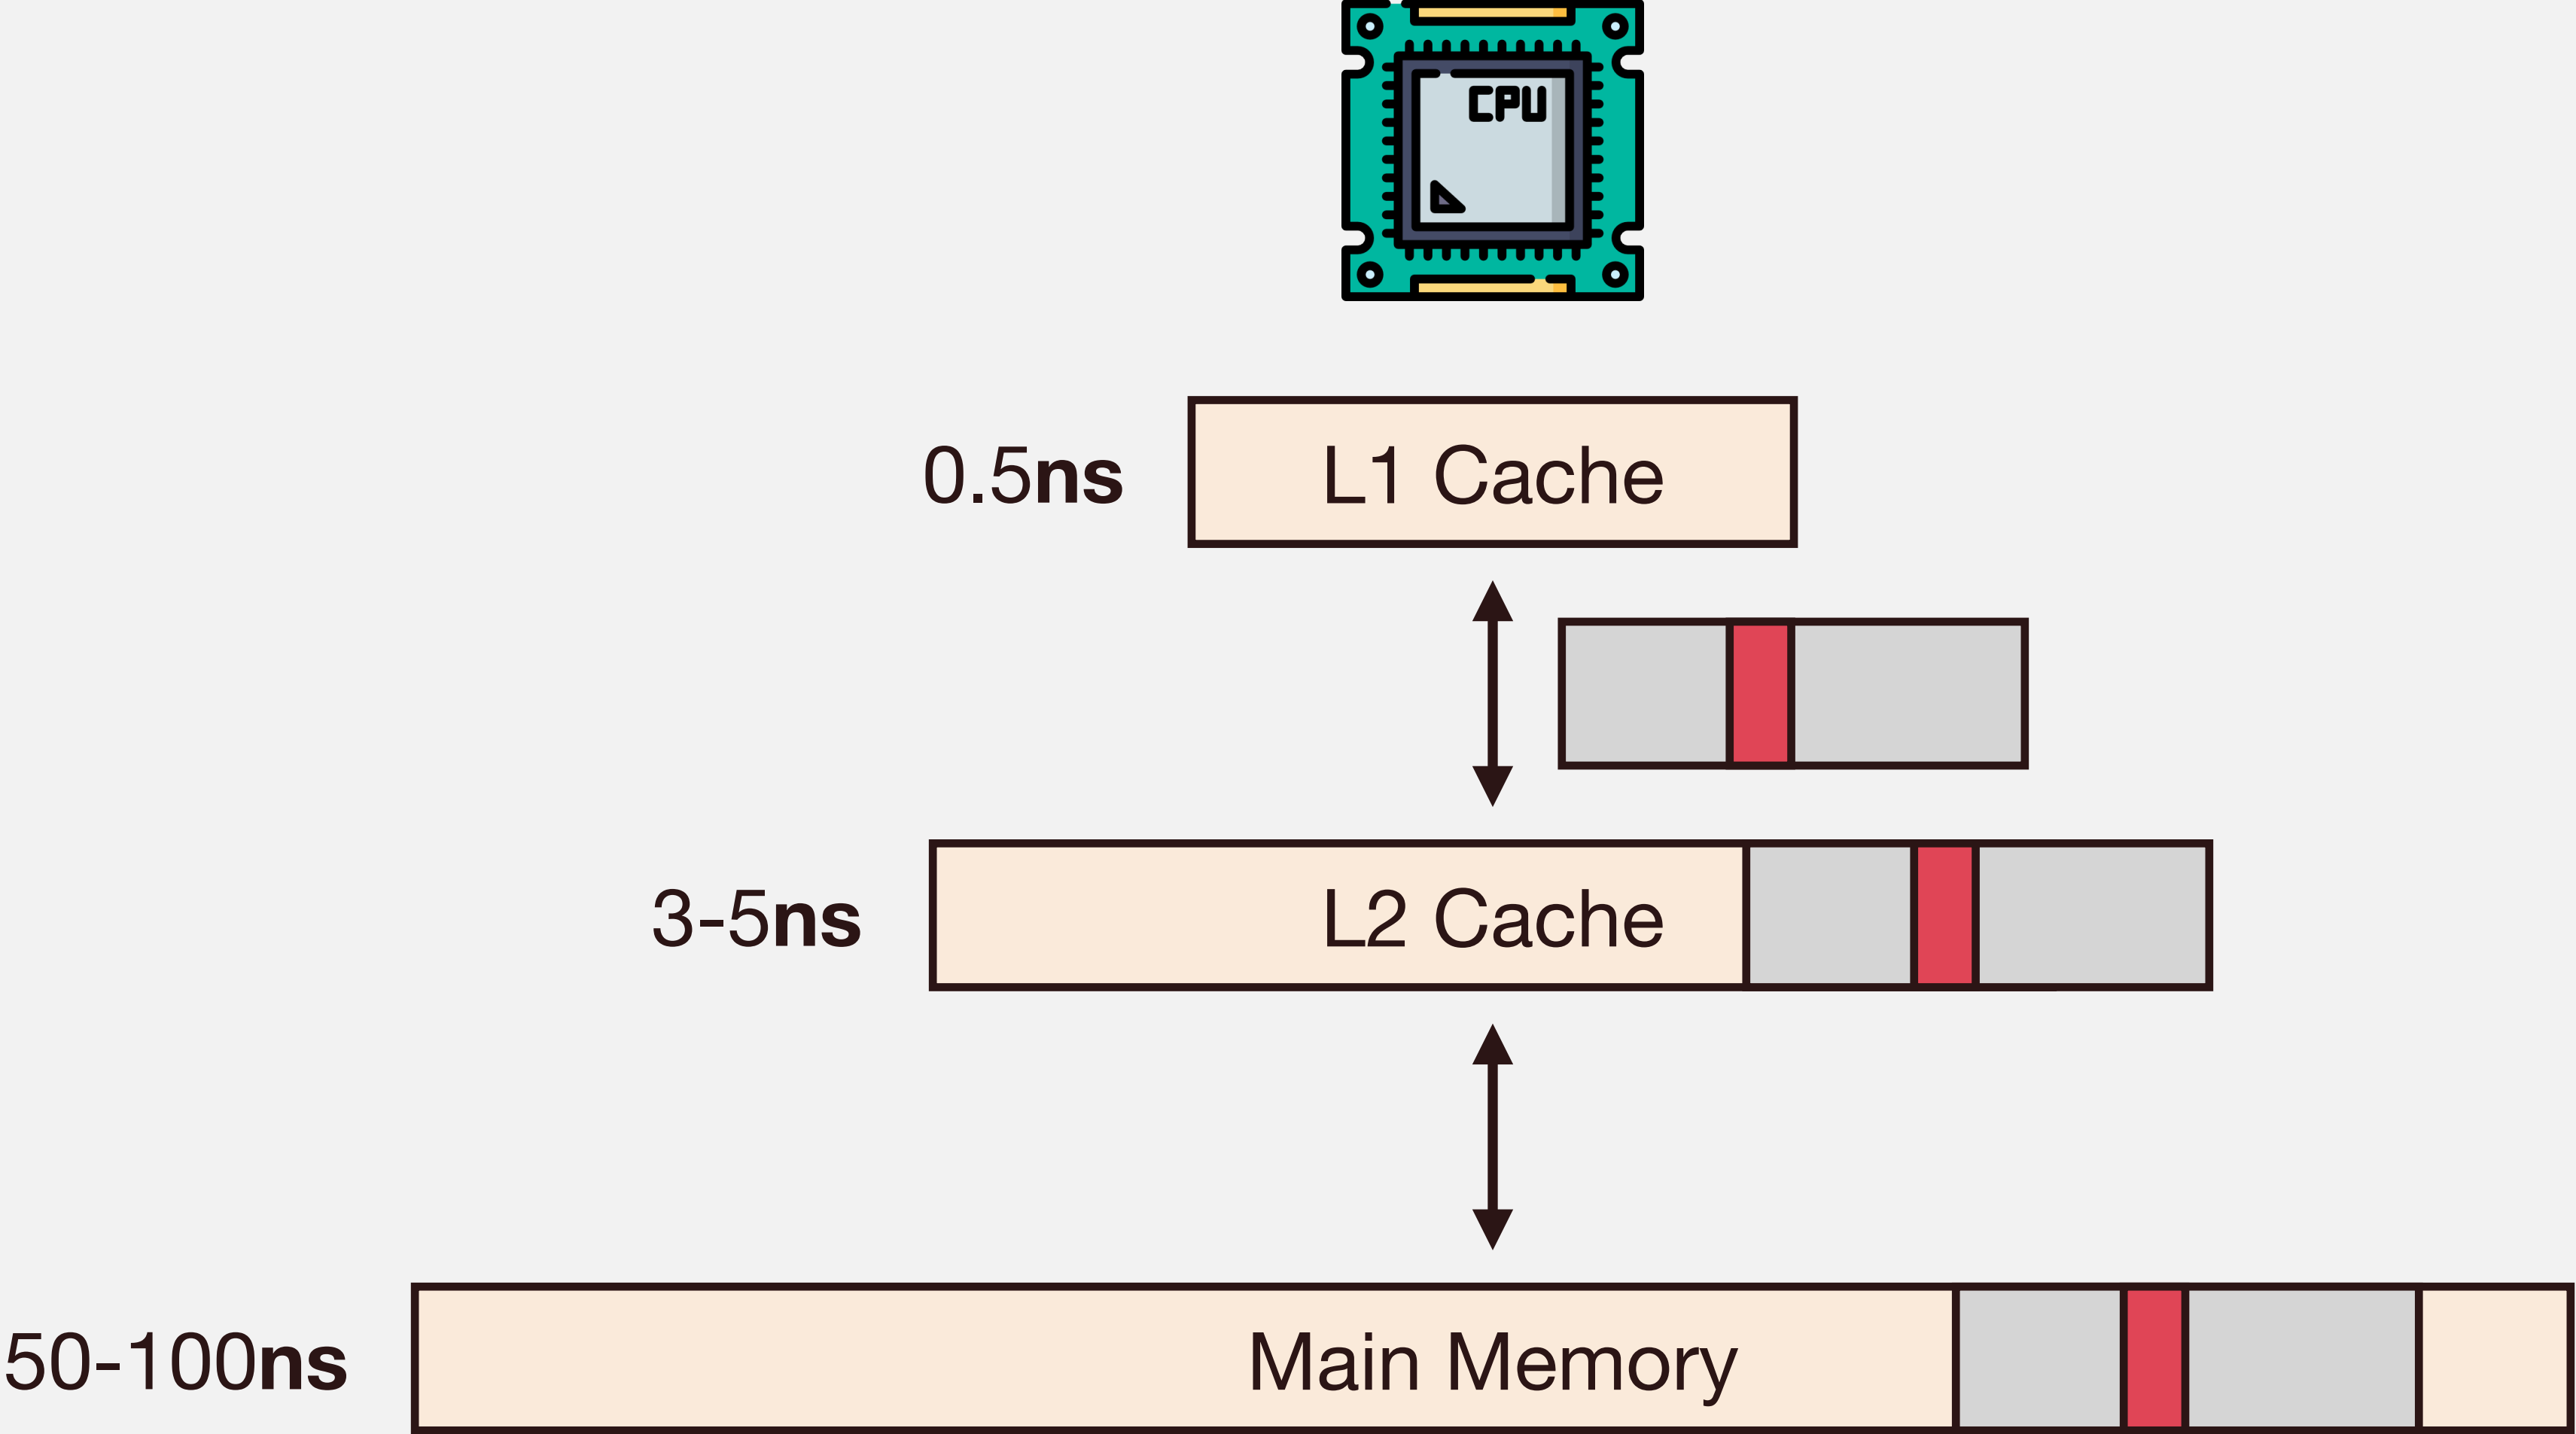
Cacheline



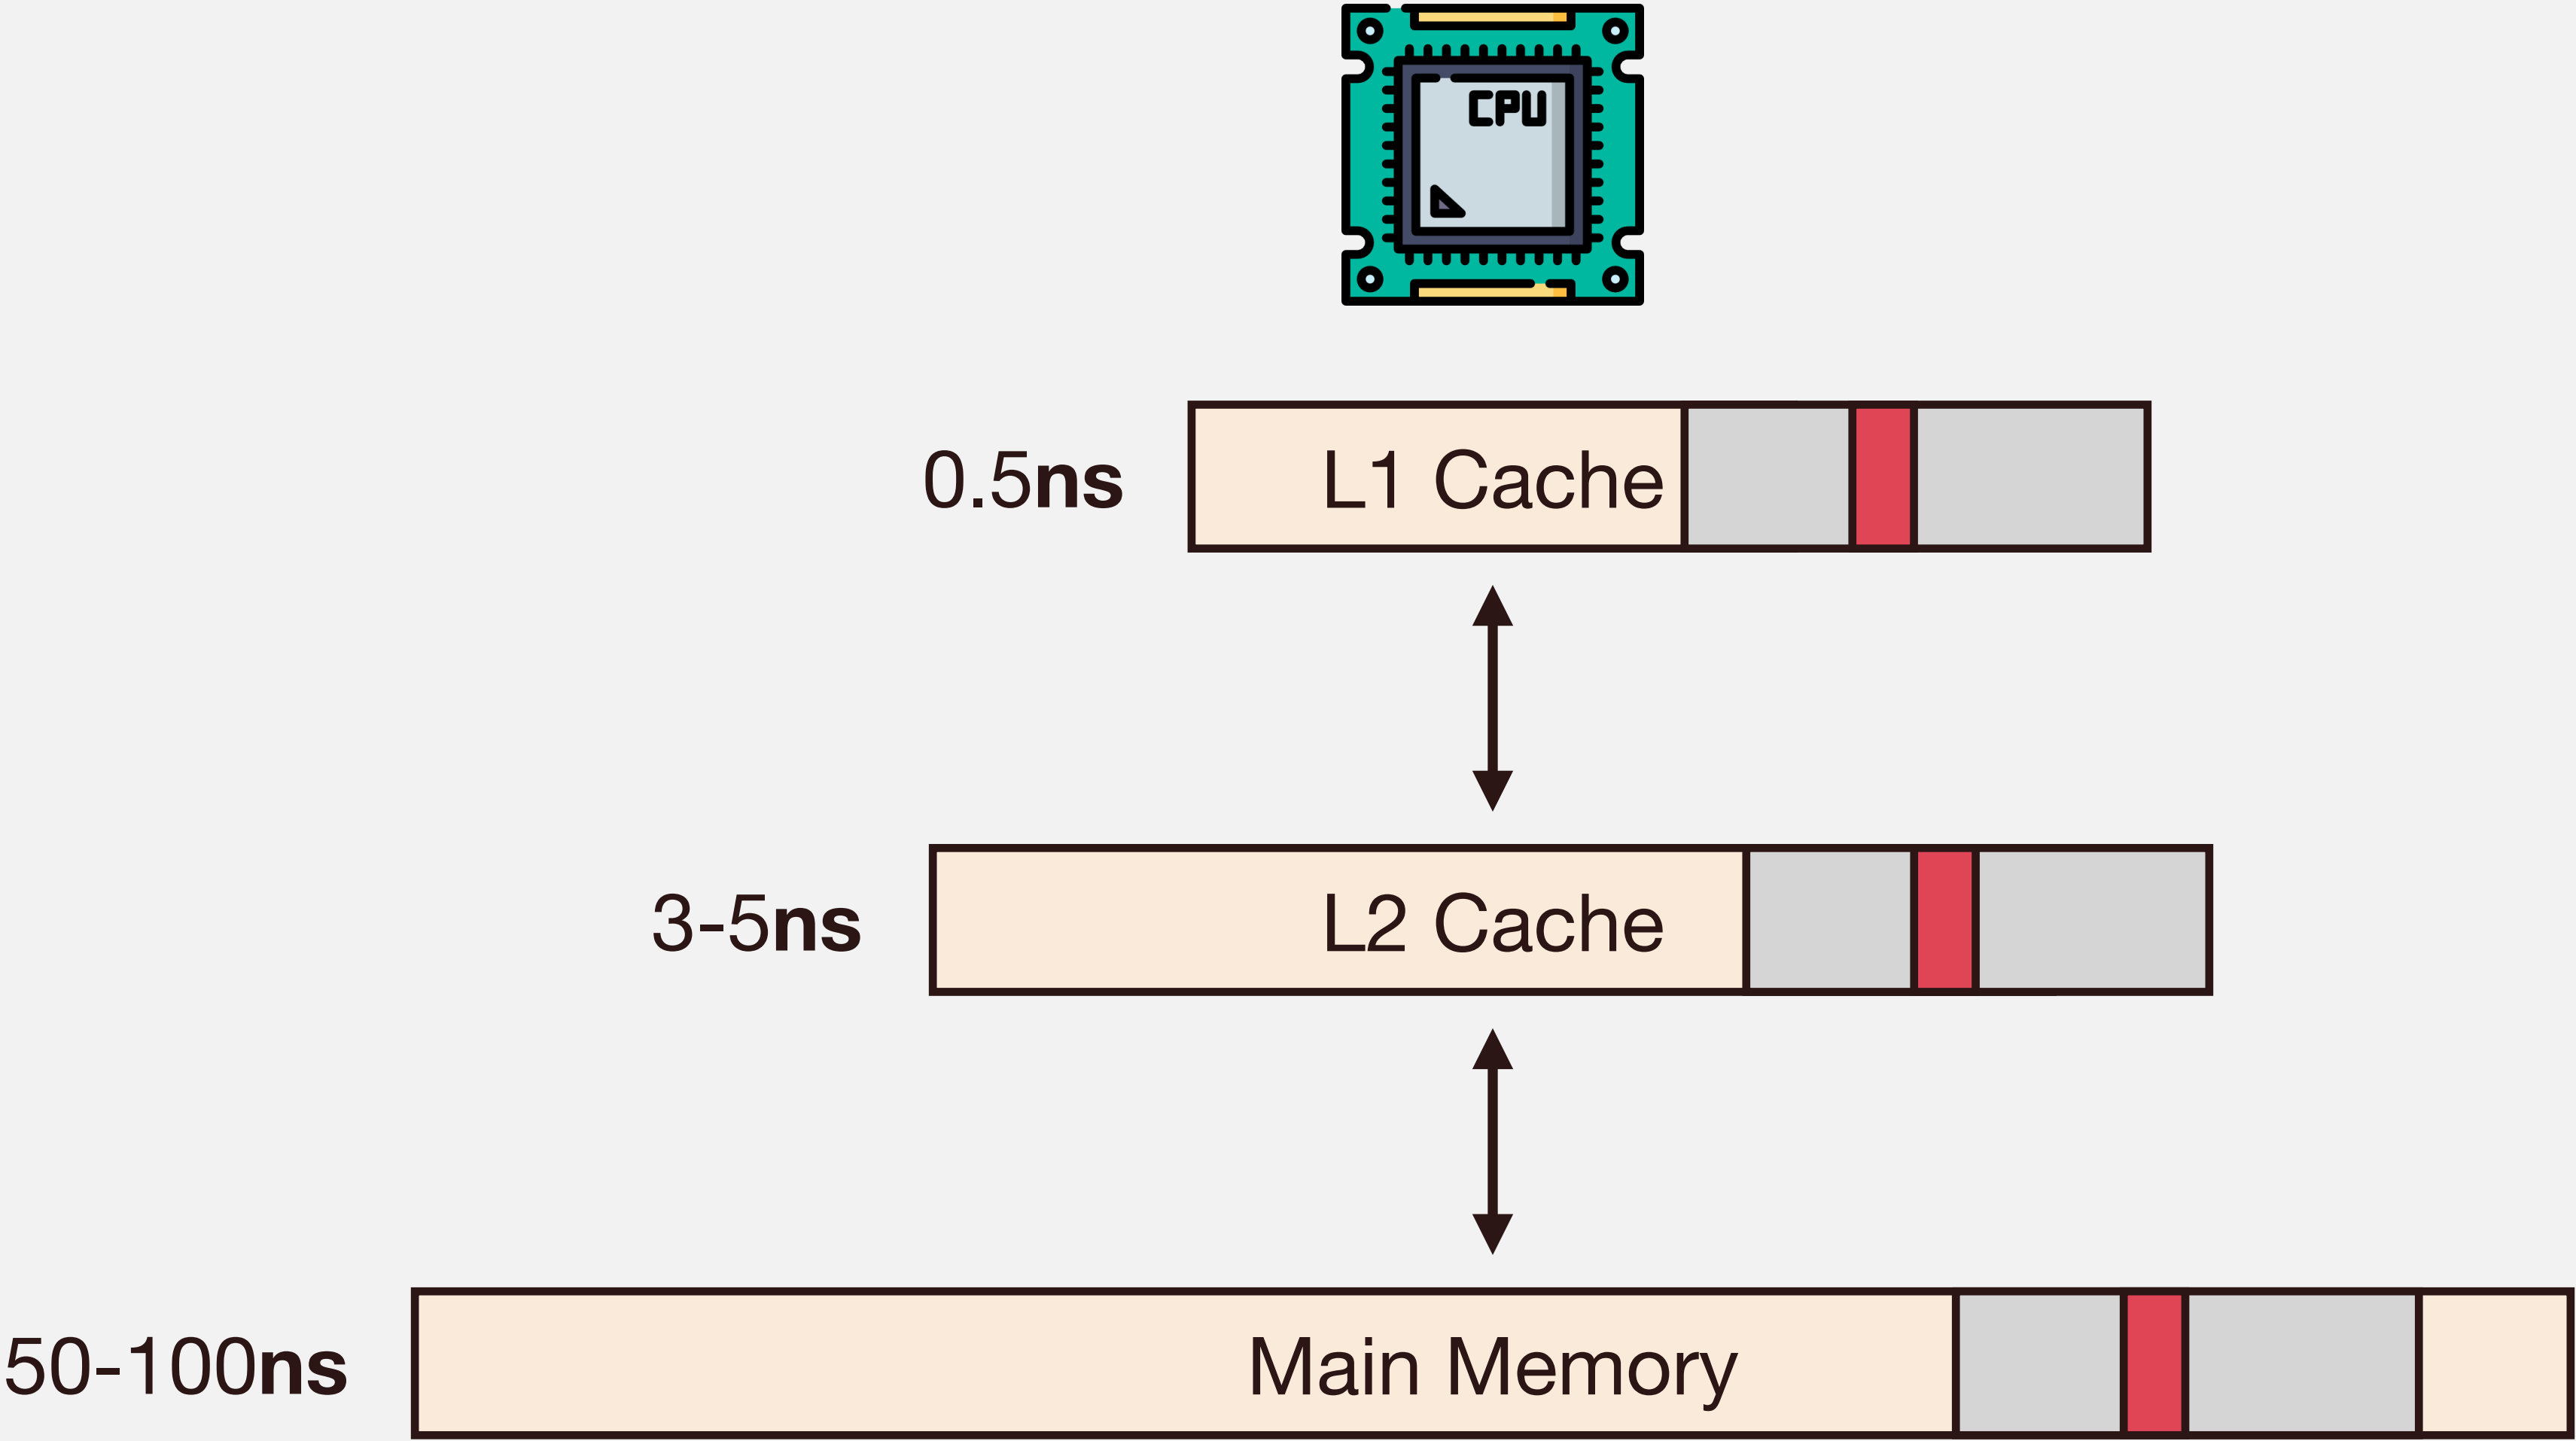
Cacheline



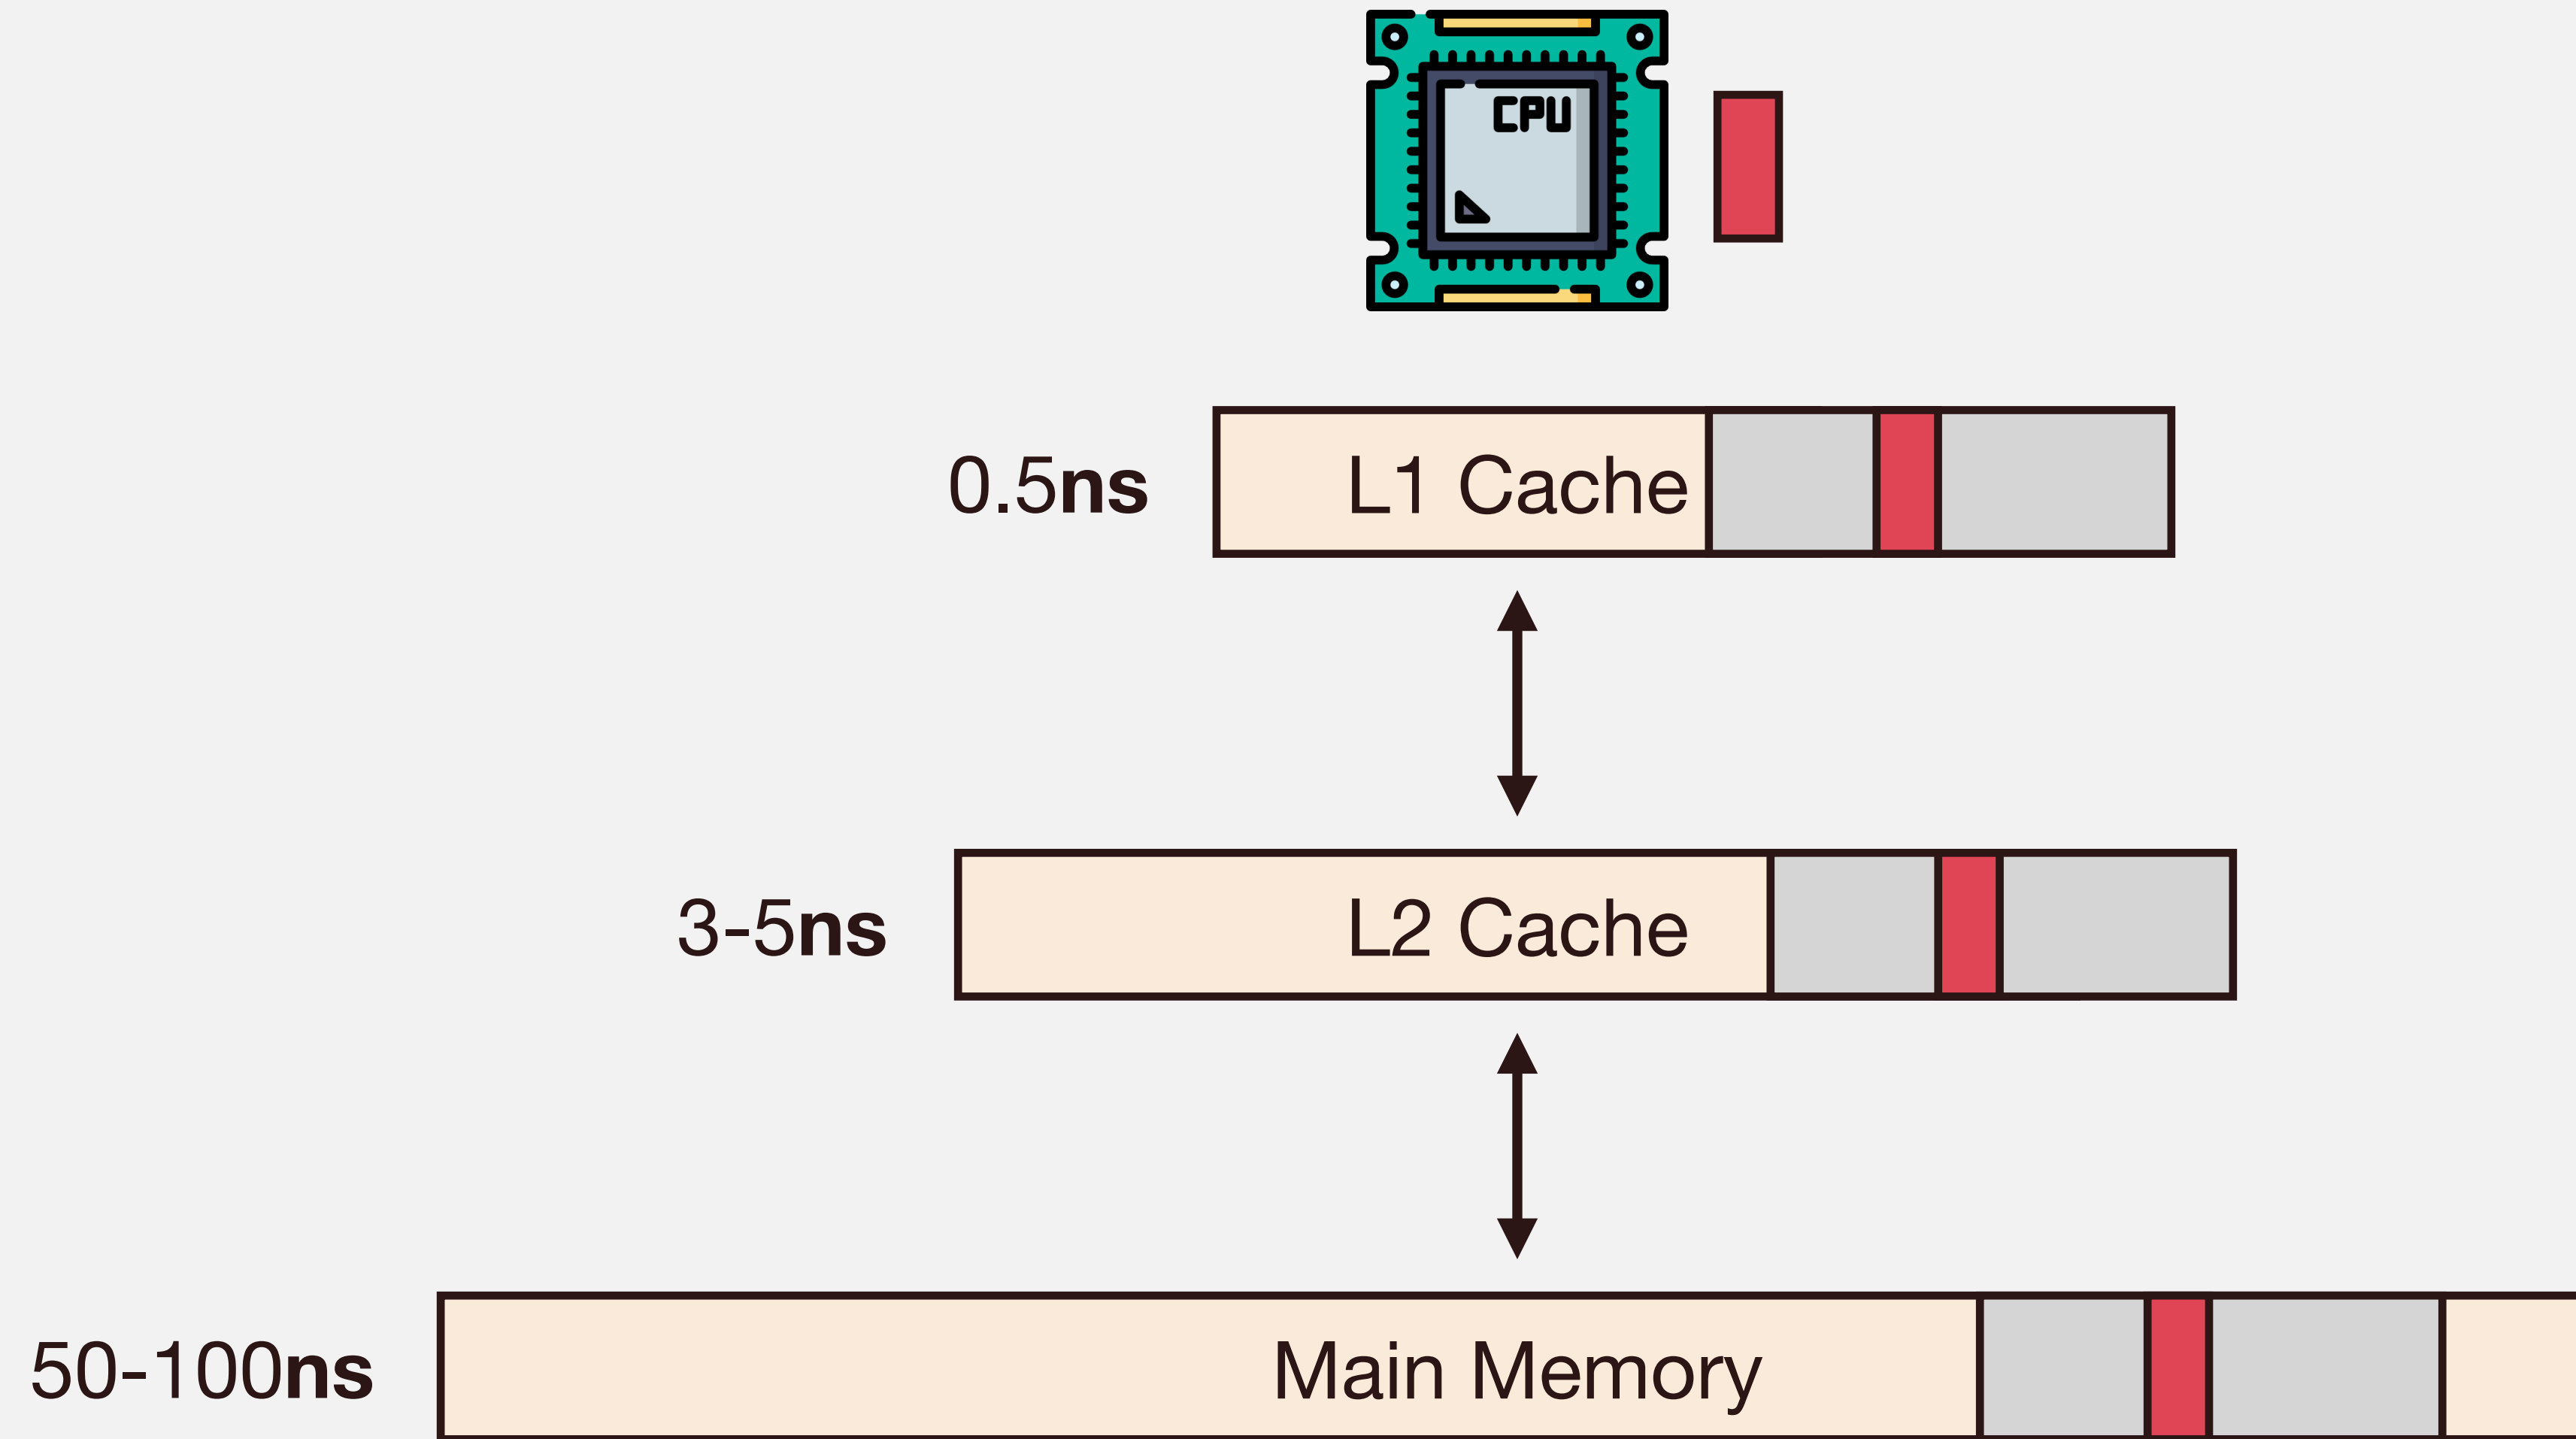
Cacheline



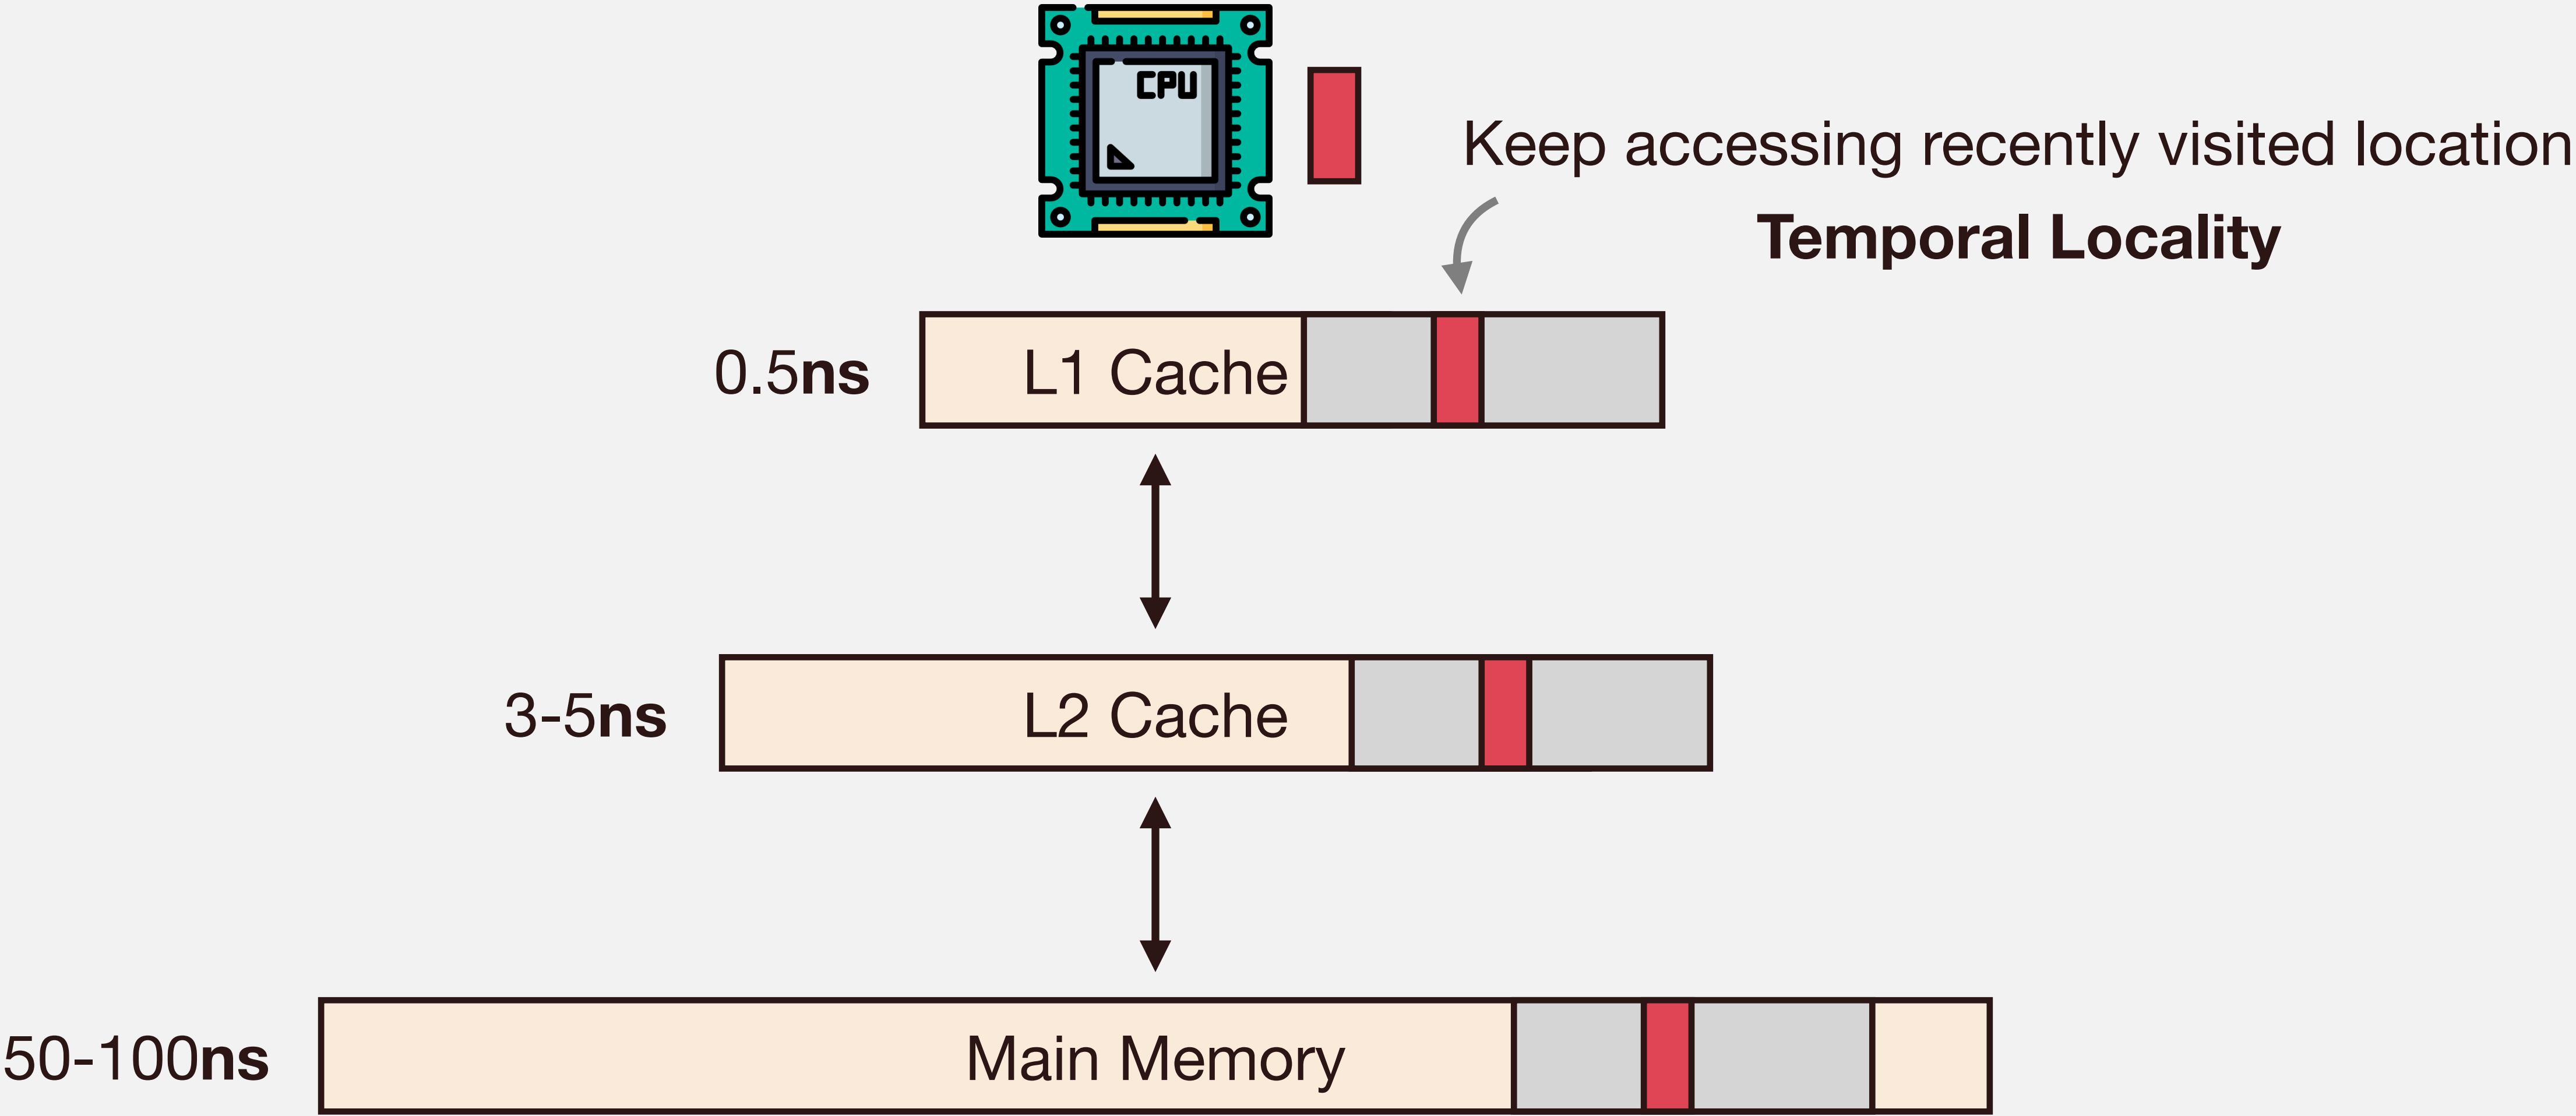
Cacheline



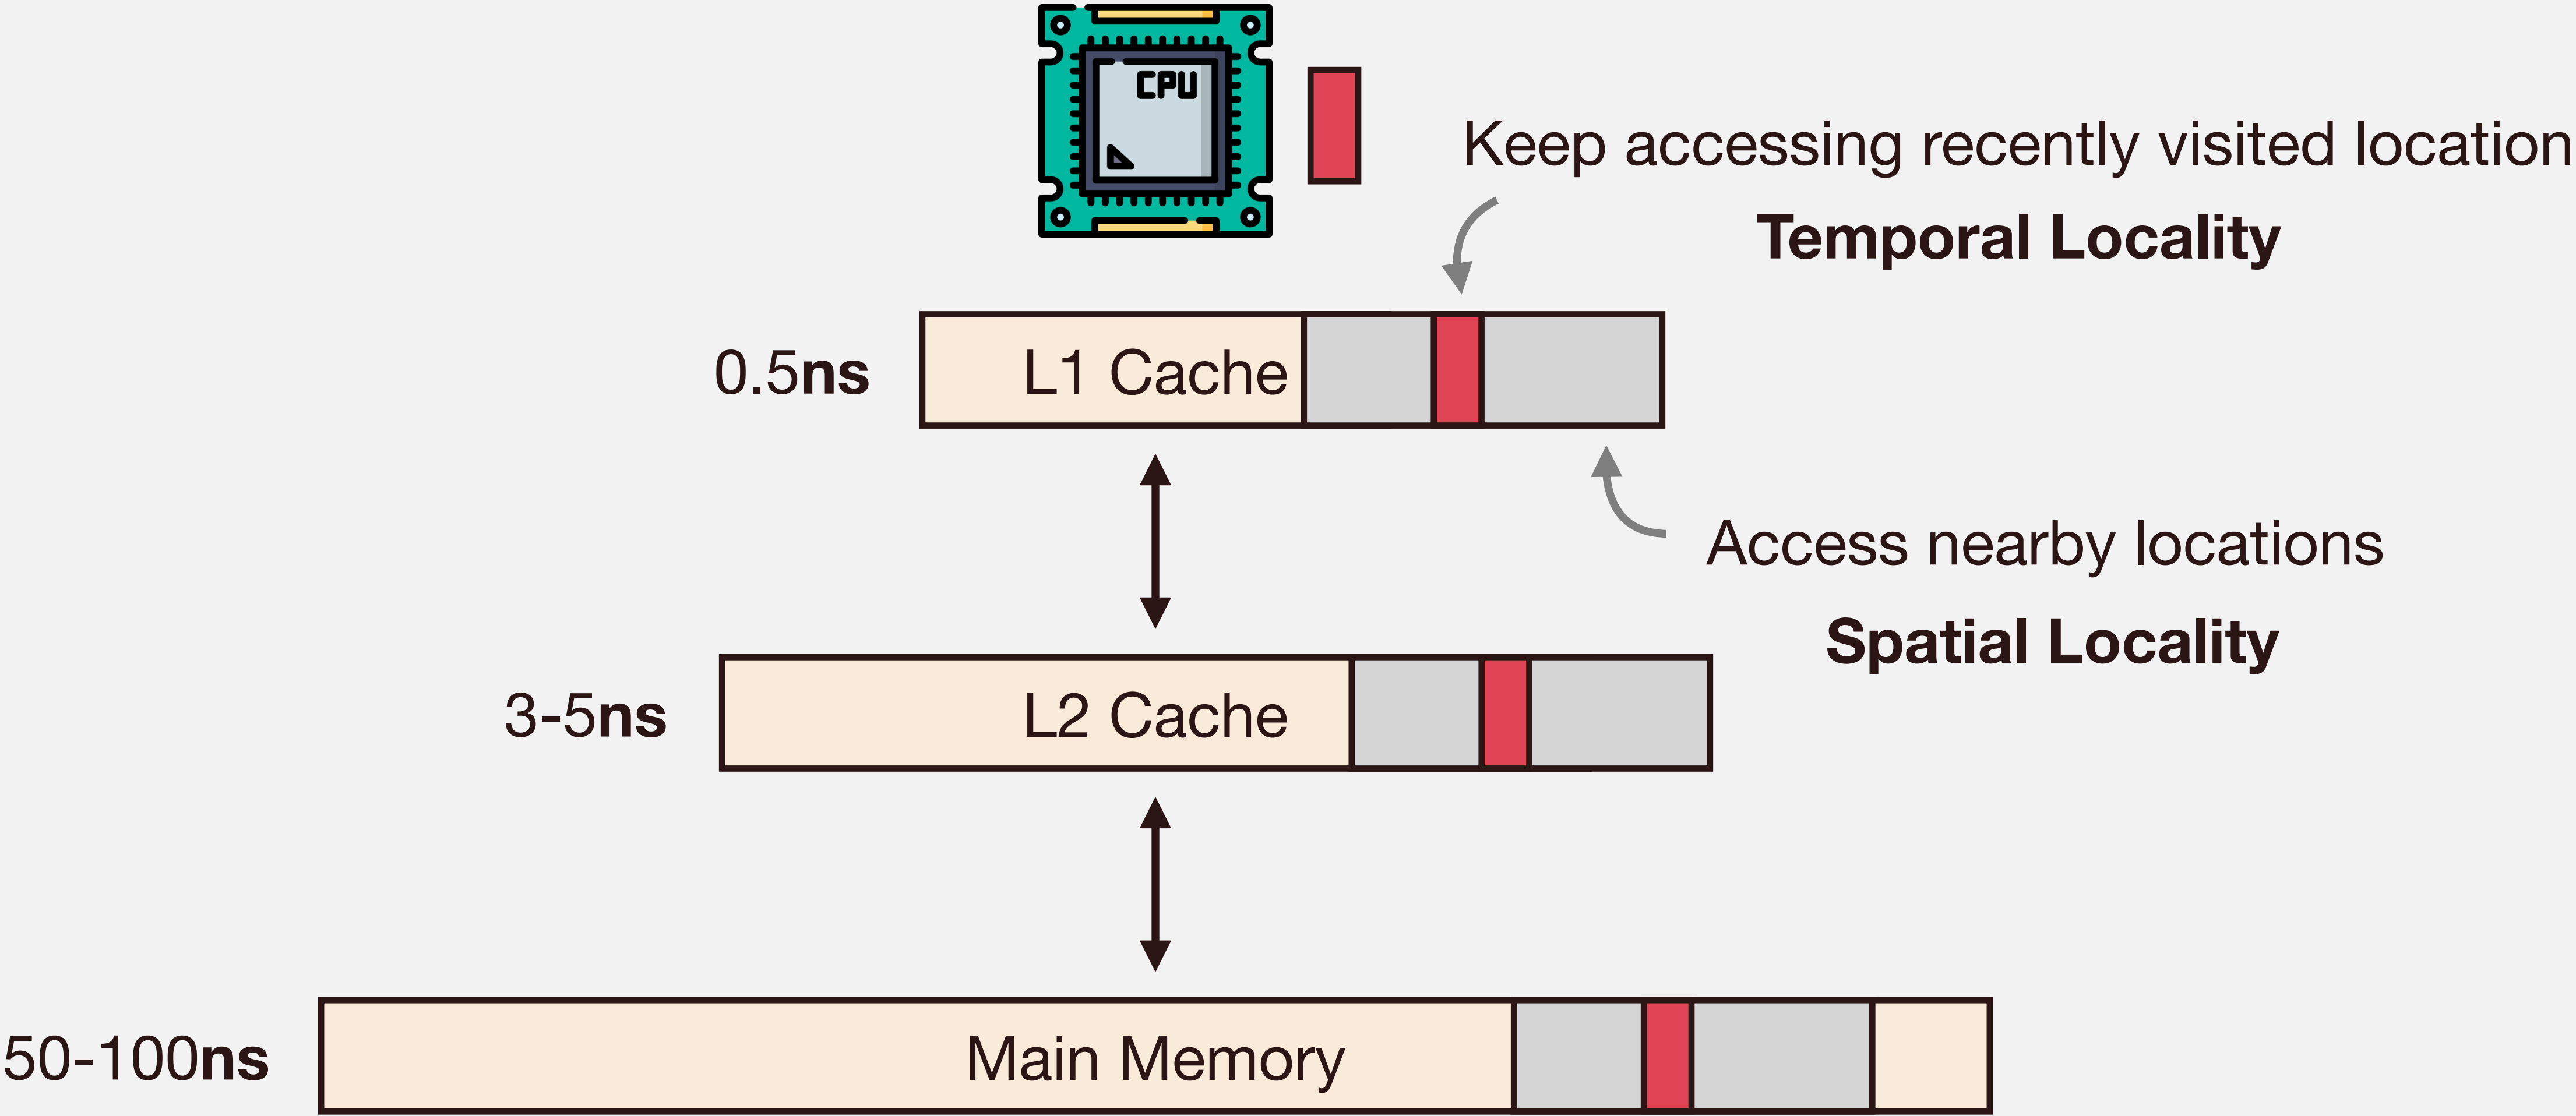
Cacheline



Cacheline



Cacheline



Low-Level Optimization Techniques

CPU

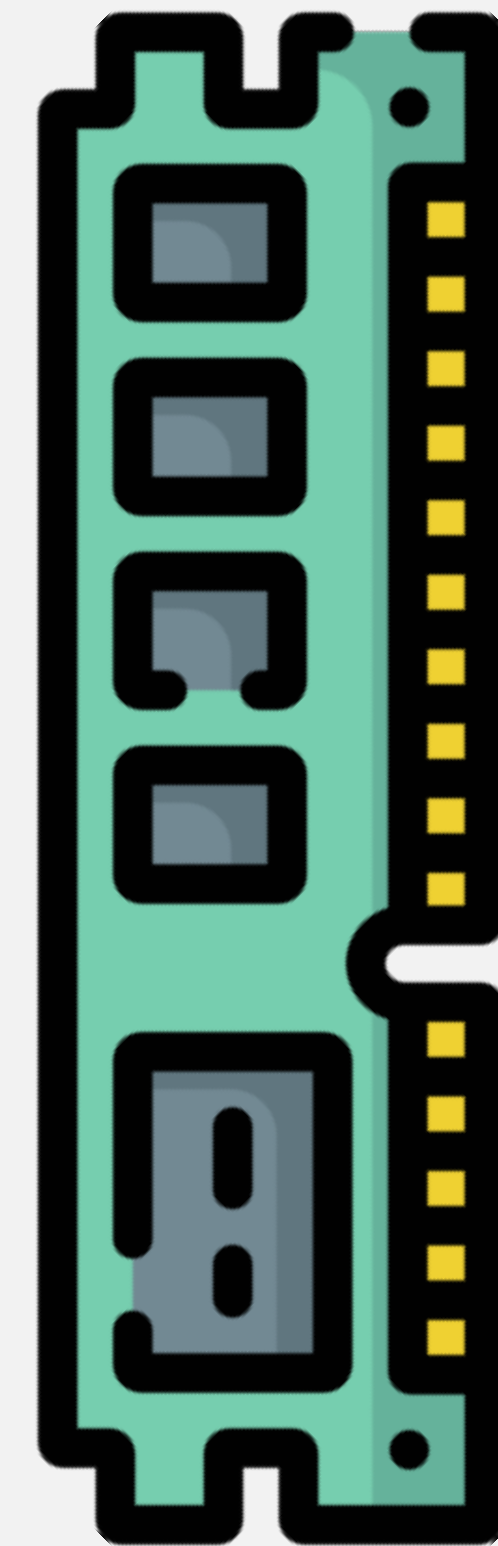
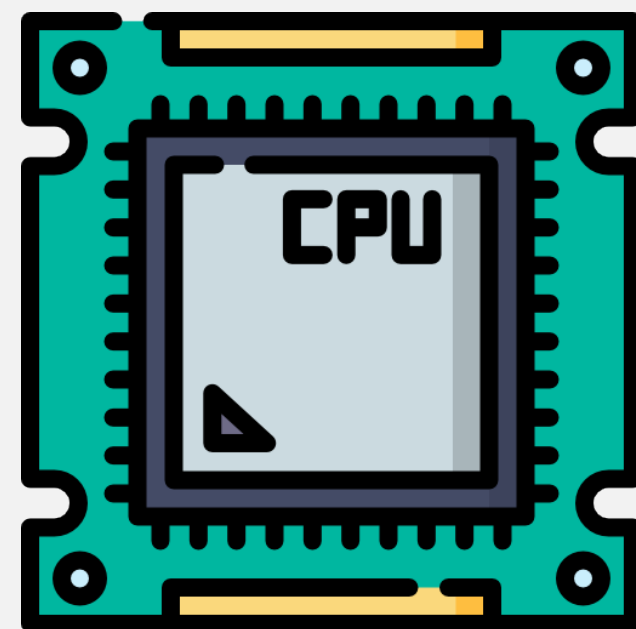
- Reduce Operator Strength
- Reduce Branch Misprediction
- Reduce Data Dependency
- SIMD

Cache

- Temporal Locality
- Spacial Locality

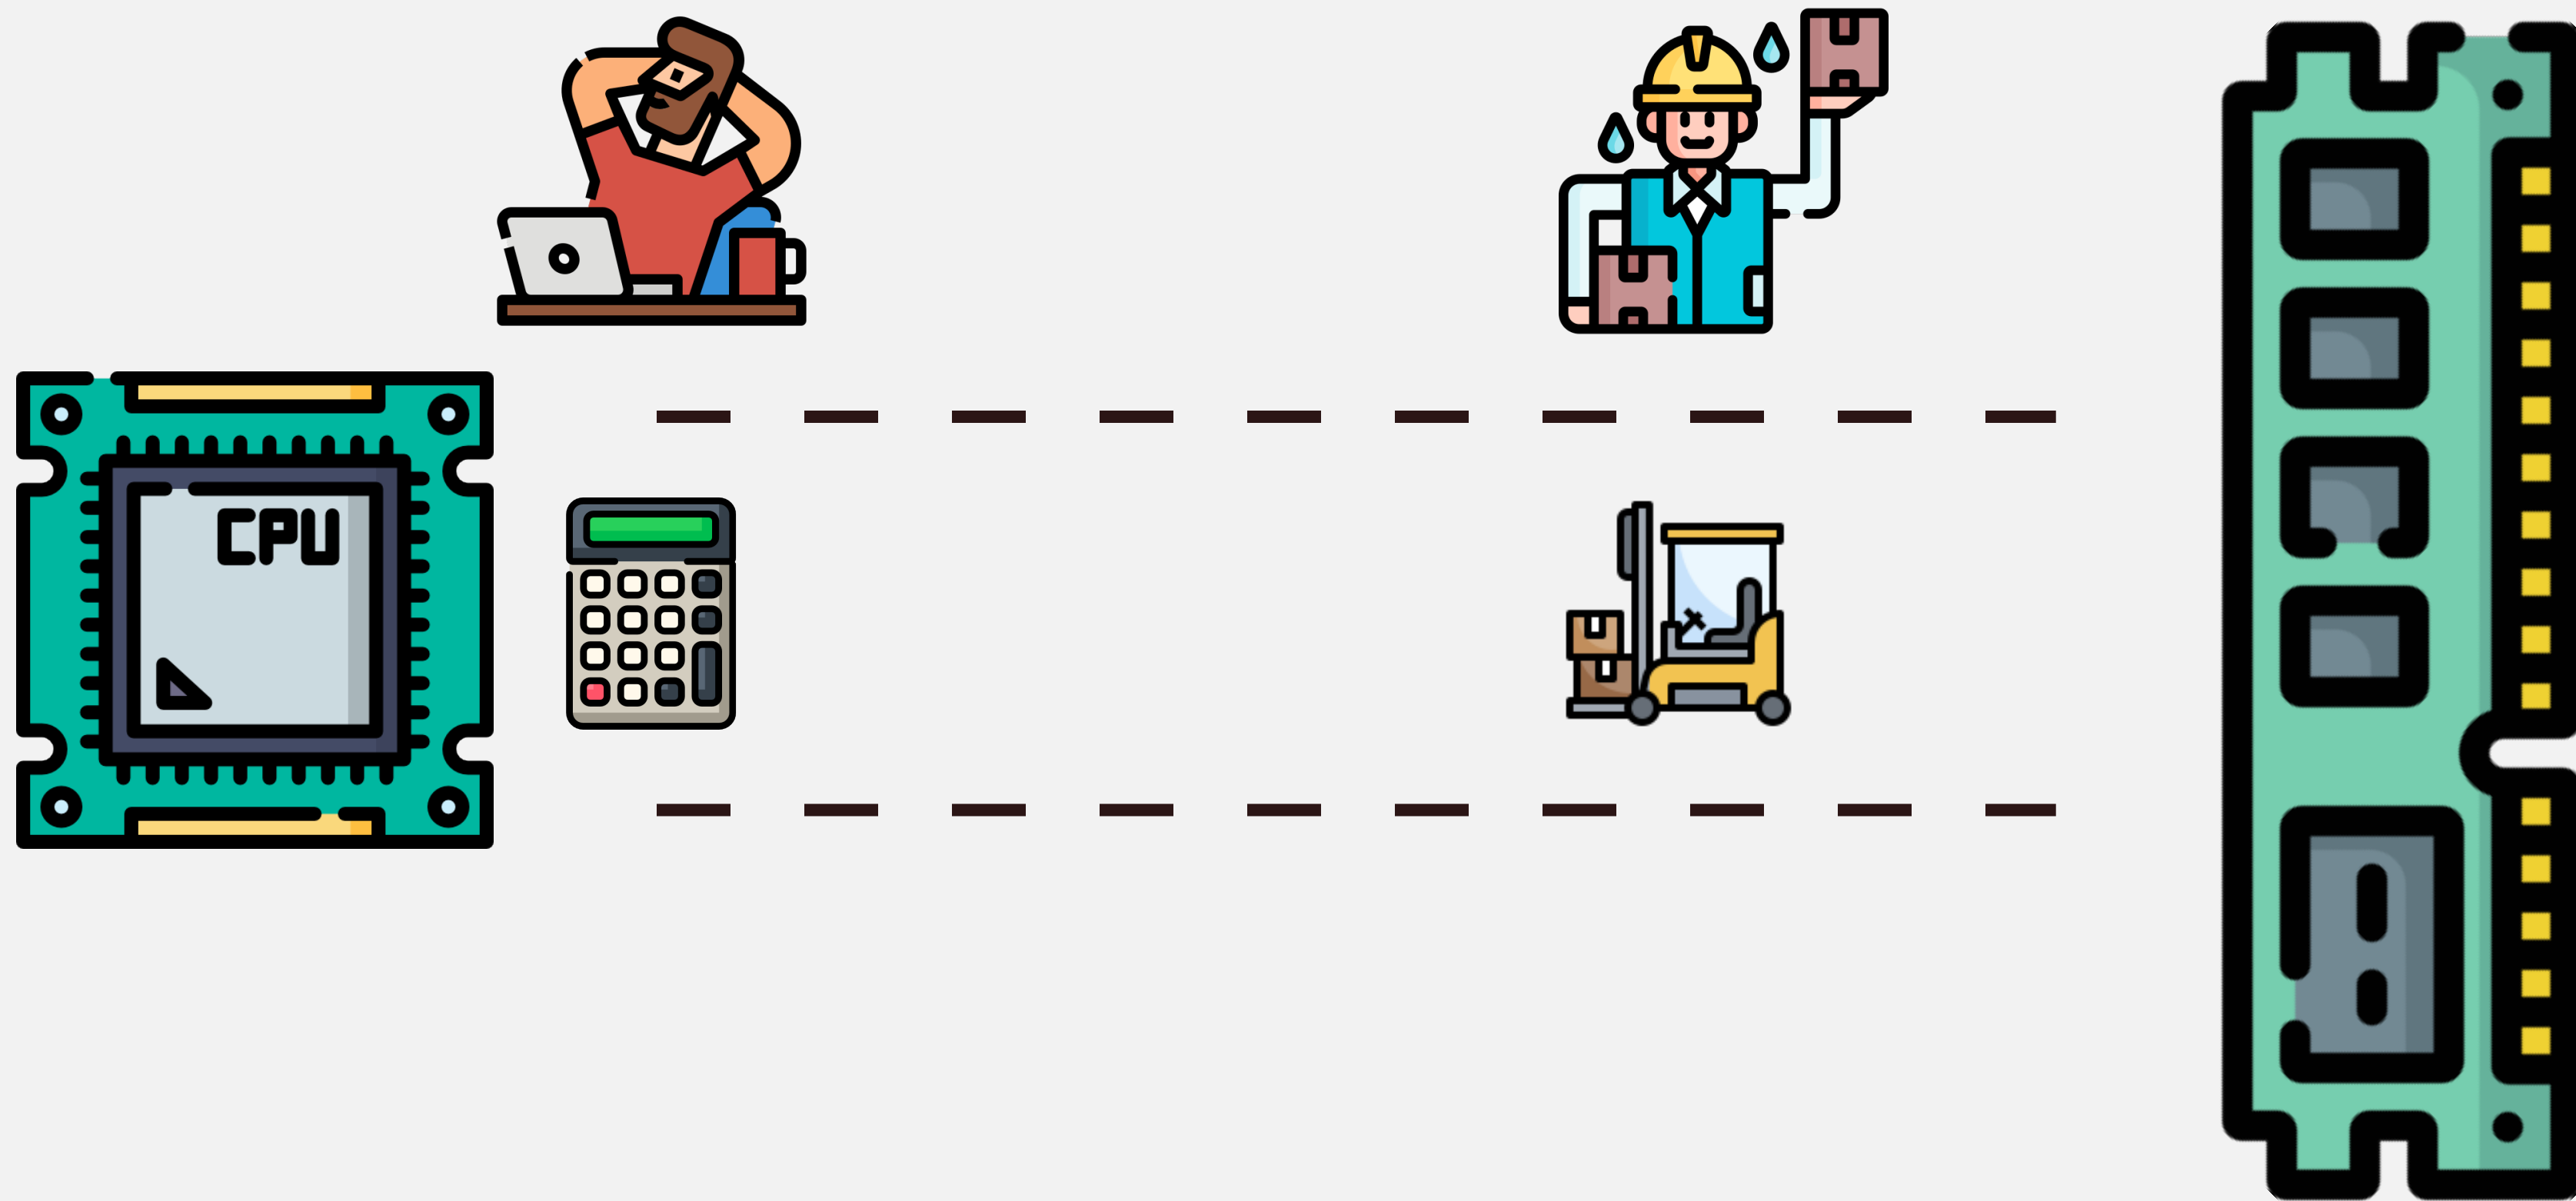
Who's the Bottleneck?

Cache / Memory

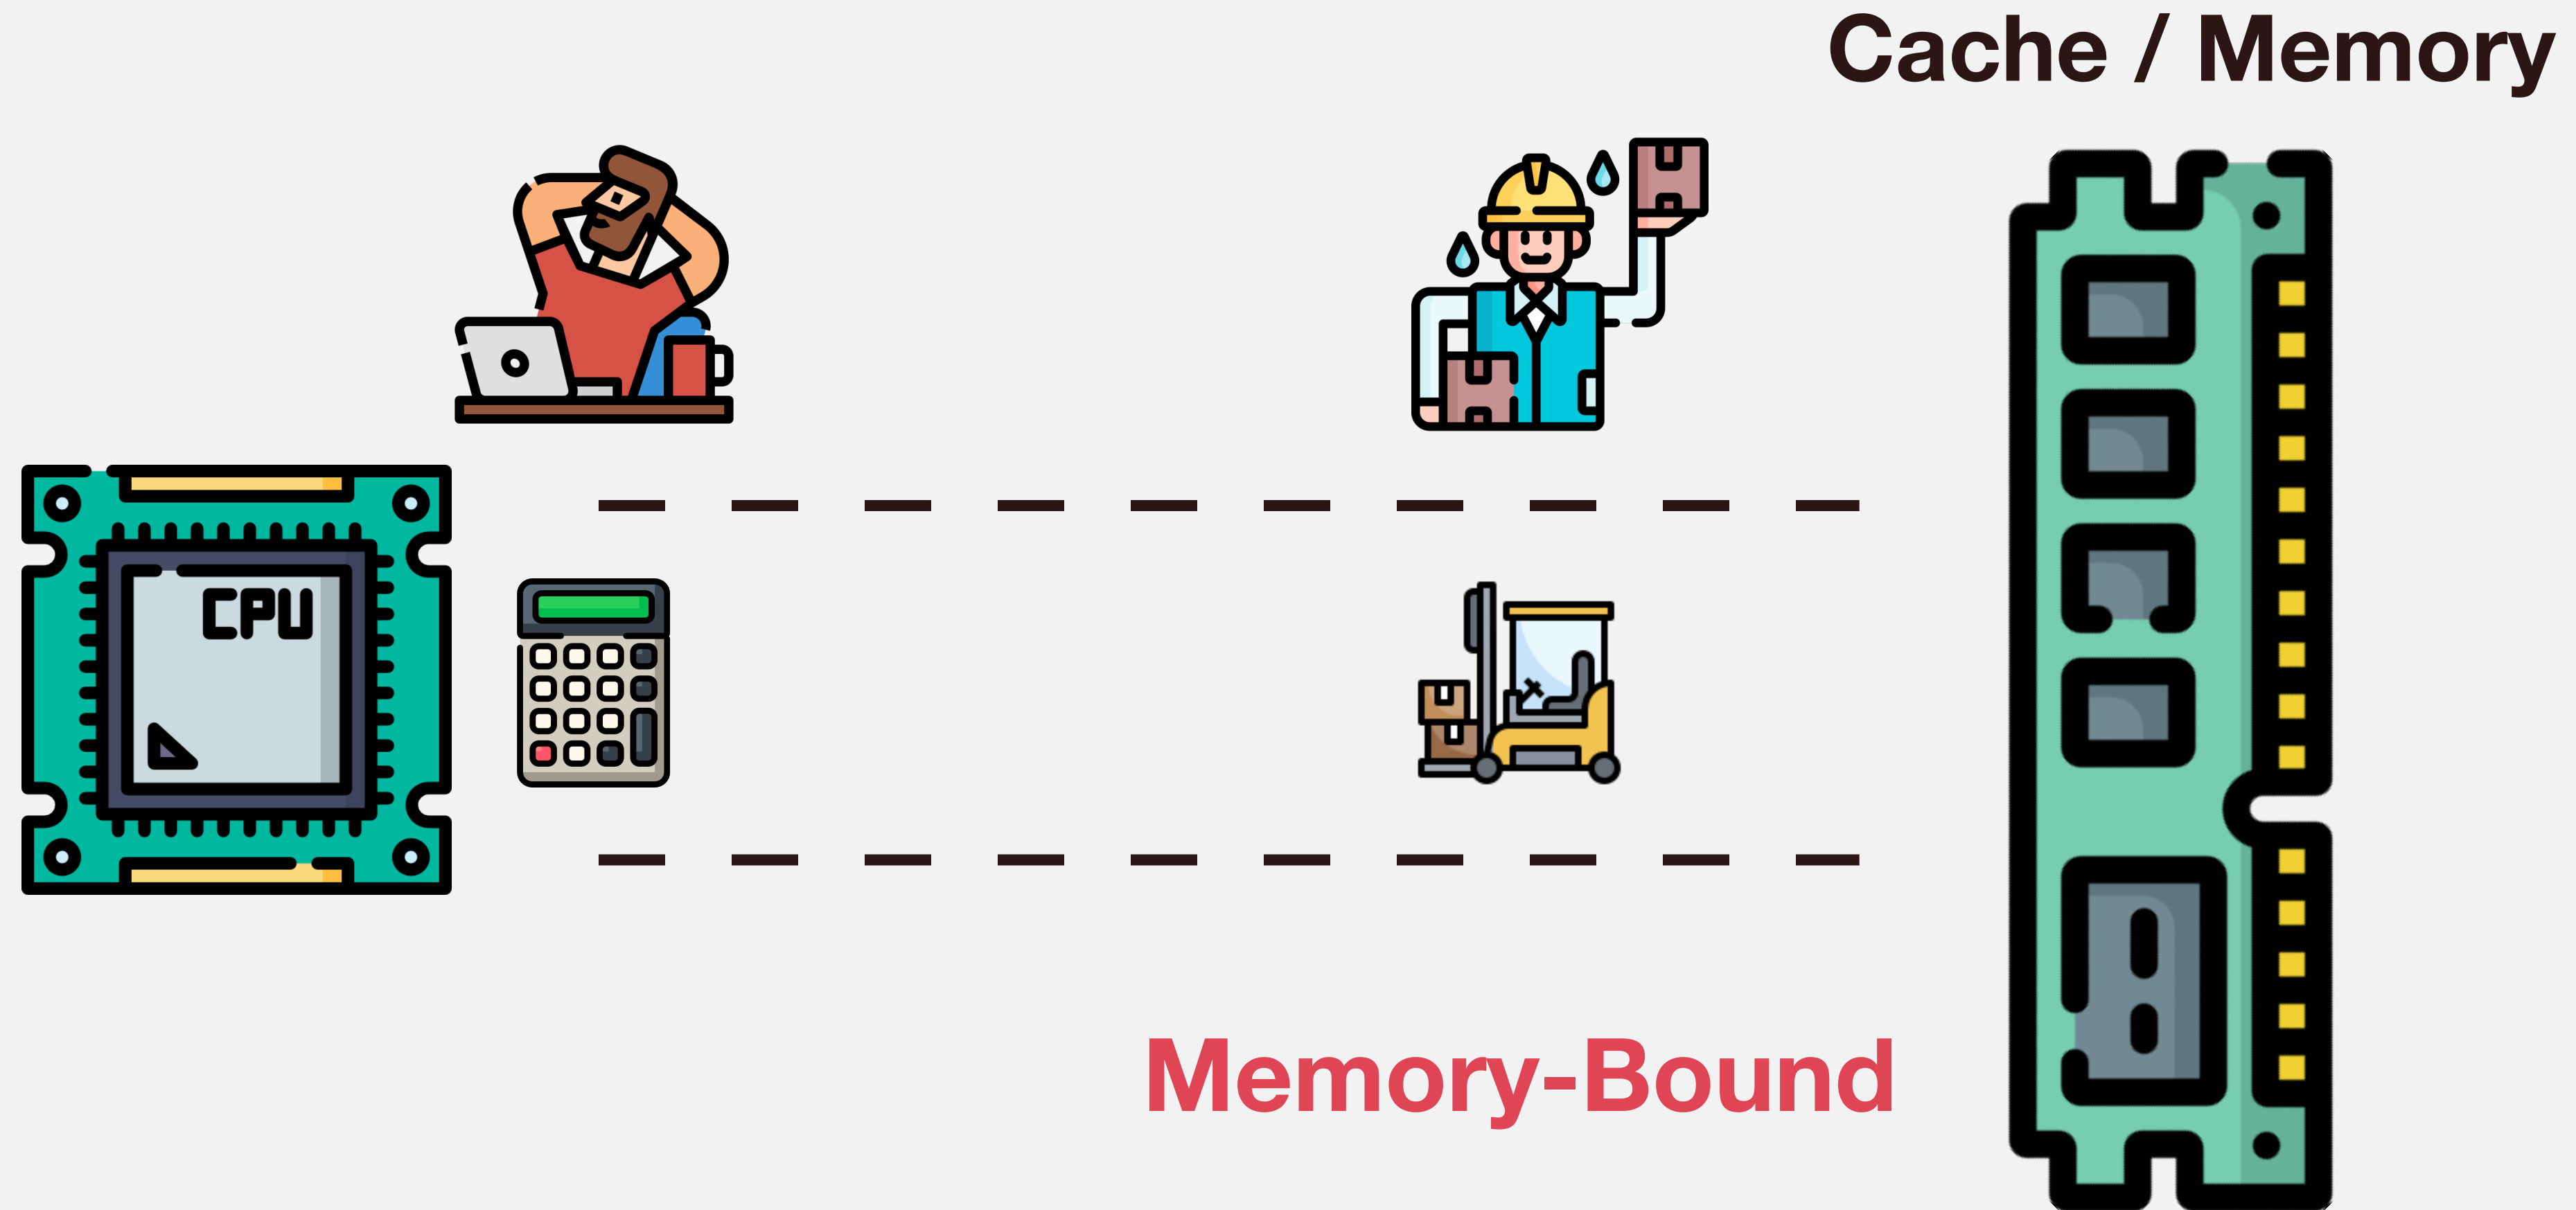


Who's the Bottleneck?

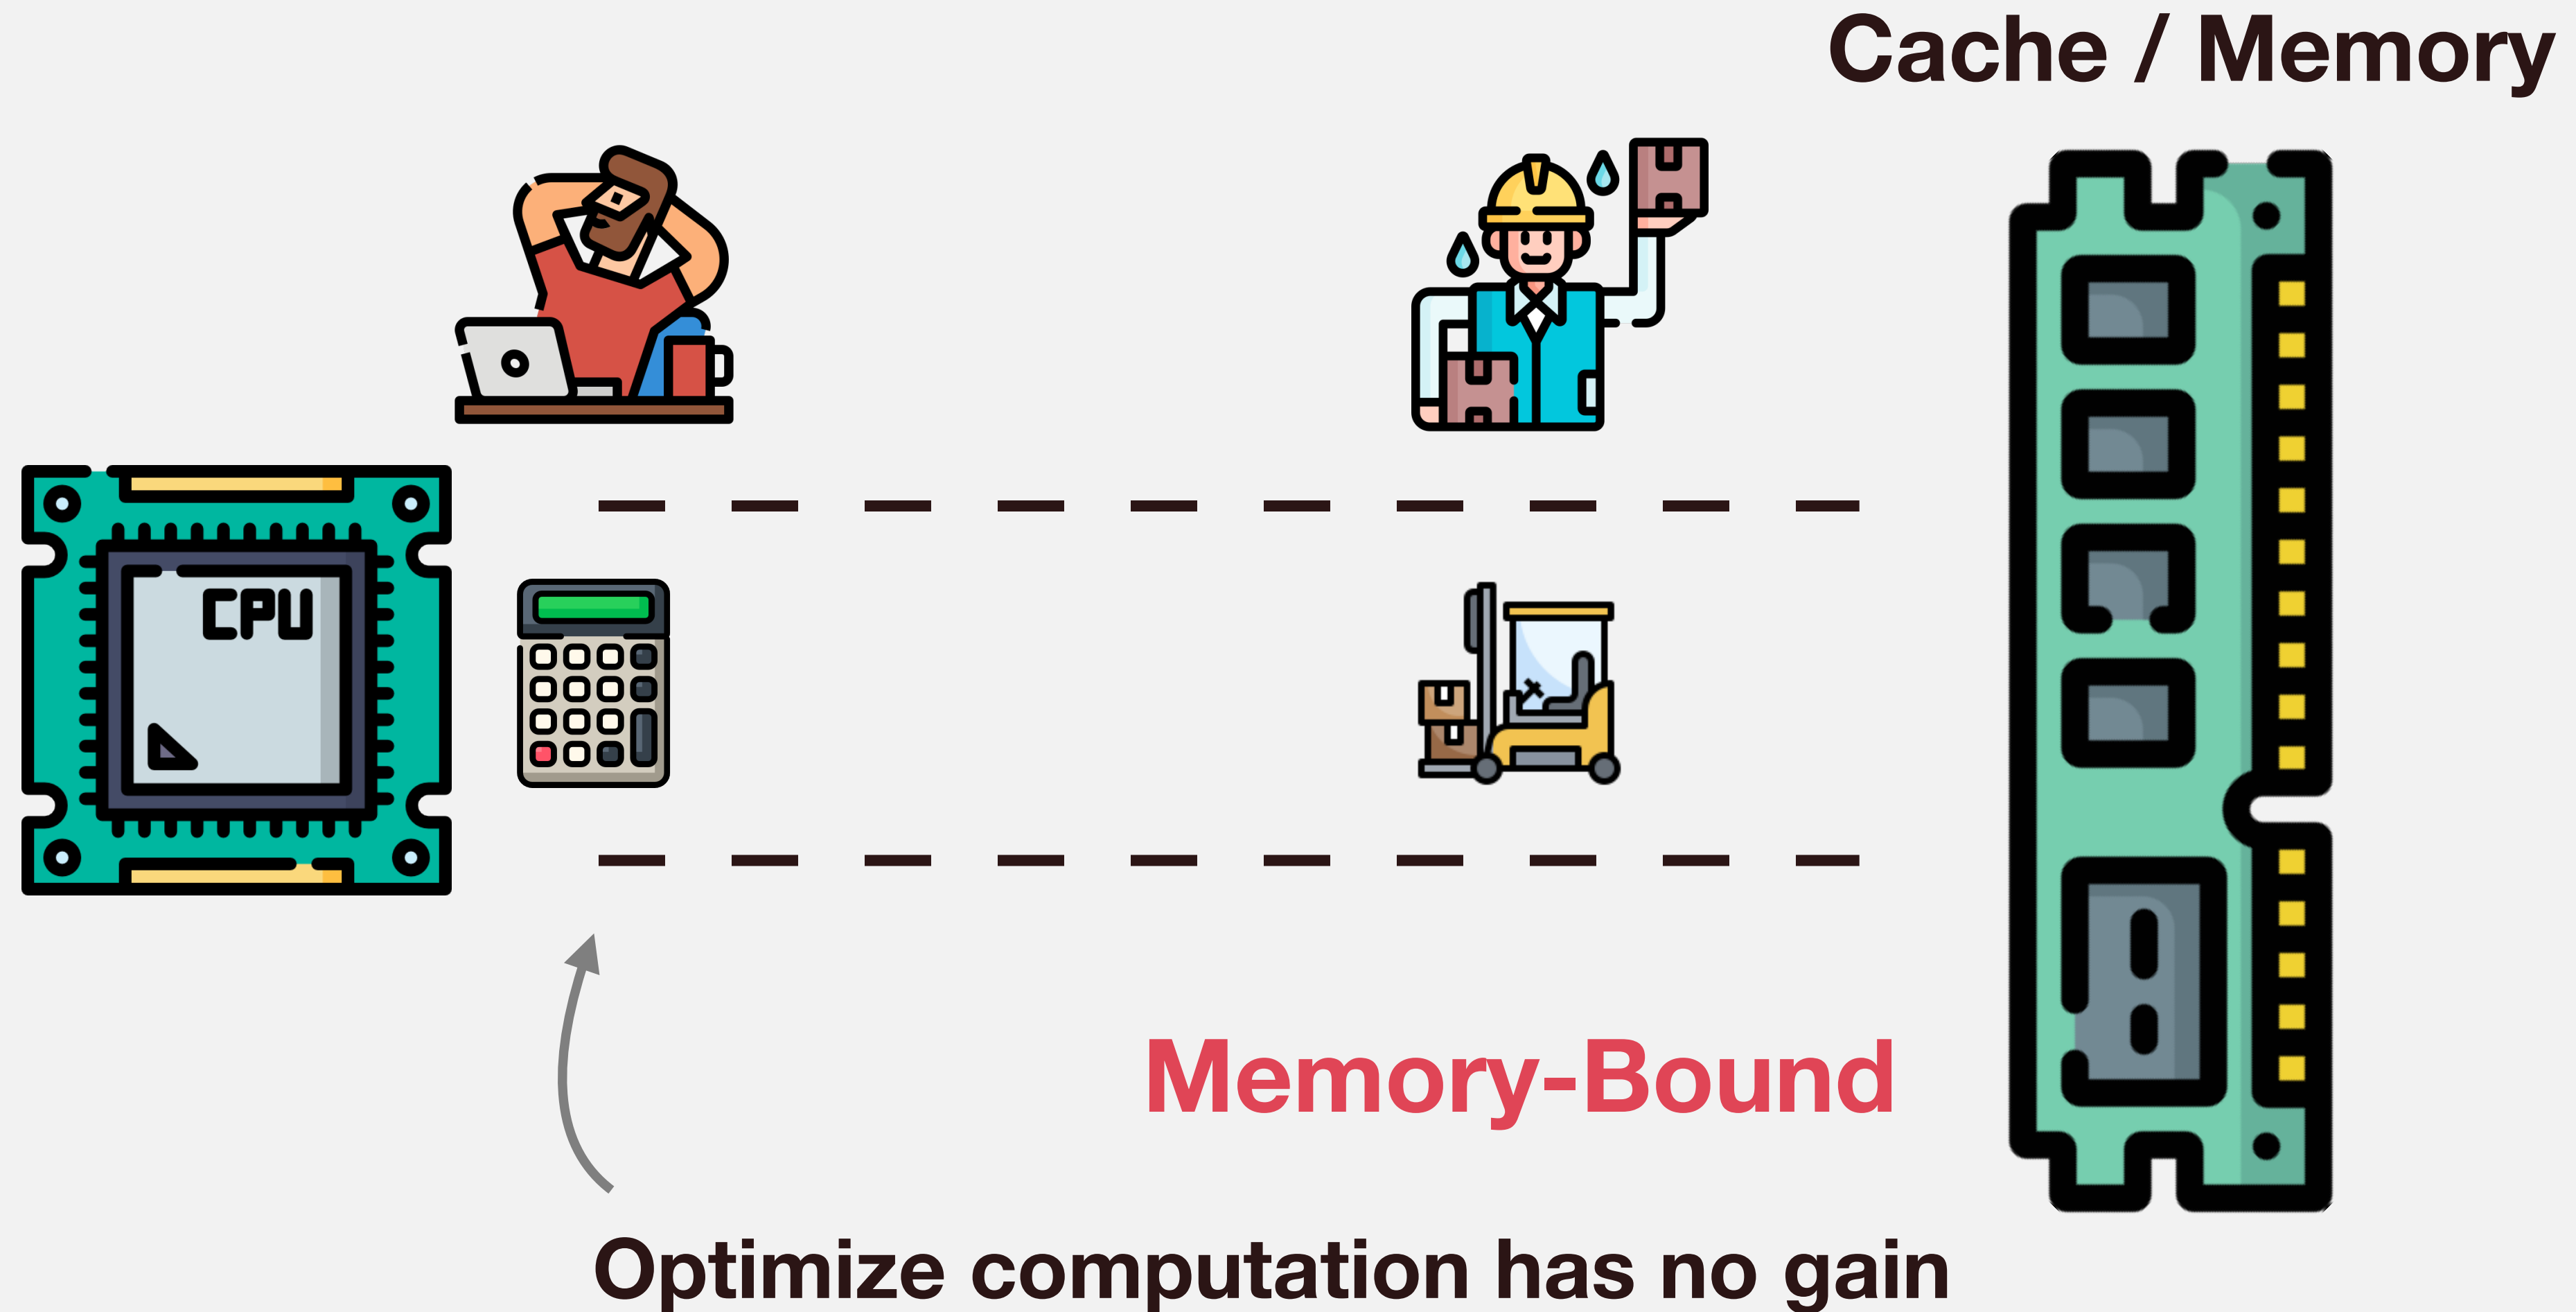
Cache / Memory



Who's the Bottleneck?

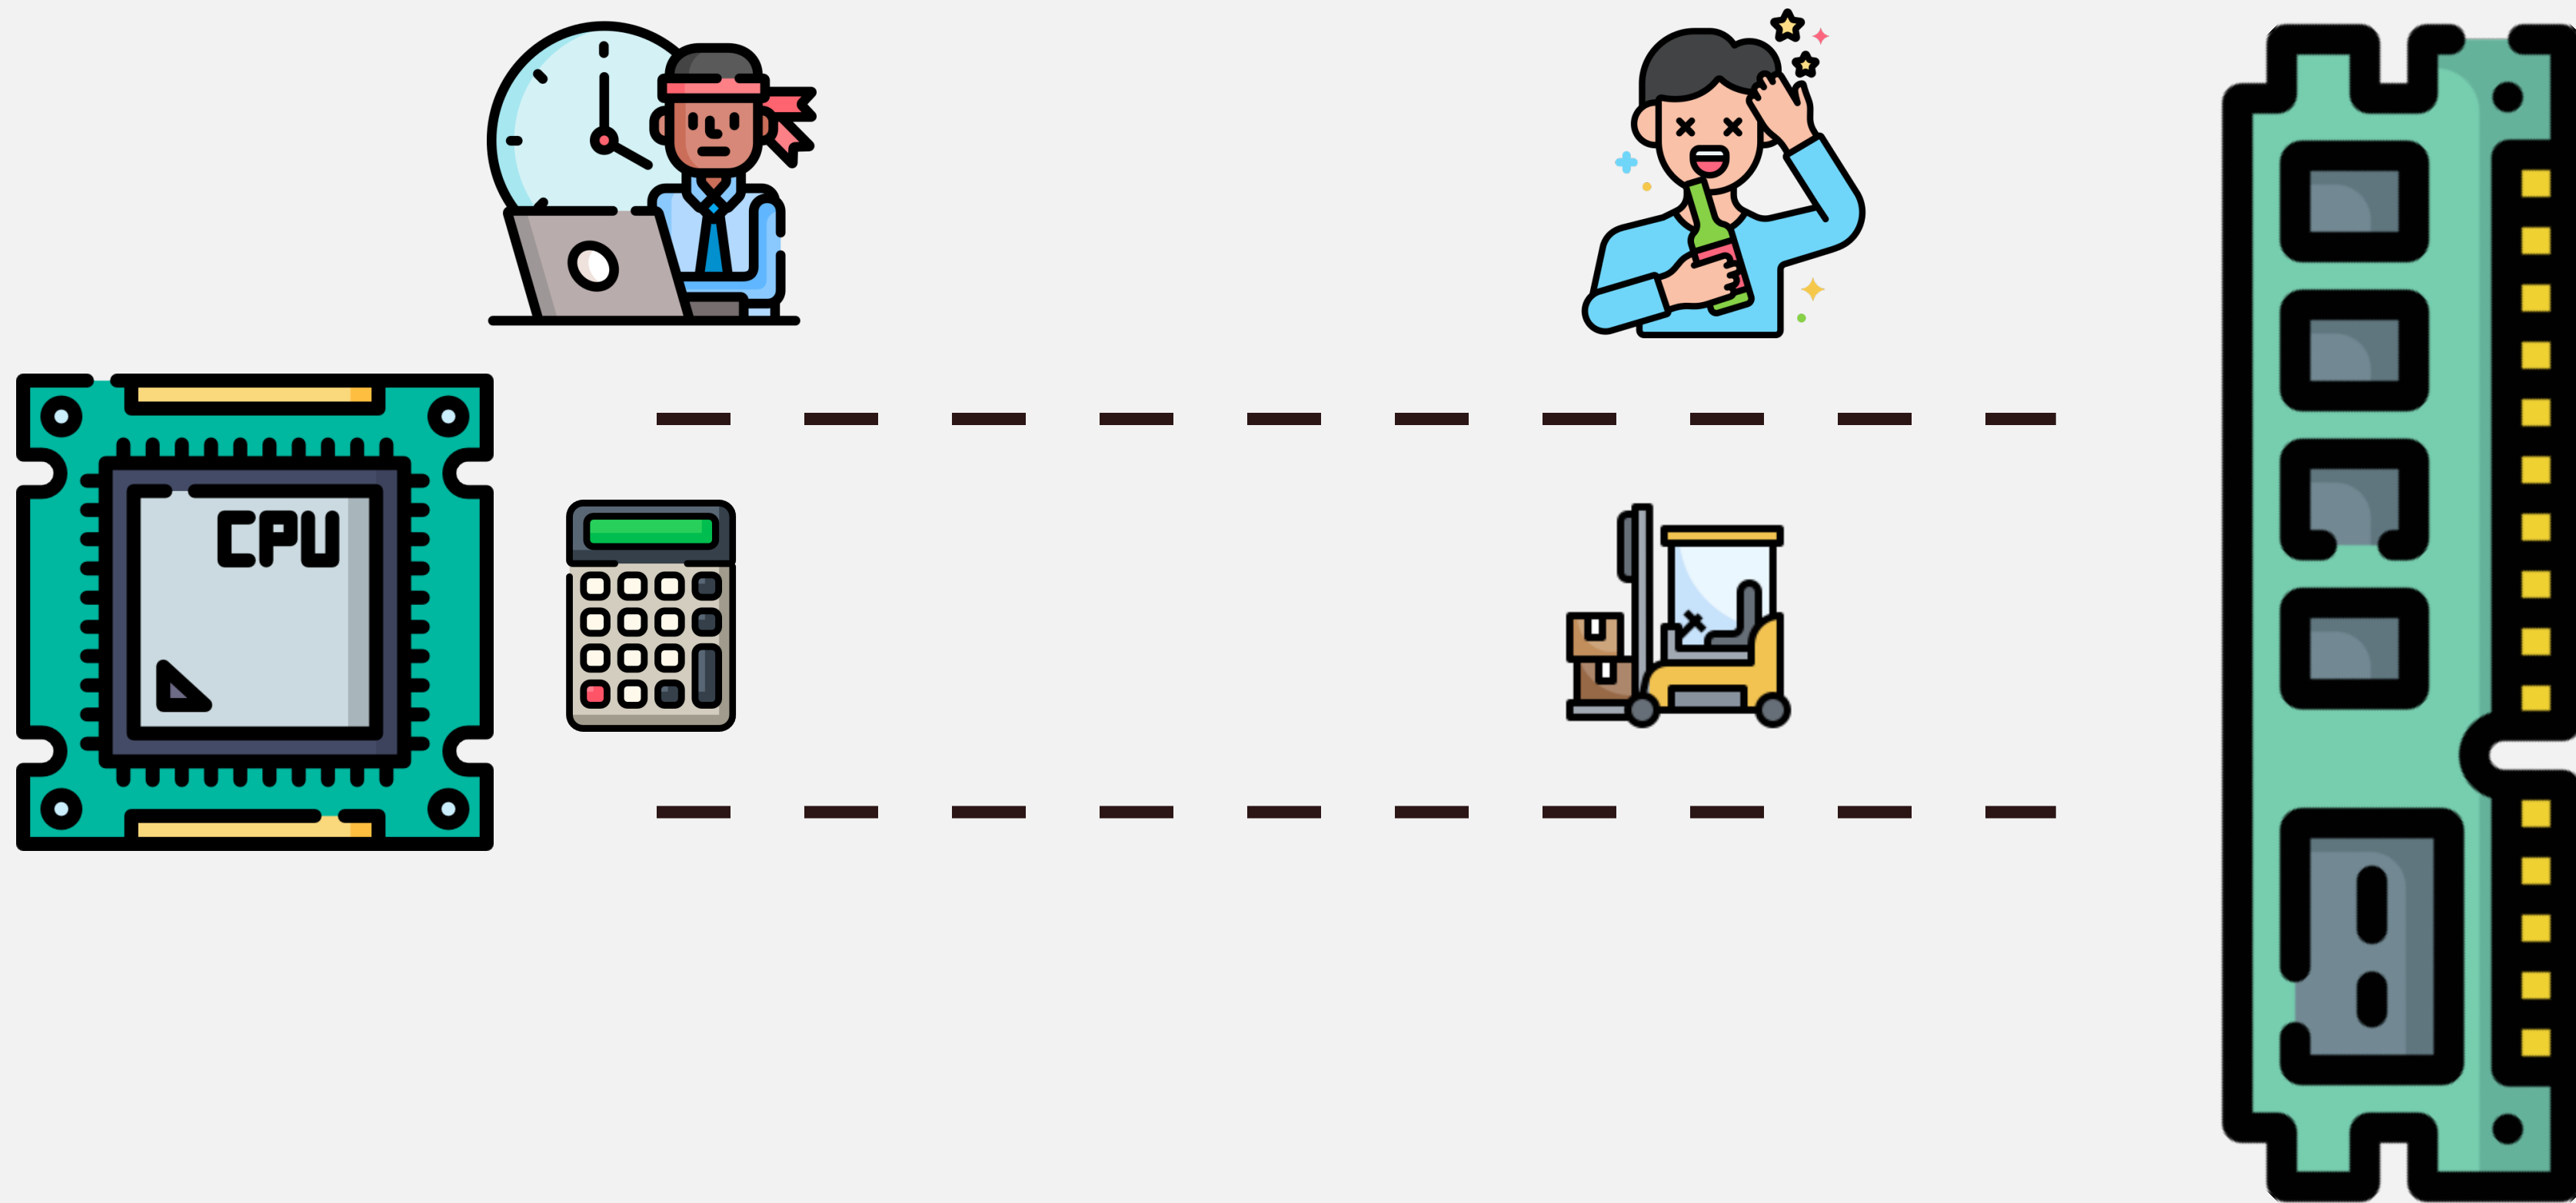


Who's the Bottleneck?

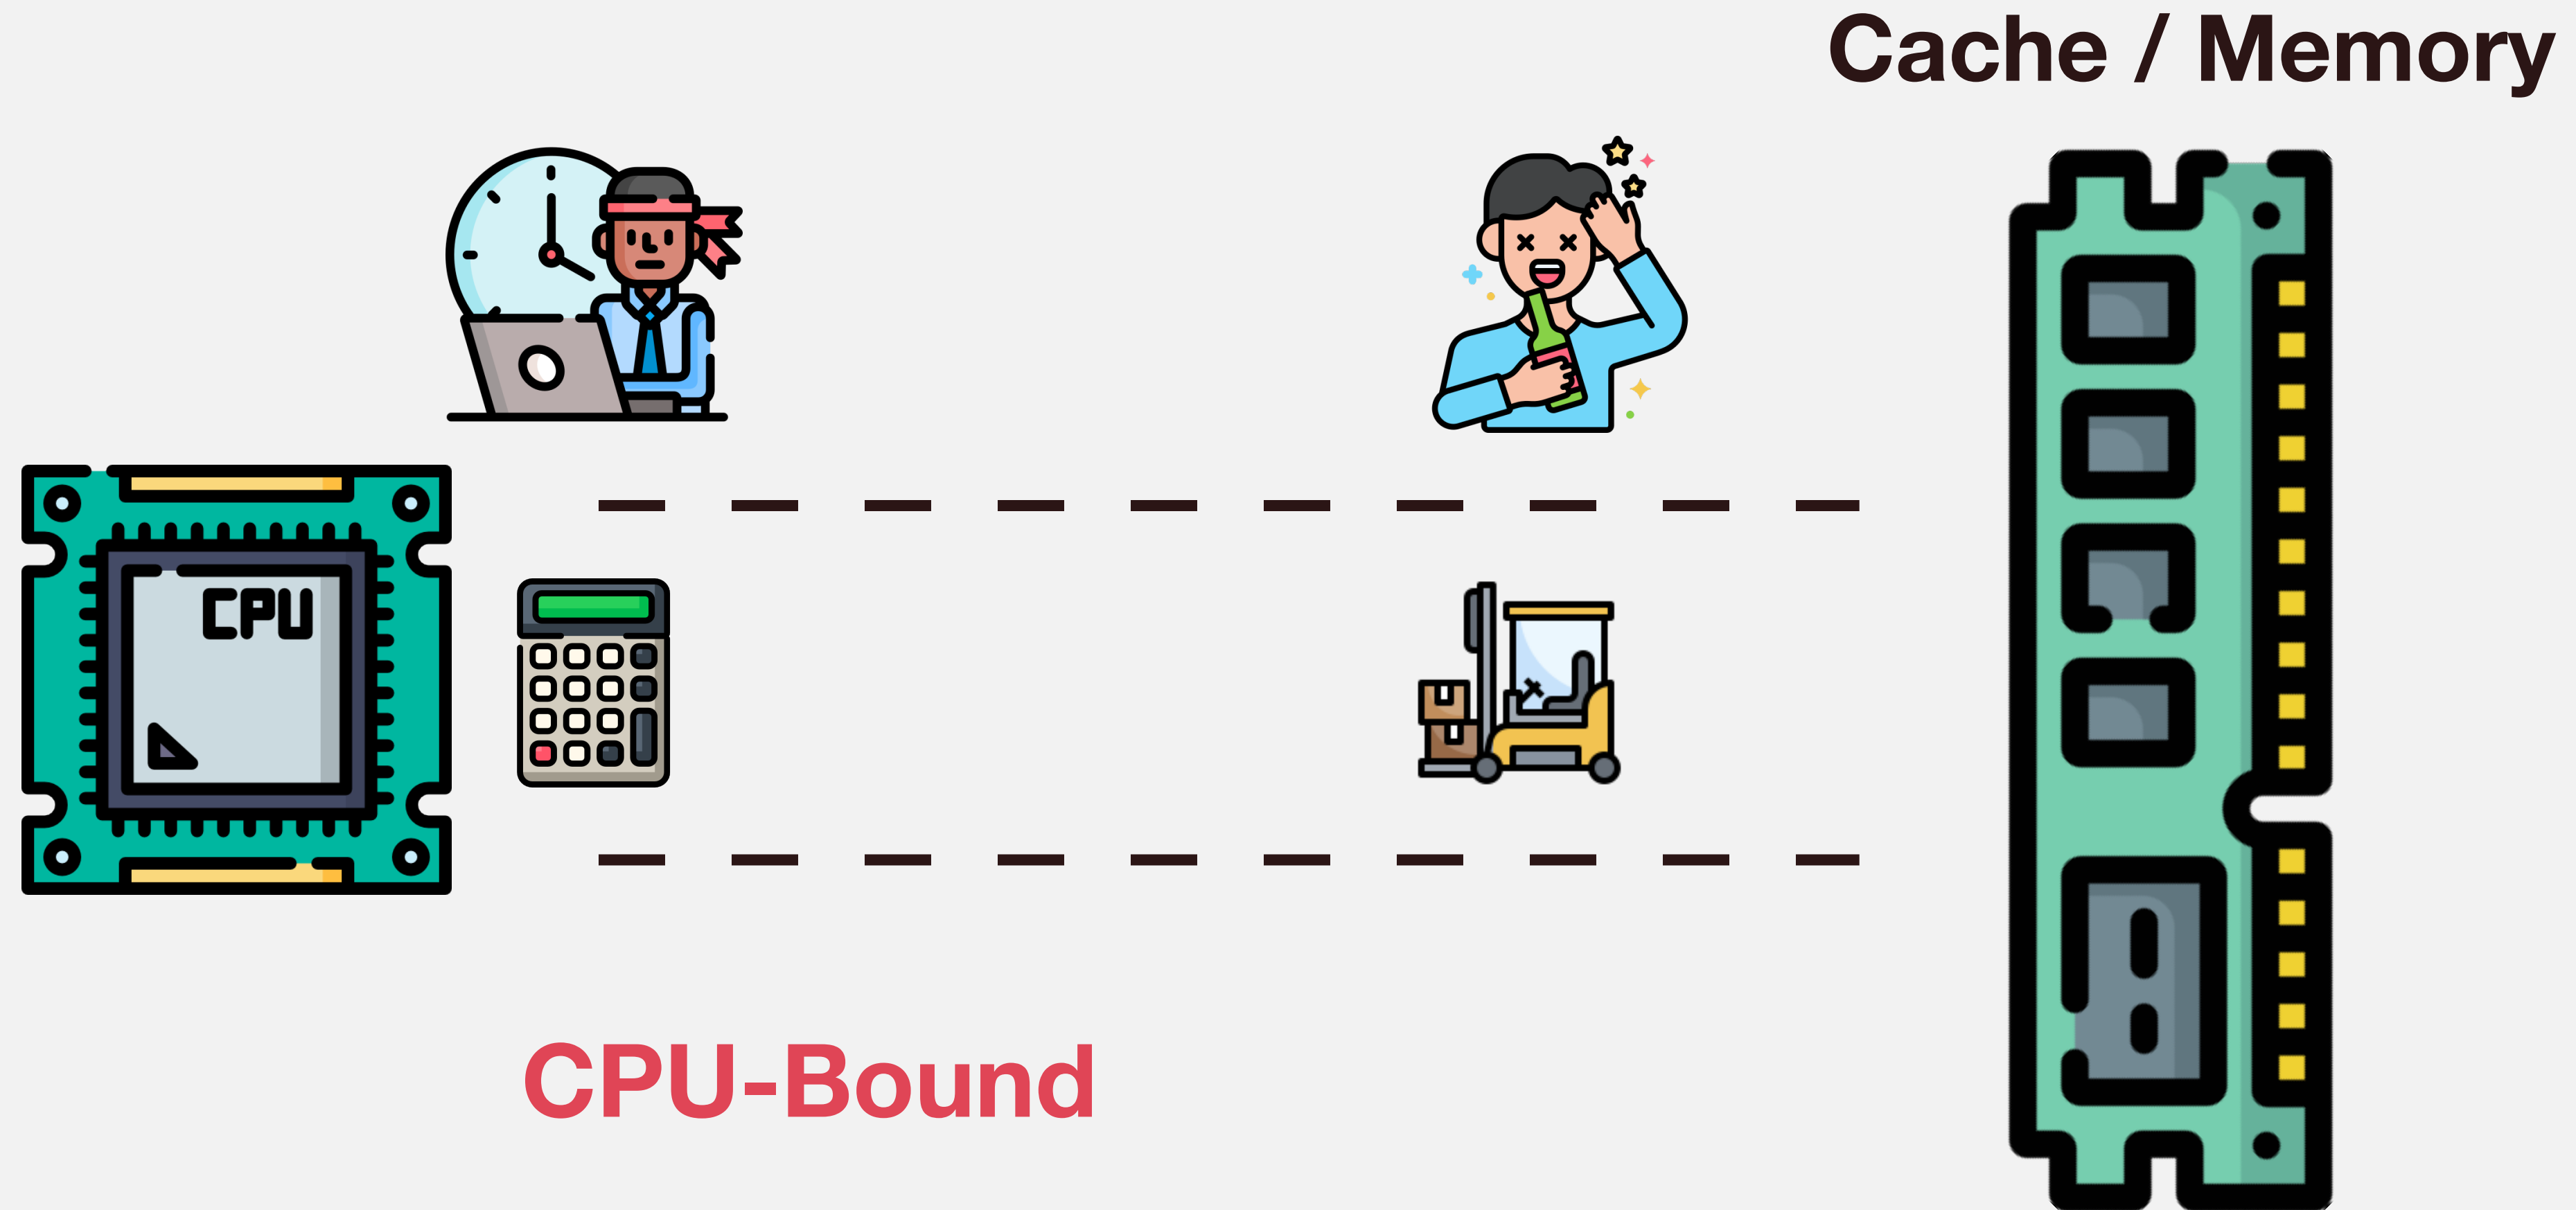


Who's the Bottleneck?

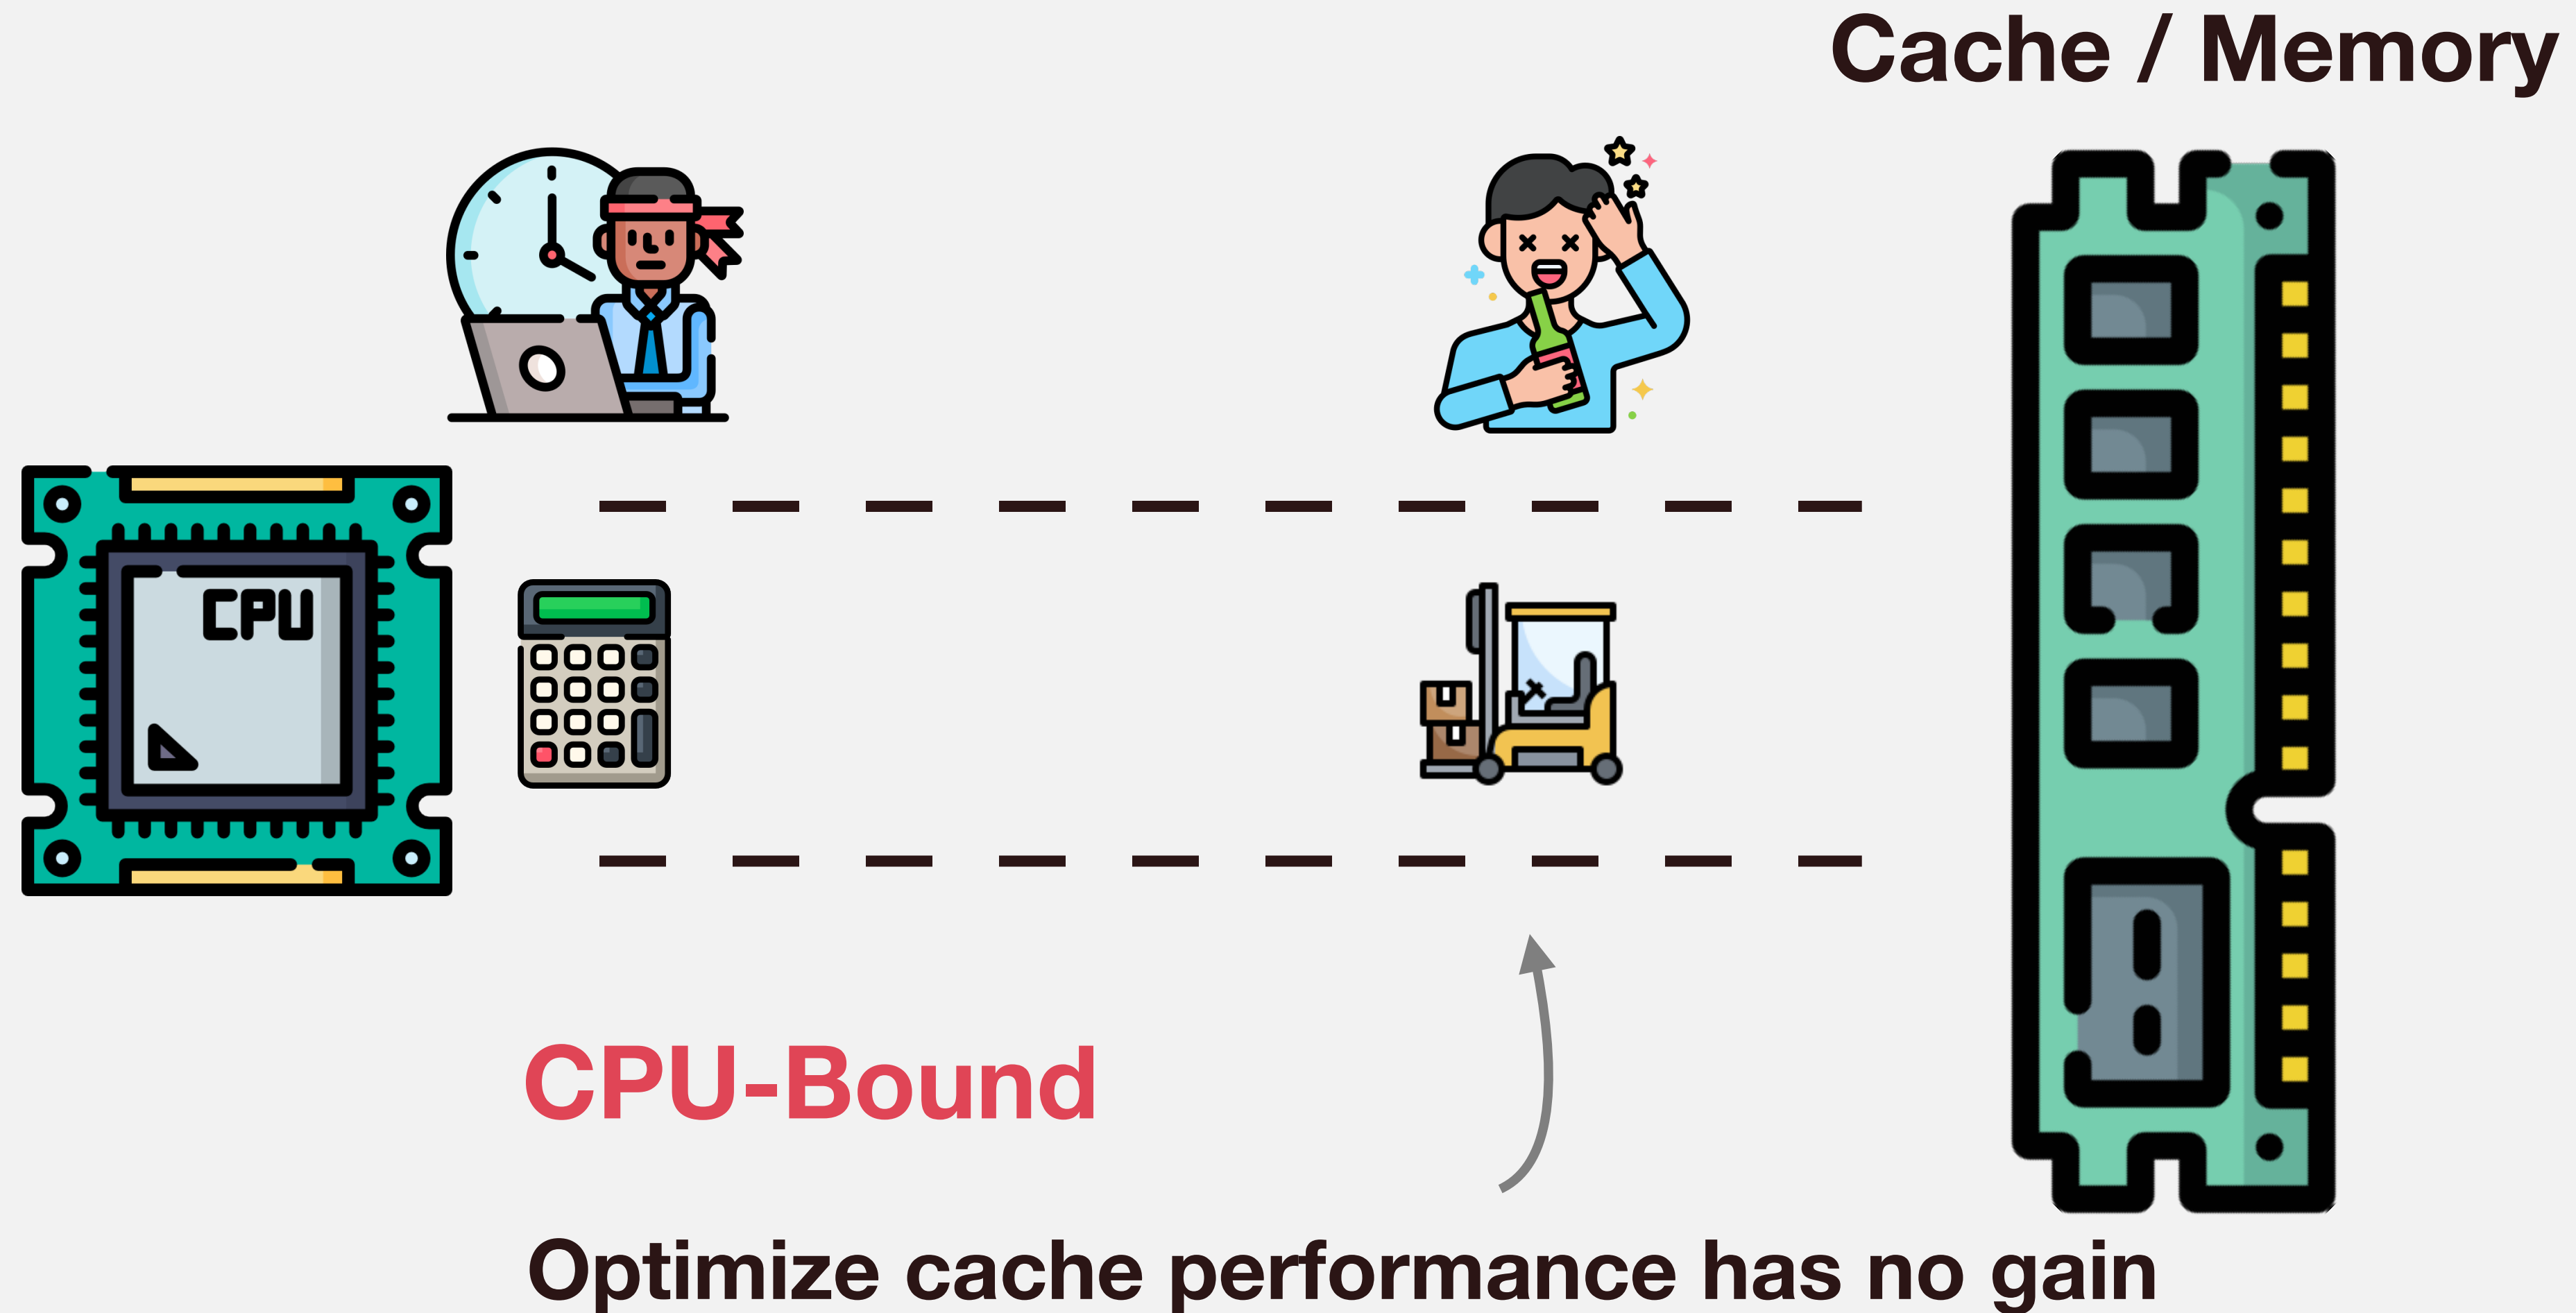
Cache / Memory



Who's the Bottleneck?



Who's the Bottleneck?



Don't Forget Amdahl's Law

$$\textit{Speedup} = \frac{1}{(1 - p) + \frac{p}{s}}$$

% of exec time benefit
from your optimization

speedup of that
portion of code

Optimizing Code Performance

→ Design efficient algorithms

Optimizing Code Performance

- Design efficient algorithms
- Eliminate unnecessary work

Optimizing Code Performance

- Design efficient algorithms
- Eliminate unnecessary work
- Leverage compiler optimization

Optimizing Code Performance

- Design efficient algorithms
- Eliminate unnecessary work
- Leverage compiler optimization
- Detect hotspots / bottleneck

Optimizing Code Performance

- Design efficient algorithms
- Eliminate unnecessary work
- Leverage compiler optimization
- Detect hotspots / bottleneck
- Computation optimization techniques

Optimizing Code Performance

- Design efficient algorithms
- Eliminate unnecessary work
- Leverage compiler optimization
- Detect hotspots / bottleneck
- Computation optimization techniques
- Cache optimization techniques

Optimizing Code Performance

- Design efficient algorithms
- Eliminate unnecessary work
- Leverage compiler optimization

- Detect hotspots / bottleneck
- Computation optimization techniques
- Cache optimization techniques

